

A Year of Artificial Intelligence

Where we make the mathematics, science, linguistics, and philosophy of artificial intelligence fun and simple.

[ALGORITHMS](#)[TODAY I LEARNED](#)[CASE STUDIES](#)[PHILOSOPHICAL](#)[META](#)[Follow](#)

New

Rohan & Lenny #3: Recurrent Neural Networks & LSTMs

The ultimate guide to machine learning's favorite child.



Rohan Kapur

Apr 13, 2017 · 85 min read




Rohan & Lenny #2: Convolutional Neural Networks

The ingenious algorithm that accelerated computer vision (and its biological roots).



Lenny Khazan
Nov 12, 2016 • 40 min read



Rohan #7: Can artificial intelligence possess human-like minds?

A philosophical question — not a technical one.



Rohan Kapur
Jun 5, 2016 • 14 min read



Rohan #6: The beautiful math behind logistic regression.

What most beginner ML courses won't teach you.



Rohan Kapur
May 21, 2016 • 4 min read



Lenny #2: Autoencoders and Word Embeddings

I want to take a quick break from my reinforcement learning endeavor (more on that soon) to talk about an interesting unsupervised learning...



Lenny Khazan
Apr 24, 2016 • 12 min read



Rohan #5: What are bias units?

This is the fifth entry in my journey to extend my knowledge of Artificial Intelligence in the year of 2016. Learn more about my motives in...



Rohan Kapur
Apr 8, 2016 • 5 min read

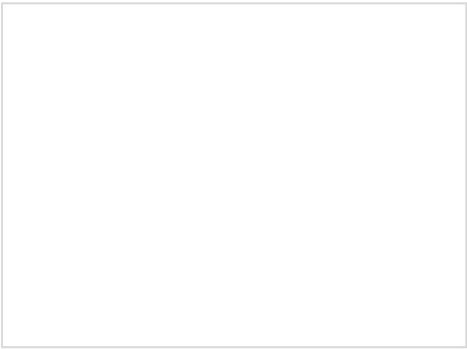


Rohan #4: The vanishing gradient problem

Oh no — an obstacle to deep learning!



Rohan Kapur
Mar 16, 2016 • 11 min read



Rohan & Lenny #1: Neural Networks & The Backpropagation Algorithm, Explained

Do you know the chain rule? Then you know the neural network backpropagation algorithm!



Rohan Kapur

Mar 4, 2016 • 30 min read



Rohan #3: Deriving the Normal Equation using matrix calculus

Understanding the analytical solution to linear regression.



Rohan Kapur

Feb 16, 2016 • 11 min read

Lenny #1: Robots + Reinforcement Learning

This is my first post for A Year of AI, a project that Rohan Kapur and I will be undertaking over the next year to learn more about the...



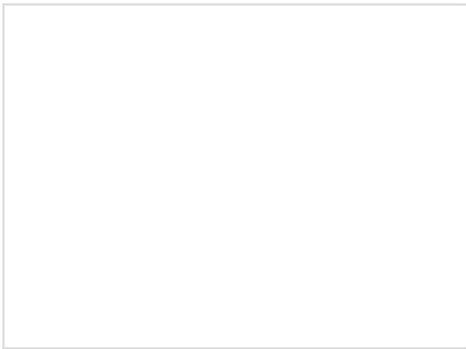
Lenny Khazan

Feb 16, 2016 • 5 min read



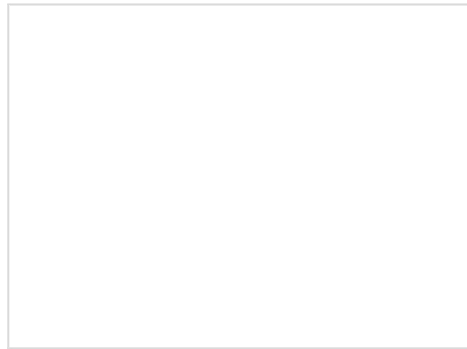
Rohan #2: Artificial intelligence, $\partial \text{Progress} / \partial \text{Time}$

A technical and philosophical analysis on the past, present, and future of artificial intelligence.



Rohan #1: Logistic regression case study — diagnosing cancer

How classification algorithms like Logistic Regression can understand the difference between malignant and



#0: 2016 is the year I venture into Artificial Intelligence.

From iPhone Apps to Neural Networks.



Rohan Kapur

Jan 11, 2016 • 4 min read



Rohan Kapur

Feb 9, 2016 · 43 min read

benign breast cancer tumors.



Rohan Kapur

Jan 22, 2016 · 34 min read

[About A Year of Artificial Intelligence](#) · [Latest Stories](#) · [Archive](#) · [About Medium](#) · [Terms](#) · [Privacy](#)