

Stochastic Gradient Descent

Minimize an objective function using a stochastic approximation of gradient descent.

Inputs

- Data: input dataset
- Preprocessor: preprocessing method(s)

Outputs

- Learner: stochastic gradient descent learning algorithm
- Model: trained model

The **Stochastic Gradient Descent** widget uses **stochastic gradient descent** that minimizes a chosen loss function with a linear function. The algorithm approximates a true gradient by considering one sample at a time, and simultaneously updates the model based on the gradient of the loss function. For regression, it returns predictors as minimizers of the sum, i.e. M-estimators, and is especially useful for large-scale and sparse datasets.

Stochastic Gradient Descent

Name 1

SGD

Algorithm 2

Classification loss function: Hinge

ϵ : 0,10

Regression loss function: Squared Loss

ϵ : 0,10

Regularization

3

Regularization method: Ridge (L2)

Regularization strength (α): 0,00001

Mixing parameter: 0,15

Learning parameters

4

Learning rate: Constant

Initial learning rate (η_0): 0,0100

Inverse scaling exponent (t): 0,2500

Number of iterations: 5

☒ Shuffle data after each iteration

Fixed seed for random shuffling: 0

5 Report 6 Apply Automatically

1. Specify the name of the model. The default name is “SGD”.

2. Algorithm parameters:

- Classification loss function:
 - **Hinge** (linear SVM)
 - **Logistic Regression** (logistic regression SGD)
 - **Modified Huber** (smooth loss that brings tolerance to outliers as well as probability estimates)
 - **Squared Hinge** (quadratically penalized hinge)
 - **Perceptron** (linear loss used by the perceptron algorithm)
 - **Squared Loss** (fitted to ordinary least-squares)
 - **Huber** (switches to linear loss beyond ϵ)
 - **Epsilon insensitive** (ignores errors within ϵ , linear beyond it)
 - **Squared epsilon insensitive** (loss is squared beyond ϵ -region).
- Regression loss function:
 - **Squared Loss** (fitted to ordinary least-squares)
 - **Huber** (switches to linear loss beyond ϵ)
 - **Epsilon insensitive** (ignores errors within ϵ , linear beyond it)
 - **Squared epsilon insensitive** (loss is squared beyond ϵ -region).

3. Regularization norms to prevent overfitting:

- None.
- **Lasso (L1)** (L1 leading to sparse solutions)
- **Ridge (L2)** (L2, standard regularizer)
- **Elastic net** (mixing both penalty norms).

Regularization strength defines how much regularization will be applied (the less we regularize, the more we allow the model to fit the data) and the mixing parameter what the ratio between L1 and L2 loss will be (if set to 0 then the loss is L2, if set to 1 then it is L1).

4. Learning parameters.

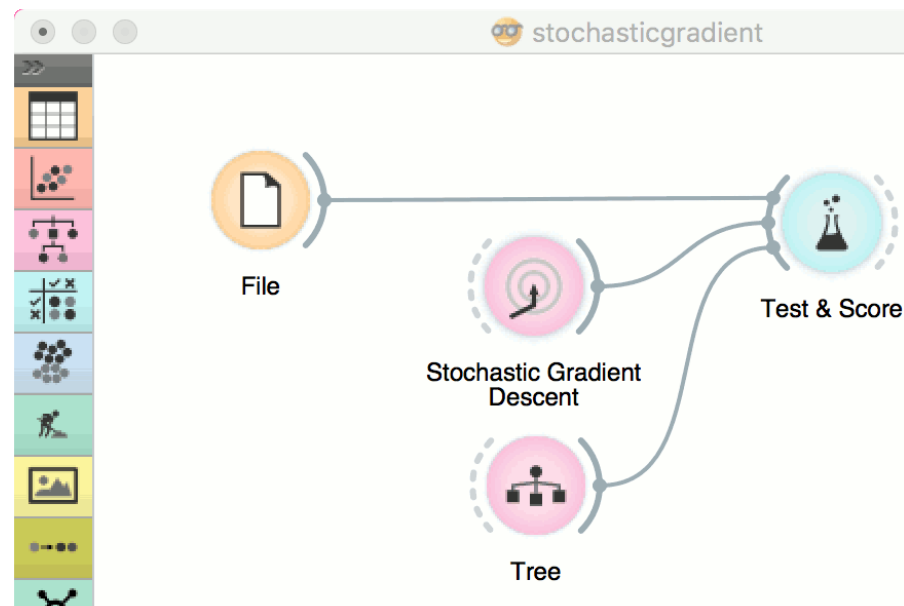
- Learning rate:
 - *Constant*: learning rate stays the same through all epochs (passes)
 - *Optimal*: a heuristic proposed by Leon Bottou
 - *Inverse scaling*: learning rate is inversely related to the number of iterations
- Initial learning rate.
- Inverse scaling exponent: learning rate decay.
- Number of iterations: the number of passes through the training data.
- If *Shuffle data after each iteration* is on, the order of data instances is mixed after each pass.
- If *Fixed seed for random shuffling* is on, the algorithm will use a fixed random seed and enable replicating the results.

5. Produce a report.

6. Press *Apply* to commit changes. Alternatively, tick the box on the left side of the *Apply* button and changes will be communicated automatically.

Examples

For the classification task, we will use *iris* dataset and test two models on it. We connected *Stochastic Gradient Descent* and *Tree* to *Test & Score*. We also connected *File* to *Test & Score* and observed model performance in the widget.



Stochastic Gradient Descent

Name

SGD

Algorithm

Classificaton loss function:

Hinge

ε: 0,10

Regression loss function:

Squared Loss

ε: 0,10

Regularization

Regularization method:

Ridge (L2)

Regularization strength (α): 0,00001

Mixing parameter: 0,15

Test & Score

Sampling

☒ Cross validation

Number of folds: 10

☒ Stratified

☐ Random sampling

Repeat train/test: 10

Training set size: 66 %

☒ Stratified

☐ Leave one out

☐ Test on train data

☐ Test on test data

Target Class

(Average over classes)

Report

Evaluation Results

Method	AUC	CA	F1	Precision	Recall
SGD	0.870	0.827	0.823	0.842	0.827
Tree	0.975	0.960	0.960	0.960	0.960

Learning parameters

Learning rate: Constant

Initial learning rate (η_0): 0,0100

Inverse scaling exponent (t): 0,2500


Number of iterations: 5

☒ Shuffle data after each iteration

☐ Fixed seed for random shuffling: 0

Report ☒ Apply Automatically

For the regression task, we will compare three different models to see which predict what kind of results. For the purpose of this example, the *housing* dataset is used. We connect the **File** widget to **Stochastic Gradient Descent**, **Linear Regression** and **kNN** widget and all four to the **Predictions** widget.



Stochastic Gradient Descent

Name: SGD


Algorithm

Classification loss function: Hinge

ϵ : 0,10

Regression loss function: Squared Loss

ϵ : 0,10



Regularization

Regularization method: Ridge (L2)Regularization strength (α): 0,00001Mixing parameter: 0,15

Learning parameters

Learning rate: Constant0,01000,25005

Iteration

Shuffling:

0Apply Automatically

Predictions

Info

Data: 506 instances.
Predictors: 3
Task: Regression

Restore Original Order

Data View

☒ Show full data set

Output

☒ Original data☒ Predictions☒ ProbabilitiesReport

	SGD	Linear Regression	kNN	MEDV	CRIM	ZN
1	28.292	30.004	21.780	24.000	0.006	18.000
2	24.281	25.026	22.900	21.600	0.027	0.000
3	27.426	30.568	25.360	34.700	0.027	0.000
4	25.272	28.607	26.060	33.400	0.032	0.000
5	24.817	27.944	27.100	36.200	0.069	0.000
6	23.460	25.256	27.100	28.700	0.030	0.000
7	23.308	23.002	20.880	22.900	0.088	12.500
8	20.959	19.536	19.100	27.100	0.145	12.500
9	15.658	11.524	18.400	16.500	0.211	12.500
10	20.186	18.920	19.480	18.900	0.170	12.500
11	20.226	18.999	19.280	15.000	0.225	12.500
12	22.214	21.587	22.000	18.900	0.117	12.500
13	21.613	20.907	24.340	21.700	0.094	12.500