

Deep Learning for Speech/Language Processing

--- machine learning & signal processing perspectives

Li Deng

Deep Learning Technology Center

Microsoft Research, Redmond, USA

Tutorial given at Interspeech, Sept 6, 2015

Thanks go to many colleagues at DLTC/MSR, collaborating universities,
and at Microsoft's engineering groups

Outline

- Part I: Basics of Machine Learning (Deep and Shallow) and of Signal Processing
- Part II: Speech
- Part III: Language

(In case you did not get link to slides, send email to:
alexander.raake@tu-ilmenau.de)

Reading Material

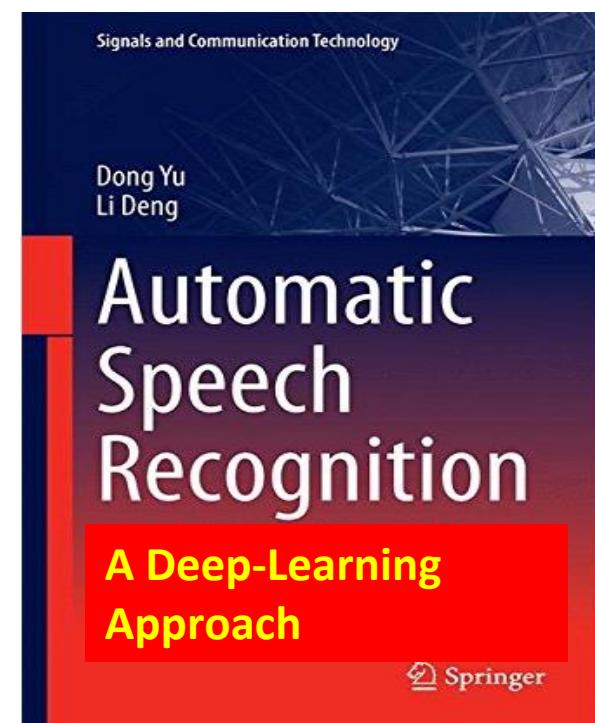
Books:

Bengio, Yoshua (2009). "[Learning Deep Architectures for AI](#)".

L. Deng and D. Yu (2014) "Deep Learning: Methods and Applications"

<http://research.microsoft.com/pubs/209355/DeepLearning-NowPublishing-Vol7-SIG-039.pdf>

D. Yu and L. Deng (2014). "Automatic Speech Recognition: A Deep Learning Approach" (Publisher: Springer).



DEEP LEARNING

An MIT Press book in preparation

Reading Material

Yoshua Bengio, Ian Goodfellow and Aaron Courville

- [Table of Contents](#)
- [Deep Learning for AI](#)
 - 1. [Part header: applied math and machine learning basics](#)
 - [Linear Algebra](#)
 - [Probability and Information Theory](#)
 - [Numerical Computation](#)
 - [Machine Learning Basics](#)
 - 2. [Part header: modern practical deep networks](#)
 - [Feedforward Deep Networks](#)
 - [Regularization](#)
 - [Numerical Optimization](#)
 - [Convolutional Networks](#)
 - [Sequence Modeling: Recurrent and Recursive Nets](#)
 - 3. [Part header: deep learning research](#)
 - [Structured Probabilistic Models: A Deep Learning Perspective](#)
 - [Monte-Carlo Methods](#)
 - [Linear Factor Models and Auto-Encoders](#)
 - [Representation Learning](#)
 - [The Manifold Perspective on Representation Learning](#)
 - [Confronting the Partition Function](#)
 - [Approximate Inference](#)
 - [Deep Generative Models](#)
- [References](#)

Robust Automatic Speech Recognition

A Bridge to Practical Applications

Jinyu Li
Li Deng
Reinhold Haeb-Umbach
Yifan Gong



ISBN: 978-0-12-802398-3

PUB DATE: October 2015



MACHINE LEARNING

A BAYESIAN AND OPTIMIZATION PERSPECTIVE

SERGIOS THEODORIDIS



Reading Material (cont'd)

Wikipedia:

https://en.wikipedia.org/wiki/Deep_learning

Papers:

G. E. Hinton, R. Salakutdinov. "[Reducing the Dimensionality of Data with Neural Networks](#)". *Science* **313**: 504–507, 2016.

G. E. Hinton, L. Deng, D. Yu, etc. "Deep Neural Networks for Acoustic Modeling in Speech Recognition: The shared views of four research groups," *IEEE Signal Processing Magazine*, pp. 82–97, November 2012.

G. Dahl, D. Yu, L. Deng, A. Acero. "[Context-Dependent Pre-Trained Deep Neural Networks for Large-Vocabulary Speech Recognition](#)". *IEEE Trans. Audio, Speech, and Language Processing*, Vol 20(1): 30–42, 2012. (plus other papers in the same special issue)

Y. Bengio, A. Courville, and P. Vincent. "Representation Learning: A Review and New Perspectives," *IEEE Trans. PAMI, special issue Learning Deep Architectures*, 2013.

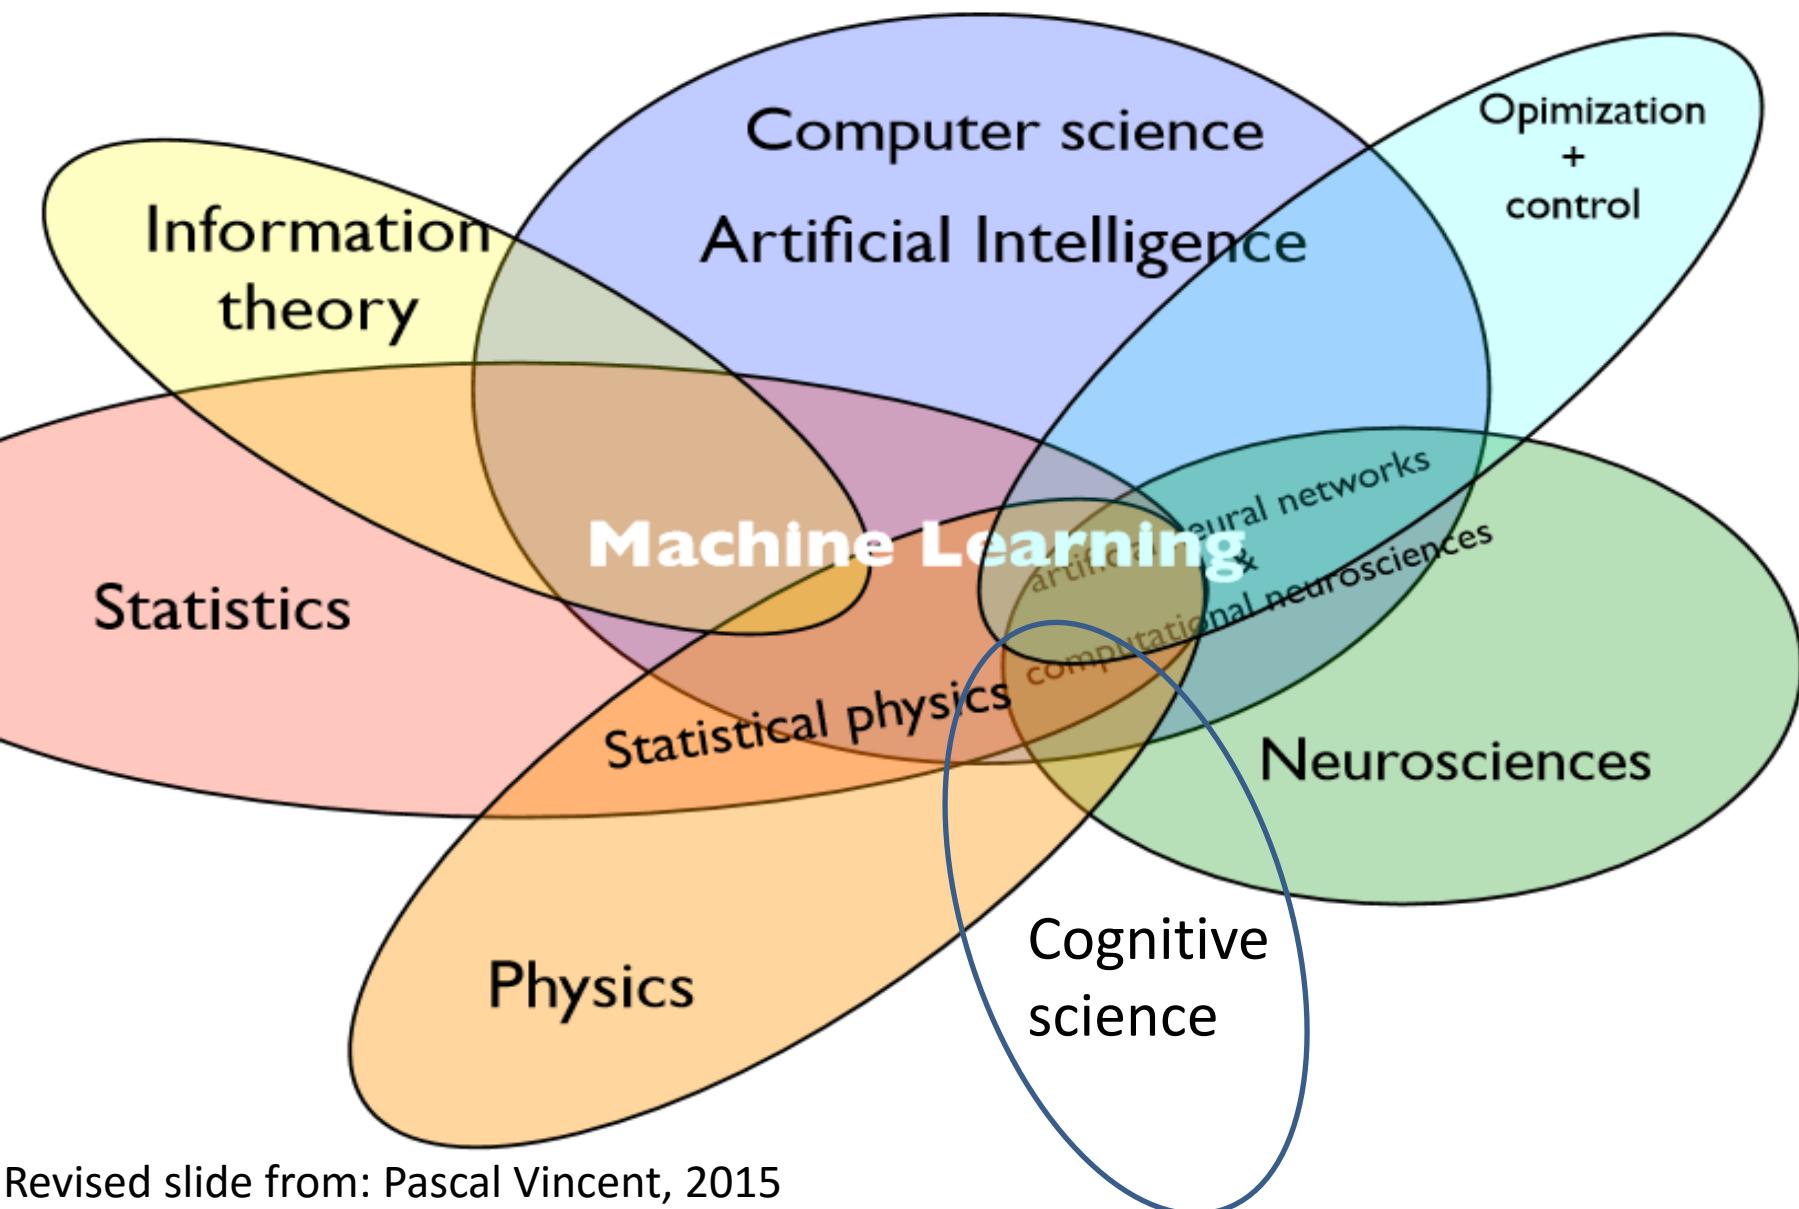
J. Schmidhuber. "Deep learning in neural networks: An overview," arXiv, October 2014.

Y. LeCun, Y. Bengio, and G. Hinton. "Deep Learning", Nature, Vol. 521, May 2015.

J. Bellegarda and C. Monz. "State of the art in statistical methods for language and speech processing," *Computer Speech and Language*, 2015

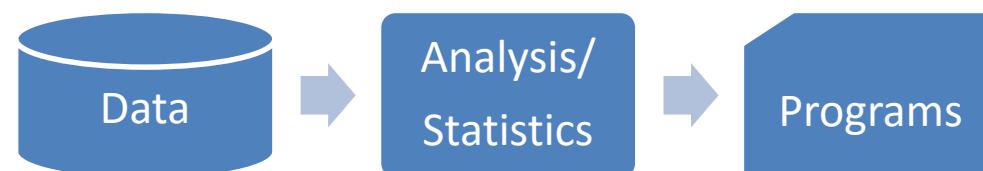
Part I: Machine Learning (Deep/Shallow) and Signal Processing

Current view of ML founding disciplines



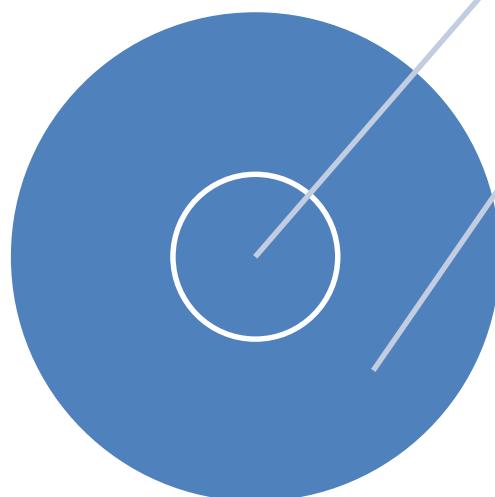
Machine Learning & Deep Learning

Machine learning

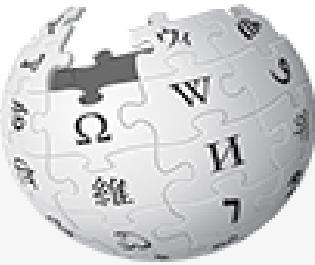


Deep learning

Machine
learning



What Is Deep Learning?



WIKIPEDIA
The Free Encyclopedia

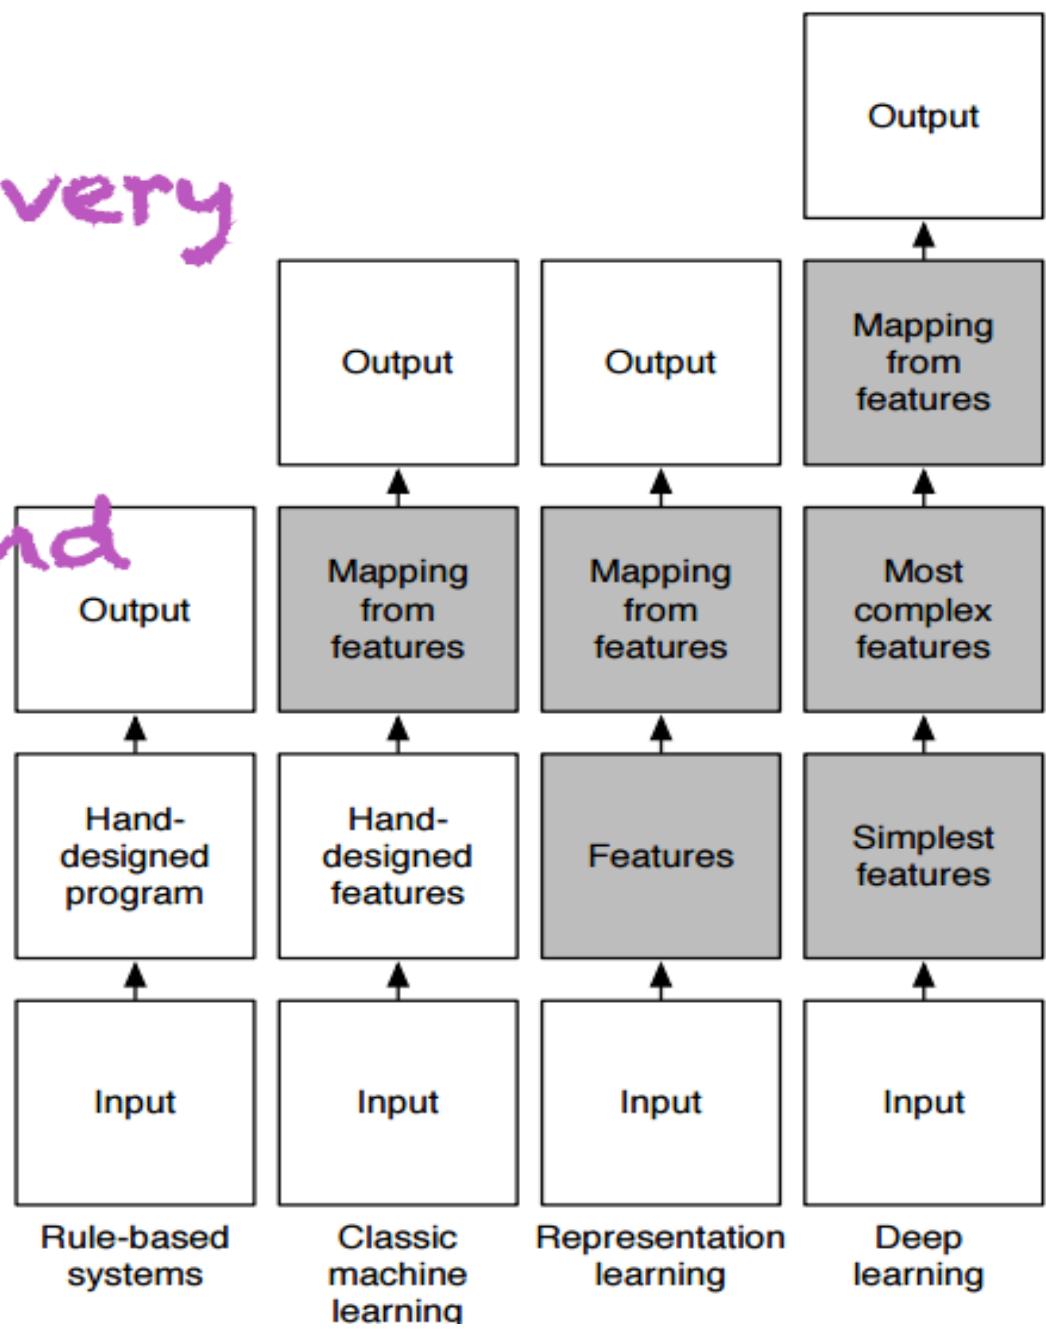
Deep learning

From Wikipedia, the free encyclopedia

Deep learning (*deep machine learning*, or *deep structured learning*, or *hierarchical learning*, or sometimes *DL*) is a branch of [machine learning](#) based on a set of [algorithms](#) that attempt to model high-level abstractions in data by using model architectures, with complex structures or otherwise, composed of multiple [non-linear transformations](#).^{[1](p198)[2][3][4]}

Automating Feature Discovery

Discovering and
representing
higher-level
abstractions



(Slide from: Yoshua Bengio)



SHALLOW

Boosting

Perceptron

AE

SVM

RBM

GMM Sparse
Coding

DecisionTree

Modified from

DEEP

Y LeCun
MA Ranzato

Neural Net

D-AE

RNN

Conv. Net

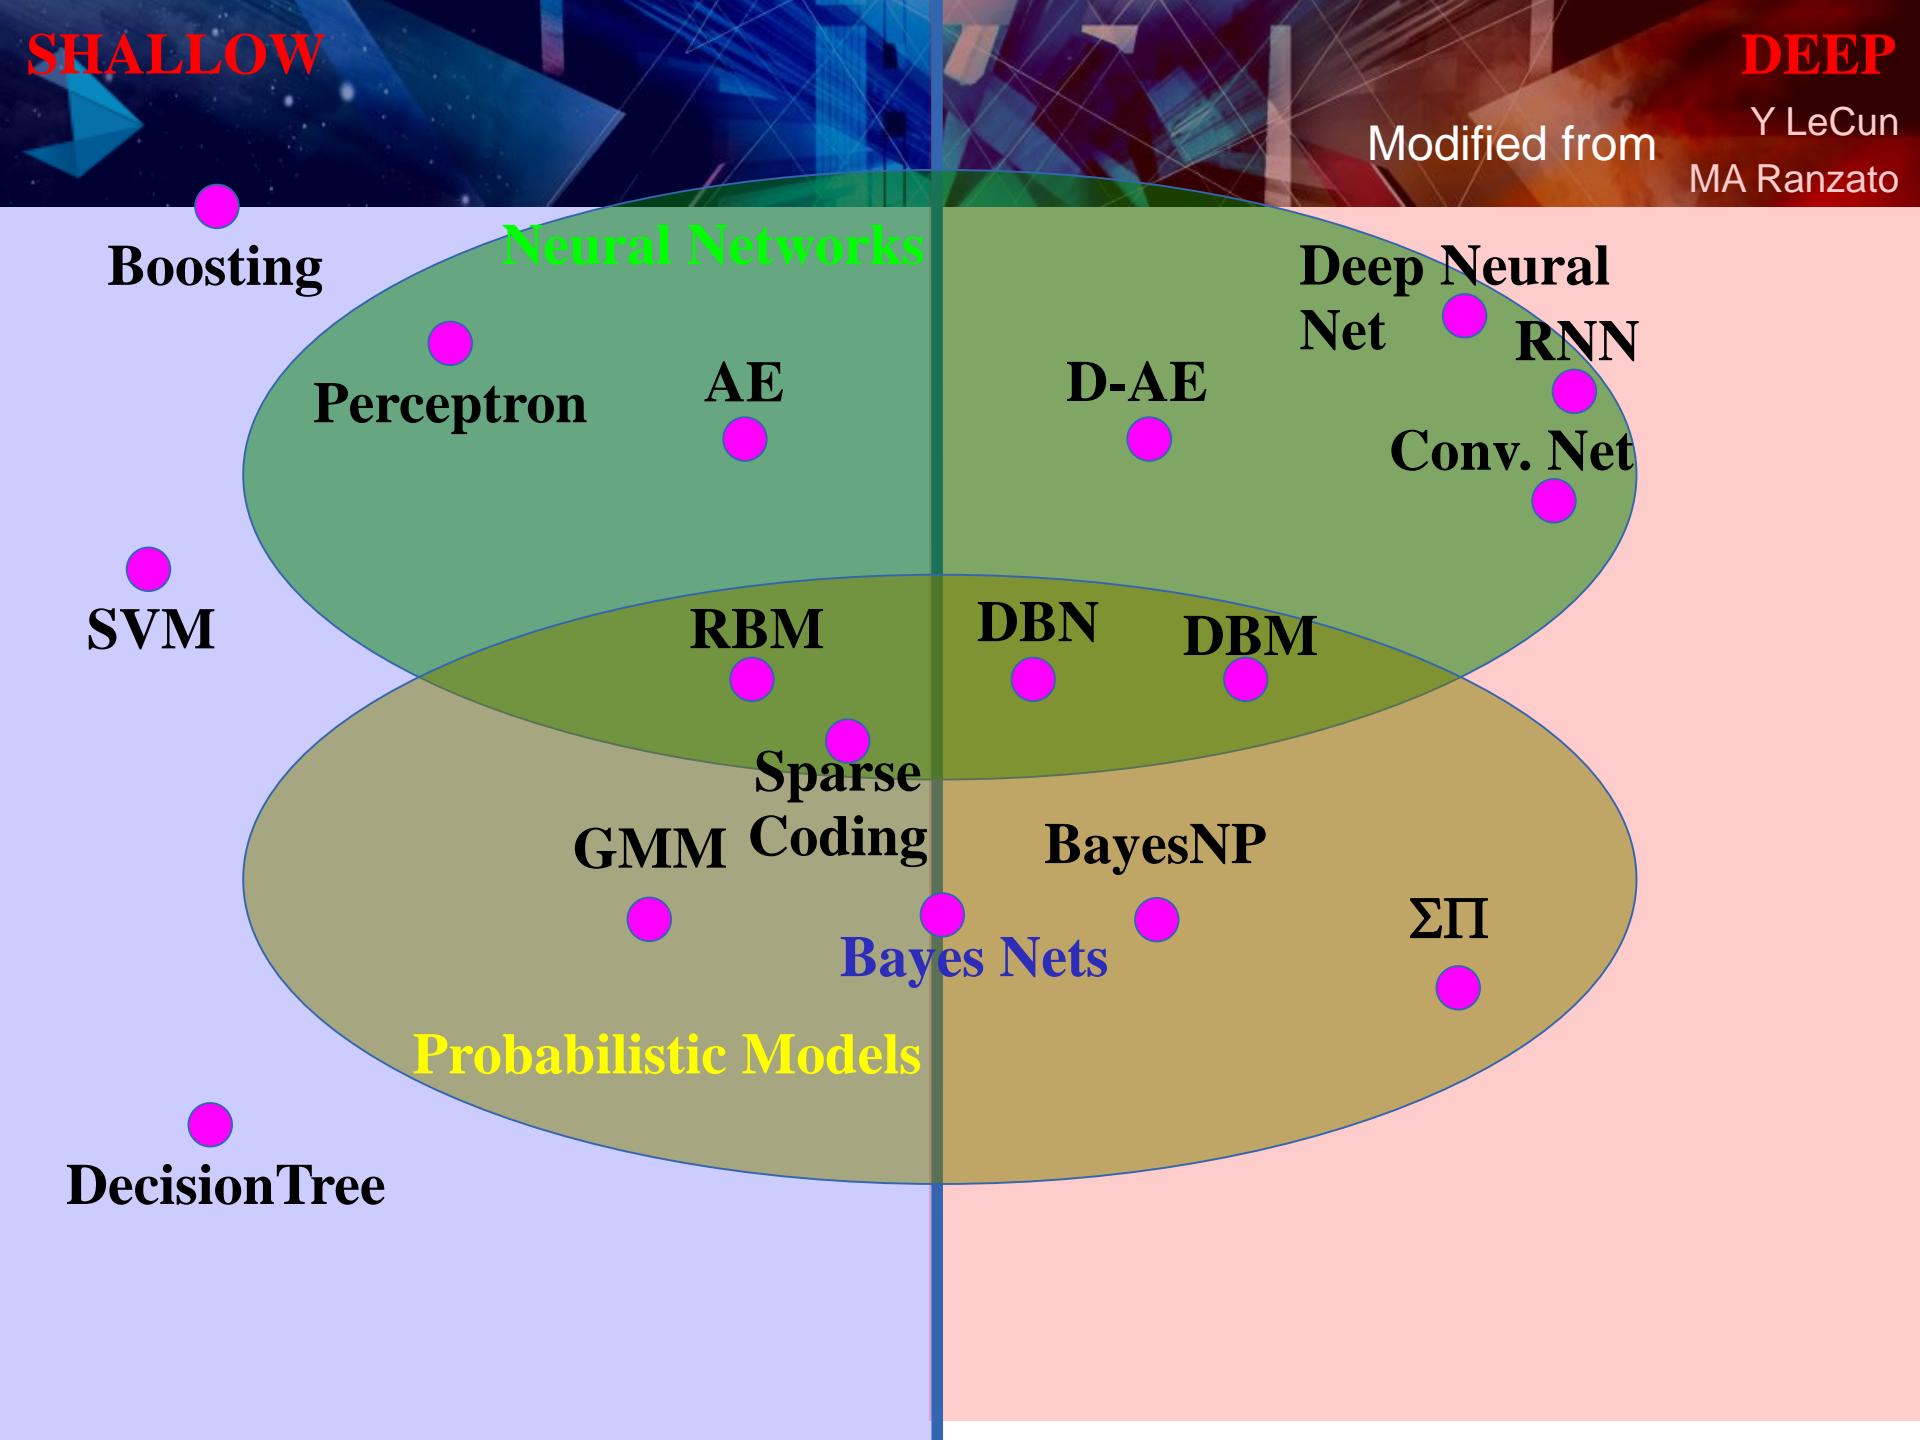
DBN

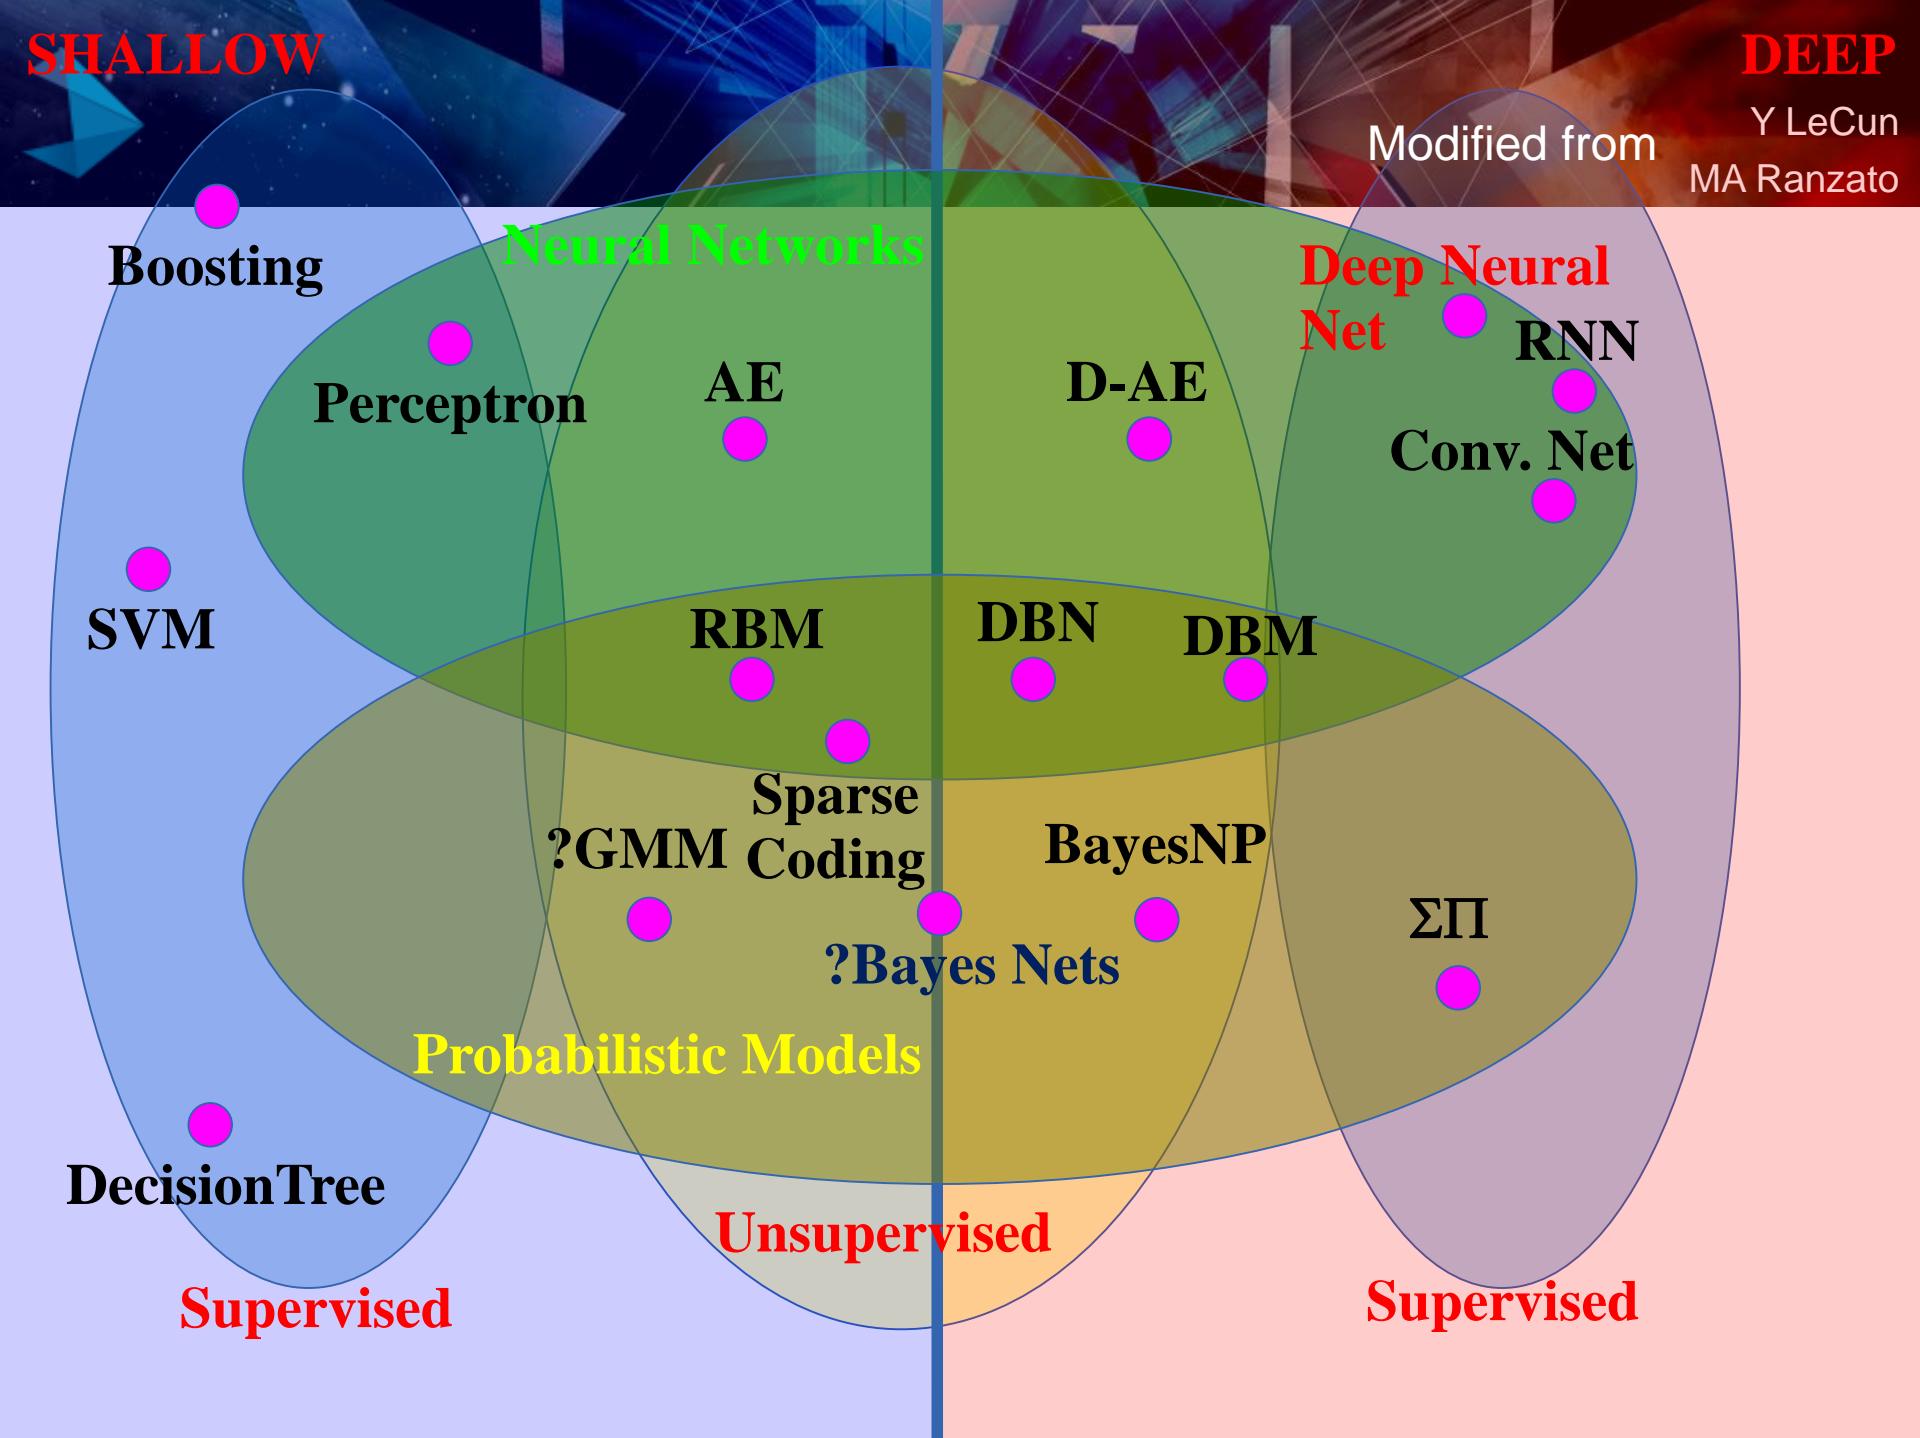
DBM

BayesNP

Bayes Nets

$\Sigma \Pi$





Signal Processing → Information Processing

Signals Processing	Audio/Music	Speech	Image/Animation/Graphics	Video	Text/Language
Coding/Compression	Audio Coding	Speech Coding	Image Coding	Video Coding	Document Compression/Summary
Communication	Voice over IP, DAB,etc		4G/5G Networks, DVB, Home Networking, etc		
Security	Multimedia watermarking, encryption, etc.				
Enhancement/Analysis	De-noising/Source separation	Speech Enhancement/Feature extraction	Image/video enhancement (Clear Type), Segmentation, feature extraction		Grammar checking, Text Parsing
Synthesis/Rendering	Computer Music	Speech Synthesis (text-to-speech)	Computer Graphics/	Video Synthesis	Natural Language Generation
User-Interface	Multi-Modal Human Computer Interaction (HCI --- Input Methods)				
Recognition	Auditory Scene Analysis (Computer audition; e.g. Melody Detection & Singer ID)	Automatic Speech/Speaker Recognition	Image Recognition	Computer Vision (e.g. 3-D object recognition)	Document Recognition
Understanding (Semantic IE)		Spoken Language Understanding	Image Understanding		Natural Language Understanding
Retrieval/Mining	Music Retrieval	Spoken Document Retrieval & Voice/Mobile Search	Image Retrieval	Video Search	Text Search (info retrieval)
translation		Speech translation			Machine translation
Social Media Apps	.	.	Photo Sharing (e.g. flickr)	Video Sharing (e.g. Youtube)	Blogs, Wiki, delicio.us

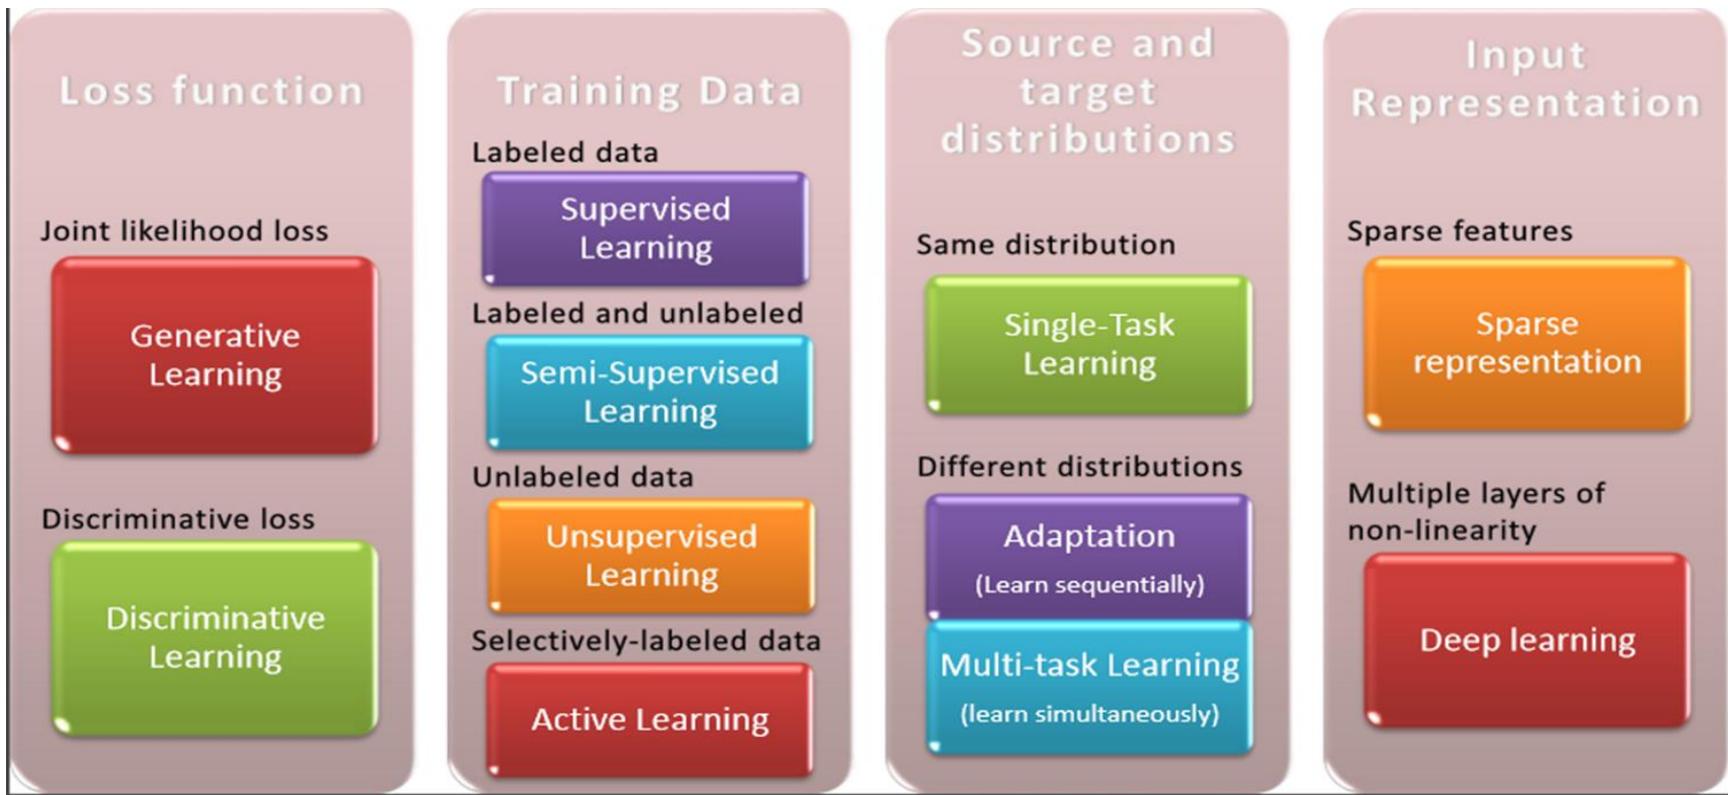
Machine Learning Basics

Machine Learning Paradigms for Speech Recognition: An Overview

Li Deng, *Fellow, IEEE*, and Xiao Li, *Member, IEEE*

Abstract—Automatic Speech Recognition (ASR) has historically been a driving force behind many machine learning (ML) techniques, including the ubiquitously used hidden Markov model, discriminative learning, structured sequence learning, Bayesian learning, and adaptive learning. Moreover, ML can and

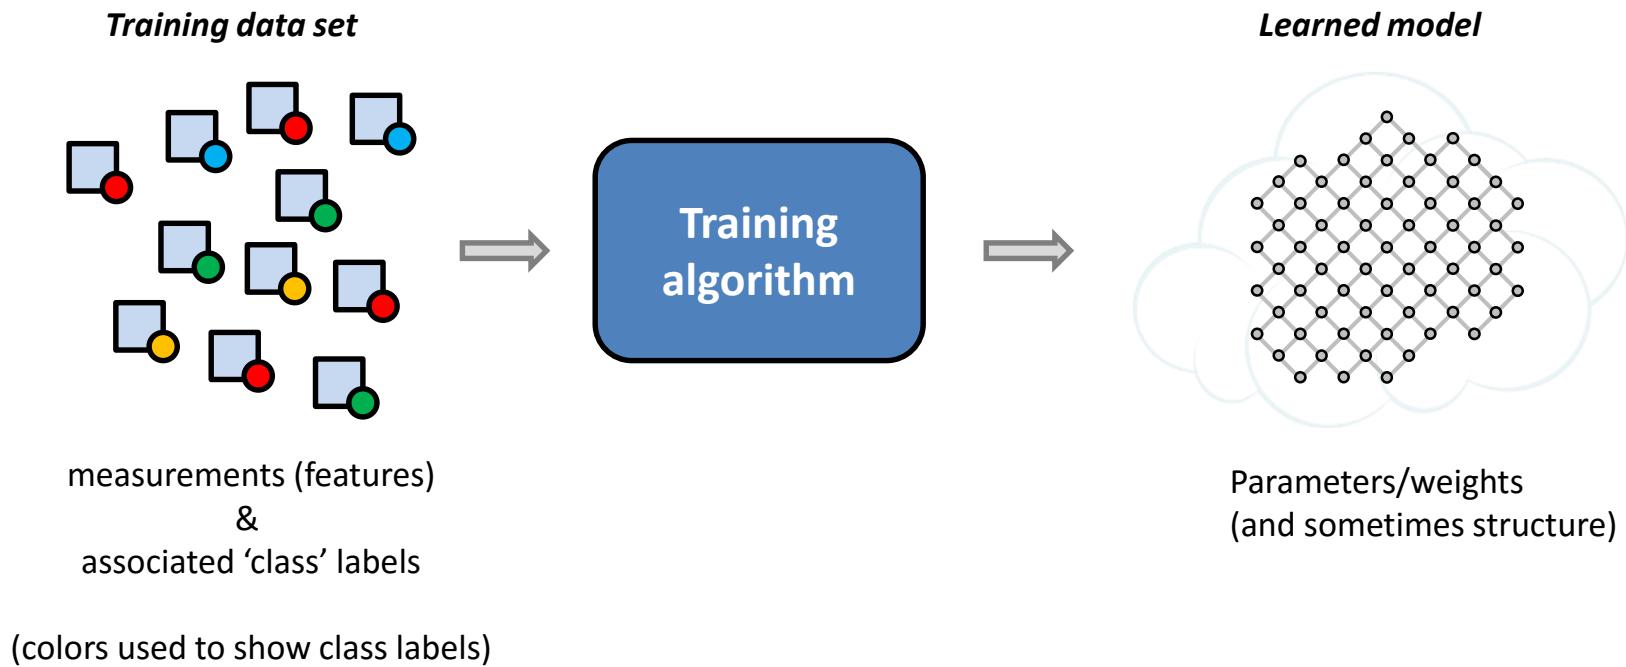
community to make assumptions about a problem, develop precise mathematical theories and algorithms to tackle the problem given those assumptions, but then evaluate on data sets that are relatively small and sometimes synthetic. ASR research, on the



	Input	Output	Decision function	Loss function
Inductive	$\{(x_i, y_i)\}_{i \in TR}$	f^{TS}		
Transductive	$\{(x_i, y_i)\}_{i \in TR}$ and $\{x_i\}_{i \in TS}$	$\{y_i\}_{i \in TS}$		
Generative			Generative models	$L = -\ln p(x, y; \theta)$
Discriminative			Generative <i>OR</i> Discriminative models	Discriminative loss (form varies)
Supervised	$\{(x_i, y_i)\}_{i \in TR}$			
Unsupervised	$\{x_i\}_{i \in U}$			
Semi-supervised	$\{(x_i, y_i)\}_{i \in TR}$ and $\{x_i\}_{i \in U}$			
Active	$\{(x_i, y_i)\}_{i \in TR}$ and $\{x_i\}_{i \in U}$ and $\{y_i\}_{i \in L \subseteq U}$			
Adaptive	f^{TR} and $\{(x_i, y_i)\}_{i \in AD}$	f^{TS}		
Multitask	$\{(x_i, y_i)\}_{i \in TR_k}$ for all k	f^k or $\{y_i\}_{i \in TS_k}$ for all k		

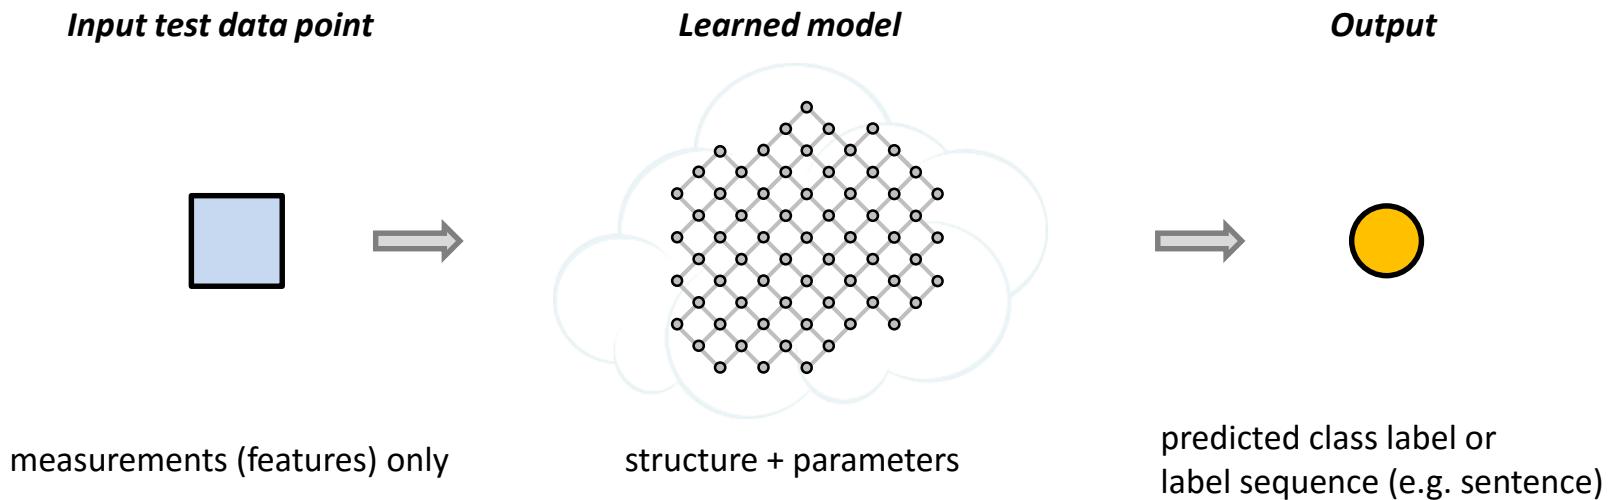
Supervised Machine Learning (classification)

Training phase (usually offline)

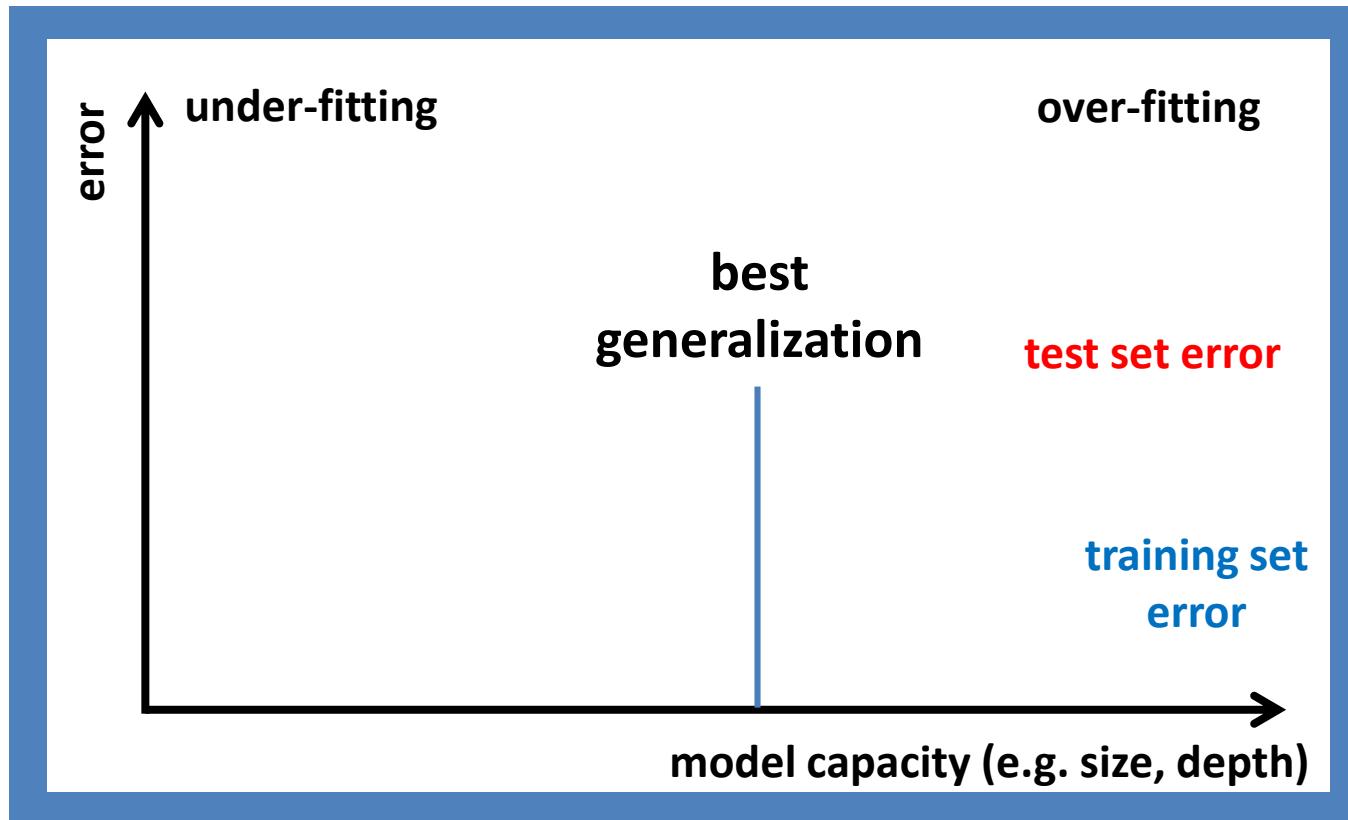


Supervised Machine Learning (classification)

Test phase (run time, online)

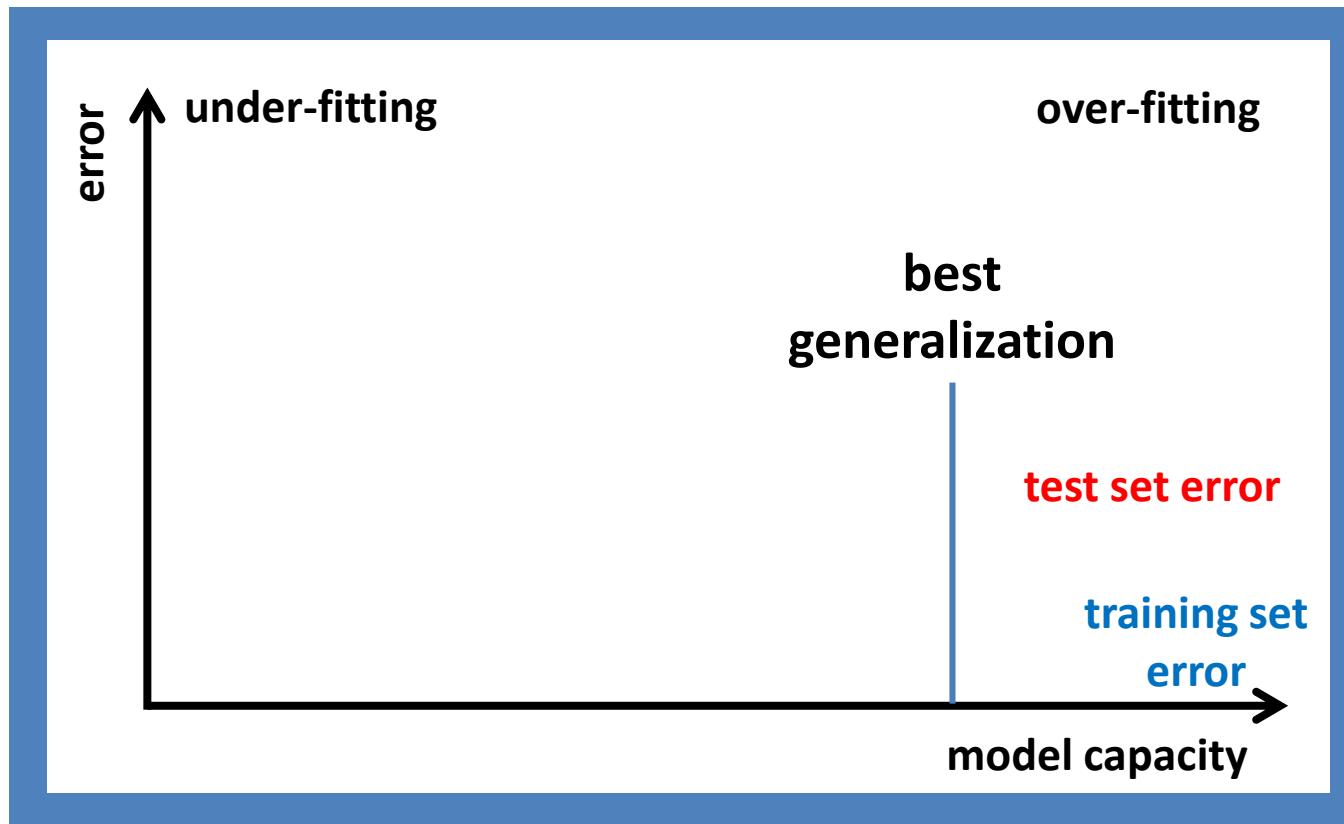


A key ML concept: Generalization



- To avoid over-fitting,
- The need for regularization (or make model simpler, or to add more training)
- → move the training objective away from (empirical) error rate

Generalization – effects of more data

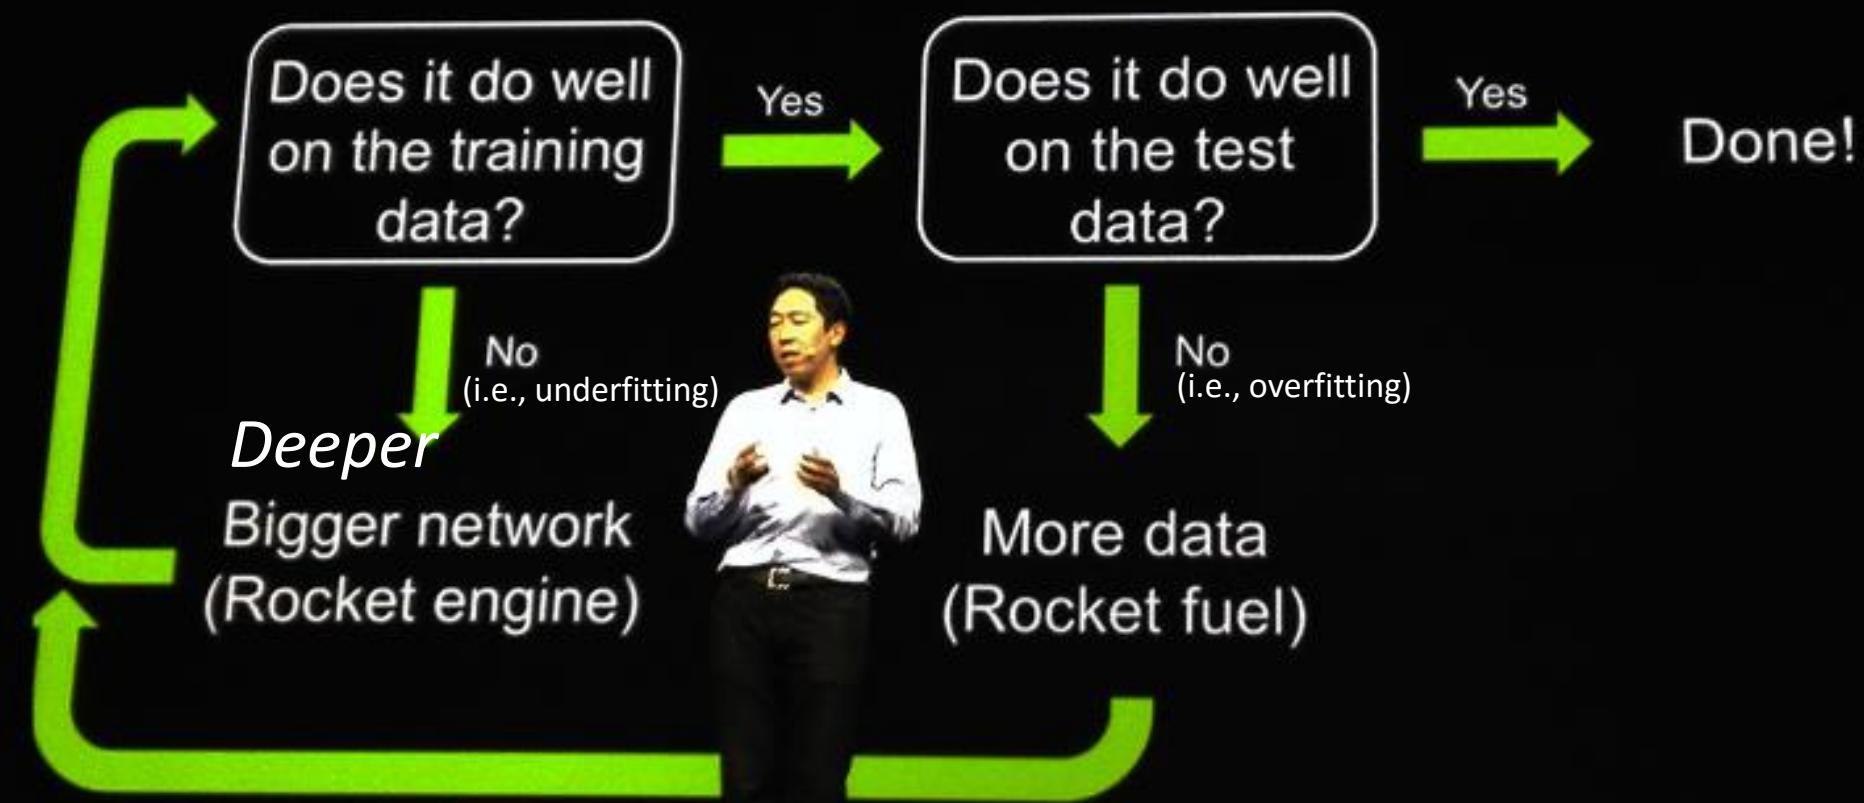


A variety of ML methods

- Decision trees/forests/jungles, Boosting
- Support Vector Machines (SVMs)
- Model-Based (Graphical models, often generative models: sparse connections w. interpretability)
 - model tailored for each new application and incorporates prior knowledge
 - Bayesian statistics exploited to ‘invert the model’ & infer variables of interest
- Neural Networks (DNN, RNN, dense connections)

These two types of methods can be made DEEP: Deep generative models and DNNs

Recipe for (Supervised) Deep Learning with Big Data



Contrast with Signal Processing Approaches

- Strong focus on sophisticated objective functions for optimization (e.g. MCE, MWE, MPE, MMI, string-level, super-string-level, ...)
- Can be regarded as “end2end” learning in ASR
- Almost always non-convex optimization (praised by deep-ML researchers)
- Weaker focus on regularization & overfitting
- Why?
- Our ASR community has been using shallow, low-capacity models for too long (e.g., GMM-HMM)
- Less need for overcoming overfitting
- Now, deep models add a new dimension for increasing model capacity
- Regularization becomes essential for DNN
- E.g. DBN pre-training, “**dropout**” method, STDP spiking neurons, etc.



Discriminative Learning for Speech Recognition

Theory and Practice

Xiaodong He
Li Deng

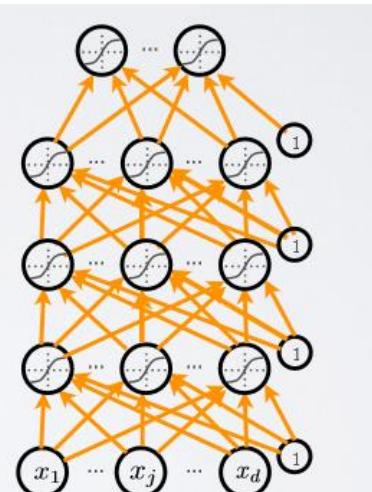
2008

CHAPTER 3

Discriminative Learning:
A Unified Objective Function

Deep Neural Net (DNN) Basics

--- why gradient vanishes & how to rescue it

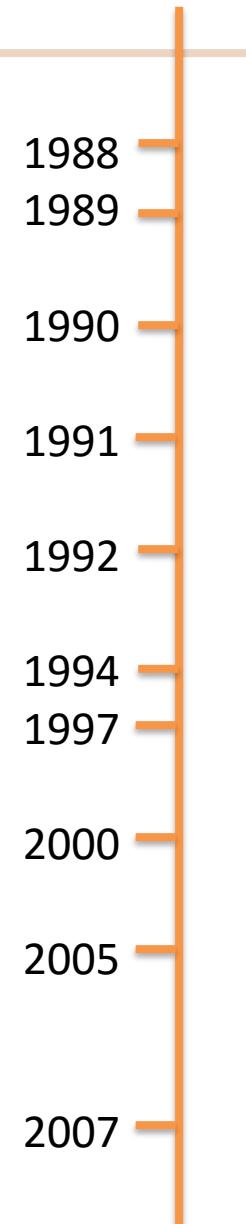


(Shallow) Neural Networks for ASR

(prior to the rise of deep learning)

Temporal & Time-Delay (1-D Convolutional) Neural Nets

- Atlas, Homma, and Marks, "An Artificial Neural Network for Spatio-Temporal Bipolar Patterns, Application to Phoneme Classification," NIPS, 1988.
- Waibel, Hanazawa, Hinton, Shikano, Lang. "Phoneme recognition using time-delay neural networks." IEEE Transactions on Acoustics, Speech and Signal Processing, 1989.



Hybrid Neural Nets-HMM

- Morgan and Bourlard. "Continuous speech recognition using MLP with HMMs," ICASSP, 1990.

Recurrent Neural Nets

- Bengio. "Artificial Neural Networks and their Application to Speech/Sequence Recognition", Ph.D. thesis, 1991.
- Robinson. "A real-time recurrent error propagation network word recognition system," ICASSP 1992.

Neural-Net Nonlinear Prediction

- Deng, Hassanein, Elmasry. "Analysis of correlation structure for a neural predictive model with applications to speech recognition," *Neural Networks*, vol. 7, No. 2, 1994.

Bidirectional Recurrent Neural Nets

- Schuster, Paliwal. "Bidirectional recurrent neural networks," IEEE Trans. Signal Processing, 1997.

Neural-Net TANDEM

- Hermansky, Ellis, Sharma. "Tandem connectionist feature extraction for conventional HMM systems." ICASSP 2000.
- Morgan, Zhu, Stolcke, Sonmez, Sivadas, Shinozaki, Ostendorf, Jain, Hermansky, Ellis, Doddington, Chen, Cretin, Bourlard, Athineos, "Pushing the envelope - aside [speech recognition]," IEEE Signal Processing Magazine, vol. 22, no. 5, 2005.

← DARPA EARS Program 2001-2004: Novel Approach I (Novel Approach II: Deep Generative Model)

Bottle-neck Features Extracted from Neural-Nets

- Grezl, Karafiat, Kontar & Cernocky. "Probabilistic and bottle-neck features for LVCSR of meetings," ICASSP, 2007.

One-hidden-layer neural networks

- Starting from (nonlinear) regression

$$y(\mathbf{x}, \mathbf{w}) = f \left(\sum_{j=1}^M w_j \phi_j(\mathbf{x}) \right)$$

- Replace each ϕ_j with a variable z_j , where

$$z_j = h(a_j) \quad a_j = \sum_{i=1}^D w_{ji}^{(1)} x_i + w_{j0}^{(1)}$$

and $h()$ is a fixed activation function

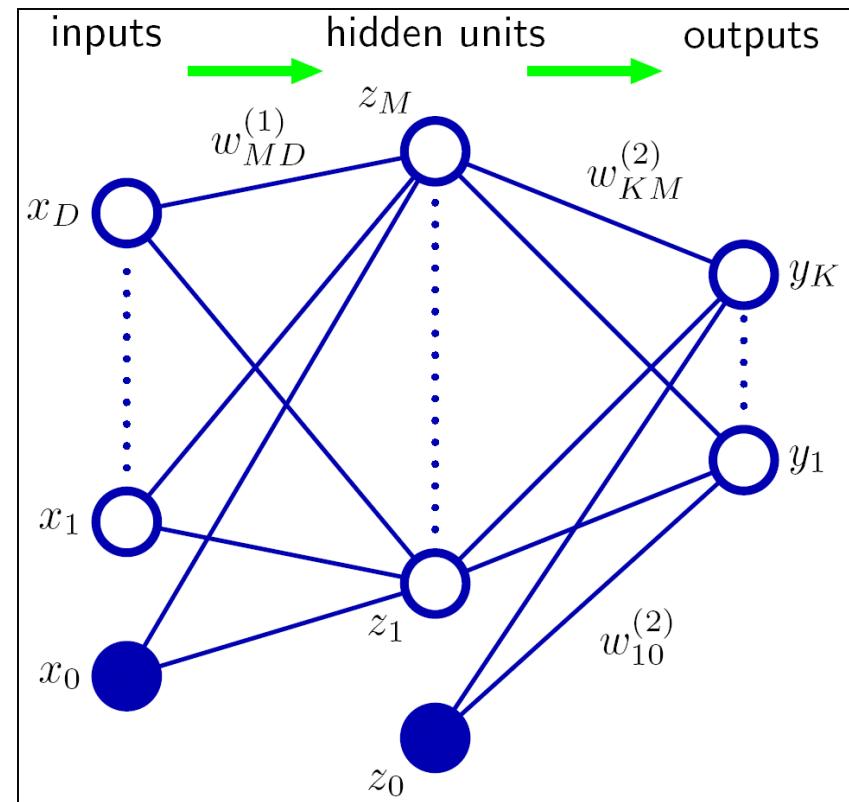
- The outputs obtained from

$$y_k = \sigma(a_k) \quad a_k = \sum_{j=1}^M w_{kj}^{(2)} z_j + w_{k0}^{(2)}$$

where $\sigma()$ is another fixed function

- In all, we have (**simplifying biases**):

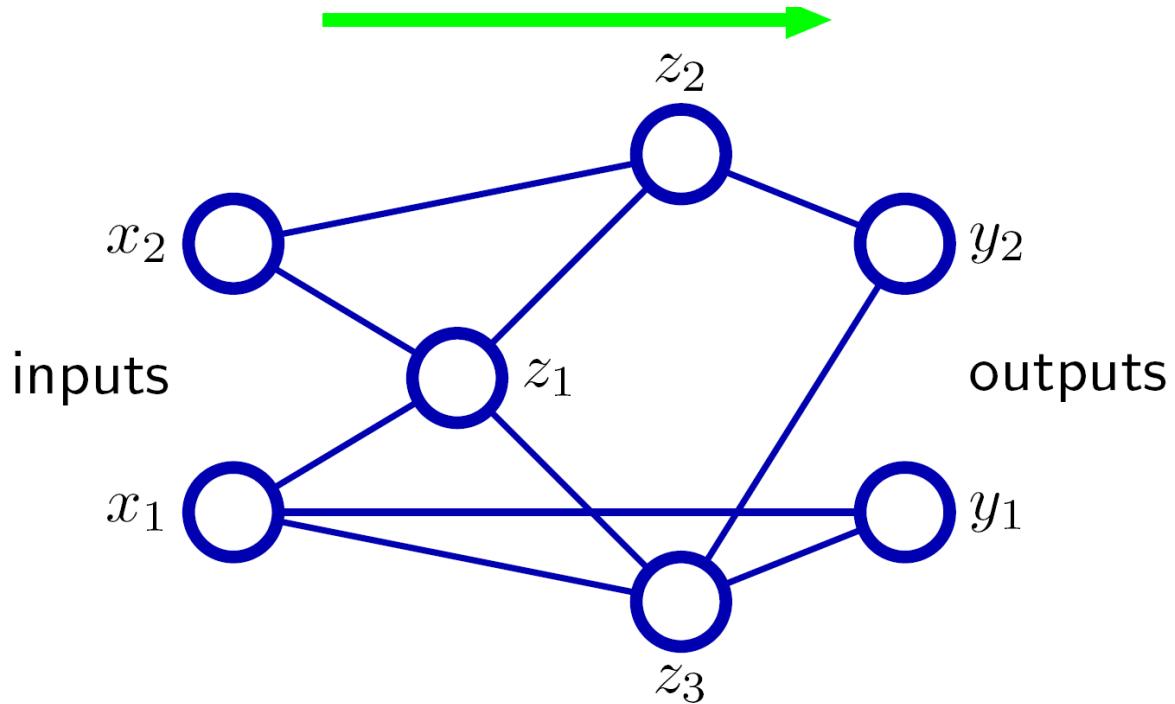
$$y_k(\mathbf{x}, \mathbf{w}) = \sigma \left(\sum_{j=0}^M w_{kj}^{(2)} h \left(\sum_{i=0}^D w_{ji}^{(1)} x_i \right) \right)$$



Multi-layer neural networks

Denote all activation functions by h

$$z_k = h \left(\sum_j w_{kj} z_j \right)$$



- The sum is over those values of j with instantiated weights w_{kj}

Unchallenged learning algorithm: Back propagation (BP)

- For regression, we consider a squared error cost function:

$$E(\mathbf{w}) = \frac{1}{2} \sum_n \sum_k (t_{nk} - y_k(\mathbf{x}_n, \mathbf{w}))^2$$

which corresponds to a Gaussian density $p(\mathbf{t}|\mathbf{x})$

- We can substitute

$$y_k(\mathbf{x}, \mathbf{w}) = \sigma \left(\sum_{j=0}^M w_{kj}^{(2)} h \left(\sum_{i=0}^D w_{ji}^{(1)} x_i \right) \right)$$

and use a general purpose optimizer to estimate \mathbf{w} ,
but it is much more efficient to exploit derivatives of E ,
the essence of BP

Learning neural networks

$$E(\mathbf{w}) = \frac{1}{2} \sum_n \sum_k (t_{nk} - y_k(\mathbf{x}_n, \mathbf{w}))^2$$

- Recall that for linear regression:

$$\frac{\partial E(\mathbf{w})}{\partial w_m} = -\sum_n (t_n - y_n) x_{nm}$$

Weight in-between error
signal and input signal Error signal Input signal

- We'll use the chain rule of differentiation to derive a similar-looking expression, where
 - Local input signals are forward-propagated from the input
 - Local error signals are back-propagated from the output

Local signals needed for learning

- For clarity, consider the error for one training case:

$$E_n = \frac{1}{2} \sum_k (t_{nk} - y_{nk})^2$$

- To compute $\partial E_n / \partial w_{ji}$, note that w_{ji} appears in only one term of the overall expression, namely

$$a_j = \sum_i w_{ji} z_i \quad \begin{array}{l} \text{if } w_{ji} \text{ is in the 1st layer,} \\ z_i \text{ is actually input } x_i \end{array}$$

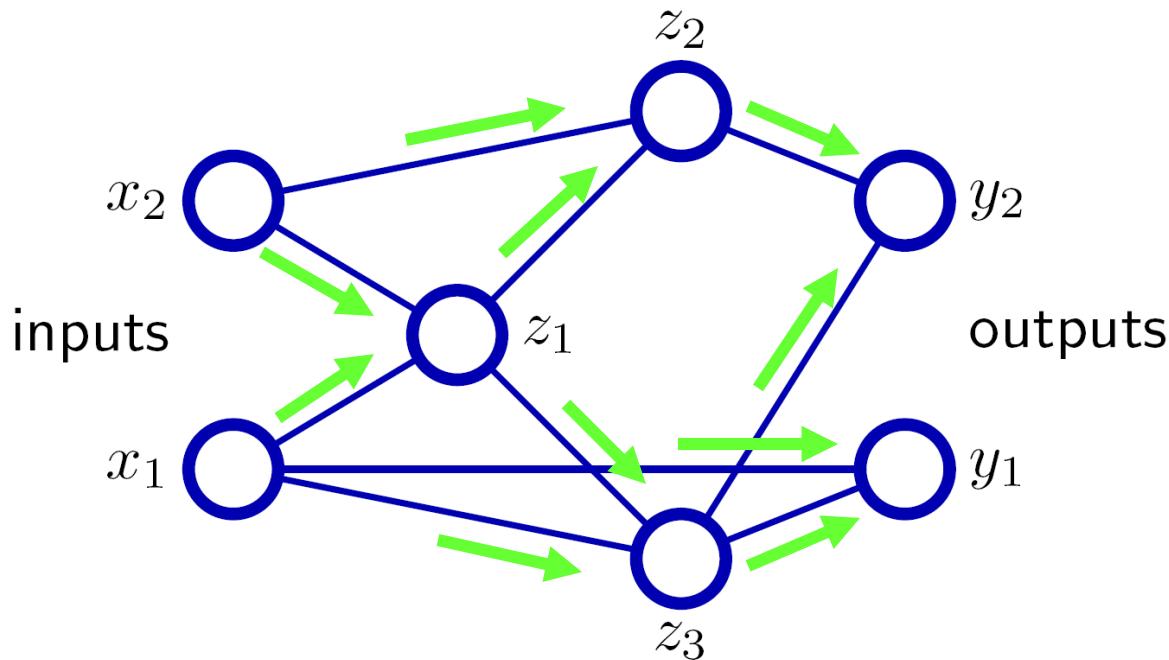
- Using the chain rule of differentiation, we have

$$\frac{\partial E_n}{\partial w_{ji}} = \underbrace{\frac{\partial E_n}{\partial a_j}}_{\text{Weight}} \frac{\partial a_j}{\partial w_{ji}} = \delta_j z_i$$



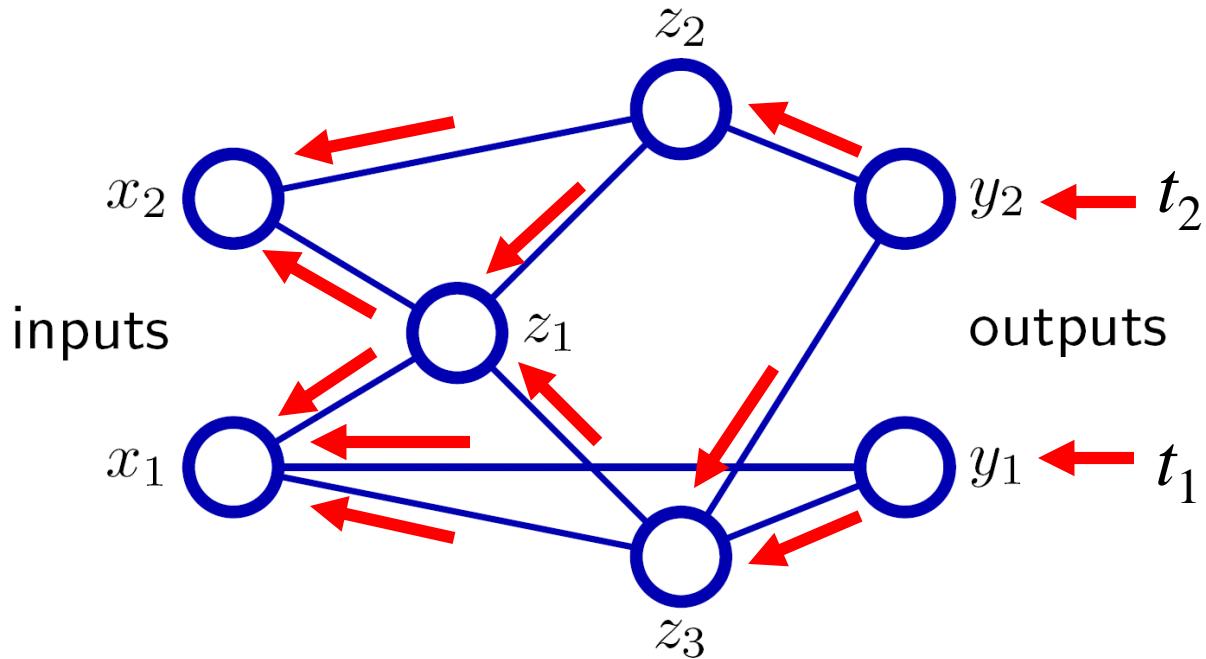
where $\delta_j \equiv \frac{\partial E_n}{\partial a_j}$

Forward-propagating local input signals



- Forward propagation gives all the a 's and z 's

Back-propagating local error signals



- Back-propagation gives all the δ 's

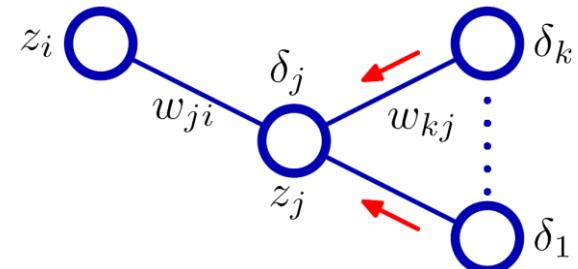
Back-propagating error signals

- To compute $\partial E_n / \partial a_j$ (i.e., δ_j), a_j (also called “logit”) appears in all those expressions $a_k = \sum_i w_{ki} h(a_i)$ that depend on a_j
- Using the chain rule, we have

$$\delta_j \equiv \frac{\partial E_n}{\partial a_j} = \sum_k \frac{\partial E_n}{\partial a_k} \frac{\partial a_k}{\partial a_j}$$

- The sum is over k s.t. unit j is connected to unit k and for each such term, $\partial a_k / \partial a_j = w_{kj} h'(a_j)$
- Noting that $\partial E_n / \partial a_k = \delta_k$, we get the back-propagation rule:

$$\delta_j = h'(a_j) \sum_k w_{kj} \delta_k$$



- For output units: $\delta_j = -h'(a_j)(t_{nk} - y_{nk})$

Putting the propagations together

- For each training case n , apply forward propagation and back-propagation to compute

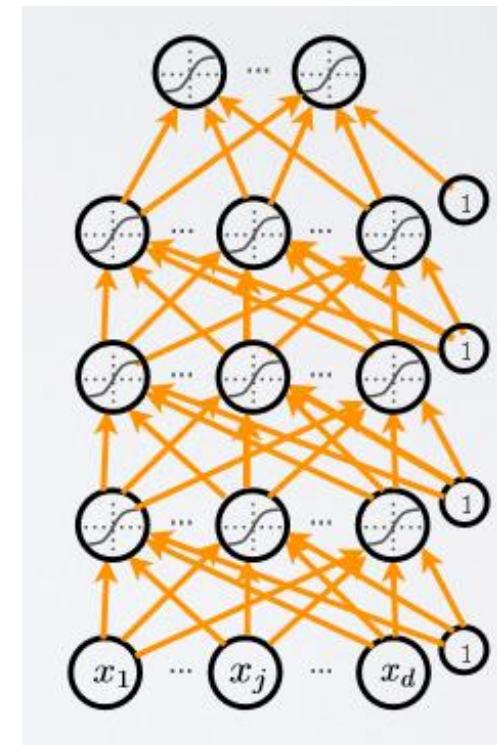
$$\frac{\partial \hat{E}_n}{\partial w_{ji}} = \delta_j z_i$$

for each weight w_{ji}

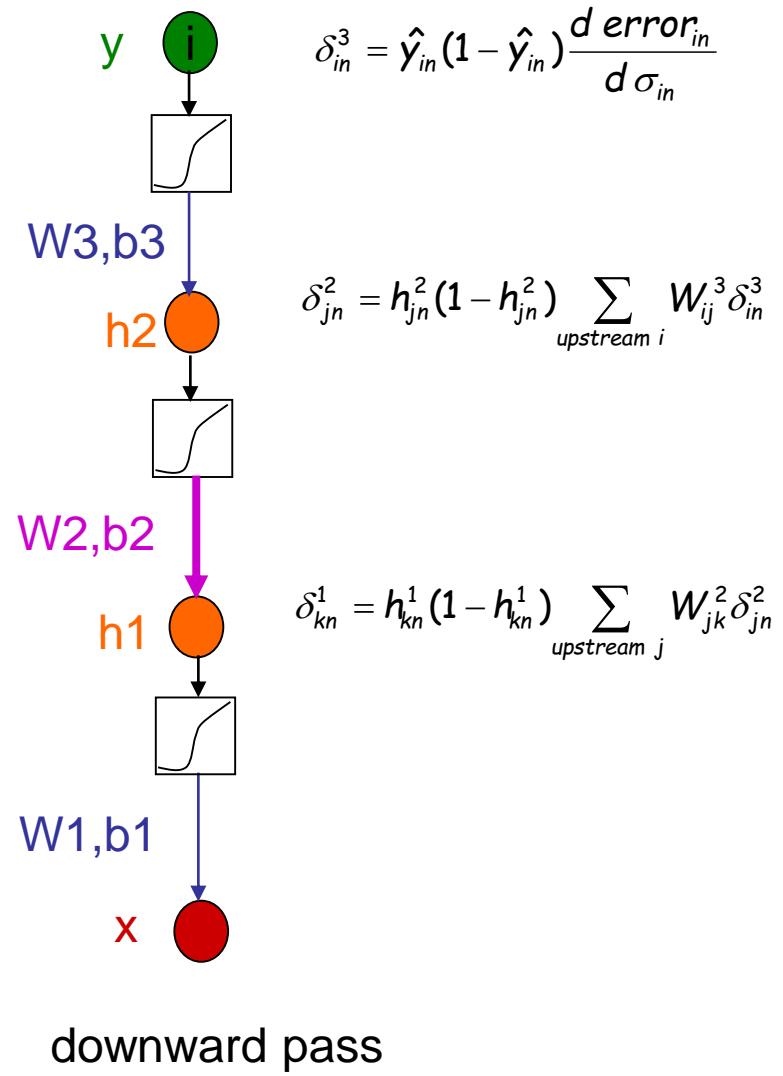
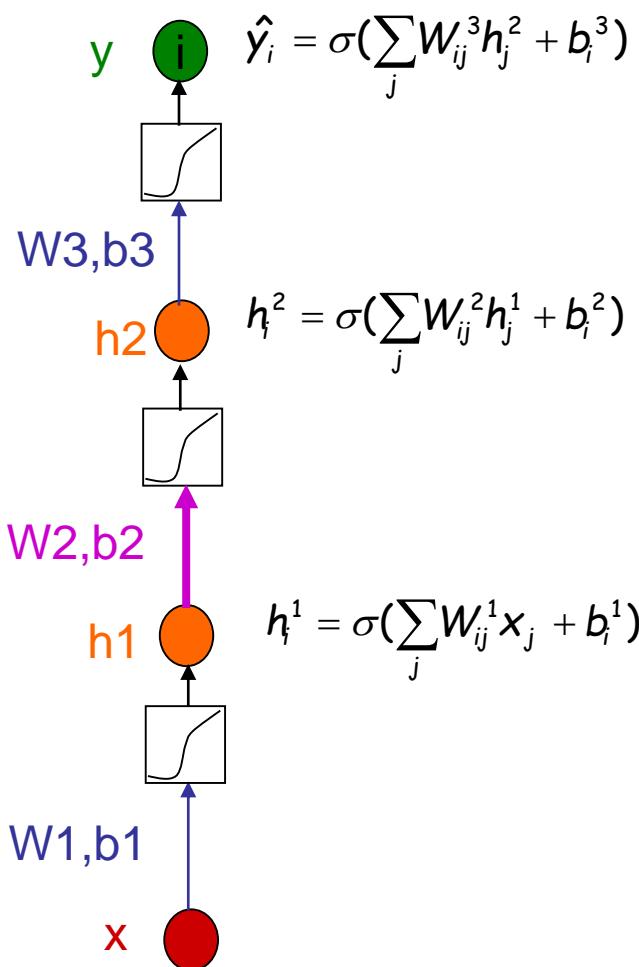
- Sum these over training cases to compute $\frac{\partial E}{\partial w_{ji}}$
- Use these derivatives for steepest descent learning (too slow for large set of training data)
- Minibatch learning: After a small set of input samples, use the above gradient to update the weights (so update more often)

Why gradients tend to vanish for DNN

- Recall BP for adjacent layer pair: $\delta_j = h'(a_j) \sum_k w_{kj} \delta_k$
- For sigmoid units: $h'(a_j) = h(a_j)(1-h(a_j))$
- If there is one hidden layer, gradients are not likely to vanish
- Problem becomes serious when nets get deep



Why gradients tend to vanish for DNN



Why gradients tend to vanish for DNN

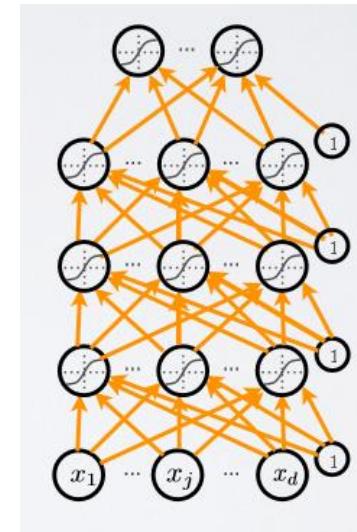
- To illustrate to problem, let's use matrix form of error BP:

$$\begin{aligned}\delta^l &= ((w^{l+1})^T \delta^{l+1}) \odot \sigma'(z^l) \\ &= ((w^{l+1})^T \left[\begin{array}{c} ((w^{l+2})^T \delta^{l+2}) \odot \sigma'(z^{l+1}) \\ \vdots \\ ((w^{l+2})^T \delta^{l+2}) \odot \sigma'(z^{l+1}) \end{array} \right]) \odot \sigma'(z^l) \\ &\quad \dots, \dots, \dots \\ &\quad \dots, \dots, \dots\end{aligned}$$

- So even if forward pass is nonlinear, **error backprop is a linear process**
- It suffers from all problems associated with linear processes
- Many terms of $\sigma(1 - \sigma)$ for sigmoid units
- In addition, many terms in the product of W's
- If any sigmoid unit saturates in either direction, the error gradient becomes zero
- If $\|W\| < 1$, the product will shrink fast for high depths
- If $\|W\| > 1$, the product may grow fast for high depths

How to rescue it

- Pre-train the DNN by generative DBN (solution by 2010)
 - Complicated process
- Discriminative pre-training (much easier to do)
- Still random initialization but with carefully set variance values; e.g.
 - Layer-dependent variance values
 - For lower layers (with more terms in the product), make the variance closer to 0 (e.g. < 0.1)
- Use **bigger training data** to reduce the chance of vanishing gradients for each epoch
- Use ReLU units: only one side of zero gradient instead of two as for sigmoid units



The power of understanding root causes!!!
(mid 2010 at MSR Redmond)

An alternative way of training NN

- Backprop takes partial derivatives: $\frac{\partial E_n}{\partial w_{ji}} = \delta_j z_i$
- If the **output layer is linear**, total (instead of partial) derivative can be computed

$$\frac{dE}{d\mathbf{W}} = \frac{\partial E}{\partial \mathbf{W}} + \frac{\partial E}{\partial \mathbf{U}^*} \frac{\partial \mathbf{U}^*}{\partial \mathbf{W}}.$$

- This is basic learning method for Deep Convex/Stacking Net (DSN), designed to be easily parallelizable by batch training
- Using the total derivative is equivalent to coordinate descent algorithm with an “infinite” step size to achieve the global optimum along the “coordinate” of updating U while fixing W.

Deep Stacking Nets

- Learn weight matrices U and W in individual modules separately.
- Given W and linear output layer, U can be expressed as explicit nonlinear function of W .
- This nonlinear function is used as the constraint in solving nonlinear least square for learning W .
- Initializing W with RBM (bottom layer)
- For higher layers, part of W is initialized with the optimized W from the immediately lower layer and part of it with random numbers

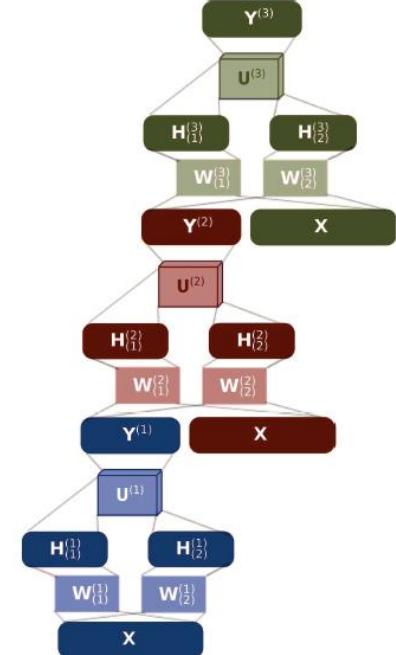
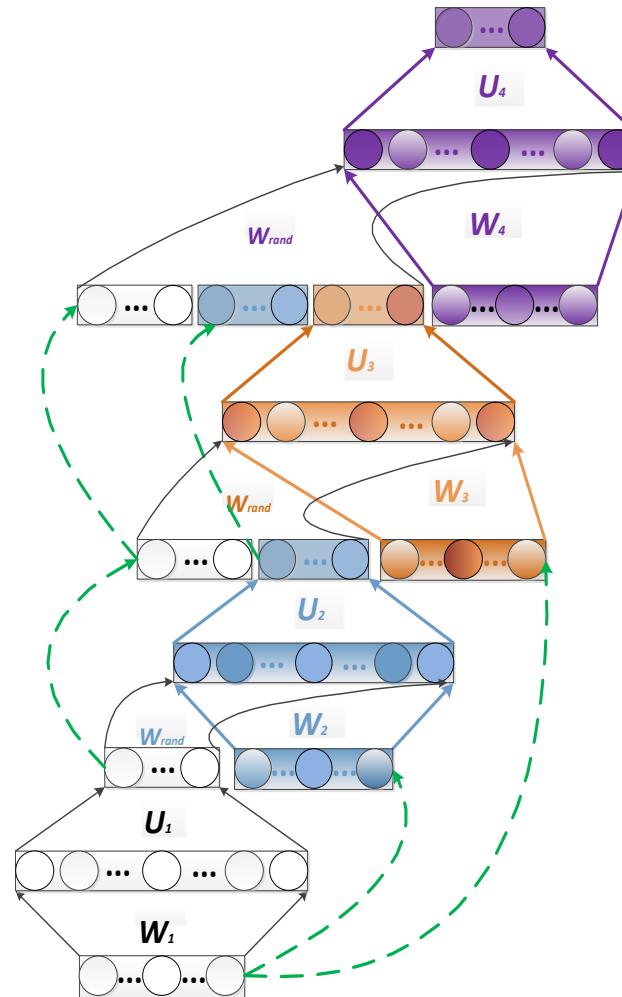


Fig. 1. An example T-DSN architecture with three stacking blocks, where each block consists of three layers, and superscript is used to indicate the block number. Inputs

A neat way of learning DSN weights

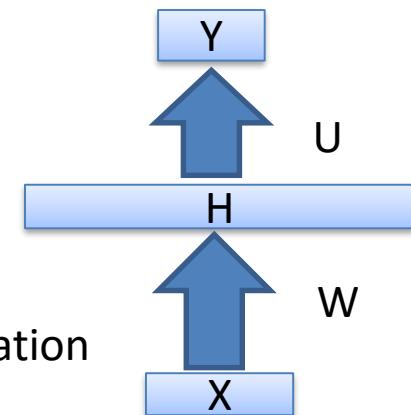
$$E = \frac{1}{2} \sum_n ||y_n - t_n||^2, \quad \text{where } y_n = U^T h_n = U^T \sigma(W^T x_n) = G_n(U, W)$$

$$\begin{aligned} \frac{\partial E}{\partial U} &= 2H(U^T H - T)^T \rightarrow U = (HH^T)^{-1} HT^T = F(W), \text{ where } h_n = \sigma(W^T x_n) \\ E &= \frac{1}{2} \sum_n ||G_n(U, W) - t_n||^2, \text{ subject to } U = F(W), \end{aligned}$$

Use of Lagrange multiplier method:

$$E = \frac{1}{2} \sum_n ||G_n(U, W) - t_n||^2 + \lambda ||U - F(W)||$$

to learn W (& then U) \rightarrow full derivation $\frac{dE}{dW}$ in closed form
 (i.e. no longer recursion on partial derivation as in backpropagation)



- Advantages found:
 - less noise in gradient than using chain rule which ignores explicit constraint $U = F(W)$
 - batch learning is effective, aiding parallel training

How the Brain May Do BackProp

- Canadian Psychology, Vol 44, pp 10-13, 2003.
-

The Ups and Downs of Hebb Synapses

GEOFFREY HINTON
University of Toronto

Abstract

Modelers have come up with many different learning rules for neural networks. When a teacher specifies the correct

This approach led to effective "error-driven" learning rules such as the Widrow-Hoff rule (Widrow & Hoff, 1960) and the perceptron convergence procedure (Rosenblatt, 1961) and it was later generalized to multilayer networks by using backpropagation of the errors to get training signals for intermediate "hidden" layers (Rumelhart, Hinton, & Williams, 1986).

- Feedback system in biological neural nets
- Key roles of STDP (Spike-Time-Dependent Plasticity) --- temporal derivative
- To provide a way to encode error derivatives

How the Brain May Do BackProp

- Backprop algorithm requires that feedforward and feedback weights are the same
- This is clearly not true for biological neural nets
- How to reconcile this discrepancy?
- Recent studies showed that use of random feedback weights in BP performs close to rigorous BP
- Implications for regularizing BP learning (like dropout, which may not make sense at 1st glance)

Random feedback weights support learning in deep neural networks

Timothy P. Lillicrap^{1*}, Daniel Cownden², Douglas B. Tweed^{3,4}, Colin J. Akerman¹

¹Department of Pharmacology, University of Oxford, Oxford, United Kingdom

²Centre for the Study of Cultural Evolution, Stockholm University, Stockholm, Sweden

³Departments of Physiology and Medicine, University of Toronto, Toronto, Canada

⁴Centre for Vision Research, York University, Toronto, Canada

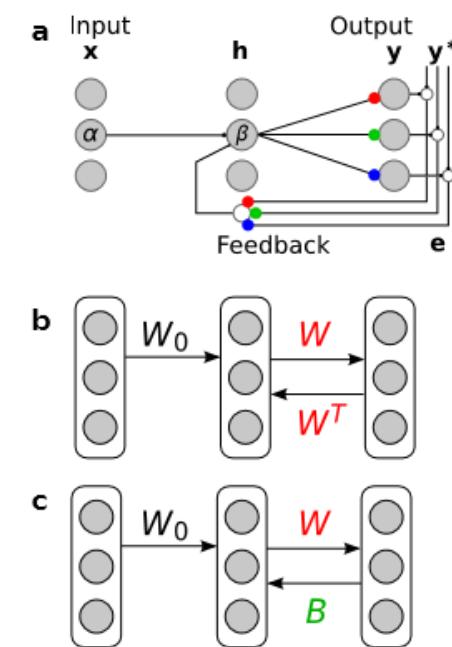
*To whom correspondence should be addressed:

timothy.lillicrap@pharm.ox.ac.uk

colin.akerman@pharm.ox.ac.uk

Abstract

The brain processes information through many layers of neurons. This deep architecture is representationally powerful^{1,2,3,4}, but it complicates learning by making it hard to identify the responsible neurons when a mistake is made^{1,5}. In machine learning, the backpropagation algorithm¹ assigns blame to a few



Recurrent Neural Net (RNN) Basics

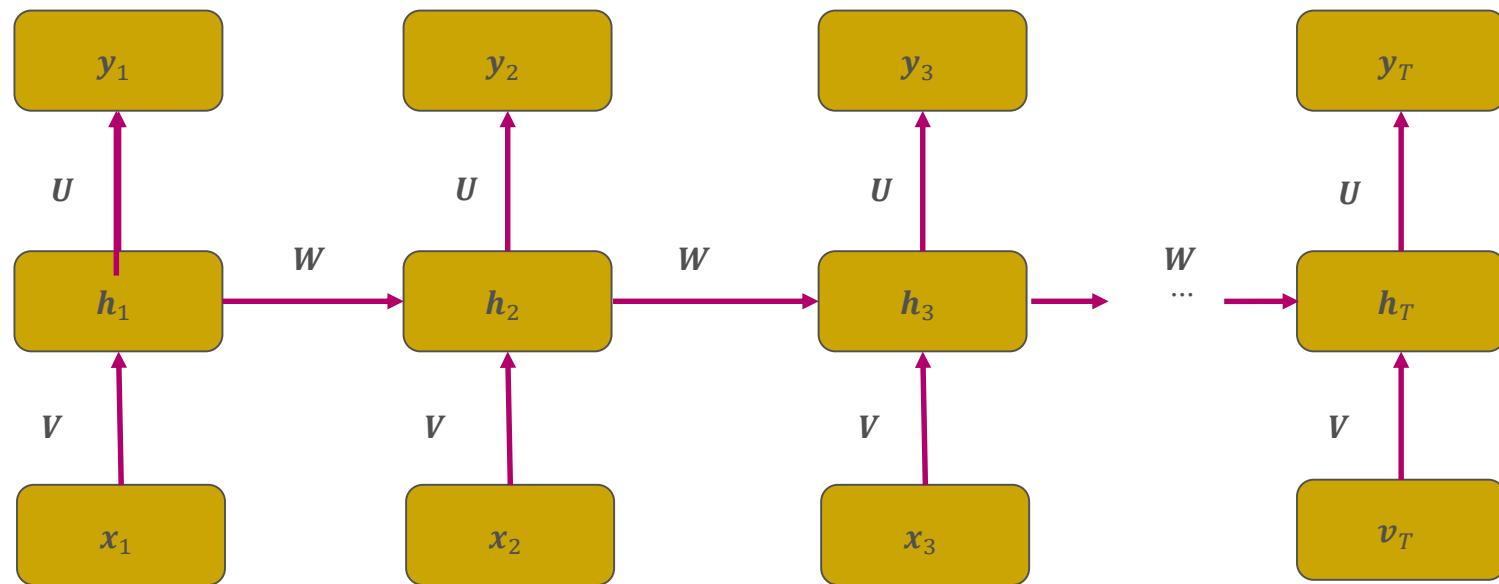
- why memory decays fast or explode (1990's)
- how to rescue it (2010's, in the new deep learning era)

Basic architecture of an RNN

h_t is the hidden layer that carries the information from time $0 \sim t$

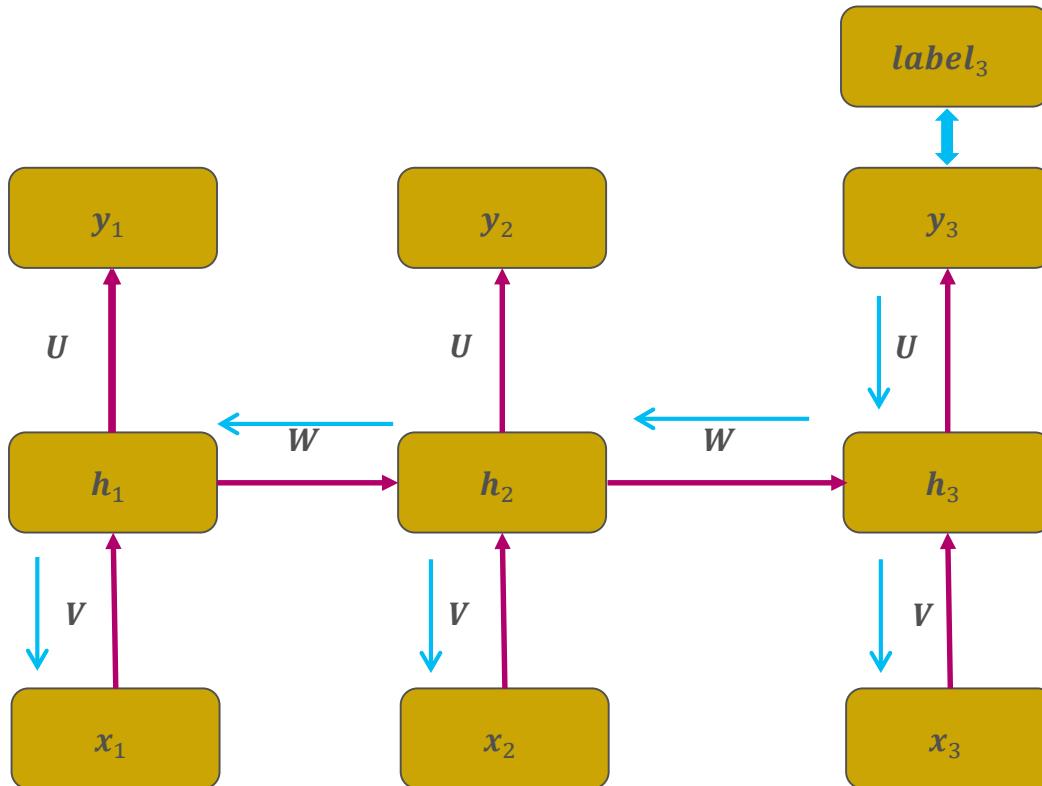
where x_t : the input word , y_t : the output tag

$$y_t = \text{SoftMax}(U \cdot h_t), \text{ where } h_t = \sigma(W \cdot h_{t-1} + V \cdot x_t)$$



Used for slot filling in SLU [Mesnil, He, Deng, Bengio, 2013; Yao, Zweig, Hwang, Shi, Yu, 2013] IEEE TASLP

Back-propagation through time (BPTT)



at time $t = 3$

1. Forward propagation
2. Generate output
3. Calculate error
4. Back propagation
5. Back prop. through time

A Good Read to appreciate the “Magic” of modern RNNs



About Hacker's guide to Neural Networks

The Unreasonable Effectiveness of Recurrent Neural Networks

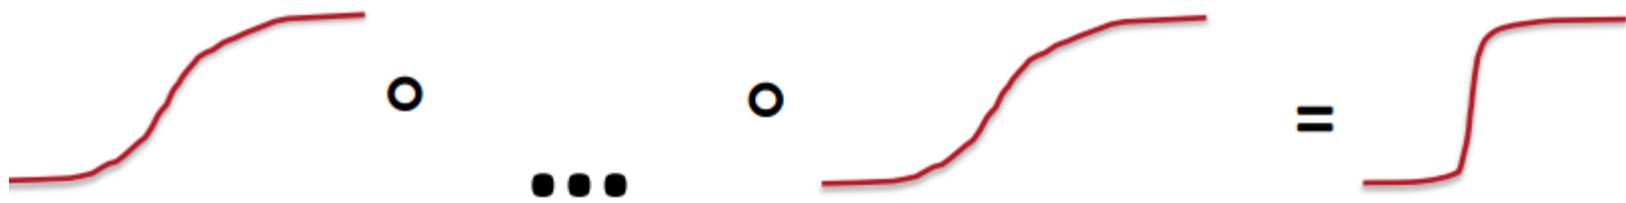
May 21, 2015

There's something magical about Recurrent Neural Networks (RNNs). I still remember when I trained my first recurrent network for [Image Captioning](#). Within a few dozen minutes of training my first baby model (with rather arbitrarily-chosen hyperparameters) started to generate very nice looking descriptions of images that were on the edge of making sense. Sometimes the ratio of how simple your model is to the quality of the results you get out of it blows past your expectations, and this was one of those times. What made this result so shocking at the time was that the common wisdom was that RNNs were supposed to be difficult to train (with more experience I've in fact reached the opposite conclusion). Fast forward about a year: I'm training RNNs all the time and I've witnessed their power and robustness many times, and yet their magical outputs still find ways of amusing me. This post is about sharing some of that magic with you.

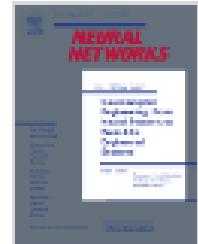
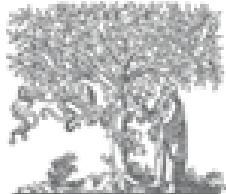
<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

Some early attempts to examine difficulties in learning RNNs

- Yoshua Bengio's Ph.D. thesis at McGill University (1991)
- Based on attractor properties of nonlinear dynamic systems
- His recent, more intuitive explanation --- in extreme of nonlinearity, discrete functions and gradients vanish or explode:



- An alternative analysis: based on perturbation analysis of nonlinear differential equations (next several slides)



Contributed article

Analysis of the correlation structure for a neural predictive model with application to speech recognition *

L. Deng , K. Hassanein, M. Elmasry

University of Waterloo Canada

- NN for nonlinear sequence prediction (like NN language model used today)
- Memory (temporal correlation) proved to be stronger than linear prediction
- No GPUs to use; very slow to train with BP; did not make NN big and deep, etc.
- Conceptually easy to make it deep using (state-space) signal processing & graphical models.. by moving away from NN...

The method we use to derive the correlation function for (2) resembles the perturbation analysis for the study of nonlinear differential equations [13]. To proceed, we construct a family of models which is parameterized by α :

$$Y_{t+1}(\alpha) = \alpha f(Y_t(\alpha)) + \epsilon_{t+1}, \quad (3)$$

Once the model is parameterized, the autoregression on the data Y_t can be removed by performing power-series expansion of the nonlinear function $f(\cdot)$:

$$Y_1(\alpha) = \epsilon_1 + \alpha f(Y_0(\alpha)),$$

$$\begin{aligned} Y_2(\alpha) &= \epsilon_2 + \alpha f(Y_1(\alpha)) \\ &= \epsilon_2 + \alpha f(\epsilon_1) + \alpha^2 f(Y_0)f'(\epsilon_1) + \frac{1}{2}\alpha^3 f^2(Y_0)f''(\epsilon_1) + \dots, \end{aligned}$$

$$\begin{aligned} Y_3(\alpha) &= \epsilon_3 + \alpha f(Y_2(\alpha)) \\ &= \epsilon_3 + \alpha f(\epsilon_2) + \alpha^2 f(\epsilon_1)f'(\epsilon_2) + \alpha^3 f(Y_0)f'(\epsilon_1)f'(\epsilon_2) + \dots, \end{aligned}$$

⋮

and in general,

$$Y_t(\alpha) = \epsilon_t + \alpha f(\epsilon_{t-1}) + \alpha^2 f(\epsilon_{t-2})f'(\epsilon_{t-1}) + \alpha^3 f(\epsilon_{t-3})f'(\epsilon_{t-2})f'(\epsilon_{t-1}) + \dots \quad (4)$$

(In the above, $f'(\cdot)$ denotes the derivative of $f(\cdot)$ with respect to its argument.)

From (4) the covariance function for model (3) is calculated to give

$$\begin{aligned} & \text{Cov}[Y_t(\alpha), Y_{t+\tau}(\alpha)] \\ & \approx \text{Cov}(\epsilon_t, \epsilon_{t+\tau}) + \alpha \text{Cov}[f(\epsilon_t), \epsilon_{t+\tau}] + \alpha \text{Cov}[\epsilon_t, f(\epsilon_{t+\tau-1})] \\ & \quad + \alpha^2 \text{Cov}[f(\epsilon_{t-2})f'(\epsilon_{t-1}), \epsilon_{t+\tau}] + \alpha^2 \text{Cov}[\epsilon_t, f(\epsilon_{t+\tau-2})f'(\epsilon_{t+\tau-1})] \\ & \quad + \alpha^3 \text{Cov}[f(\epsilon_{t-2})f'(\epsilon_{t-1}), f(\epsilon_{t+\tau-1})] \\ & \quad + \alpha^3 \text{Cov}[f(\epsilon_{t-1}), f(\epsilon_{t+\tau-2})f'(\epsilon_{t+\tau-1})] \\ & \quad + \alpha^4 \text{Cov}[f(\epsilon_{t-2})f'(\epsilon_{t-1}), f(\epsilon_{t+\tau-2})f'(\epsilon_{t+\tau-1})]. \end{aligned} \quad (5)$$

Among the eight terms in (5), the first, second, fourth, and sixth terms are zero for $\tau \geq 0$. This is due to the IID assumption for ϵ_t and to the fact that $f(\cdot)$ is a static function containing no memory. The fifth term, $\text{Cov}[\epsilon_t, f(\epsilon_{t+\tau-2})f'(\epsilon_{t+\tau-1})]$, is non-zero only for $\tau = 1$ and $\tau = 2$. The seventh and the eighth terms are non-zero only for $\tau = 1$. Likewise, any higher order terms of α in the covariance function which are omitted due to cutoff in the power-series expansion of $Y_t(\alpha)$ would contain non-zero values only for small time lags.

We conclude from the above analysis that prediction of a time series with a single nonlinear term alone does not produce long-term temporal correlations in the model's output.

III.3. Joint prediction with nonlinear and linear terms

In this section we investigate correlation properties of the data generated from the stationary time series model

$$Y_{t+1} = \phi Y_t + f(Y_t) + \epsilon_{t+1}, \quad t = 1, 2, \dots, T, \quad (6)$$

We now decompose the stationary random process $Y_t(\alpha)$ into its stationary component processes by representing it as a power-series expansion on α

$$Y_{t+1}(\alpha) = Y_{t+1,0} + \alpha Y_{t+1,1} + \frac{1}{2!} \alpha^2 Y_{t+1,2} + \frac{1}{3!} \alpha^3 Y_{t+1,3} + \dots \quad (8)$$

In order to identify the component processes $Y_{t,i}$, $i = 0, 1, 2, \dots$, we substitute (8) into (7) and approximate the nonlinear function $f(\cdot)$ by truncating its power-series expansion. This gives

$$\begin{aligned} Y_{t+1}(\alpha) &\approx \phi(Y_{t,0} + \alpha Y_{t,1} + \frac{1}{2!} \alpha^2 Y_{t,2} + \frac{1}{3!} \alpha^3 Y_{t,3}) \\ &\quad + \alpha[f(Y_{t,0}) + f'(Y_{t,0})(\alpha Y_{t,1} + \frac{1}{2!} \alpha^2 Y_{t,2} + \frac{1}{3!} \alpha^3 Y_{t,3})] + \epsilon_{t+1} \\ &= (\phi Y_{t,0} + \epsilon_{t+1}) + \alpha[\phi Y_{t,1} + f(Y_{t,0})] + \alpha^2[\frac{1}{2} \phi Y_{t,2} + f'(Y_{t,0}) Y_{t,1}] \\ &\quad + \alpha^3[\frac{1}{6} \phi Y_{t,3} + \frac{1}{2} f'(Y_{t,0}) Y_{t,2}] + \dots \end{aligned} \quad (9)$$

By equating the coefficients of α^i in (8) and in (9), we obtain the following recursive relations among the component processes $Y_{t,k}$, $k = 0, 1, 2, \dots$:

$$\begin{aligned} Y_{t+1,0} &= \phi Y_{t,0} + \epsilon_{t+1}, \\ Y_{t+1,1} &= \phi Y_{t,1} + f(Y_{t,0}), \\ Y_{t+1,2} &= \phi Y_{t,2} + 2f'(Y_{t,0}) Y_{t,1}, \\ Y_{t+1,3} &= \phi Y_{t,3} + 3f'(Y_{t,0}) Y_{t,2}, \\ &\vdots \end{aligned} \quad (10)$$

According to (10), we can proceed to derive the autocovariance function for $Y_t(\alpha)$ denoted by

$$\gamma = \text{Cov}[Y_t(\alpha), Y_{t+\tau}(\alpha)].$$

Using (8) and truncating the expansion up to the first order, we have

$$\gamma \approx Cov[Y_{t,0} + \alpha Y_{t,1}, Y_{t+\tau,0} + \alpha Y_{t+\tau,1}]. \quad (11)$$

Use of the stationarity property of $Y_{t,0}$ and $Y_{t,1}$ leads to

$$\begin{aligned}\gamma = & \phi^2\gamma + \alpha^2Cov[f(Y_{t-1,0}), f(Y_{t+\tau-1,0})] \\ & + \phi\alpha Cov[Y_{t-1,0} + \alpha Y_{t-1,1}, f(Y_{t+\tau-1,0})] \\ & + \phi\alpha Cov[Y_{t+\tau-1,0} + \alpha Y_{t+\tau-1,1}, f(Y_{t-1,0})].\end{aligned}$$

Re-arranging terms and using the stationarity property of $Y_{t,0}$ and $Y_{t,1}$ again give γ which is equal to

$$\frac{1}{(1 - \phi^2)} \{ \alpha^2 Cov[f(Y_{t-1,0}), f(Y_{t+\tau-1,0})] + 2\phi\alpha Cov[Y_{t,0} + \alpha Y_{t,1}, f(Y_{t+\tau,0})] \}. \quad (12)$$

$Y_{t,0}$, the zero-th order expansion of $Y_t(\alpha)$, is a linear process and its properties are well understood (Section III.1). To obtain the desired

$$\begin{aligned}\gamma = & \frac{1}{(1 - \phi^2)} \{ \alpha^2 Cov[f(Y_{t-1,0}), f(Y_{t+\tau-1,0})] + 2\phi\alpha Cov[Y_{t,0}, f(Y_{t+\tau,0})] \\ & + 2\phi\alpha^2 \sum_{i=0}^{t-1} \phi^i Cov[f(Y_{t-i-1,0}), f(Y_{t+\tau,0})] \}.\end{aligned}$$

The first two terms in the above expression are exponentially declining as a function of time lag τ because the component processes involved are just static functions of linear processes. The remaining summation, however, would in general decay more slowly because of the many contributing terms.

Why gradients vanish/explode for RNN

- The easiest account is to follow the analysis for DNN
- except that “depth” of RNN is much larger: the length of input sequence
- Especially serious for speech sequence (not as bad for text input)
- Tony Robinson group was the only one that made RNN work for TIMIT phone recognition (1994)

How to rescue it

- Echo state nets (“lazy” approach)
 - avoiding problems by not training input & recurrent weights
 - H. Jaeger. “Short term memory in echo state networks”, 2001
 - But if you figure out smart ways to train them, you get much better results
-

Learning Input and Recurrent Weight Matrices in Echo State Networks

Hamid Palangi
University of British Columbia
Vancouver, BC, Canada
hamidp@ece.ubc.ca

Li Deng
Microsoft Research
Redmond, WA, USA
deng@microsoft.com

Rabab K Ward
University of British Columbia
Vancouver, BC, Canada
rababw@ece.ubc.ca

Abstract

The traditional echo state network (ESN) is a special type of a temporally deep model, the recurrent network (RNN), which carefully designs the recurrent matrix and fixes both the recurrent and input matrices in the RNN. The ESN also adopts the linear output (or readout) units to simplify the learning of the only output matrix in the RNN. In this paper, we devise a special technique that takes advantage of the linearity in the output units in the ESN to learn the input and

(NIPS-WS, 2013)

How to rescue it

- By Better optimization
 - Hessian-free method
 - Primal-dual method

Deep learning via Hessian-free optimization

James Martens

University of Toronto, Ontario, M5S 1A1, Canada

JMARTENS@CS.TORONTO.EDU

Abstract

We develop a 2nd-order optimization method based on the “Hessian-free” approach, and apply it to training deep auto-encoders. Without using pre-training, we obtain results superior to those reported by [Hinton & Salakhutdinov \(2006\)](#) on the same tasks they considered. Our method is practical, easy to use, scales nicely to very large

nately, there has yet to be a demonstration that any of these methods are effective on deep learning problems that are known to be difficult for gradient descent.

Much of the recent work on applying 2nd-order methods to learning has focused on making them practical for large datasets. This is usually attempted by adopting an “on-line” approach akin to the one used in stochastic gradient descent (SGD). The only demonstrated advantages of these meth-

A Primal-Dual Method for Training Recurrent Neural Networks Constrained by the Echo-State Property

Jianshu Chen

Department of Electrical Engineering
University of California
Los Angeles, CA 90034, USA
cjs09@ucla.edu

Li Deng

Machine Learning Group
Microsoft Research
Redmond, WA 98052, USA
deng@microsoft.com

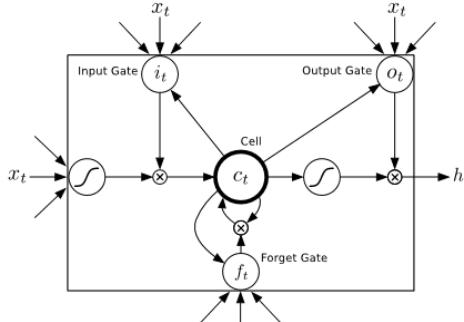
Abstract

We present an architecture of a recurrent neural network (RNN) with a fully-connected deep neural network (DNN) as its feature extractor. The RNN is equipped with both causal temporal prediction and non-causal look-ahead, via auto-regression (AR) and moving-average (MA), respectively. The focus of this paper is a primal-dual training method that formulates the learning of the RNN as a formal optimization problem with an inequality constraint that provides a sufficient condition for the stability of the network dynamics. Experimental results demonstrate the effectiveness of this new method, which achieves 18.86% phone

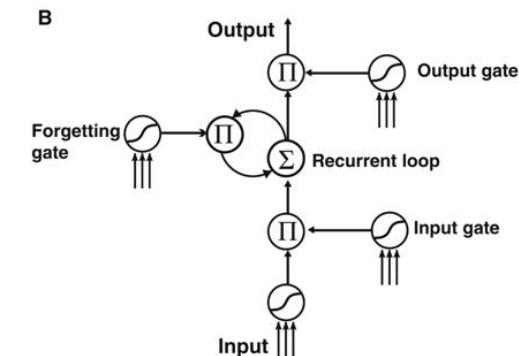
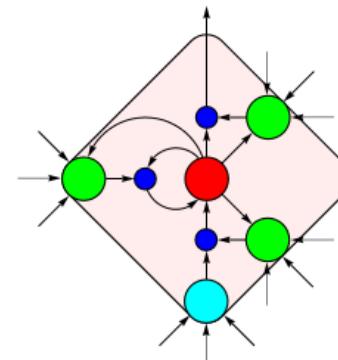
How to rescue it

- Use of LSTM (long short-term memory) cells
 - Sepp Hochreiter & Jürgen Schmidhuber (1997). "Long short-term memory" *Neural Computation* 9 (8): 1735–1780.
 - Many earlier-to-read materials, especially after 2013
- The best way so far to train RNNs well
 - Increasingly popular in speech/language processing
 - Attracted big attention from ASR community at ICASSP-2013's DNN special session (Graves et al.)
 - Huge progress since then

Many ways to show an LSTM cell



$$\begin{aligned} i_t &= \sigma(W_{xi}x_t + W_{hi}h_{t-1} + W_{ci}c_{t-1} + b_i) \\ f_t &= \sigma(W_{xf}x_t + W_{hf}h_{t-1} + W_{cf}c_{t-1} + b_f) \\ c_t &= f_t c_{t-1} + i_t \tanh(W_{xc}x_t + W_{hc}h_{t-1} + b_c) \\ o_t &= \sigma(W_{xo}x_t + W_{ho}h_{t-1} + W_{co}c_t + b_o) \\ h_t &= o_t \tanh(c_t) \end{aligned}$$

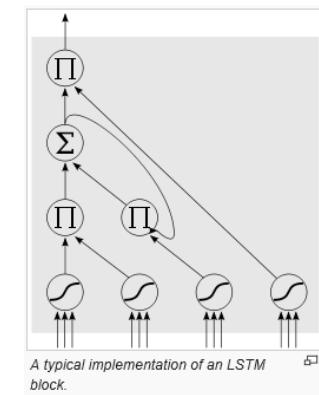


- ▶ Long Short-Term Memory (LSTM)^{52,53}
- ▶ Provides vanishing gradient problem solution
- ▶ Input, Output and Forget⁵⁴ gates flow information through
- ▶ Linear memory cell (called Constant Error Carousel – CEC)

⁵²Hochreiter:95fki207r.

⁵³Hochreiter:97lstm.

⁵⁴Gers:99a.



Many ways to show an LSTM cell

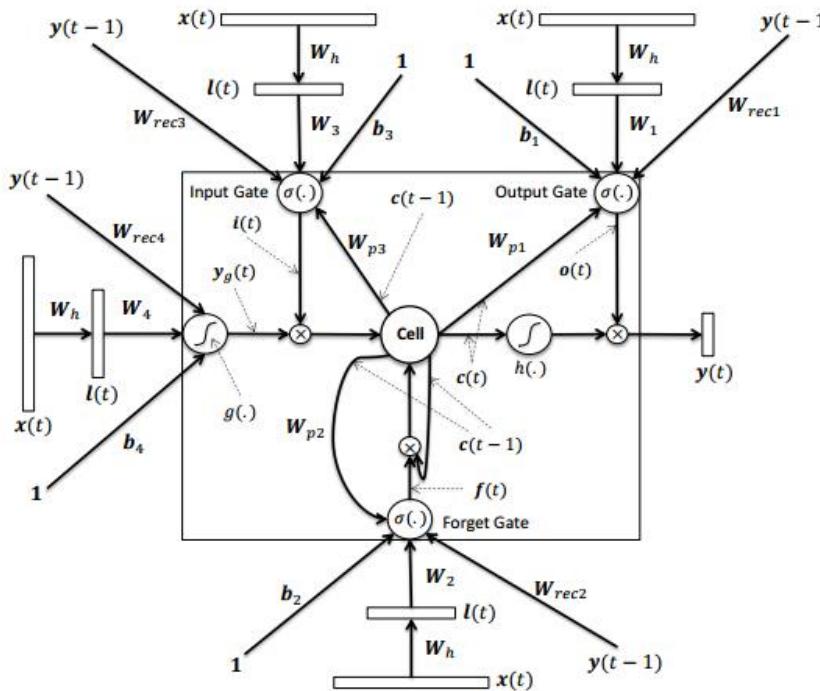


Fig. 2. The basic LSTM architecture used for sentence embedding

Deep Sentence Embedding Using Long Short-Term Memory Networks

Hamid Palangi, Li Deng, Yelong Shen, Jianfeng Gao, Xiaodong He, Jianshu Chen, Xinying Song,
Rabab Ward

Abstract—This paper develops a model that addresses sentence embedding, a hot topic in current natural language processing research, using recurrent neural networks (RNN) with Long Short-Term Memory (LSTM) cells. The proposed LSTM-RNN model sequentially takes each word in a sentence, extracts its information, and embeds it into a semantic vector. Due to its ability to capture long term memory, the LSTM-RNN accumulates increasingly richer information as it goes through the sentence, and when it reaches the last word, the hidden layer of the network

embedding is learned using sentence pairs. As a result, sentence embedding can better discover salient words and topics in a sentence, and thus is more suitable for tasks that require computing semantic similarities between text strings. By mapping texts into a unified semantic representation, the embedding vector can be further used for different language processing applications, such as machine translation [1], sentiment analysis [2], and information retrieval [3]. In machine translation

Jul 2015

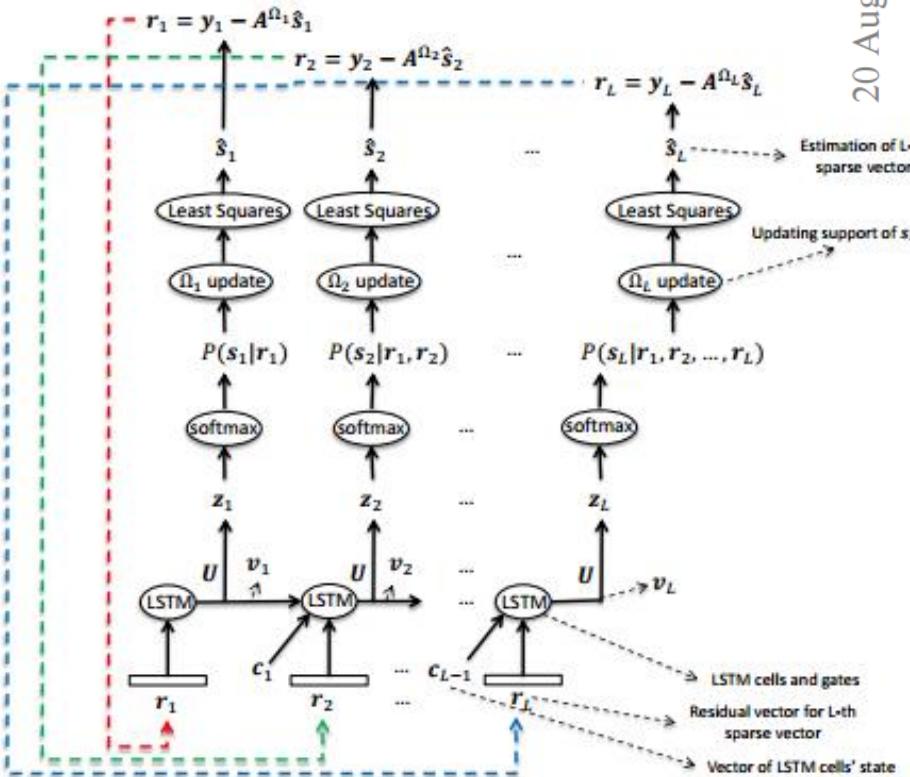
$$\begin{aligned}
 y_g(t) &= g(W_4 l(t) + W_{rec4} y(t-1) + b_4) \\
 i(t) &= \sigma(W_3 l(t) + W_{rec3} y(t-1) + W_{p3} c(t-1) + b_3) \\
 f(t) &= \sigma(W_2 l(t) + W_{rec2} y(t-1) + W_{p2} c(t-1) + b_2) \\
 c(t) &= f(t) \circ c(t-1) + i(t) \circ y_g(t) \\
 o(t) &= \sigma(W_1 l(t) + W_{rec1} y(t-1) + W_{p1} c(t) + b_1) \\
 y(t) &= o(t) \circ h(c(t))
 \end{aligned} \tag{2}$$

Many ways to show an LSTM cell

Distributed Compressive Sensing: A Deep Learning Approach

Hamid Palangi, Rabab Ward, Li Deng

20 Aug 2015



Abstract—We address the problem of compressed sensing with Multiple Measurement Vectors (MMVs) when the structure of sparse vectors in different channels depend on each other. *The sparse vectors are not necessarily joint sparse*. We capture this dependency by computing the conditional probability of each entry of each sparse vector to be non-zero given “residuals” of all previous sparse vectors. To compute these probabilities, we propose to use Long Short-Term Memory (LSTM) [1], a bottom up data driven model for sequence modelling. To compute model parameters we minimize a cross entropy cost function. We propose a greedy solver that uses above probabilities at

where $\mathbf{y} \in \mathbb{R}^{M \times 1}$ is the known measured vector and $\Phi \in \mathbb{R}^{M \times N}$ is a random measurement matrix. An important assumption needed by the decoder to uniquely recover \mathbf{x} given \mathbf{y} and Φ , is that \mathbf{x} is sparse in a given basis Ψ . This means that

$$\mathbf{x} = \Psi \mathbf{s} \quad (2)$$

where \mathbf{s} is K -sparse, i.e., \mathbf{s} has at most K non-zero elements. The basis Ψ can be complete; i.e., $\Psi \in \mathbb{R}^{N \times N}$, or over-complete; i.e., $\Psi \in \mathbb{R}^{N \times N_1}$ where $N < N_1$ (compressed sensing for over-complete dictionaries is

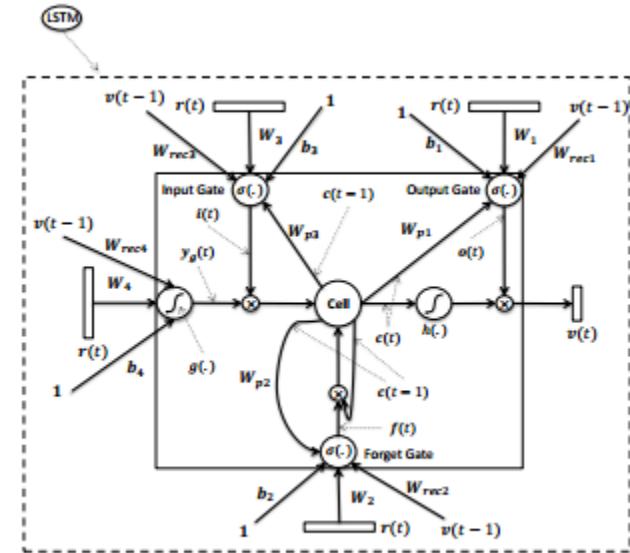


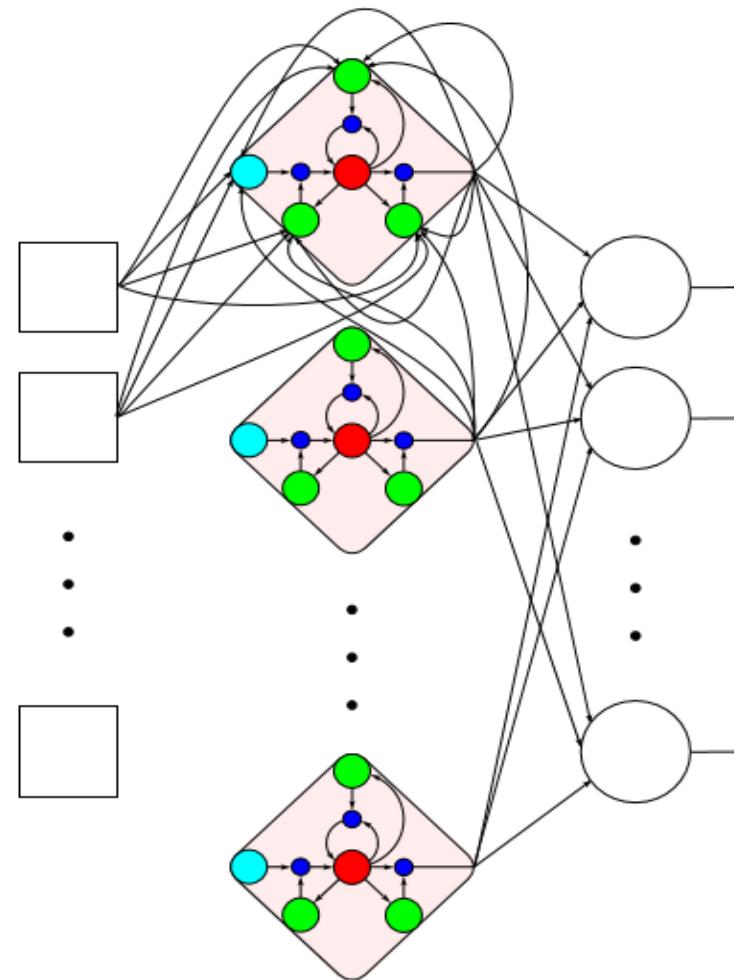
Fig. 2. Block diagram of the Long Short-Term Memory (LSTM).

Huh? . . .

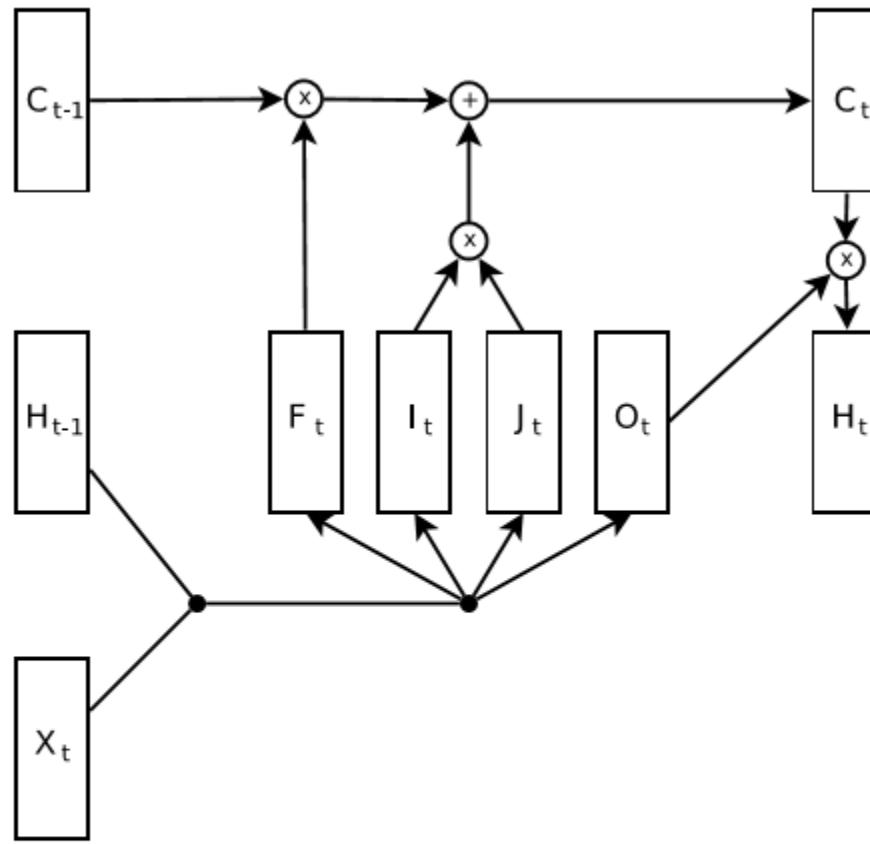


Don't worry, we'll come back to this shortly

LSTM Cells in an RNN



LSTM cell unfolding over time



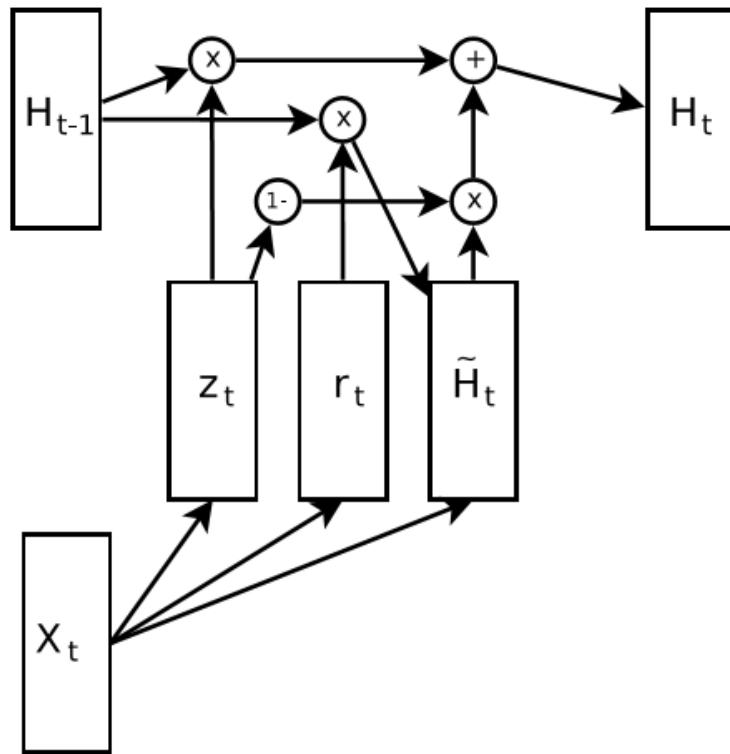
$$\begin{aligned}
 i_t &= \tanh(W_{xi}x_t + W_{hi}h_{t-1} + b_i) \\
 j_t &= \text{sigm}(W_{xj}x_t + W_{hj}h_{t-1} + b_j) \\
 f_t &= \text{sigm}(W_{xf}x_t + W_{hf}h_{t-1} + b_f) \\
 o_t &= \tanh(W_{xo}x_t + W_{ho}h_{t-1} + b_o) \\
 c_t &= c_{t-1} \odot f_t + i_t \odot j_t \\
 h_t &= \tanh(c_t) \odot o_t
 \end{aligned}$$

Figure 1. The LSTM architecture. The value of the cell is increased by $i_t \odot j_t$, where \odot is element-wise product. The LSTM's output is typically taken to be h_t , and c_t is not exposed. The forget gate f_t allows the LSTM to easily reset the value of the cell.

(Jozefowics, Zaremba, Sutskever,
ICML 2015)

Research Gated Recurrent Unit (GRU)

(simpler than LSTM; no output gates)

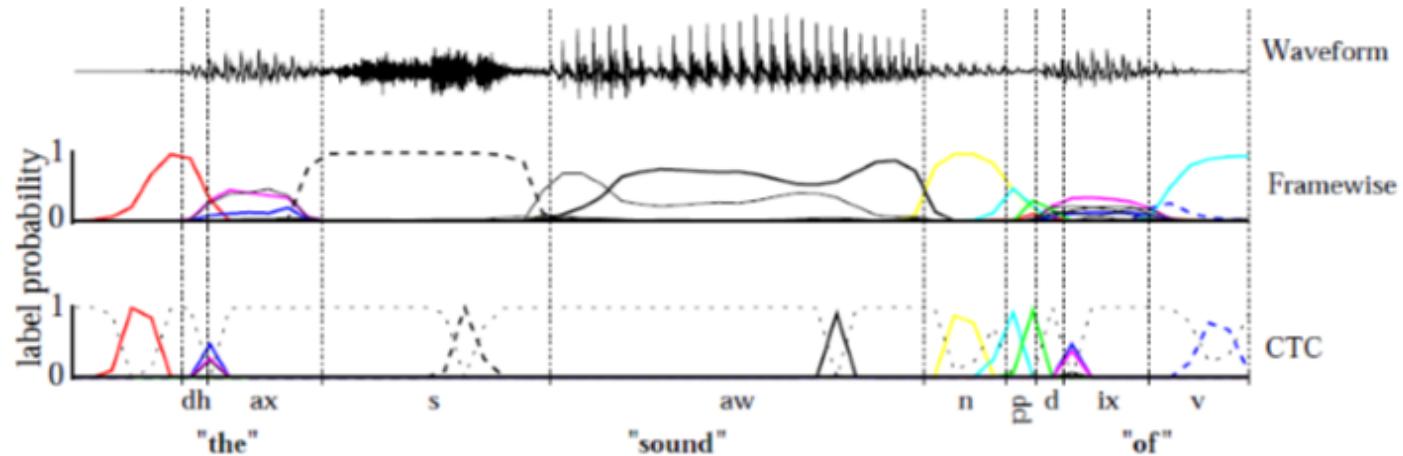


$$\begin{aligned}
 r_t &= \text{sigm}(W_{xr}x_t + W_{hr}h_{t-1} + b_r) \\
 z_t &= \text{sigm}(W_{xz}x_t + W_{hz}h_{t-1} + b_z) \\
 \tilde{h}_t &= \tanh(W_{xh}x_t + W_{hh}(r_t \odot h_{t-1}) + b_h) \\
 h_t &= z_t \odot h_{t-1} + (1 - z_t) \odot \tilde{h}_t
 \end{aligned}$$

Figure 2. The Gated Recurrent Unit. Like the LSTM, it is hard to tell, at a glance, which part of the GRU is essential for its functioning.

(Jozefowics, Zaremba, Sutskever, ICML 2015; Google Kumar et al., arXiv, July, 2015; Metamind)

Connectionist Temporal Classification



- ▶ Connectionist Temporal Classification (CTC)⁶⁰
- ▶ Automatically transforms frame labels to classification segments
- ▶ Provides labels for all time steps
- ▶ Allows to back-propagate the error from sparse labels

⁶⁰Graves:06icml.

Alex Graves

Supervised Sequence Labelling with Recurrent Neural Networks



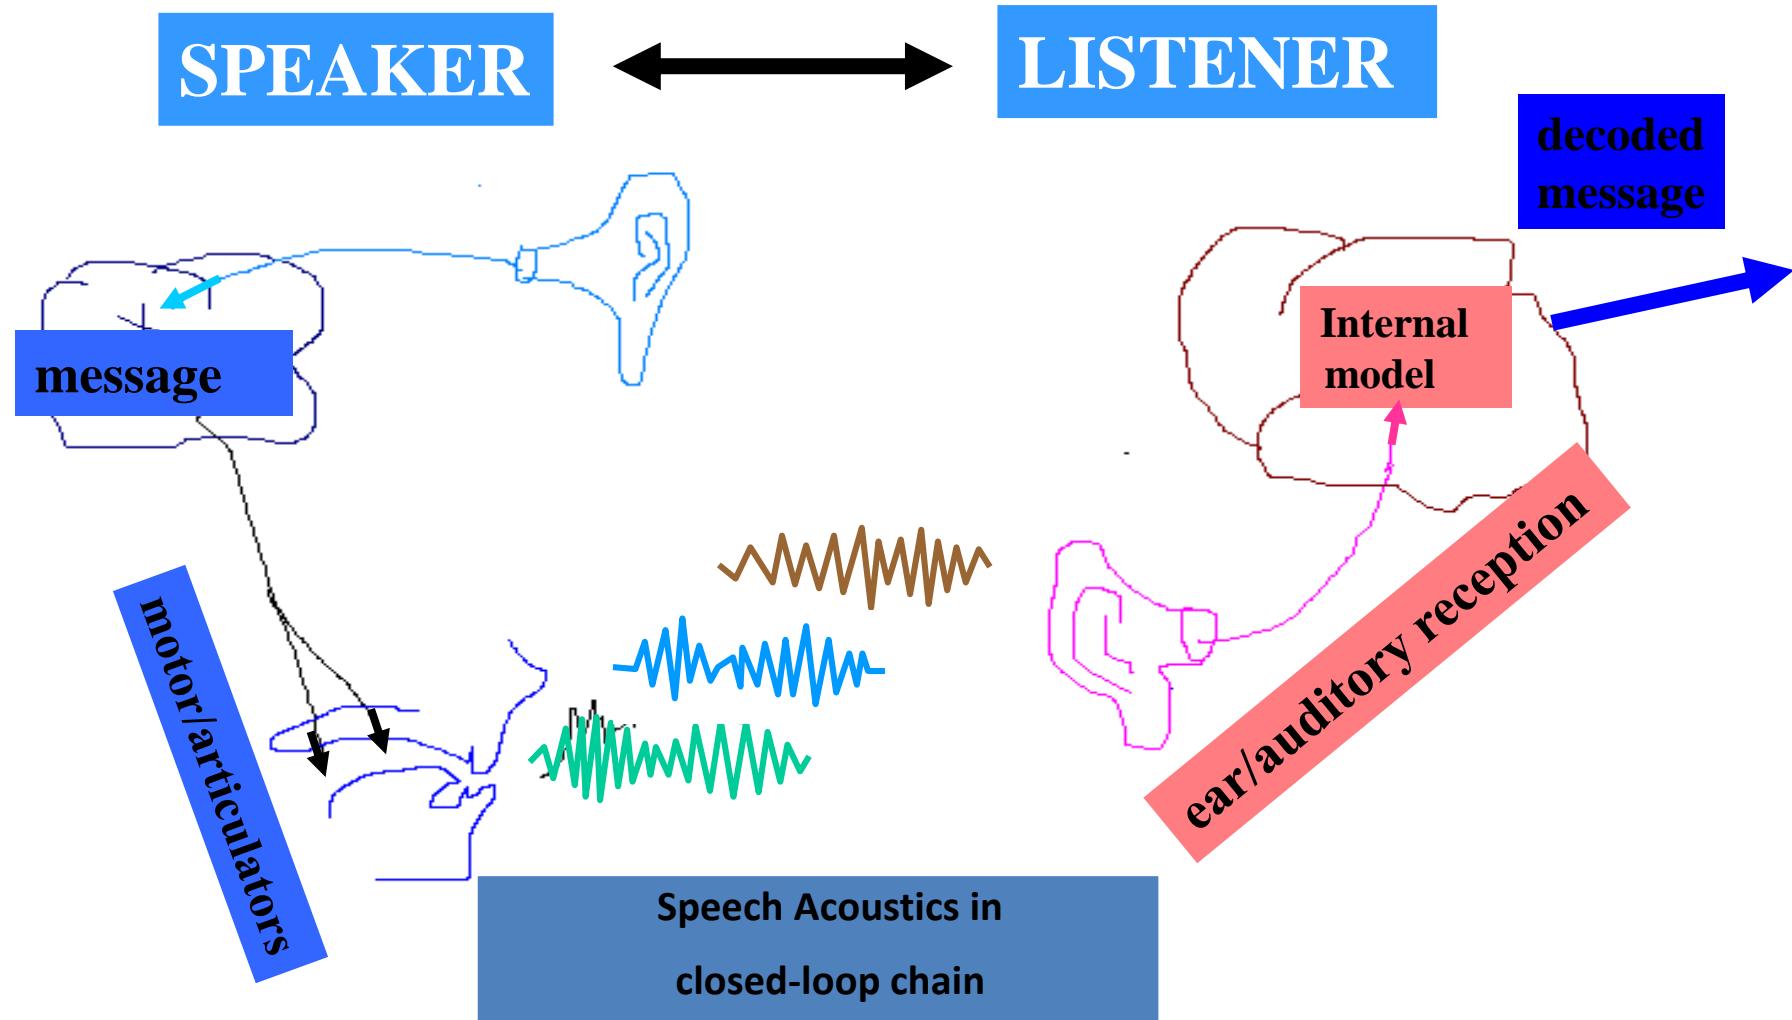
Springer

Part II: Speech

Deep/Dynamic Structure in Human Speech Production and Perception

(part of my tutorial at 2009 NIPS WS)

Production & Perception: Closed-Loop Chain

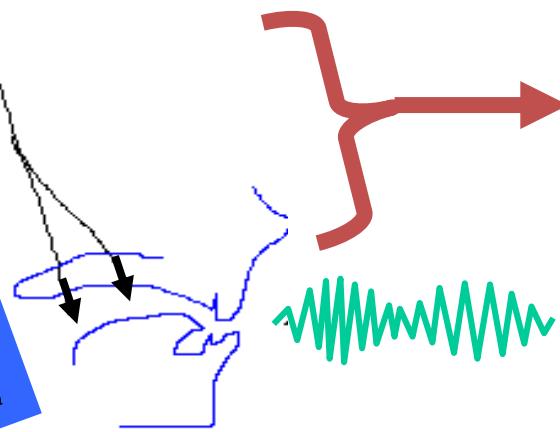


Encoder: Two-Stage Production Mechanisms

SPEAKER



motor/articulators



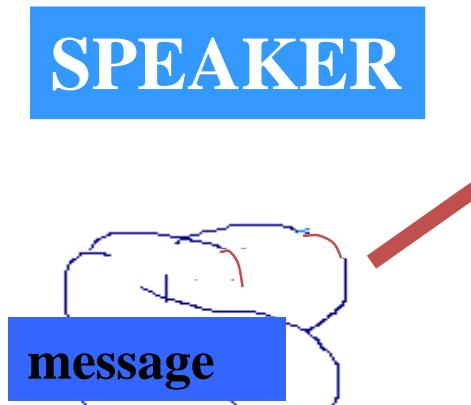
Phonology (higher level):

- Symbolic encoding of linguistic message
- Discrete representation by phonological features
- Loosely-coupled multiple feature tiers
- Overcome beads-on-a-string phone model
- Theories of distinctive features, feature geometry & articulatory phonology
- Account for partial/full sound deletion/modification in casual speech

Phonetics (lower level):

- Convert discrete linguistic features to continuous acoustics
- Mediated by motor control & articulatory dynamics
- Mapping from articulatory variables to VT area function to acoustics
- Account for co-articulation and reduction (target undershoot), etc.

Encoder: Phonological Modeling



Computational phonology:

- Represent pronunciation variations as constrained factorial Markov chain
- Constraint: from articulatory phonology
- Language-universal representation

ten themes

/ t ε n θ i: m z /

Labial-closure

Alveolar closure

Alveolar closure

dental constr.

Alveolar constr.

Mid / Front

High / Front

Nasality

Nasality

GLO:

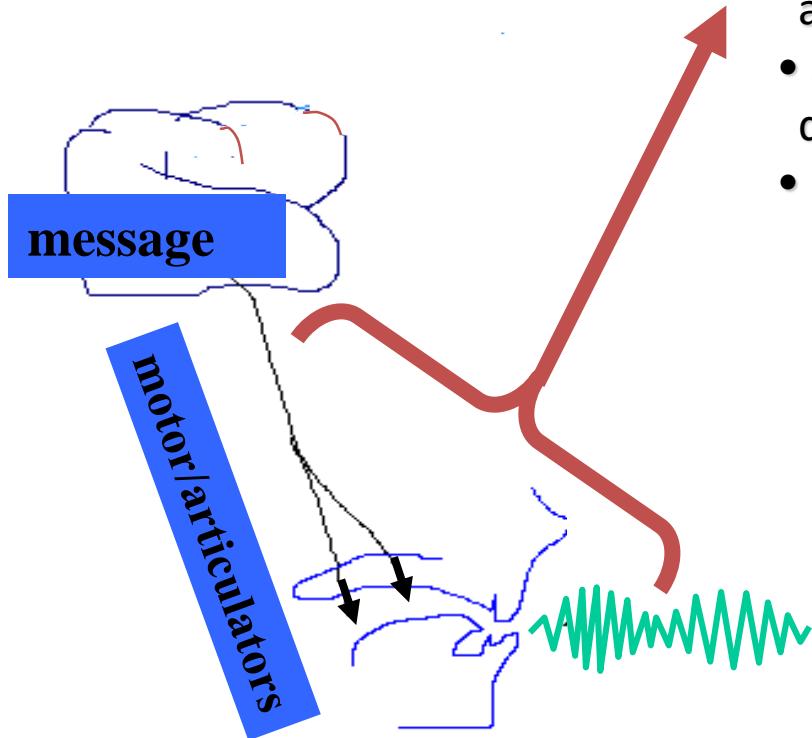
Aspiration

Voicing

Voicing

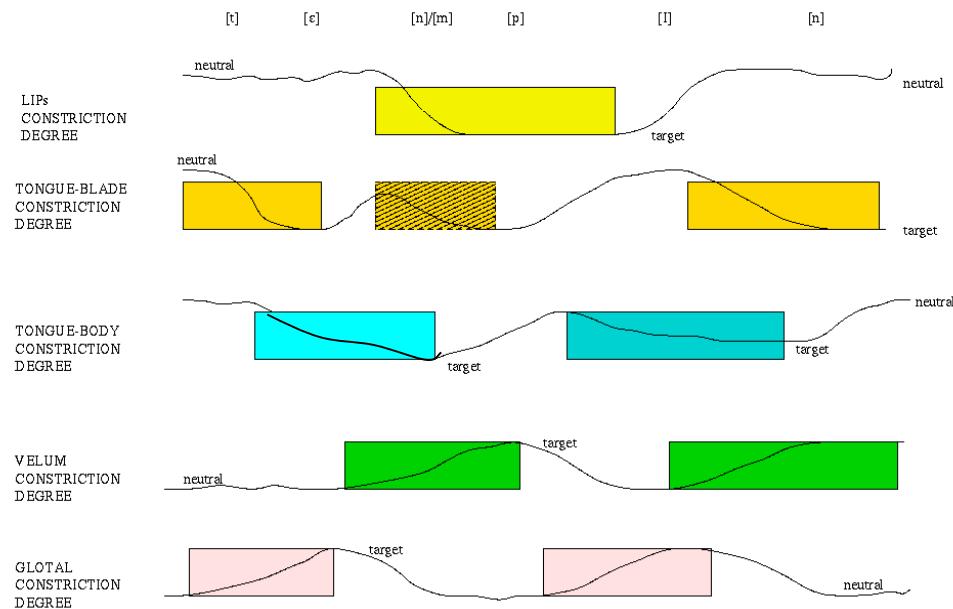
Encoder: Phonetic Modeling

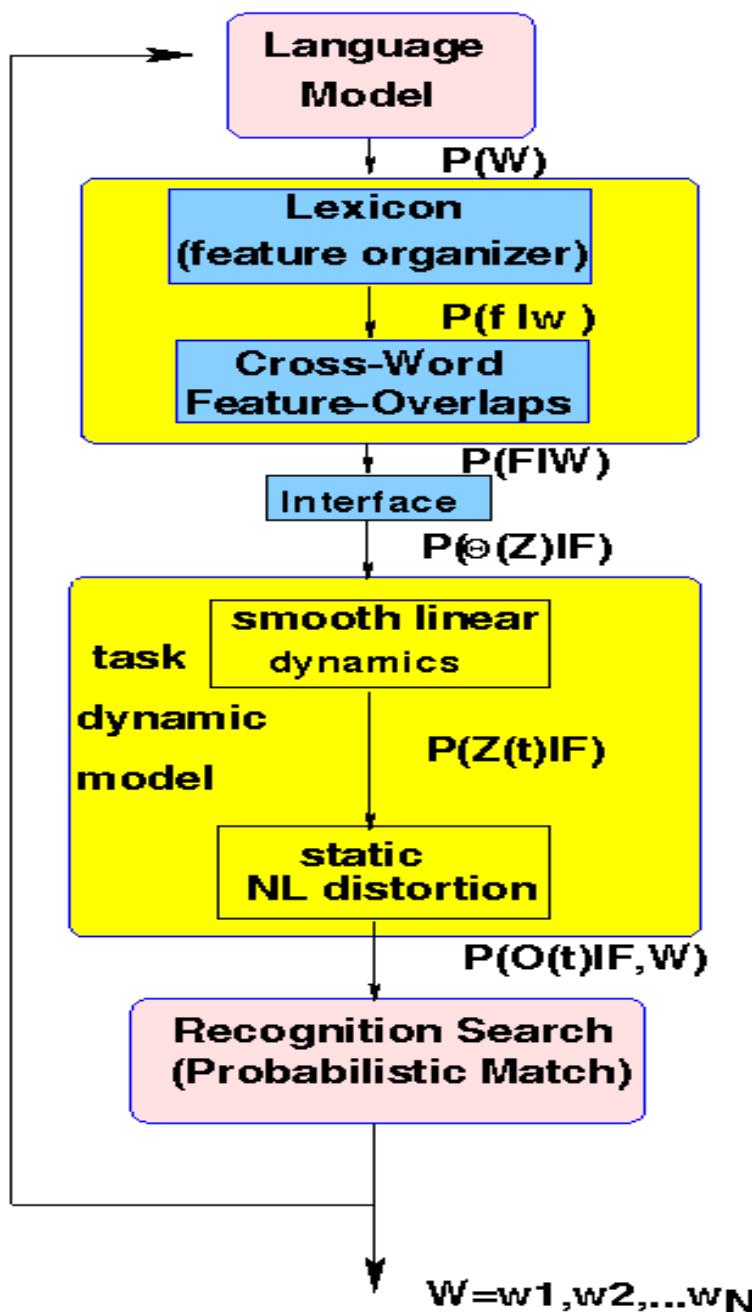
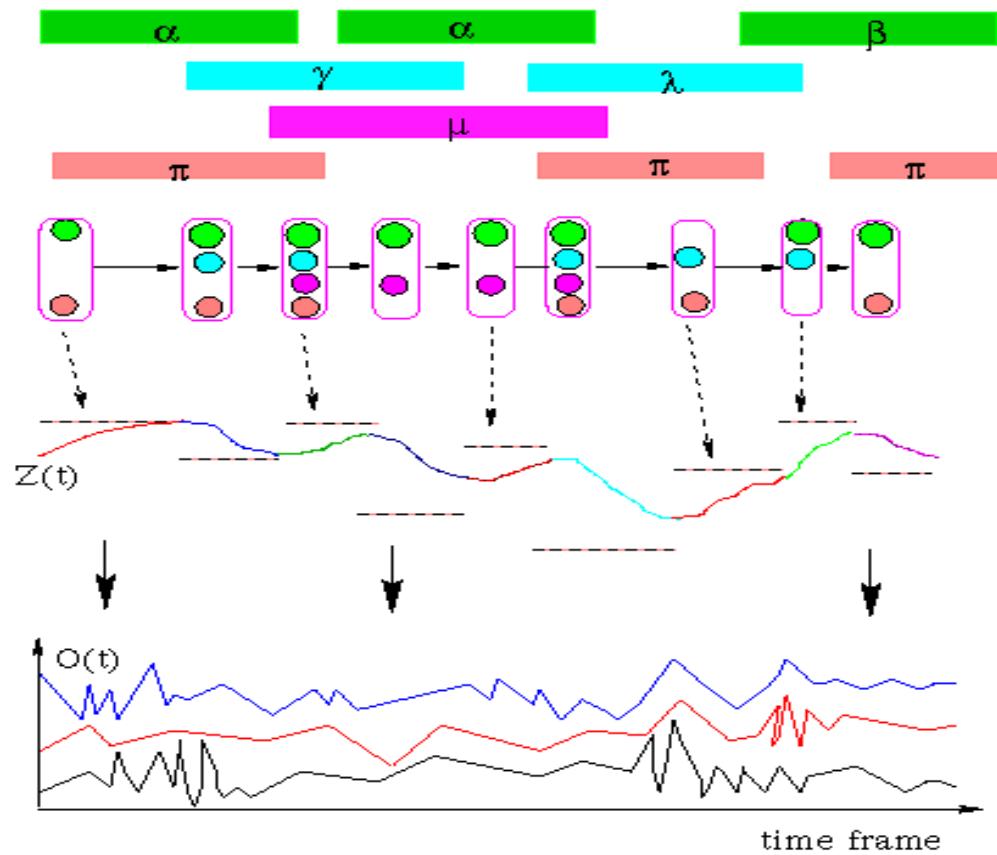
SPEAKER



Computational phonetics:

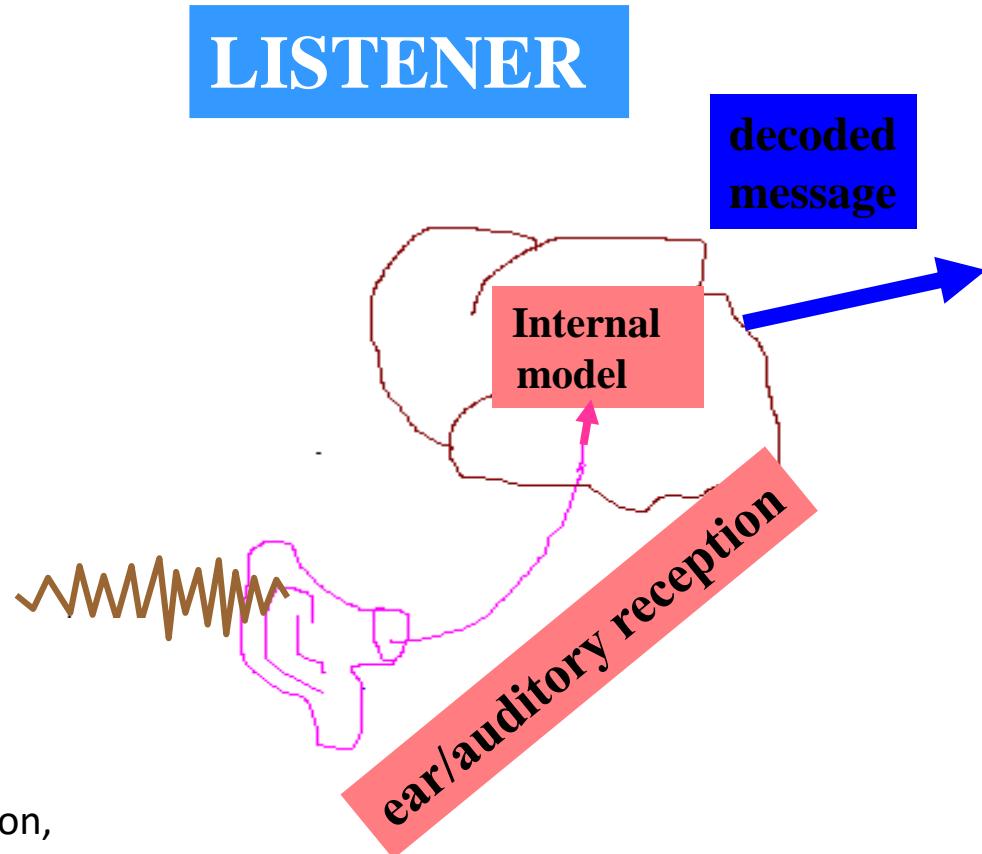
- Segmental factorial HMM for sequential target in articulatory or vocal tract resonance domain
- Switching trajectory model for target-directed articulatory dynamics
- Switching nonlinear state-space model for dynamics in speech acoustics
- Illustration:




 $W=w_1, w_2, \dots, w_N$

 $W=w_1, w_2, \dots, w_N$

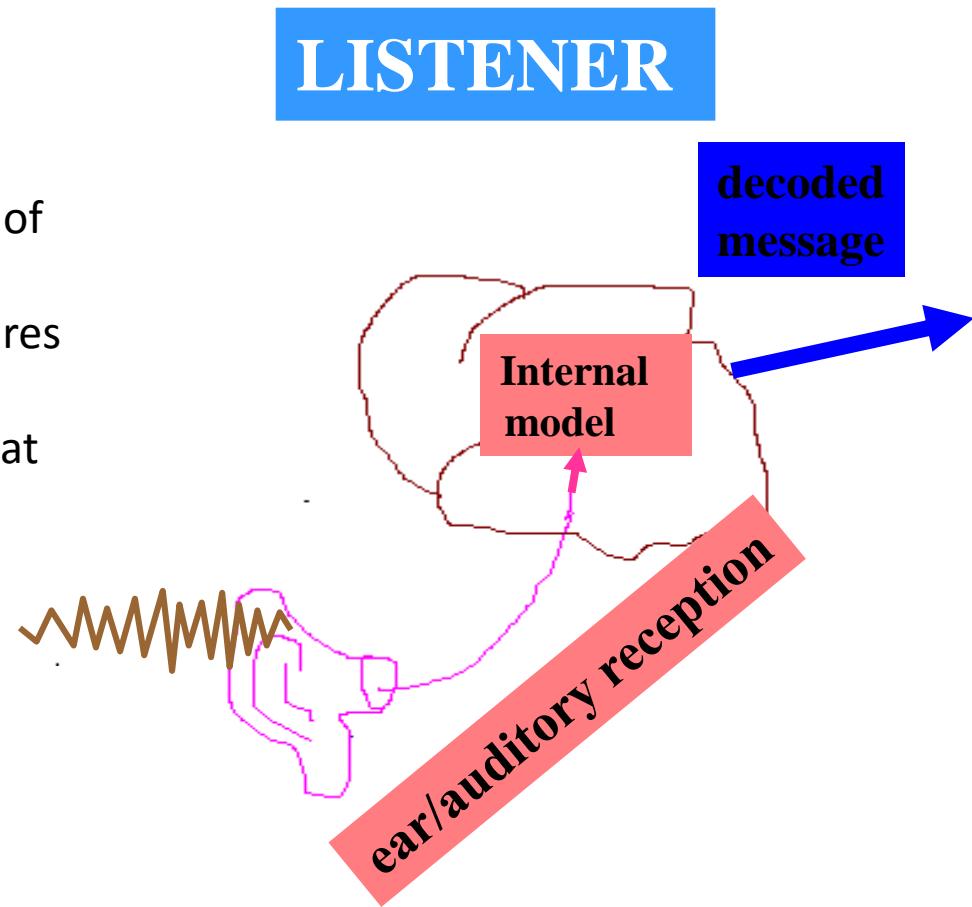
Decoder I: Auditory Reception

- Convert speech acoustic waves into efficient & robust auditory representation
- This processing is largely independent of phonological units
- Involves processing stages in cochlea (ear), cochlear nucleus, SOC, IC,..., all the way to A1 cortex
- Principal roles:
 - 1) combat environmental acoustic distortion;
 - 2) detect relevant speech features
 - 3) provide temporal landmarks to aid decoding
- Key properties:
 - 1) Critical-band freq scale, logarithmic compression,
 - 2) adapt freq selectivity, cross-channel correlation,
 - 3) sharp response to transient sounds
 - 4) modulation in independent frequency bands,
 - 5) binaural noise suppression, etc.



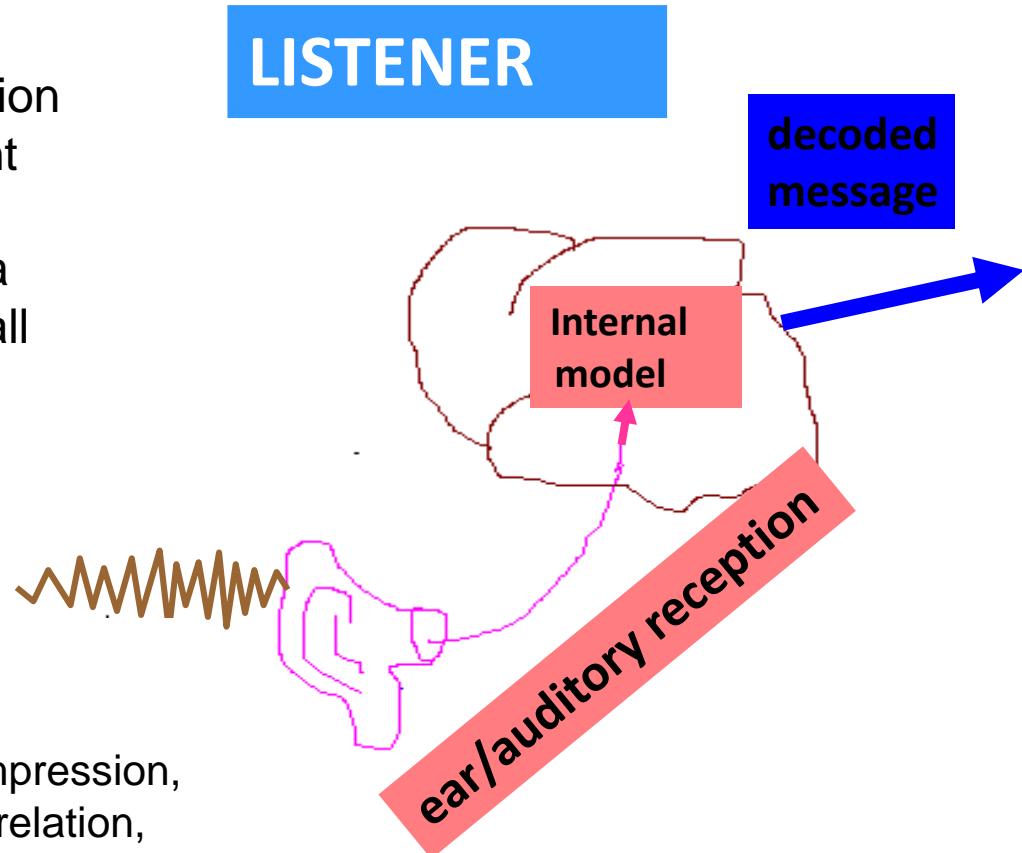
Decoder II: Cognitive Perception

- Cognitive process: recovery of linguistic message
- Relies on
 - 1) “Internal” model: structural knowledge of the encoder (production system)
 - 2) Robust auditory representation of features
 - 3) Temporal landmarks
- Child speech acquisition process is one that gradually establishes the “internal” model
- Strategy: analysis by synthesis
- i.e., Probabilistic inference on (deeply) hidden linguistic units using the internal model
- No motor theory: the above strategy requires no articulatory recovery from speech acoustics



Human Speech Perception (decoder)

- Convert speech acoustic waves into efficient & robust auditory representation
- This processing is largely independent of phonological units
- Involves processing stages in cochlea (ear), cochlear nucleus, SOC, IC,..., all the way to A1 cortex
- Two principal roles:
 - 1) combat environmental acoustic distortion;
 - 2) provide temporal landmarks to aid decoding
- Key properties:
 - 1) Critical-band freq scale, logarithmic compression,
 - 2) adapt freq selectivity, cross-channel correlation,
 - 3) sharp response to transient sounds (CN),
 - 4) modulation in independent frequency bands,
 - 5) binaural noise suppression, etc.



Types of Speech Perception Theories

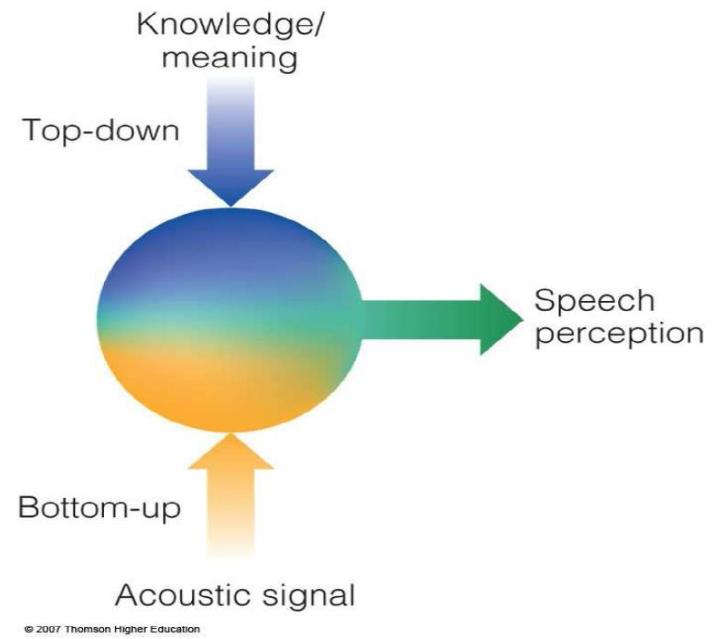
- Active vs. Passive
- Bottom up vs./and Top Down
- Autonomous vs. Interactive

Active vs. Passive

- **Active theories** suggests that speech perception and production are closely related
 - Listener knowledge of how sounds are produced facilitates recognition of sounds
- **Passive theories** emphasizes the sensory aspects of speech perception
 - Listeners utilize internal filtering mechanisms
 - Knowledge of vocal tract characteristics plays a minor role, for example when listening in noise conditions

Bottom up ~~vs.~~ & Top Down

- Top-down processing works with knowledge a listener has about a language, context, experience, etc.
 - Listeners use stored information about language and the world to make sense of the speech
- Bottom-up processing works in the absence of a knowledge base providing top-down information
 - listeners receive auditory information, convert it into a neural signal and process the phonetic feature information



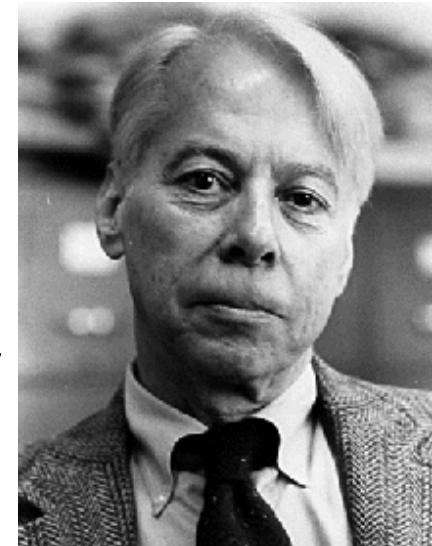
© 2007 Thomson Higher Education

Specific Speech Perception Theories

- Motor Theory
- Acoustic Invariance Theory
- Direct Realism
- Trace Model (based on neural nets)
- Cohort Theory
- Fuzzy Logic Model of Perception
- Native Language Magnet Theory

Motor Theory

- Postulates speech is perceived by reference to how it is produced
 - when perceiving speech, listeners access their own knowledge of how phonemes are articulated
 - Articulatory gestures (such as rounding or pressing the lips together) are units of perception that directly provide the listener with phonetic information



Liberman, Cooper, Shankweiler, & Studdert-Kennedy, 1967

Acoustic Invariance Theory

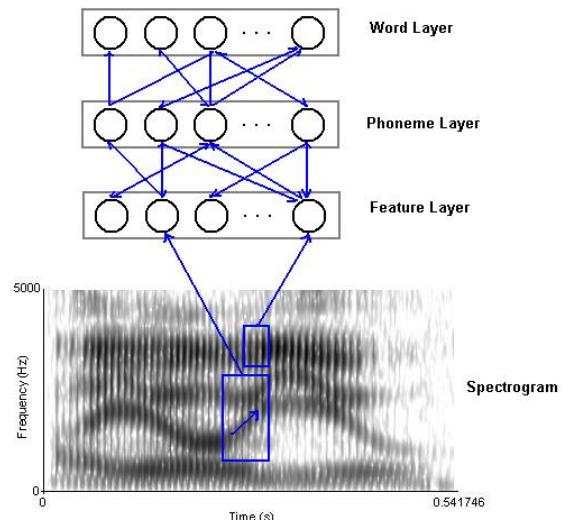
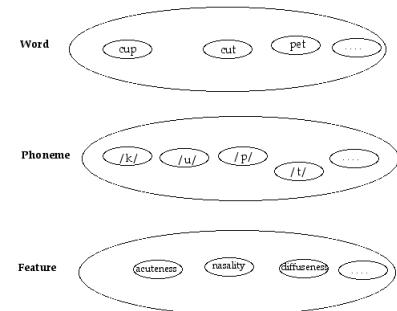
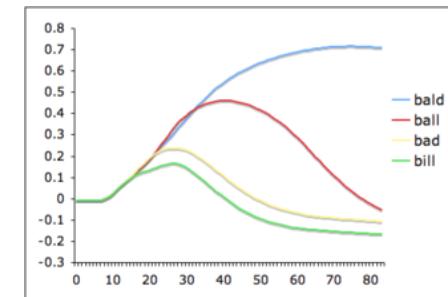
- Listeners inspect the incoming signal for the so-called acoustic landmarks which are particular events in the spectrum carrying information about gestures which produced them.
- Gestures are limited by the capacities of humans' articulators and listeners are sensitive to their auditory correlates, the lack of invariance simply does not exist in this model.
- The acoustic properties of the landmarks constitute the basis for establishing the distinctive features.
- Bundles of the distinctive features uniquely specify phonetic segments (phonemes, syllables, words).



Stevens, K.N. (2002). ["Toward a model of lexical access based on acoustic landmarks and distinctive features"](#) ([PDF](#)). *Journal of the Acoustical Society of America* 111 (4): 1872–1891.

TRACE Model

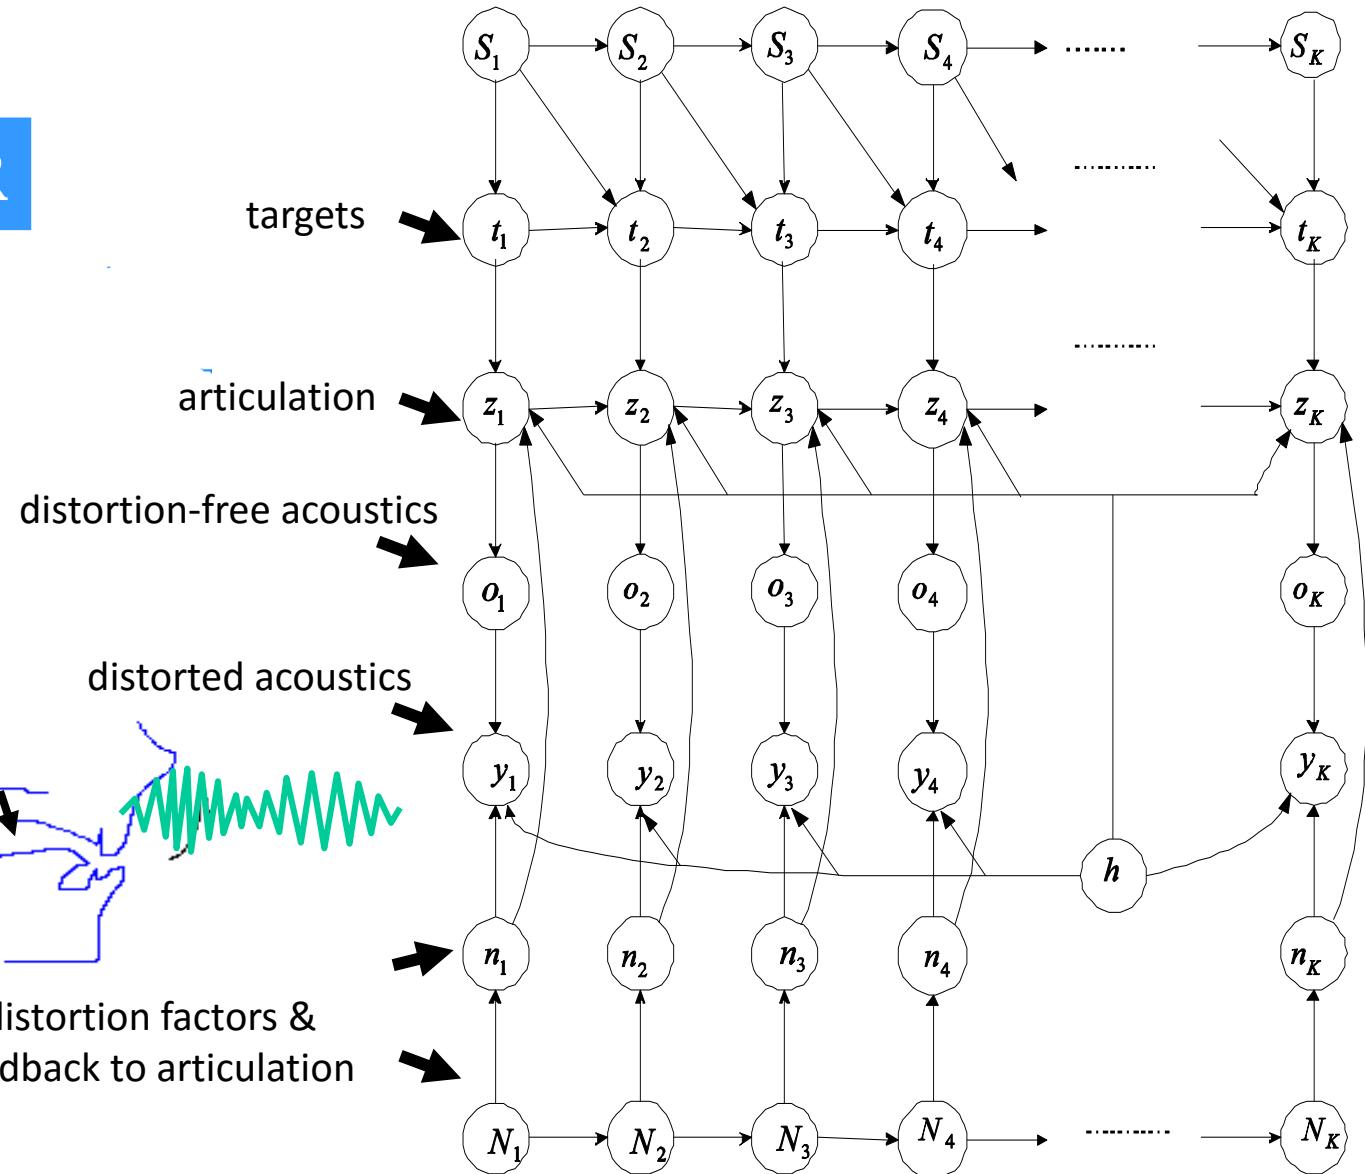
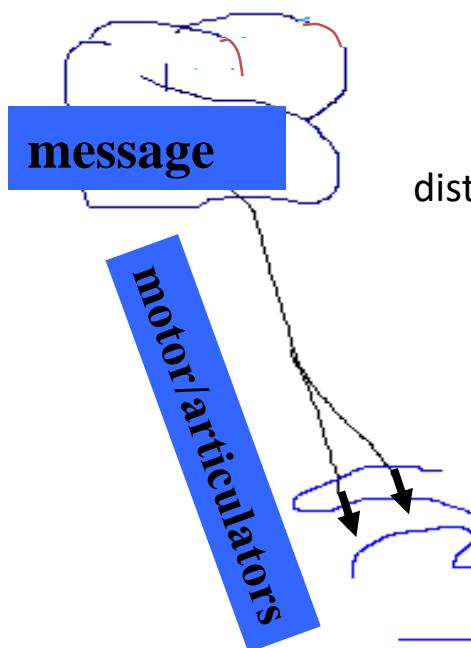
- For example, a listener hears the beginning of *bald*, and the words bald, ball, bad, bill become active in memory. Then, soon after, only bald and ball remain in competition (bad, bill have been eliminated because the vowel sound doesn't match the input).
 - Soon after, bald is recognized.
- TRACE simulates this process by representing the temporal dimension of speech, allowing words in the lexicon to vary in activation strength, and by having words compete during processing.



A Deep/Generative Model of Speech Production/Perception

--- Perception as “*variational inference*”

SPEAKER



Deep Generative Models, Variational Interference/Learning, & Applications to Speech

Deep Learning

\approx

Neural Networks, Deep

in space & time (recurrent LSTM), & deep RNN

+

Generative Models, Deep

in space & time (dynamic), & deep/hidden dynamic models

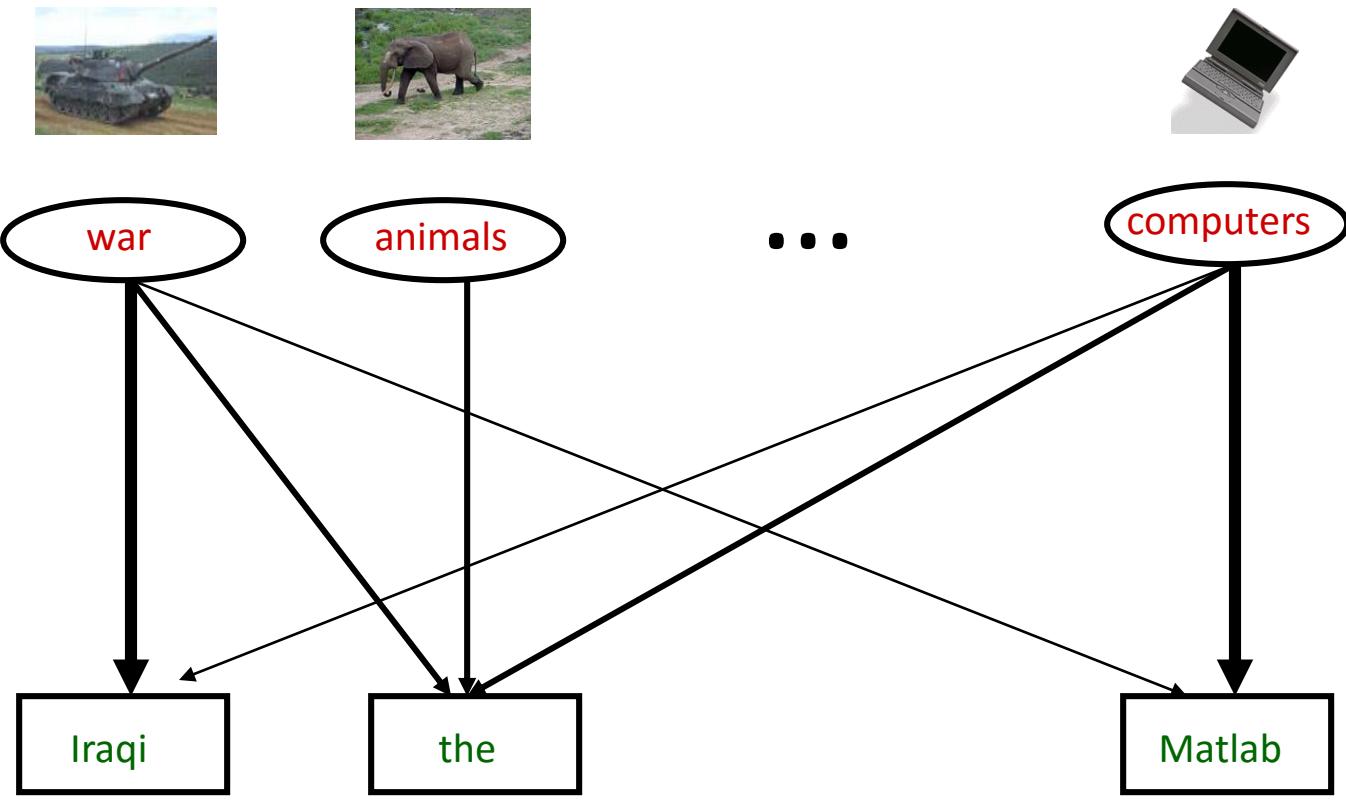
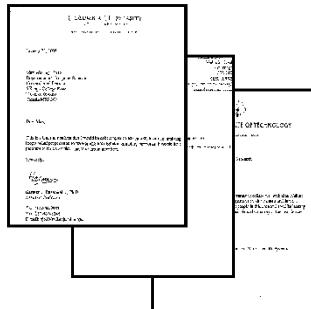
+

..., ..., ...

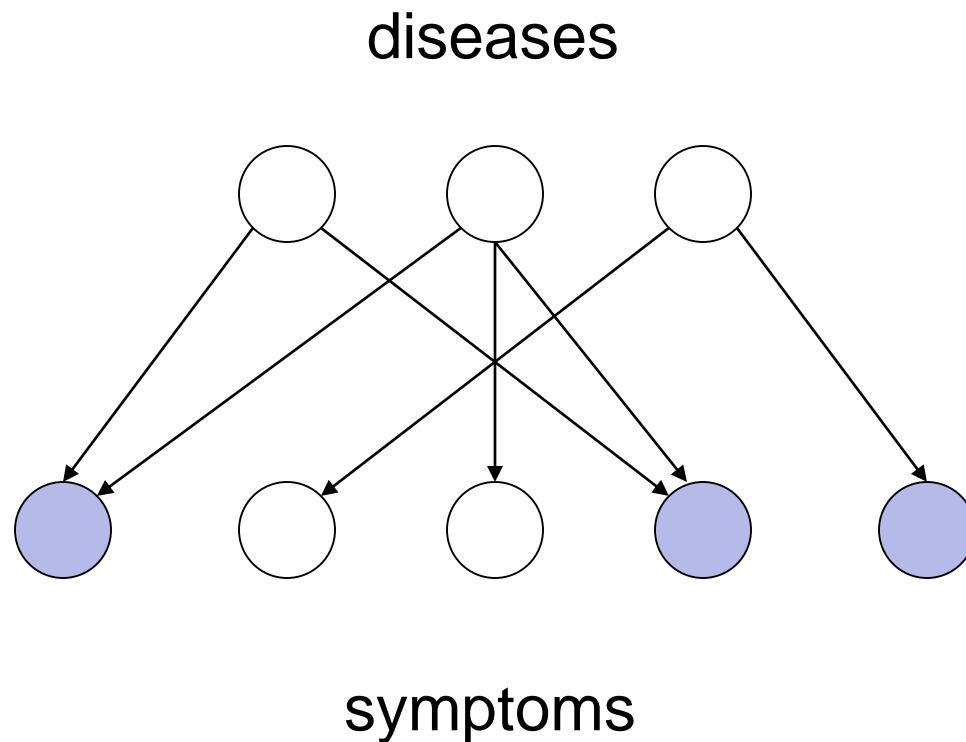
	Deep Neural Nets	Deep Generative Models
Structure	Graphical; info flow: bottom-up	Graphical; info flow: top-down
Incorp constraints & domain knowledge	Hard	Easy
Semi/ <i>unsupervised</i>	Harder or impossible	<i>Easier, at least possible</i>
Interpretation	Harder	Easy (generative “story” on data and hidden variables)
Representation	Distributed	Localist (mostly); can be distributed also
Inference/decode	Easy	Harder (but note recent progress)
Scalability/compute	Easier (regular computes/GPU)	Harder (but note recent progress)
Incorp. uncertainty	Hard	Easy
Empirical goal	Classification, feature learning, ...	Classification (via Bayes rule), latent variable inference...
Terminology	Neurons, activation/gate functions, weights ...	Random vars, stochastic “neurons”, potential function, parameters ...
Learning algorithm	A single, unchallenged, algorithm -- BackProp	A major focus of open research, many algorithms, & more to come
Evaluation	On a black-box score – end performance	On almost every intermediate quantity
Implementation	Hard (but increasingly easier)	Standardized but insights needed
Experiments	Massive, real data	Modest, often simulated data
Parameterization	Dense matrices	Sparse (often PDFs); can be dense

Example: (Shallow) Generative Model

“TOPICS”
as hidden layer

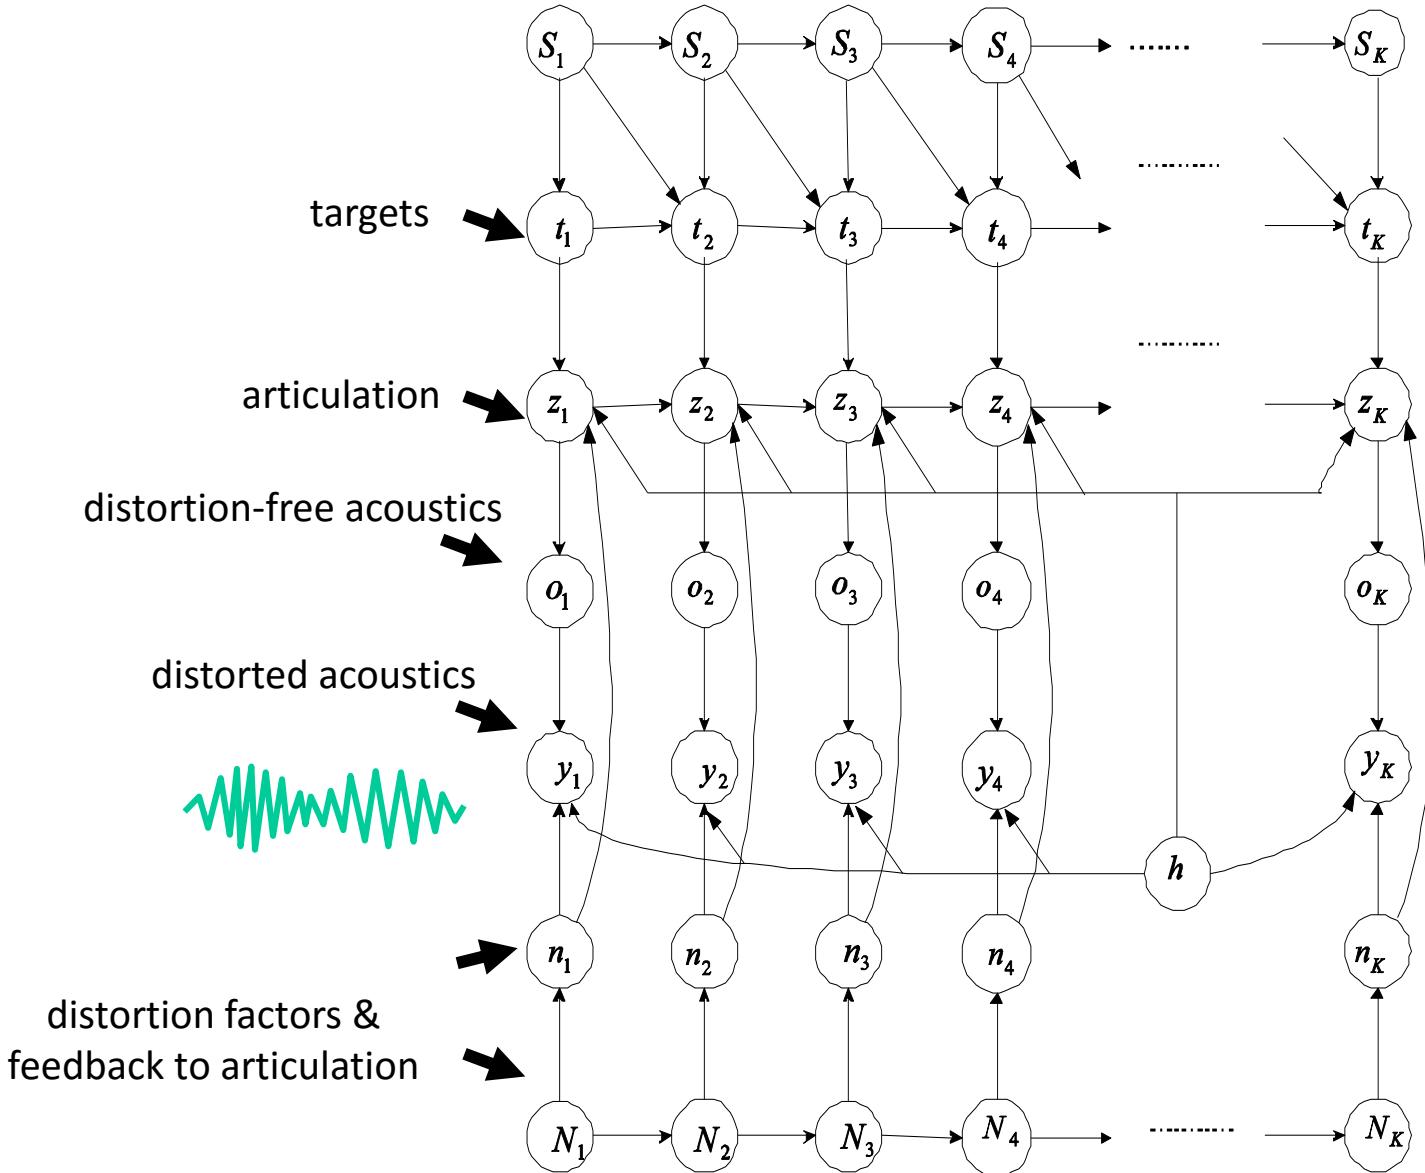


Another example: Medical Diagnosis



Inference problem:
What is the most
probable disease
given the
symptoms?

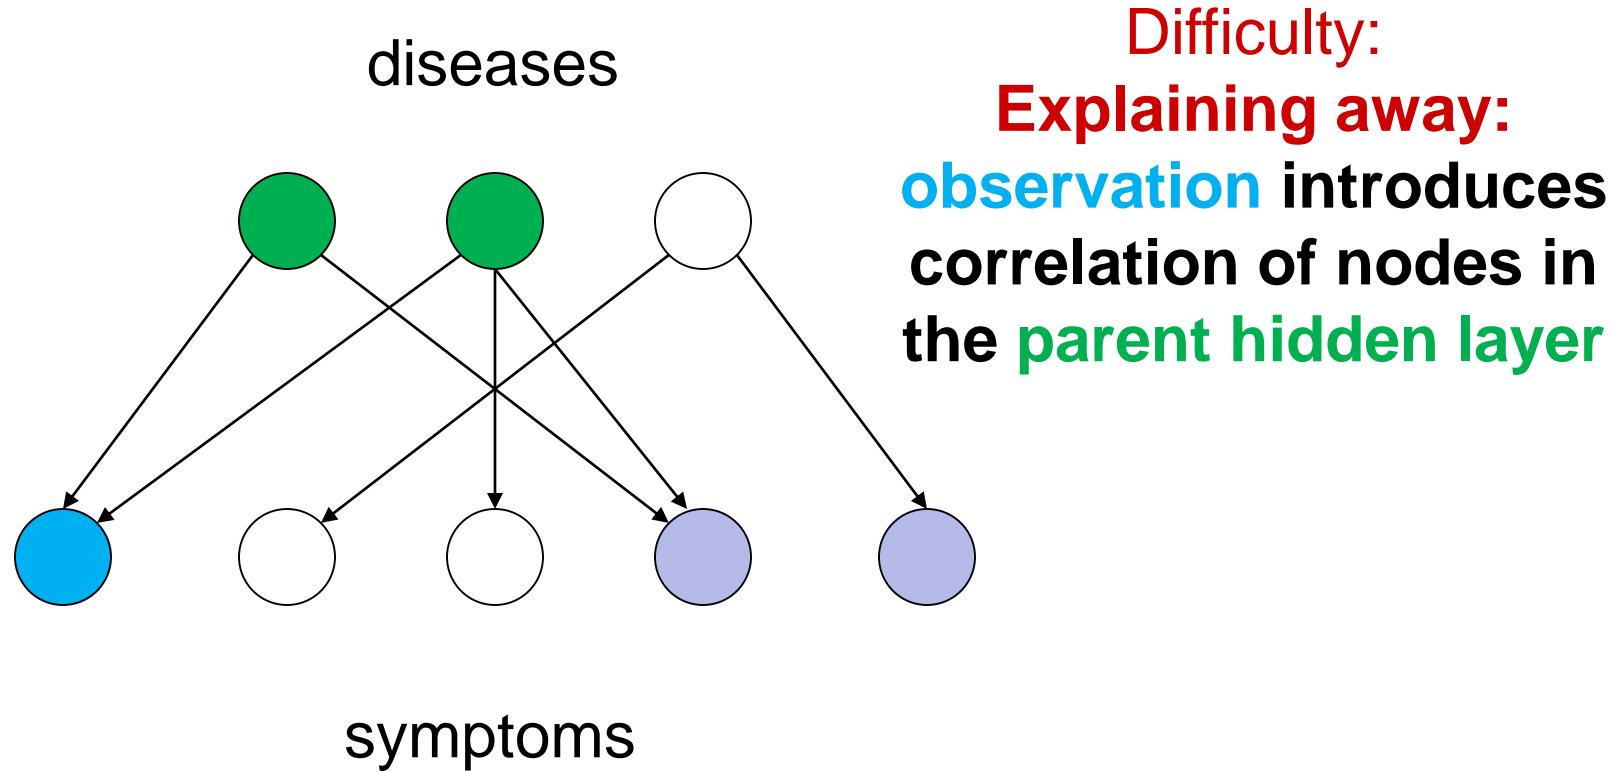
Example: (Deep) Generative Model



Deep Generative/Graphical Model Inference

- Key issues:
 - **Representation:** syntax and semantics (directed/undirected, variables/factors,...)
 - **Inference:** computing probabilities and most likely assignments/explanations
 - **Learning:** of model parameters based on observed data. *Relies on inference!*
- Inference is NP-hard (incl. approximation hardness)
- Exact inference: works for very limited subset of models/structures
 - E.g., chains or low-treewidth trees
- Approximate inference: highly computationally intensive
 - Deterministic: **Variational**, loopy belief propagation, expectation propagation
 - Numerical sampling (Monte Carlo): Gibbs sampling
- Variational learning:
 - EM algorithm
 - E-step uses variational inference (recent new advances in ML)

Variational inference/learning is not trivial: Example



Variational EM

Step 1: Maximize the bound with respect to Q

$$(\text{E step}): Q^{(k+1)} = \arg \max_Q L(Q, \theta^{(k)})$$

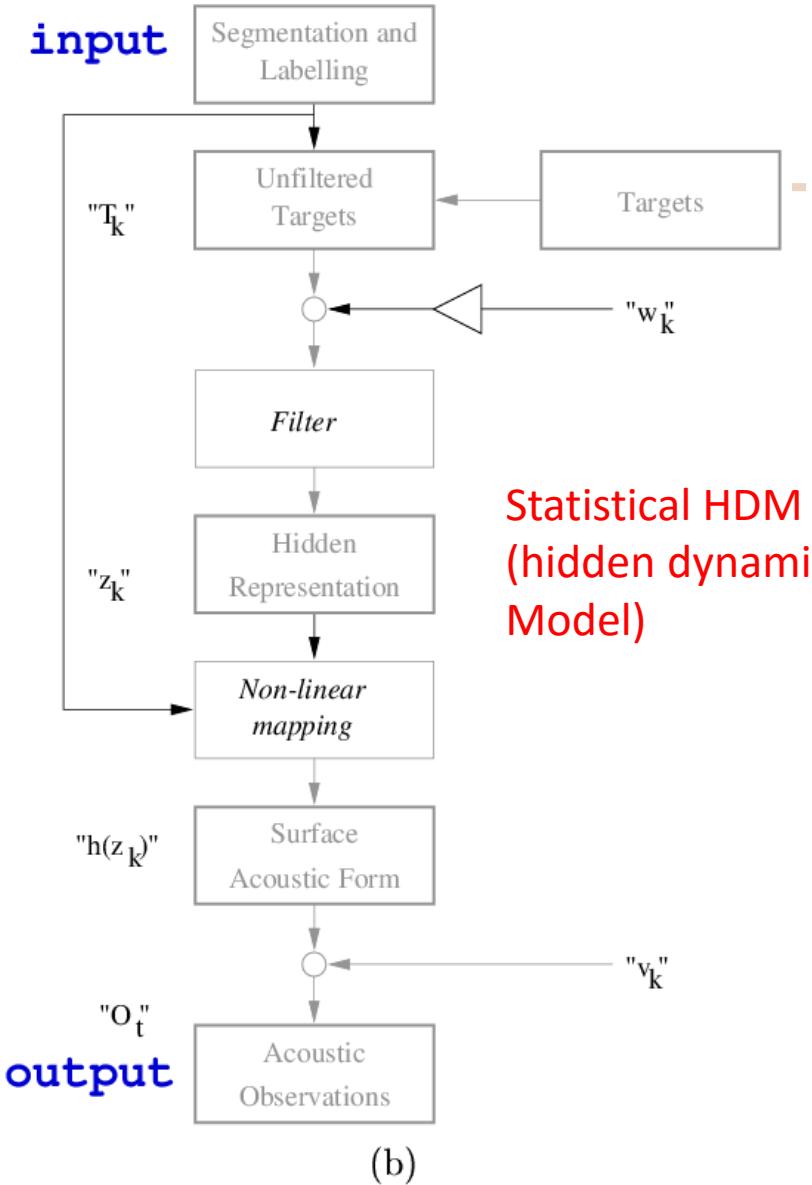
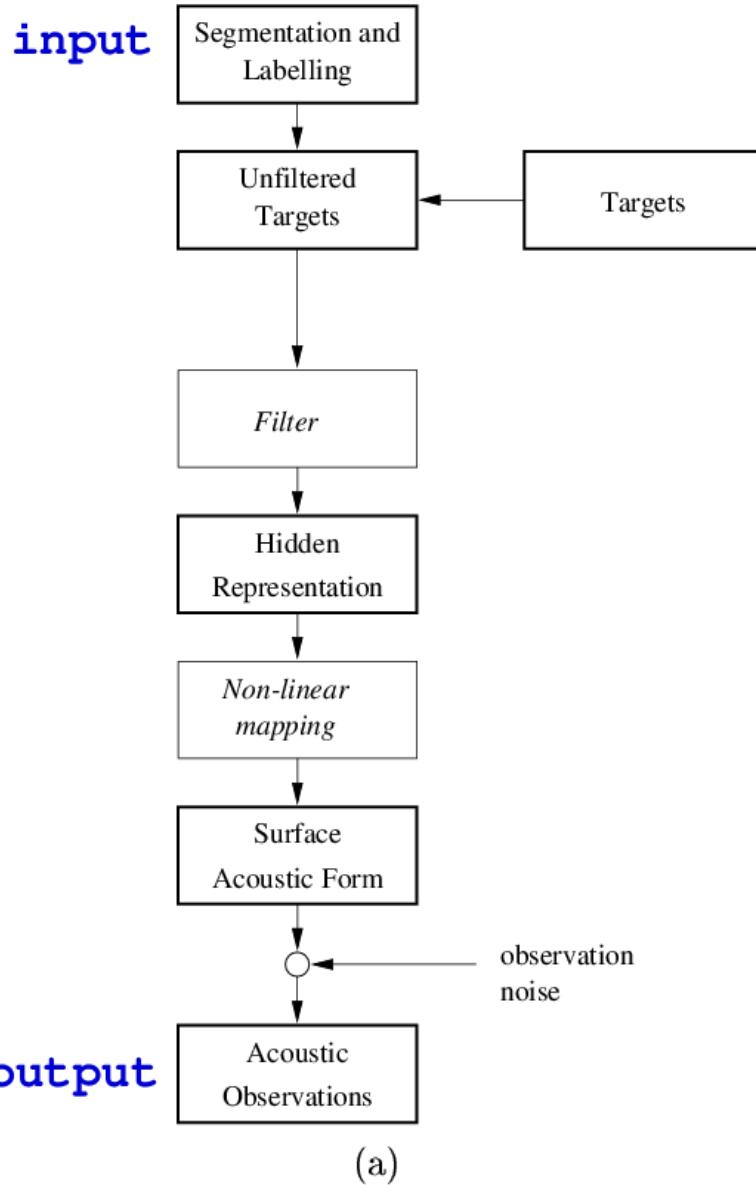
(Q approximates true posterior, often by factorizing it)

Step 2: Fix Q, maximize with respect to

$$(\text{M step}): \theta^{(k+1)} = \arg \max_\theta L(Q^{(k+1)}, \theta)$$

Note in traditional EM, Q is precise; e.g. posteriors computed by forward/backward algorithm for HMMs

Deterministic HDM



A MULTIMODAL VARIATIONAL APPROACH TO LEARNING AND INFERENCE IN SWITCHING STATE SPACE MODELS

Leo J. Lee^{1,2}, Hagai Attias², Li Deng² and Paul Fieguth³

University of Waterloo

¹Electrical & Computer Engineering
³Systems Design Engineering
 Waterloo, ON, N2L 3G1
 Canada

²Microsoft Corporation
 Microsoft Research
 One Microsoft Way
 Redmond, WA 98052-6339
 USA

Auxiliary function:

$$\mathcal{F}[q] = \sum_{s_{1:N}} \int d\mathbf{x}_{1:N} q(s_{1:N}, \mathbf{x}_{1:N}) \cdot [\log p(\mathbf{y}_{1:N}, \mathbf{x}_{1:N}, s_{1:N}) - \log q(s_{1:N}, \mathbf{x}_{1:N})]$$

In the variational approach we approximate the exact posterior $p(s_{1:N}, \mathbf{x}_{1:N} | \mathbf{y}_{1:N})$ by a distribution with a tractable structure, denoted by q . Here we choose the following partially factorized structure shown graphically in Fig. 1:

$$\begin{aligned} p(s_{0:N}, \mathbf{x}_{0:N} | \mathbf{y}_{1:N}) &\approx q(s_{0:N}, \mathbf{x}_{0:N} | \mathbf{y}_{1:N}) \\ &= \prod_{n=1}^N q(\mathbf{x}_n | s_n) q(s_n | s_{n-1}) \cdot q(\mathbf{x}_0 | s_0) q(s_0). \end{aligned} \quad (5)$$

E-step: sufficient statistics. As usual, the variational equations above are coupled, with the equations for $\rho_{s,n}$, $\Gamma_{s,n}$ depend on $\eta_{s's,n}$, $\gamma_{s,n}$ and vice versa. These equations are solved iteratively starting from a random or more suitable initialization if available. The solution is the set of sufficient statistics

$$\varphi = \{\rho_{s,n}, \Gamma_{s,n}, \eta_{ss',n}, \gamma_{s,n}\} \quad (16)$$

which are moments of the variational posterior.

M-step: parameter estimation. Given the sufficient statistics φ , the derivation of the M-step is achieved by taking derivatives of \mathcal{F} w.r.t. the model parameters (details omitted).

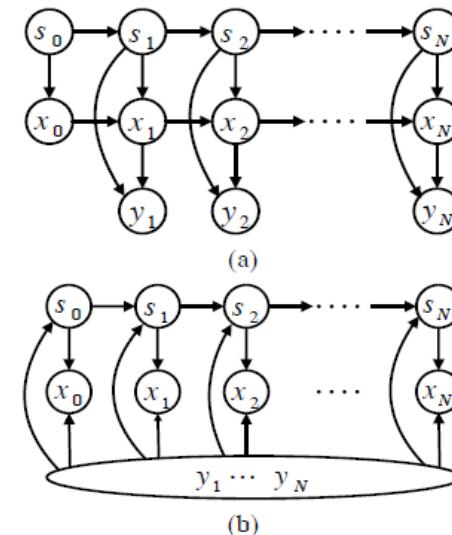
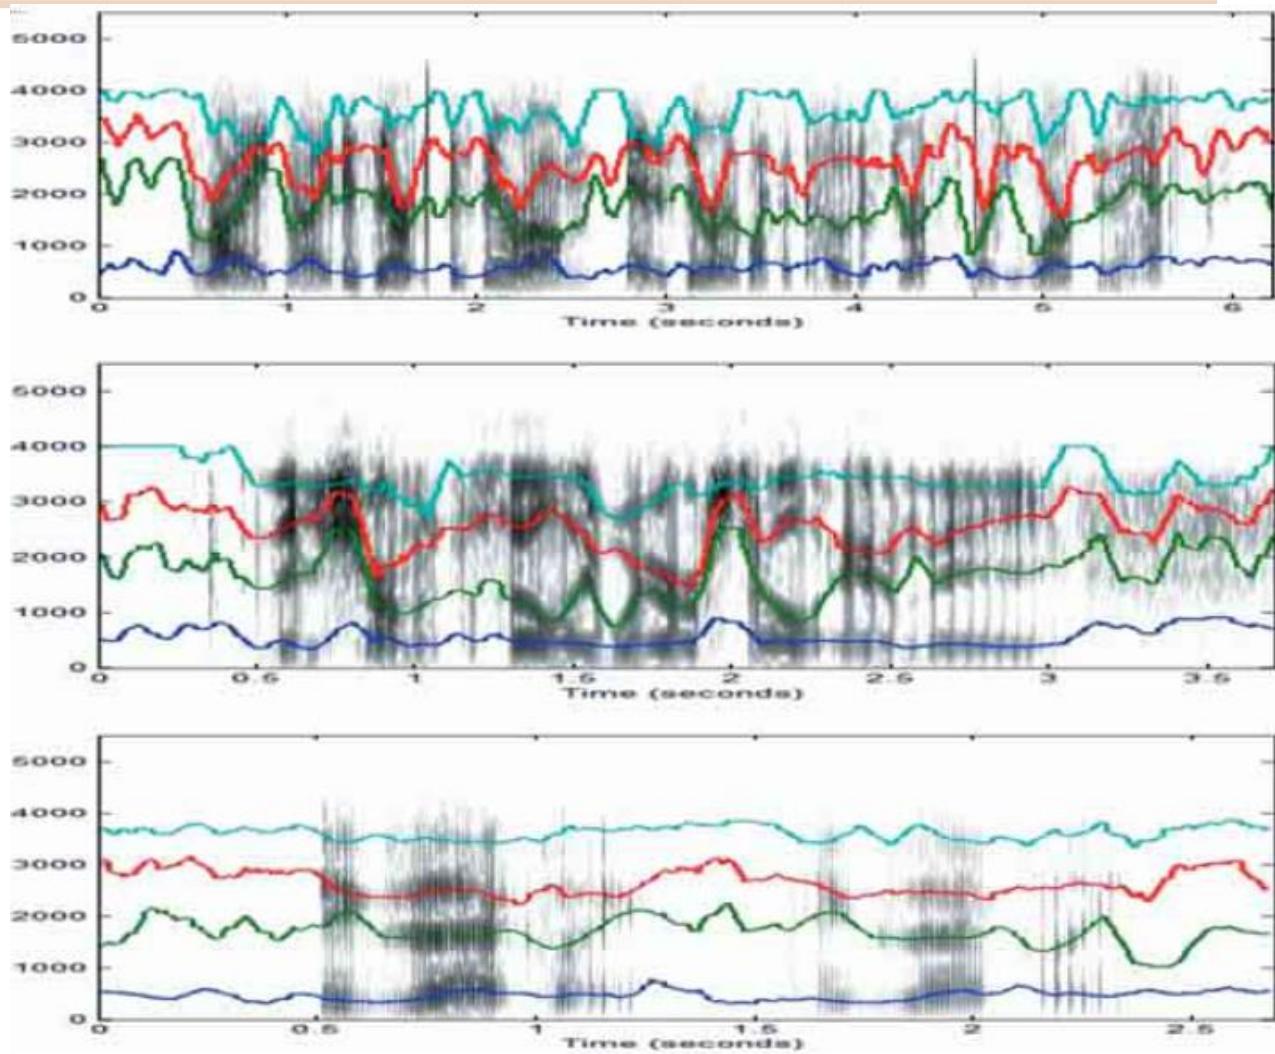


Fig. 1. The model (a) and the variational posterior (b) represented as Bayesian networks.

Surprisingly Good Inference Results for Continuous Hidden States

- By-product: accurately tracking dynamics of resonances (formants) in vocal tract (TIMIT & SWBD).
- Best formant tracker by then in speech analysis; used as basis to form a formant database as “ground truth”
- We thought we solved the ASR problem, except
- “Intractable” for decoding



Deng & Huang, [Challenges in Adopting Speech Recognition](#), Communications of the ACM, vol. 47, pp. 69-75, 2004.

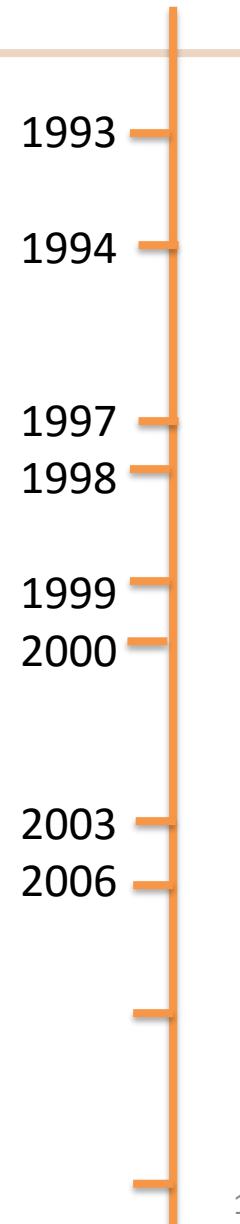
Deng, Cui, Pruvencok, Huang, Momen, Chen, Alwan, [A Database of Vocal Tract Resonance Trajectories for Research in Speech](#), ICASSP, 2006.

Deep Generative Models in Speech Recognition

(prior to the rising of deep learning)

Segment & Nonstationary-State Models

- Digalakis, Rohlicek, Ostendorf. "ML estimation of a stochastic linear system with the EM alg & application to speech recognition," IEEE T-SAP, 1993
- Deng, Aksmanovic, Sun, Wu, "Speech recognition using HMM with polynomial regression functions as nonstationary states," IEEE T-SAP, 1994.



Hidden Dynamic Models (HDM)

- Deng, Ramsay, Sun. "Production models as a structural basis for automatic speech recognition," Speech Communication, vol. 33, pp. 93–111, 1997.
- Bridle et al. "An investigation of segmental hidden dynamic models of speech coarticulation for speech recognition," Final Report Workshop on Language Engineering, Johns Hopkins U, 1998.
- Picone et al. "Initial evaluation of hidden dynamic models on conversational speech," ICASSP, 1999.
- Deng and Ma. "Spontaneous speech recognition using a statistical co-articulatory model for the vocal tract resonance dynamics," JASA, 2000.

Structured Hidden Trajectory Models (HTM)

- Zhou, et al. "Coarticulation modeling by embedding a target-directed hidden trajectory model into HMM," ICASSP, 2003. **← DARPA EARS Program 2001-2004: Novel Approach II**
- Deng, Yu, Acero. "Structured speech modeling," *IEEE Trans. on Audio, Speech and Language Processing*, vol. 14, no. 5, 2006.

Switching Nonlinear State-Space Models

- Deng. "Switching Dynamic System Models for Speech Articulation and Acoustics," in Mathematical Foundations of Speech and Language Processing, vol. 138, pp. 115 - 134, Springer, 2003.
- Lee et al. "A Multimodal Variational Approach to Learning and Inference in Switching State Space Models," ICASSP, 2004.

Other Deep Generative Models (developed outside speech)

LETTER

Communicated by Yann Le Cun

- Sigmoid belief nets & wake/sleep alg. (1992)
- Deep belief nets (DBN, 2006);
→ Start of deep learning
- Totally non-obvious result:
Stacking many RBMs (undirected)
→ not Deep Boltzmann Machine (**DBM**, undirected)
→ but a DBN (directed, generative model)
- Excellent in generating images & speech synthesis
- Similar type of deep generative models to HDM
- But simpler: no temporal dynamics
- With very different parameterization
- Most intriguing of DBN: inference is easy
(i.e. no need for approximate variational Bayes)
← "Restriction" of connections in RBM
- Pros/cons analysis → Hinton coming to MSR 2009

A Fast Learning Algorithm for Deep Belief Nets

Geoffrey E. Hinton

hinton@cs.toronto.edu

Simon Osindero

osindero@cs.toronto.edu

Department of Computer Science, University of Toronto, Toronto, Canada M5S 3G4

Yee-Whye Teh

tehyw@comp.nus.edu.sg

*Department of Computer Science, National University of Singapore,
Singapore 117543*

We show how to use "complementary priors" to eliminate the explaining-away effects that make inference difficult in densely connected belief nets that have many hidden layers. Using complementary priors, we derive a fast, greedy algorithm that can learn deep, directed belief networks one layer at a time, provided the top two layers form an undirected associative memory. The fast, greedy algorithm is used to initialize a slower learning procedure that fine-tunes the weights using a contrastive version of the wake-sleep algorithm. After fine-tuning, a network with three hidden layers forms a very good generative model of the joint distribution of handwritten digit images and their labels. This generative model gives better digit classification than the best discriminative learning algorithms. The low-dimensional manifolds on which the digits lie are modeled by long ravines in the free-energy landscape of the top-level associative memory, and it is easy to explore these ravines by using the directed connections to display what the associative memory has in mind.

This is a very different kind of deep generative model

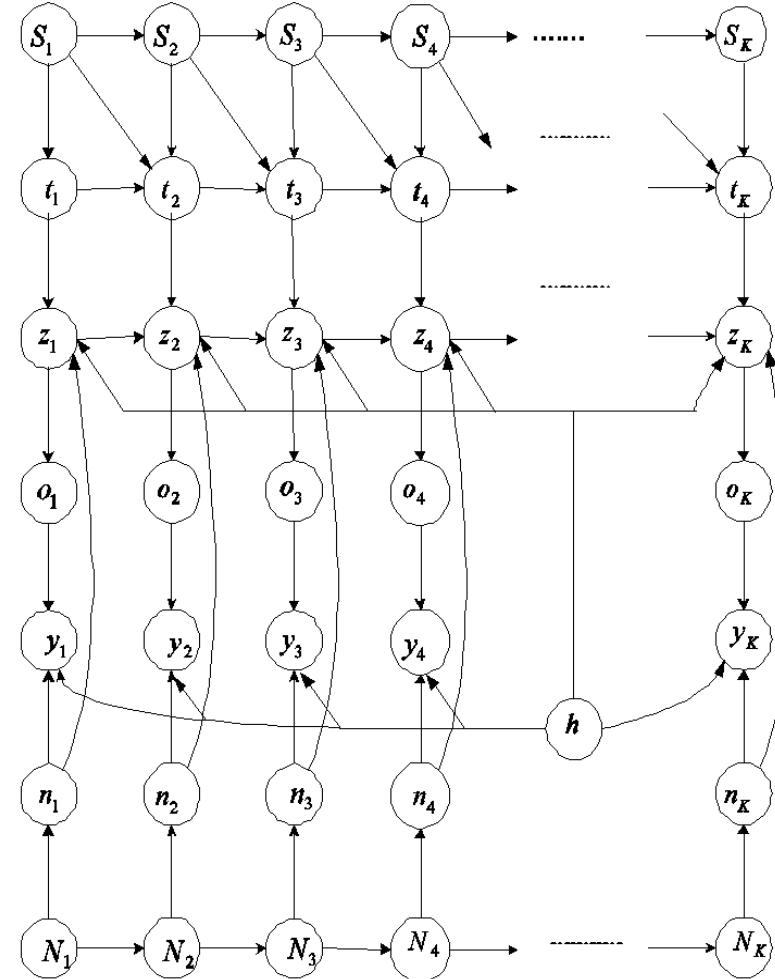
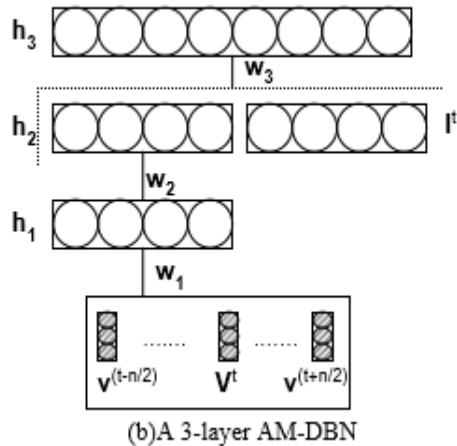
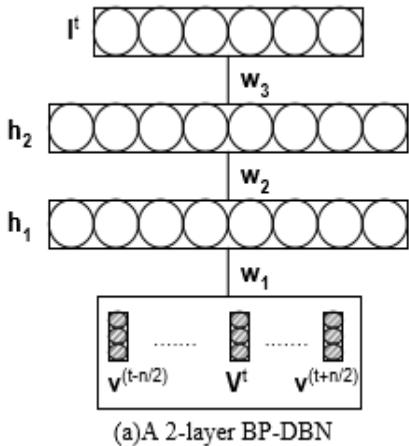


Table 4: Reported results on TIMIT core test set

Method	PER
Stochastic Segmental Models [28]	36%
Conditional Random Field [29]	34.8%
Large-Margin GMM [30]	33%
CD-HMM [4]	27.3%
Augmented conditional Random Fields [4]	26.6%
Recurrent Neural Nets [31]	26.1%
Bayesian Triphone HMM [32]	25.6%
Monophone HTMs [33]	24.8%
Heterogeneous Classifiers [34]	24.40%
Deep Belief Networks(DBNs)	23.0%

(Mohamed, Dahl, Hinton, 2009, 2012)

(after adding Backprop to the generative DBN)

(Deng et al., 2006; Deng & Yu, 2007)

Structured Speech Modeling

Li Deng, *Fellow, IEEE*, Dong Yu, *Member, IEEE*, and Alex Acero, *Fellow, IEEE*

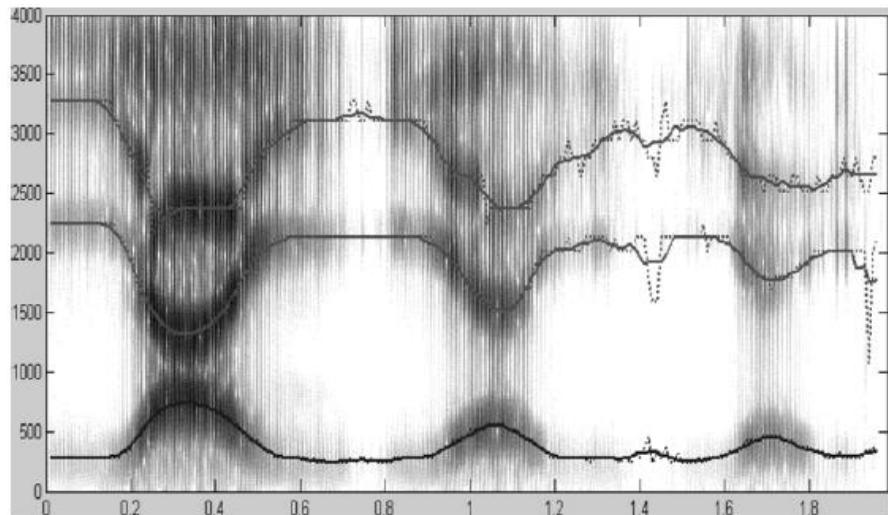
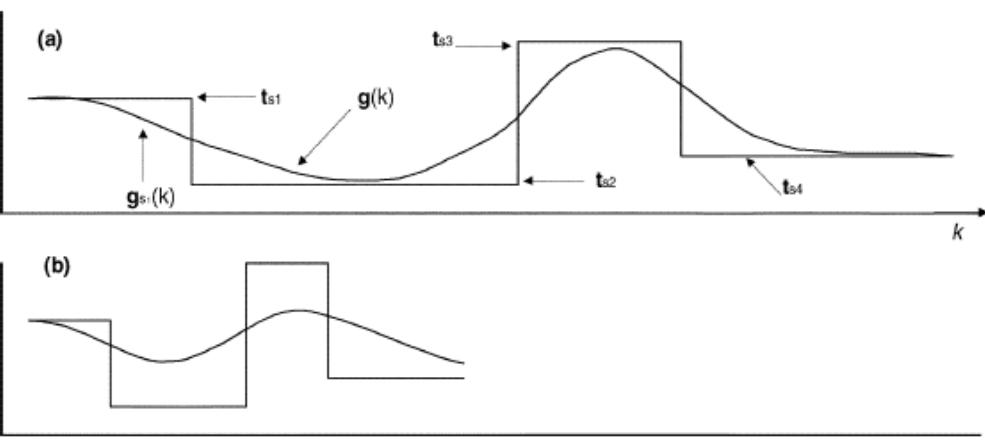


Fig. 1. Illustrations of the various VTR quantities in model Stage-I in an utterance with four phone segments. (a) and (b) are for the same four VI targets and their filtered results, but the durations of the four segments are shorter in (b) than in (a).

TABLE II
COMPARISONS OF HMM AND HTM PERFORMANCES (PERCENT CORRECT)
WITHIN EACH OF FOUR BROAD PHONE CLASSES

	Sonorants	Stops	Fricatives	Closures
Occurrences	3814	889	1252	1578
HMM	64.05	72.10	75.64	88.72
HTM	72.42	76.27	75.74	90.94

- DBN/DNN made many new errors on short, undershoot vowels
-- 11 frames contain too much "noise"

Early Successes of Deep Learning in Speech Recognition

Academic-Industrial Collaboration (2009,2010)

- I invited Geoff Hinton to work with me at MSR, Redmond
- **Well-timed** academic-industrial collaboration:
 - ASR industry searching for new solutions when “principled” deep generative approaches could not deliver
 - Academia developed deep learning tools (e.g. **DBN** 2006) looking for applications
 - Add Backprop to deep generative models (DBN) → DNN (hybrid generative/discriminative)
 - Advent of GPU computing (Nvidia CUDA library released 2007/08)
 - Big training data in speech recognition were already available



[NIPS Home](#)

[Overview](#)

[Conference Videos](#)

[Workshop Videos](#)

[Program Highlights](#)

[Tutorials](#)

[Conference Sessions](#)

[Workshops](#)

[Publication Models](#)

[Demonstrations](#)

[Mini Symposia](#)

[Accepted Papers](#)

[Dates](#)

[Committees](#)

[Sponsors](#)

[Awards](#)

[Board](#)

[Li Deng, Dong Yu, Geoffrey Hinton](#)

[Microsoft Research; Microsoft Research; University of Toronto](#)

Deep Learning for Speech Recognition and Related Applications

7:30am - 6:30pm Saturday, December 12, 2009

Location: Hilton: Cheakamus

Abstract: Over the past 25 years or so, speech recognition technology has been dominated by a "shallow" architecture --- hidden Markov models (HMMs). Significant technological success has been achieved using complex and carefully engineered variants

**Invitee 1: give me one week
to decide ...,...**

**Not worth my time to fly to
Vancouver for this...**

There has been virtually no effective communication between machine learning researchers and speech recognition researchers who are both advocating the use of deep architecture and learning. One goal of the proposed workshop is to bring together these two groups of researchers to review the progress in both fields and to identify promising and synergistic research directions for potential future cross-fertilization and collaboration.

Mohamed, Dahl, Hinton, Deep belief networks for phone recognition, NIPS 2009 Workshop on Deep Learning, 2009

Yu, Deng, Wang, Learning in the Deep-Structured Conditional Random Fields, NIPS 2009 Workshop on Deep Learning, 2009

..., ..., ...

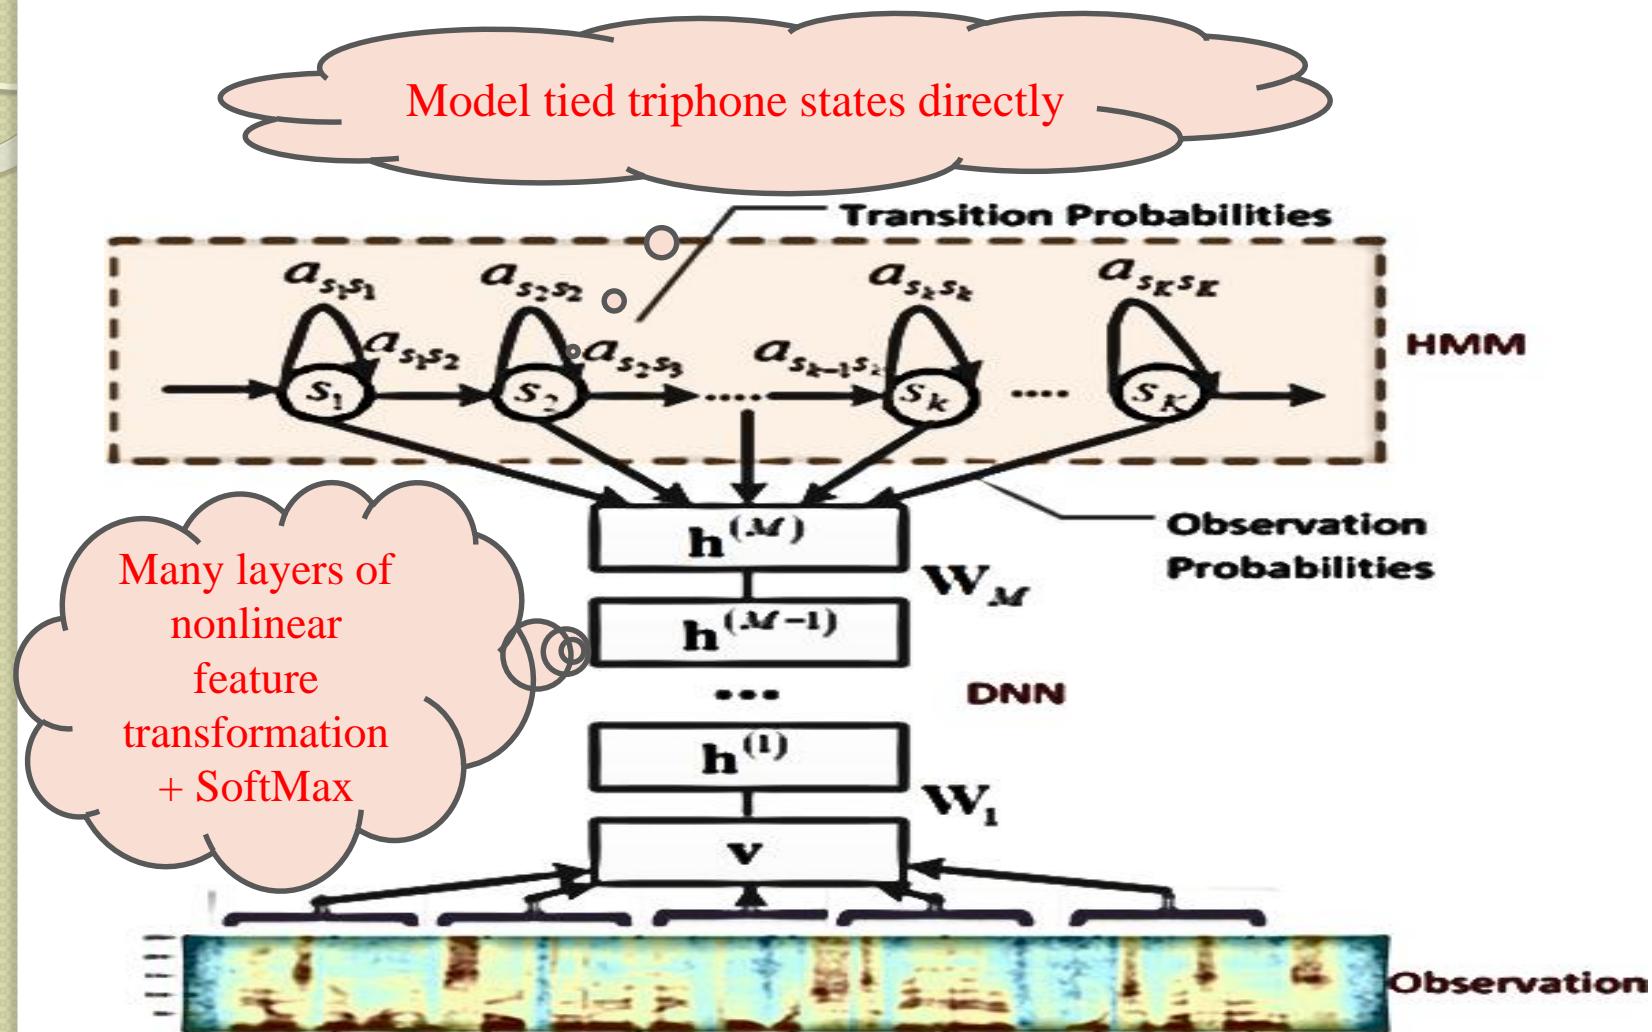
Expanding DNN at Industry Scale

- **Scale DNN's success to large speech tasks (2010-2011)**
 - Grew output neurons from context-independent phone states (100-200) to context-dependent ones (1k-30k) → CD-DNN-HMM for Bing Voice Search and then to SWBD tasks
 - Motivated initially by saving huge MSFT investment in the speech decoder software infrastructure
 - CD-DNN-HMM also gave much higher accuracy than CI-DNN-HMM
 - Earlier NNs made use of context only as appended inputs, not coded directly as outputs
 - Discovered that ***with large training data Backprop works well without DBN pre-training*** by understanding why gradients often vanish (patent filed for “discriminative pre-training” 2011)
- **Engineering for large speech systems:**
 - Combined expertise in DNN (esp. with GPU implementation) **and** speech recognition
 - Collaborations among MSRR, MSRA, academic researchers

- Yu, Deng, Dahl, [Roles of Pre-Training and Fine-Tuning in Context-Dependent DBN-HMMs for Real-World Speech Recognition](#), in *NIPS Workshop on Deep Learning*, 2010.
- Dahl, Yu, Deng, Acero, [Large Vocabulary Continuous Speech Recognition With Context-Dependent DBN-HMMS](#), in *Proc. ICASSP*, 2011.
- Dahl, Yu, Deng, Acero, [Context-Dependent Pre-trained Deep Neural Networks for Large Vocabulary Speech Recognition](#), in *IEEE Transactions on Audio, Speech, and Language Processing* (2013 IEEE SPS Best Paper Award), vol. 20, no. 1, pp. 30-42, January 2012.
- Seide, Li, Yu, "[Conversational Speech Transcription Using Context-Dependent Deep Neural Networks](#)", Interspeech 2011, pp. 437-440.
- Hinton, Deng, Yu, Dahl, Mohamed, Jaitly, Senior, Vanhoucke, Nguyen, Sainath, Kingsbury, [Deep Neural Networks for Acoustic Modeling in SpeechRecognition](#), in *IEEE Signal Processing Magazine*, vol. 29, no. 6, pp. 82-97, November 2012
- Sainath, T., Kingsbury, B., Ramabhadran, B., Novak, P., and Mohamed, A. "Making deep belief networks effective for large vocabulary continuous speech recognition," *Proc. ASRU*, 2011.
- Sainath, T., Kingsbury, B., Soltau, H., and Ramabhadran, B. "Optimization Techniques to Improve Training Speed of Deep Neural Networks for Large Speech Tasks," *IEEE Transactions on Audio, Speech, and Language Processing*, vol.21, no.11, pp.2267-2276, Nov. 2013.
- Jaitly, N., Nguyen, P., Senior, A., and Vanhoucke, V. "Application of Pretrained Deep Neural Networks to Large Vocabulary Speech Recognition," *Proc. Interspeech*, 2012.

DNN-HMM

(replacing GMM only; longer MFCC/filter-back windows w. no transformation)



DNN vs. Pre-DNN Prior-Art

- **Table:** TIMIT Phone recognition (3 hours of training)

Features	Setup	Error Rates
Pre-DNN	Deep Generative Model	24.8%
DNN	5 layers x 2048	23.4%

~10% relative improvement

- **Table:** Voice Search SER (24-48 hours of training)

Features	Setup	Error Rates
Pre-DNN	GMM-HMM with MPE	36.2%
DNN	5 layers x 2048	30.1%

~20% relative improvement

- **Table:** SwitchBoard WER (309 hours training)

Features	Setup	Error Rates
Pre-DNN	GMM-HMM with BMMI	23.6%
DNN	7 layers x 2048	15.8%

~30% relative Improvement

For DNN, the more data, the better!



Deep learning
technology enabled
speech-to-speech
translation



The New York Times

Scientists See Promise in Deep-Learning Programs

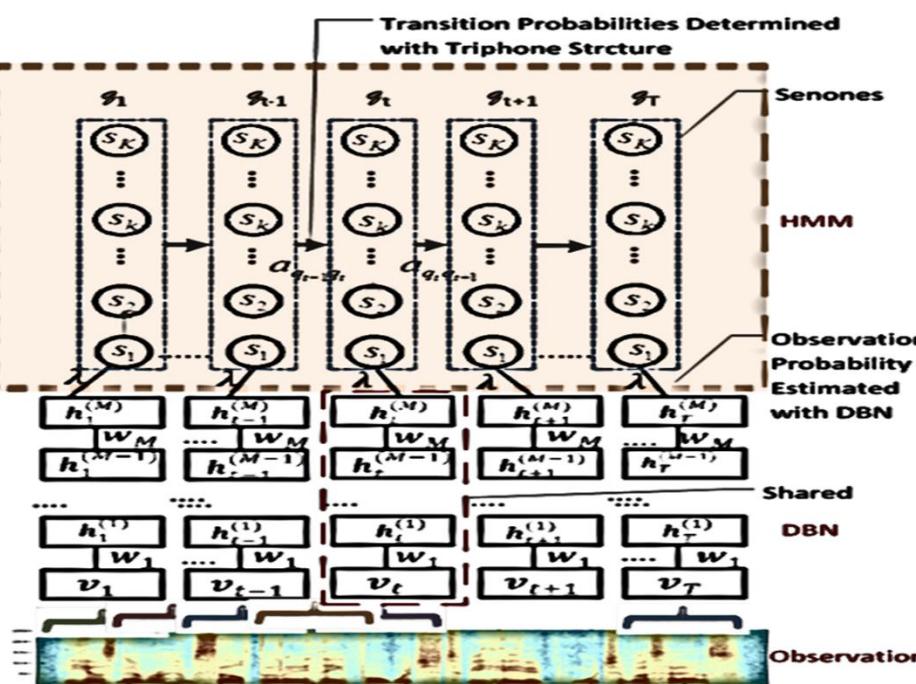
John Markoff

November 23, 2012

Rick Rashid in Tianjin, China, October, 25, 2012



A voice recognition program translated a speech given by Richard F. Rashid, Microsoft's top scientist, into Mandarin Chinese.

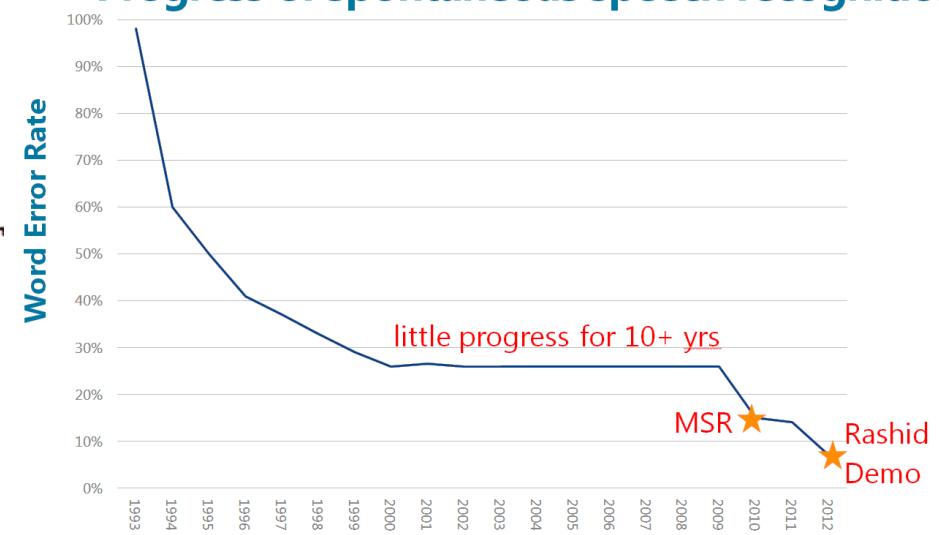


CD-DNN-HMM

Dahl, Yu, Deng, and Acero, "Context-Dependent Pre-trained Deep Neural Networks for Large Vocabulary Speech Recognition," *IEEE Trans. ASLP*, Jan. 2012 (also ICASSP 2011)

Seide et al, Interspeech, 2011.

Progress of spontaneous speech recognition

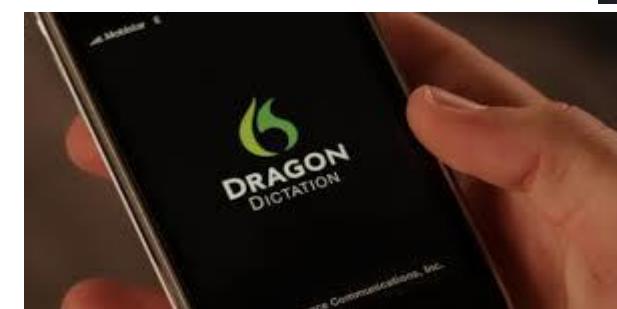


After no improvement for 10+ years
by the research community...
...MSR reduced error from ~23% to
<13% (and under 7% for Rick
Rashid's S2S demo in 2012)!

Impact of deep learning in speech technology



Skype to get 'real-time' translator



Analy



Deep Learning

With massive amounts of computational power, machines can now recognize objects and translate speech in real time. Artificial intelligence is finally getting smart.

Temporary Social Media

Messages that quickly self-destruct could enhance the privacy of online communications and make people freer to be spontaneous.

Prenatal DNA Sequencing

Reading the DNA of fetuses will be the next frontier of the genomic revolution. But do you really want to know about the genetic problems or musical aptitude of your unborn child?

Additive Manufacturing

Skeptical about 3-D printing? GE, the world's largest manufacturer, is on the verge of using the technology to make jet parts.

Baxter: The Blue-Collar Robot

Rodney Brooks's newest creation is easy to interact with, but the complex innovations behind the robot show just how hard it is to get along with people.

Memory Implants

A maverick neuroscientist believes he has deciphered the code by which the brain forms long-term memories. Next: testing a prosthetic implant for people suffering from long-term memory loss.

Smart Watches

The designers of the Pebble watch realized that a mobile phone is more useful if you don't have to take it out of your pocket.

Ultra-Efficient Solar Power

Doubling the efficiency of a solar cell would completely change the economics of renewable energy. Nanotechnology just might make it possible.

Big Data from Cheap Phones

Collecting and analyzing information from simple cell phones can provide surprising insights into how people move about and behave – and even help us understand the spread of diseases.

Supergrids

A new high-power circuit breaker could finally make highly efficient DC power grids practical.



Enabling Cross-Lingual Conversations in Real Time

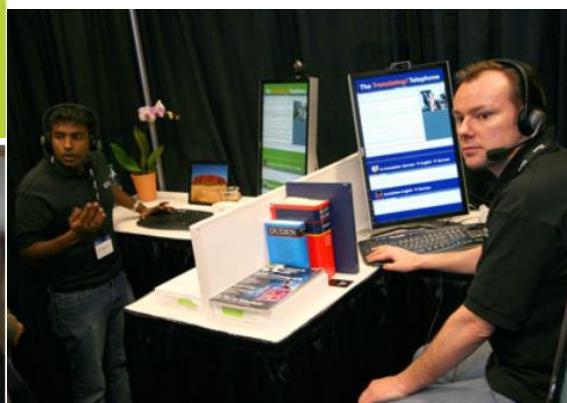
Microsoft Research

May 27, 2014 5:58 PM PT



ROBERT MCMILLAN BUSINESS 12.17.14 1:19 PM

View milestones
on the path to
Skype Translator
#speech2speech



Taking a cue from science fiction, Microsoft demos 'universal translator'

By Jacopo Prisco, for CNN
Updated 12:35 PM ET, Thu October 16, 2014



Deep Learning in the News



EXCLUSIVE

Facebook, Google in 'Deep Learning' Arms Race

Yann LeCun, an NYU artificial intelligence researcher who now works for Facebook. Photo: Josh Valcarcel/WIRED



WIRED

NEWS BULLETIN

Google Beat Facebook for DeepMind

Google Acquires Artificial Intelligence Startup DeepMind For More Than \$500M

Posted Jan 26, 2014 by Catherine Shu (@catherineshu)



The New York Tim

Monday, June 25, 2012 Last Update: 11:50 PM ET

o Hired to Make AI a Reality WEEKS

YOU TOLD US 31 NAMES WE NEED TO KNOW IN 2011
MARCH 10, 2011 | BY JEFFREY WAGNER



TECH 2/19/2015 @ 1:06PM | 6,586 views

Microsoft's Deep Learning Project Outperforms Humans In Image Recognition

[+ Comment Now](#) [+ Follow Comments](#)



Anthony Wing Kosner
Contributor

[FOLLOW](#)

TECH 12/29/2014 @ 11:37AM | 75,350 views

Tech 2015: Deep Learning And Machine Intelligence Will Eat The World

[+ Comment Now](#) [+ Follow Comments](#)

Despite what [Stephen Hawking](#) or Elon Musk say, [hostile Artificial Intelligence](#) is not going to destroy the world anytime soon. What is





Deep learning

Yann LeCun, Yoshua Bengio & Geoffrey Hinton

Affiliations | Corresponding author

Nature **521**, 436–444 (28 May 2015) | doi:10.1038/nature14539

Received 25 February 2015 | Accepted 01 May 2015 | Published online

“This joint paper from the major speech recognition laboratories was the first major industrial application of deep learning.”



[VOLUME 29 NUMBER 6 NOVEMBER 2012]

LOUD AND CLEAR FUNDAMENTAL TECHNOLOGIES IN MODERN SPEECH RECOGNITION



[Geoffrey Hinton, Li Deng, Dong Yu, George E. Dahl, Abdel-rahman Mohamed, Navdeep Jaitly, Andrew Senior, Vincent Vanhoucke, Patrick Nguyen, Tara N. Sainath, and Brian Kingsbury]

Deep Neural Networks for Acoustic Modeling in Speech Recognition

[The shared views of four research groups]

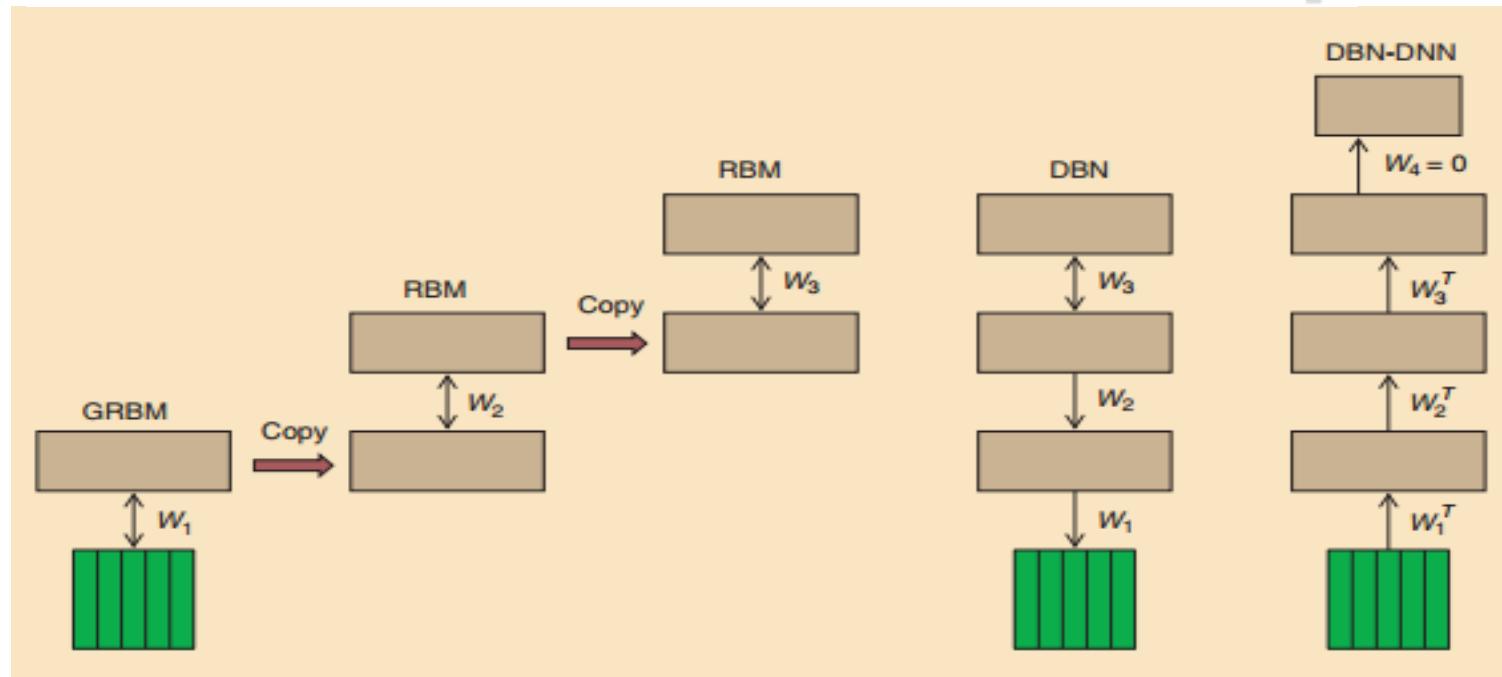


PHOTO.COM/SUCHA/LAFAT

DNN: (Fully-Connected) Deep Neural Networks

“DNN for acoustic modeling in speech recognition,” in *IEEE SPM*, Nov. 2012

Geoffrey Hinton, Li Deng, Dong Yu, George E. Dahl, Abdel-rahman Mohamed, Navdeep Jaitly,
Andrew Senior, Vincent Vanhoucke, Patrick Nguyen, Tara N. Sainath, and Brian Kingsbury



First train a stack of N models each of which has one hidden layer. Each model in the stack treats the hidden variables of the previous model as data.

Then compose them into a single Deep Belief Network (DBN).

Then add outputs and train the DNN with backprop.

Practical guide to address some common deep learning ASR issues

ASR Issues	Solutions
How to reduce the runtime without accuracy loss?	SVD
How to do speaker adaptation with low footprints?	SVD-based adaptation
How to be robust to noise?	Variable component CNN
How to reduce accuracy gap between large and small DNN?	Teacher-student learning using output posterior
How to deal with large variety of data?	DNN factorization, mixed band training
How to enable languages with limited training data?	Multi-lingual DNN

(Slide from: Jinyu Li)

More Recent Development of Deep Learning for Speech

Deep Learning
Methods and Applications

Li Deng and Dong Yu

Signals and Communication Technology

Dong Yu
Li Deng

Automatic
Speech
Recognition

A Deep-Learning
Approach

Springer

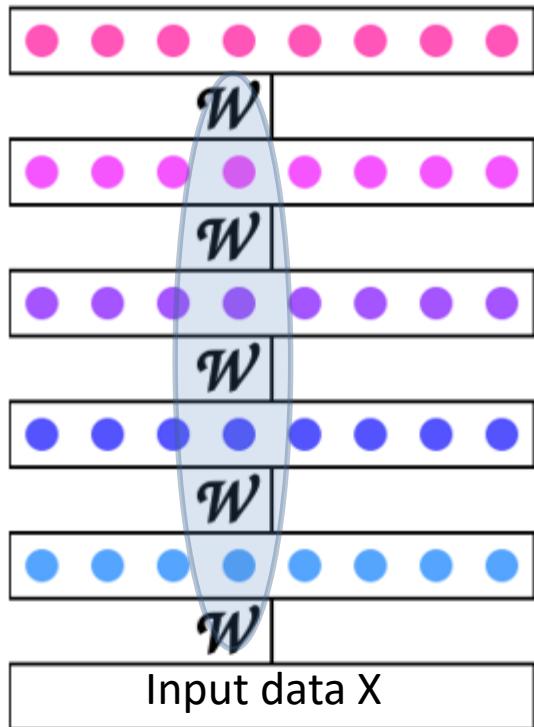
Chapter 7

Selected Applications in Speech and Audio Processing

7.1 Acoustic modeling for speech recognition

As discussed in Section 2, speech recognition is the very first successful application of deep learning methods at an industry scale. This success is a result of close academic-industrial collaboration, initiated at Microsoft Research, with the involved researchers identifying and acutely attending to the industrial need for large-scale deployment [68, 89, 109, 161, 323, 414]. It is also a result of carefully exploiting the strengths of the deep learning and the then-state-of-the-art speech recognition technology, including notably the highly efficient decoding techniques.

Innovation: Better Optimization



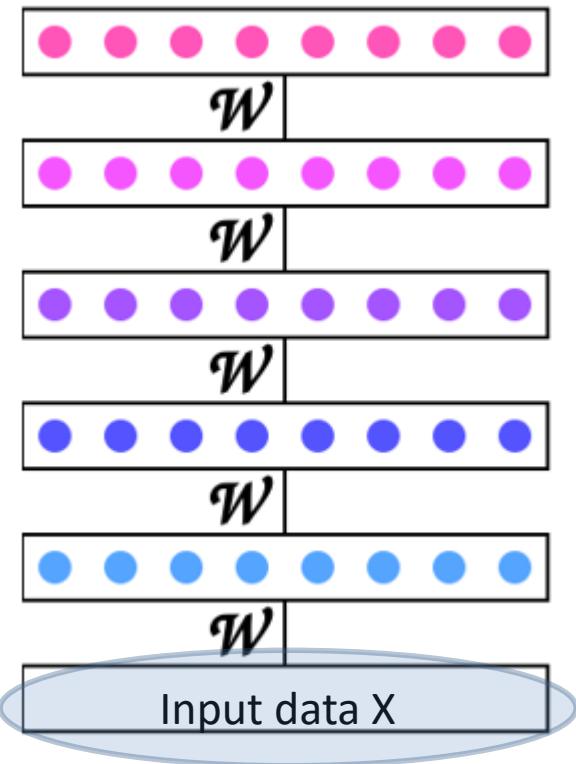
- **Sequence discriminative training for DNN:**

- Mohamed, Yu, Deng: “Investigation of full-sequence training of deep belief networks for speech recognition,” *Interspeech*, 2010.
- Kingsbury, Sainath, Soltau. “Scalable minimum Bayes risk training of DNN acoustic models using distributed hessian-free optimization,” *Interspeech*, 2012.
- Su, Li, Yu, Seide. “Error back propagation for sequence training of CD deep networks for conversational speech transcription,” ICASSP, 2013.
- Vesely, Ghoshal, Burget, Povey. “Sequence-discriminative training of deep neural networks, *Interspeech*, 2013.

- **Distributed asynchronous SGD**

- Dean, Corrado,...Senior, Ng. “Large Scale Distributed Deep Networks,” NIPS, 2012.
- Sak, Vinyals, Heigold, Senior, McDermott, Monga, Mao. “Sequence Discriminative Distributed Training of Long Short-Term Memory Recurrent Neural Networks,” *Interspeech*, 2014.

Innovation: Towards Raw Inputs



- **Bye-Bye MFCCs (no more cosine transform, Mel-scaling?)**

- Deng, Seltzer, Yu, Acero, Mohamed, Hinton. "Binary coding of speech spectrograms using a deep auto-encoder," *Interspeech*, 2010.
- Mohamed, Hinton, Penn. "Understanding how deep belief networks perform acoustic modeling," ICASSP, 2012.
- Li, Yu, Huang, Gong, "Improving wideband speech recognition using mixed-bandwidth training data in CD-DNN-HMM" SLT, 2012
- Deng, J. Li, Huang, Yao, Yu, Seide, Seltzer, Zweig, He, Williams, Gong, Acero. "Recent advances in deep learning for speech research at Microsoft," ICASSP, 2013.
- Sainath, Kingsbury, Mohamed, Ramabhadran. "Learning filter banks within a deep neural network framework," ASRU, 2013.

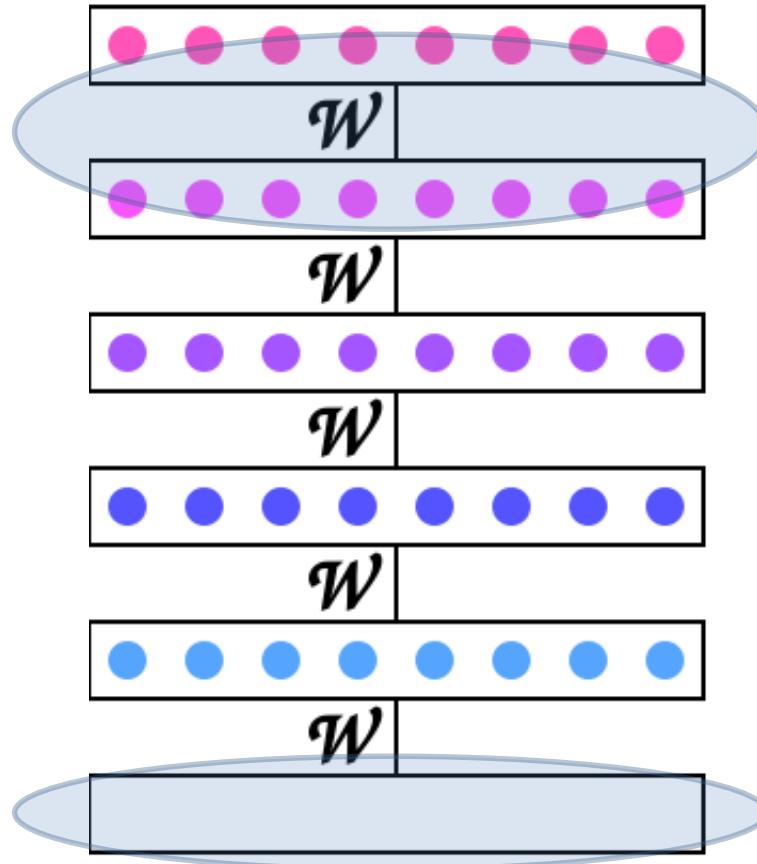
- **Bye-Bye Fourier transforms?**

- Jaithi and Hinton. "Learning a better representation of speech sound waves using RBMs," ICASSP, 2011.
- Tuske, Golik, Schluter, Ney. "Acoustic modeling with deep neural networks using raw time signal for LVCSR," Interspeech, 2014.
- Golik et al, "Convolutional NNs for acoustic modeling of raw time signals in LVCSR," Interspeech, 2015.
- Sainath et al. "Learning the Speech Front-End with Raw Waveform CLDNNs," Interspeech, 2015

- **DNN as hierarchical nonlinear feature extractors:**

- Seide, Li, Chen, Yu. "Feature engineering in context-dependent deep neural networks for conversational speech transcription, ASRU, 2011.
- Yu, Seltzer, Li, Huang, Seide. "Feature learning in deep neural networks - Studies on speech recognition tasks," ICLR, 2013.
- Yan, Huo, Xu. "A scalable approach to using DNN-derived in GMM-HMM based acoustic modeling in LVCSR," Interspeech, 2013.
- Deng, Chen. "Sequence classification using high-level features extracted from deep neural networks," ICASSP, 2014.

Innovation: Transfer/Multitask Learning & Adaptation



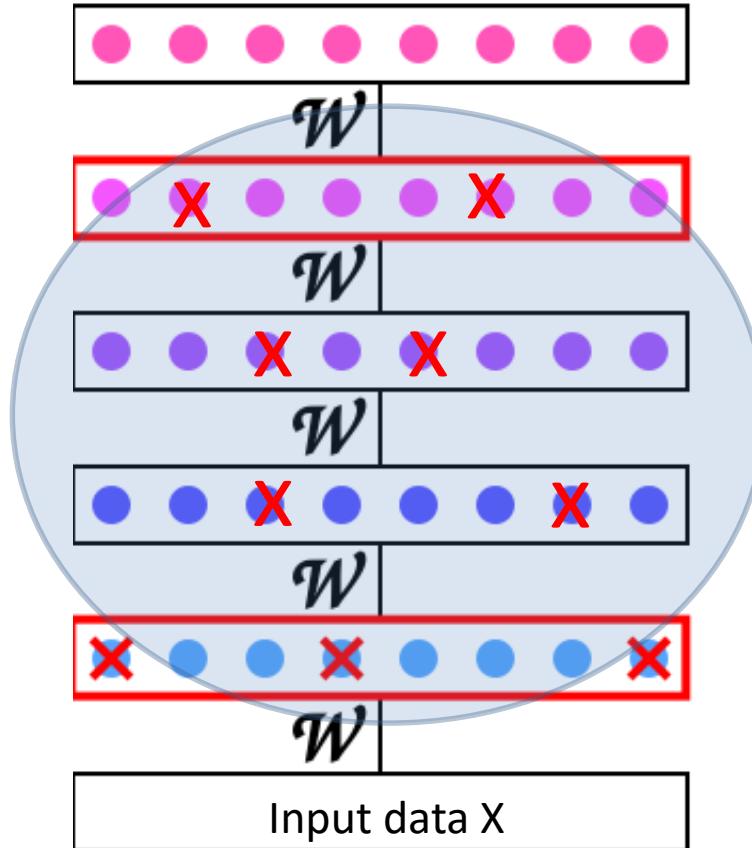
Multi-lingual acoustic modeling

Adaptation to speakers & environments (i-vectors)

Mixed-bandwidth acoustic modeling

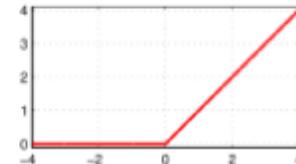
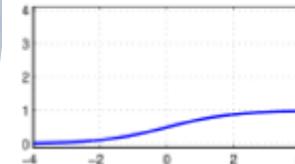
- Too many references to list & organize

Innovation: Better regularization & nonlinearity



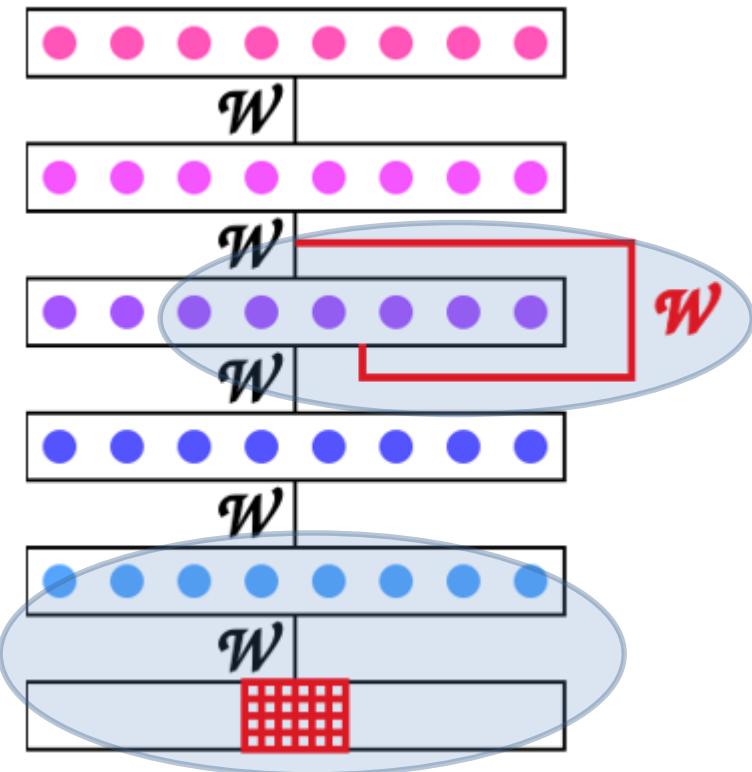
Sparsity in hidden representations

logistic \rightarrow ReLU , MaxOut,



Dropout

Innovation: Better architectures



- **Recurrent Nets (bi-directional RNN/LSTM) and Conv Nets (CNN) are superior to fully-connected DNNs**
- Sak, Senior, Beaufays. "LSTM Recurrent Neural Network architectures for large scale acoustic modeling," [Interspeech](#), 2014.
- Soltau, Saon, Sainath. "Joint Training of Convolutional and Non-Convolutional Neural Networks," [ICASSP](#), 2014.

Innovation: Ensemble Deep Learning

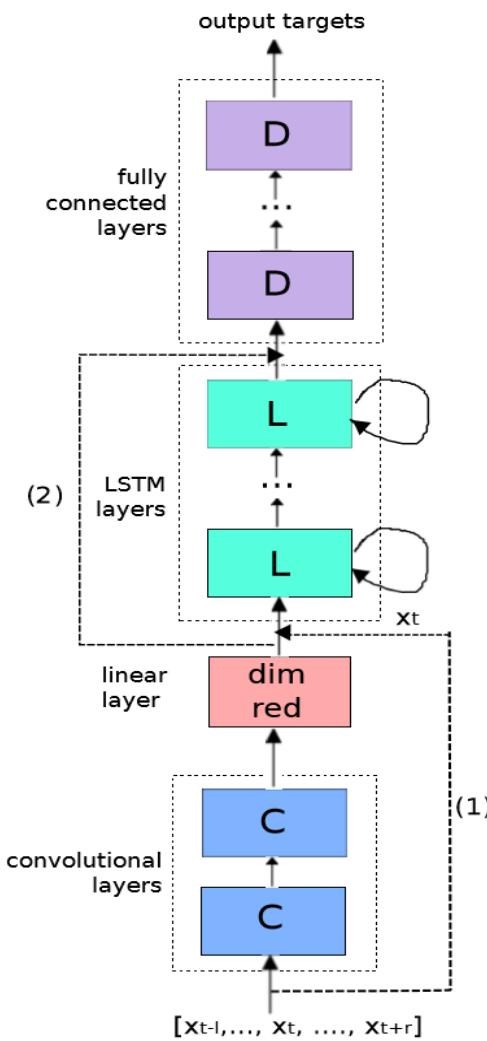
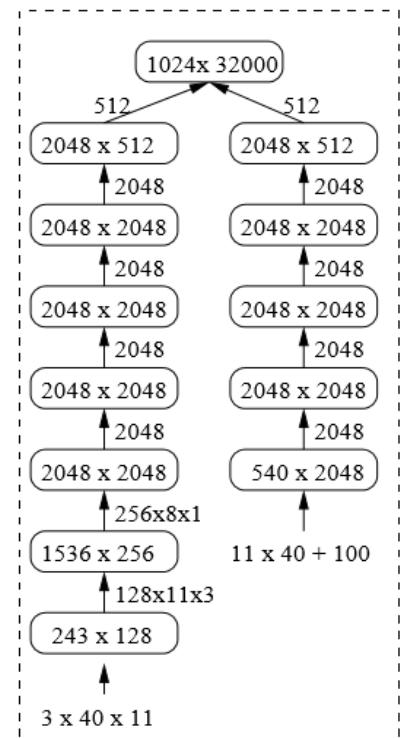
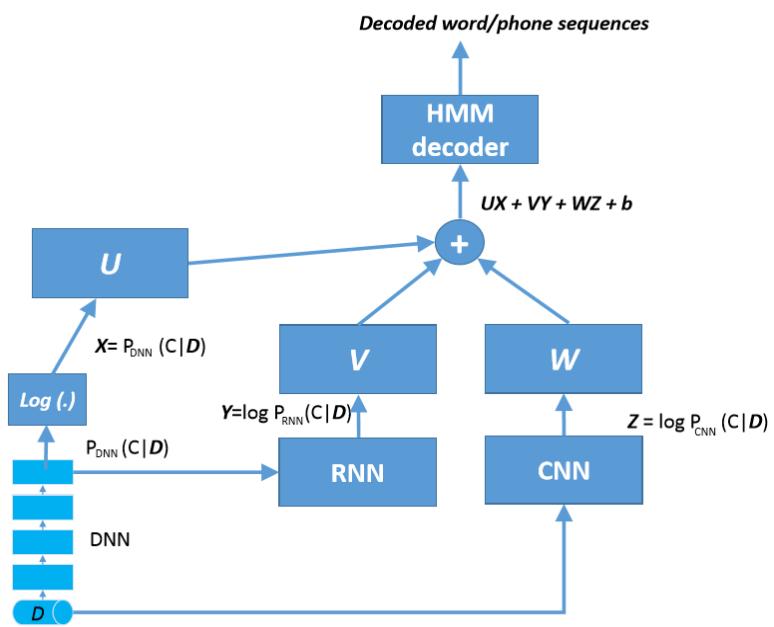


Fig. 1. CLDNN Architecture

- **Ensembles of RNN/LSTM, DNN, & Conv Nets (CNN) give huge gains:**
- T. Sainath, O. Vinyals, A. Senior, H. Sak. "Convolutional, Long Short-Term Memory, Fully Connected Deep Neural Networks," ICASSP 2015.
- L. Deng and John Platt, [Ensemble Deep Learning for Speech Recognition](#), Interspeech, 2014.
- G. Saon, H. Kuo, S. Rennie, M. Picheny. "The IBM 2015 English conversational telephone speech recognition system," arXiv, May 2015. (8% WER on SWB-309h)



Joint CNN/DNN

Innovation: Better learning objectives/methods

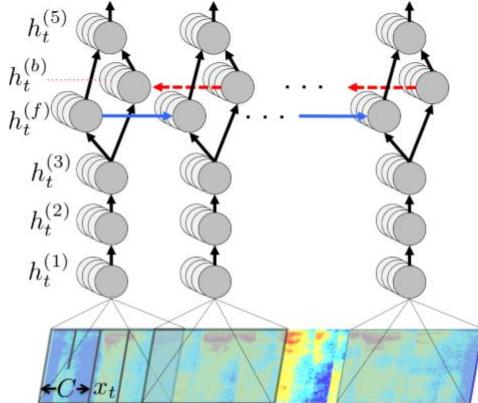
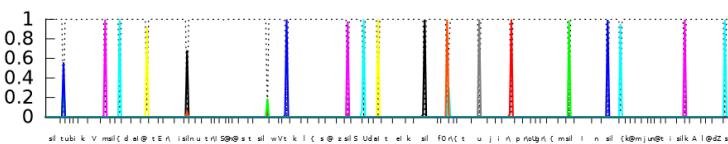


Figure 1: Structure of our RNN model and notation.

- Use of **CTC** as a new objective in RNN/LSTM with end2end learning drastically **simplifies** ASR systems
 - Predict graphemes or words directly; no pron. dictionaries; no CD; no decision trees
 - Use of “Blank” symbols may be equivalent to a special HMM state tying scheme
- CTC/RNN has NOT replaced HMM (left-to-right)
- Relative 8% gain by CTC has been shown by a very limited number of labs

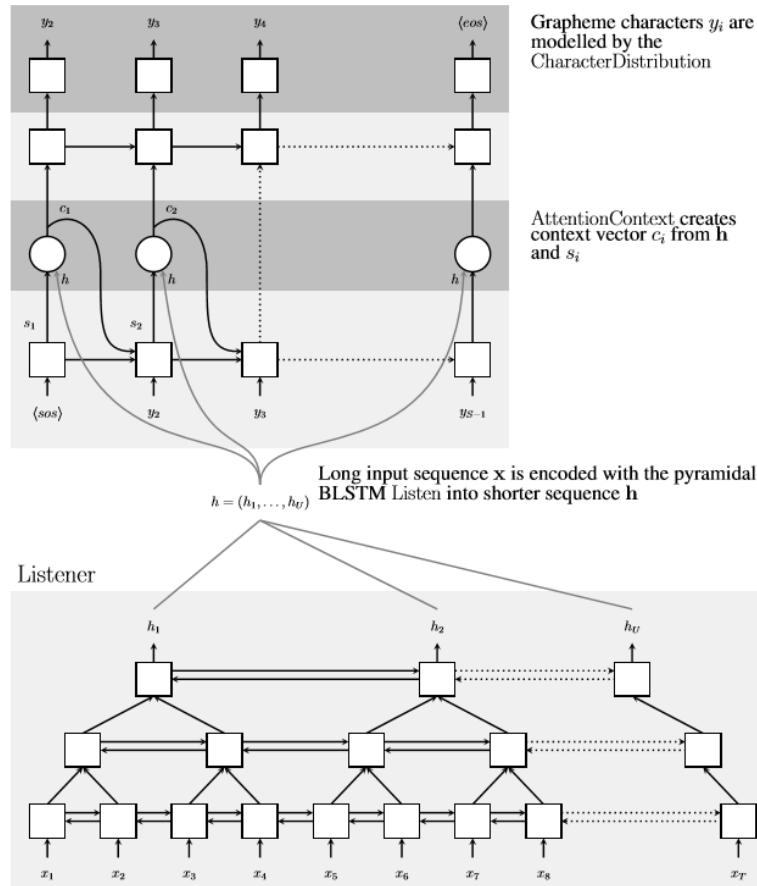


(b) 90k vocabulary

Figure 5: ‘To become a dietary nutritionist what classes should I take for a two year program in a community college’

- A. Graves and N. Jaitly. “Towards End-to-End Speech Recognition with Recurrent Neural Networks,” ICML, 2014.
- A. Hannun, A. Ng et al. “DeepSpeech: Scaling up End-to-End Speech Recognition,” arXiv Nov. 2014.
- A. Maas et al. “Lexicon-Free Conversational ASR with NN,” NAACL, 2015
- H. Sak et al. “Learning Acoustic Frame Labeling for ASR with RNN,” ICASSP, 2015
- H. Sak, A. Senior, K. Rao, F. Beaufays. “Fast and Accurate Recurrent Neural Network Acoustic Models for Speech Recognition,” 133 Interspeech, 2015

Innovation: A new paradigm for speech recognition



- **Seq2seq learning with attention mechanism (borrowed from NLP-MT)**

- W. Chan, N. Jaitly, Q. Le, O. Vinyals. "Listen, attend, and spell," arXiv, 2015.
- J. Chorowski, D. Bahdanau, D. Serdyuk, K. Cho, Y. Bengio. "Attention-Based Models for Speech Recognition," arXiv, 2015.

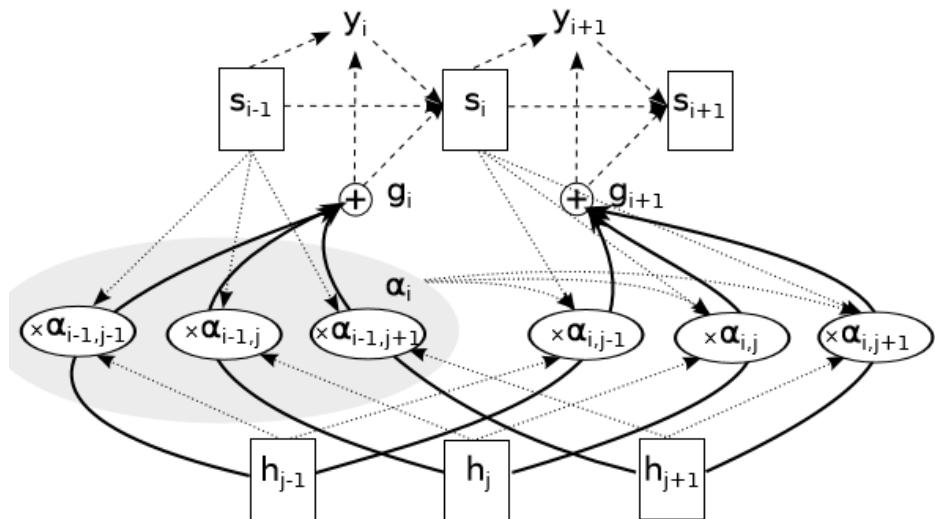


Figure 1: Listen, Attend and Spell (LAS) model: the listener is a pyramidal BLSTM encoding our input sequence x into high level features h , the speller is an attention-based decoder generating the y characters from h .

A Perspective on Recent Innovations of ASR

- All above deep learning innovations are based on supervised, discriminative learning of DNN and recurrent variants
 - Capitalizing on big, **labeled** data
 - Incorporating **monotonic-sequential** structure of speech (non-monotonic for language, later)
 - Hard to incorporate many other aspects of speech knowledge with (e.g. speech distortion model)
 - Hard to do semi- and **unsupervised** learning
- Deep generative modeling may overcome such difficulties

Li Deng and Roberto Togneri, [Chapter 6: Deep Dynamic Models for Learning Hidden Representations of Speech Features](#), pp. 153-196, Springer, December 2014.

Li Deng and Navdeep Jaitly, Chapter 2: [Deep discriminative and generative models for pattern recognition](#), ~30 pages, in *Handbook of Pattern Recognition and Computer Vision*: 5th Edition, World Scientific Publishing, Jan 2016.

	Deep Neural Nets	Deep Generative Models
Structure	Graphical; info flow: bottom-up	Graphical; info flow: top-down
Incorp constraints & domain knowledge	Hard	Easy
Unsupervised	Harder or impossible	<i>Easier, at least possible</i>
Interpretation	Harder	Easier (generative “story” on data and hidden variables)
Representation	Distributed	Localist (mostly); can be distributed also
Inference/decode	Easy	Harder (but note recent progress)
Scalability/compute	Easier (regular computes/GPU)	Harder (but note recent progress)
Incorp. uncertainty	Hard	Easy
Empirical goal	Classification, feature learning, ...	Classification (via Bayes rule), latent variable inference...
Terminology	Neurons, activation/gate functions, weights ...	Random vars, stochastic “neurons”, potential function, parameters ...
Learning algorithm	A single, unchallenged, algorithm -- BackProp	A major focus of open research, many algorithms, & more to come
Evaluation	On a black-box score – end performance	On almost every intermediate quantity
Implementation	Hard (but increasingly easier)	Standardized but insights needed
Experiments	Massive, real data	Modest, often simulated data
Parameterization	Dense matrices	Sparse (often PDFs); can be dense

Example 1: Interpretable deep learning using deep topic models (NIPS-2015)

End-to-end Learning of Latent Dirichlet Allocation by Mirror-Descent Back Propagation

Jianshu Chen*, Ji He[†], Yelong Shen*, Lin Xiao*, Xiaodong He*, Jianfeng Gao*,
Xinying Song* and Li Deng*

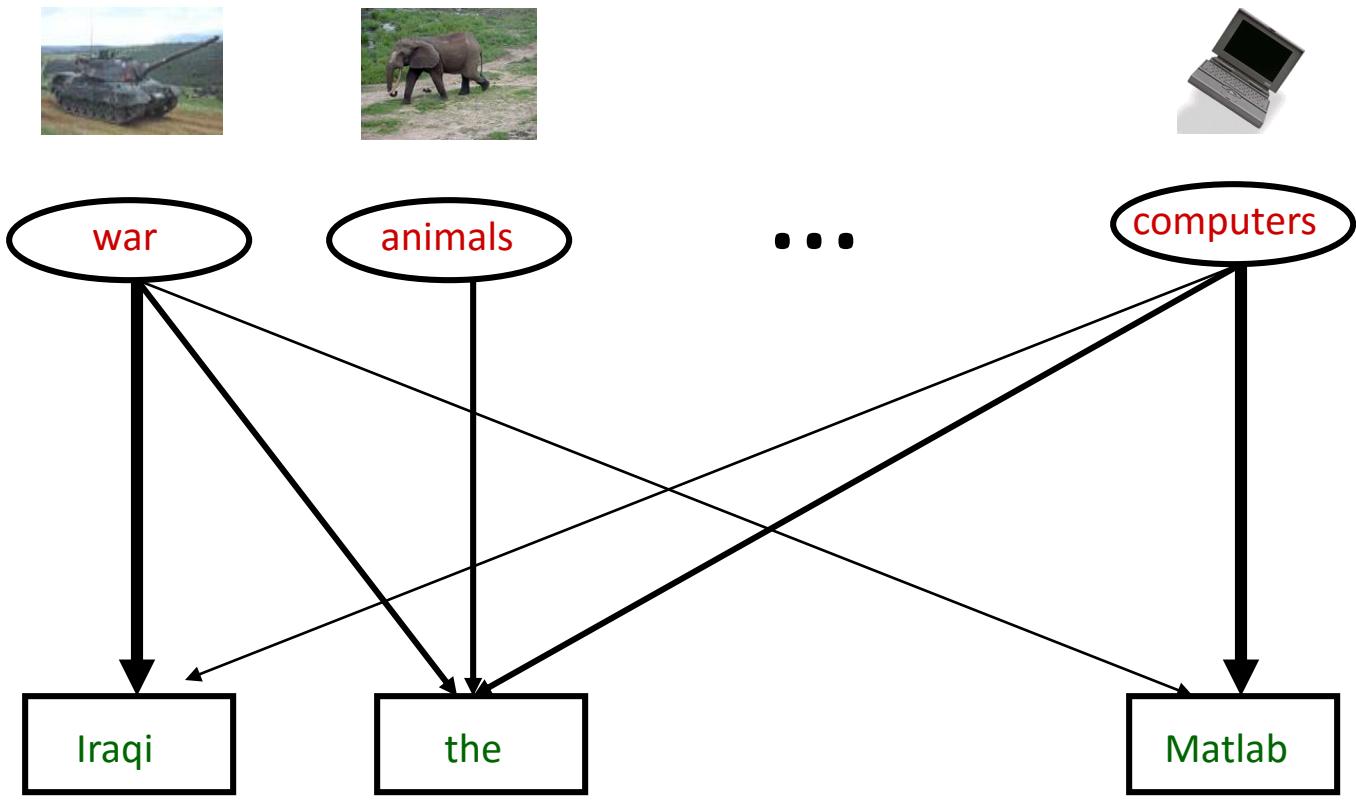
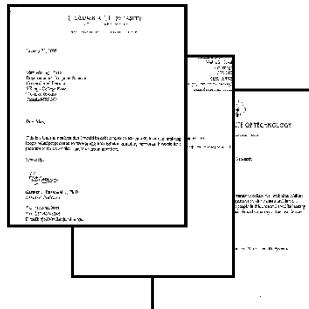
*Microsoft Research, Redmond, WA 98052, USA,
{jianshuc,yeshen,lin.xiao,xiaohe,jfgao,xinson,deng}@microsoft.com
†Department of Electrical Engineering, University of Washington, Seattle, WA 98195, USA,
jvking@uw.edu

Abstract

We develop a fully discriminative learning approach for supervised Latent Dirichlet Allocation (LDA) model, which maximizes the posterior probability of the prediction variable given the input document. Different from traditional variational learning or Gibbs sampling approaches, the proposed learning method applies (i) the mirror descent algorithm for exact maximum a posterior inference and (ii)

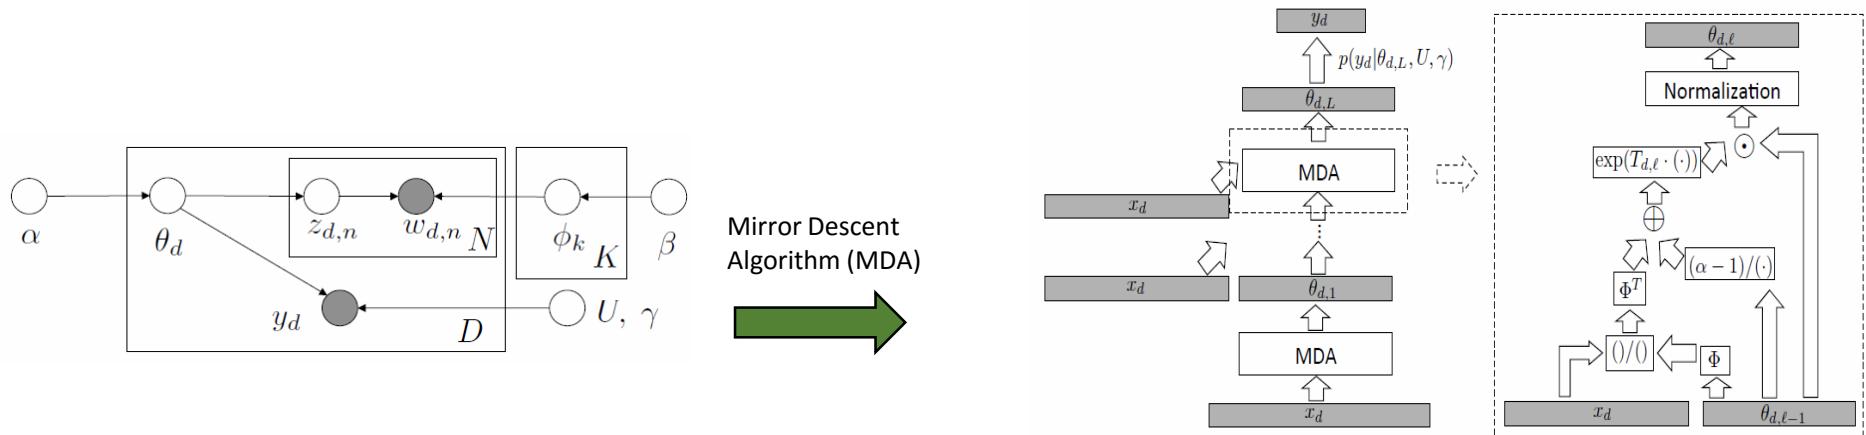
Recall: (Shallow) Generative Model

“TOPICS”
as hidden layer



Interpretable deep learning: Deep generative model

- Constructing interpretable DNNs based on generative topic models
- End-to-end learning by mirror-descent backpropagation
- to maximize posterior probability $p(y|x)$
 - y : output (win/loss), and x : input feature vector



Example 2: Unsupervised learning using deep generative model (ACL, 2013)

- Distorted character string **Images**→**Text**
- Easier than unsupervised **Speech**→**Text**
- 47% error reduction over Google's open-source OCR system

Unsupervised Transcription of Historical Documents

Taylor Berg-Kirkpatrick Greg Durrett Dan Klein

Computer Science Division

University of California at Berkeley

{tberg, gdurrett, klein}@cs.berkeley.edu

Abstract

We present a generative probabilistic model, inspired by historical printing processes, for transcribing images of documents from the printing press era. By jointly modeling the text of the document and the noisy (but regular) process of rendering glyphs, our unsupervised sys-

Motivated me to
think about
unsupervised ASR
and NLP

- (a) **a small milk saucepan ;**
- (b) **the Death of the Deceased,**
- (c) **rule along in silence**

Figure 1: Portions of historical documents with (a) unknown font, (b) uneven baseline, and (c) over-inking.

Power: Character-level LM & generative modeling for unsupervised learning

- “Image” data are naturally “generated” by the model quite accurately (like “computer graphics”)
- I had the same idea for unsupervised generative Speech-to-Text in 90’s
- Not successful because 1) Deep generative models were too simple for generating speech waves
2) Inference/learning methods for deep generative models not mature then
3) Computers were too slow

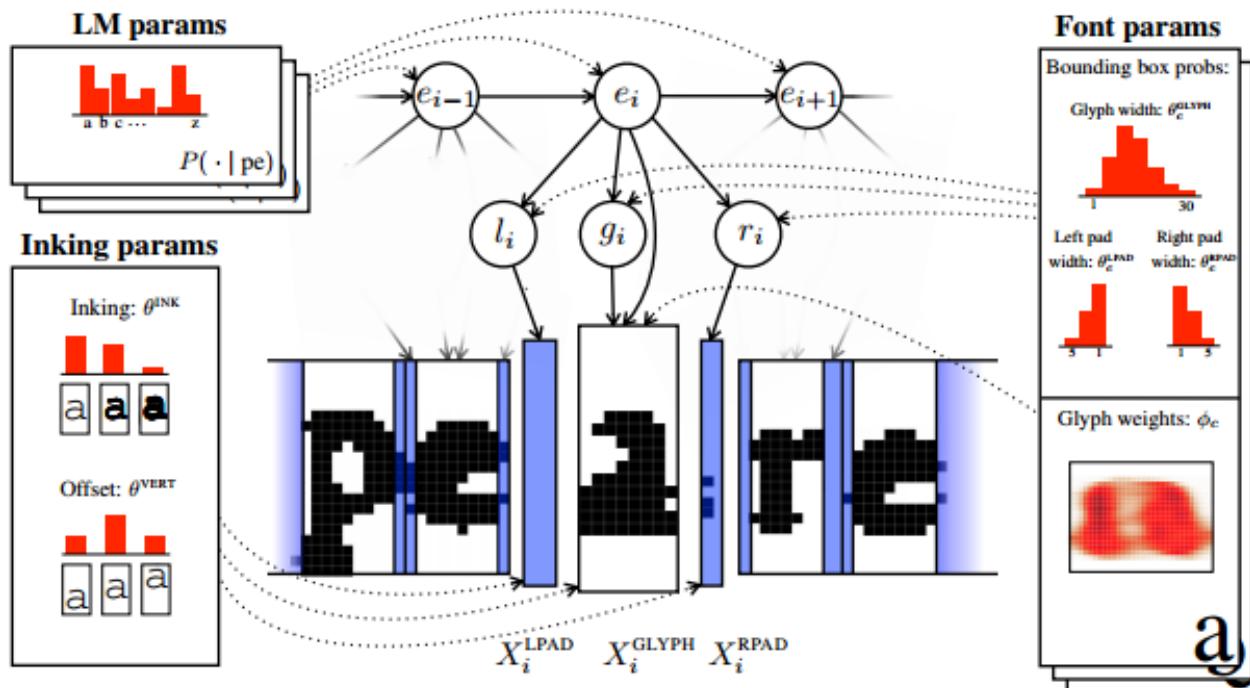
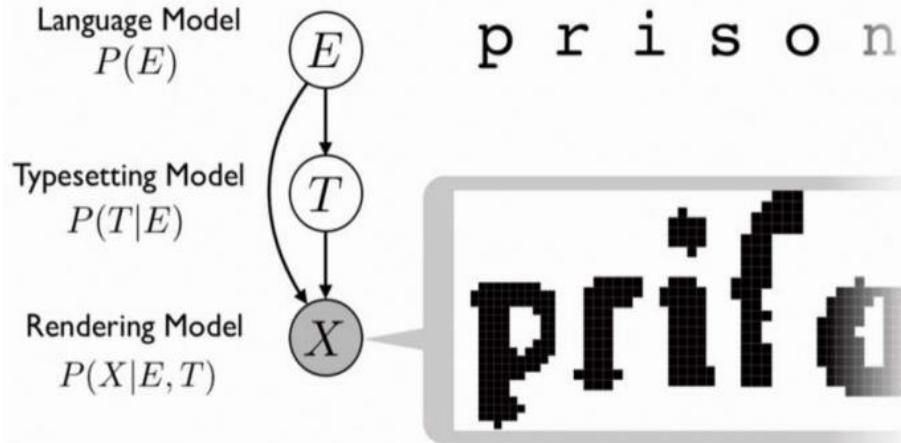


Figure 3: Character tokens e_i are generated by the language model. For each token index i , a glyph bounding box width g_i , left padding width l_i , and a right padding width r_i , are generated. Finally, the pixels in each glyph bounding box X_i^{GLYPH} are generated conditioned on the corresponding character, while the pixels in left and right padding bounding boxes, X_i^{LPAD} and X_i^{RPAD} , are generated from a background distribution.

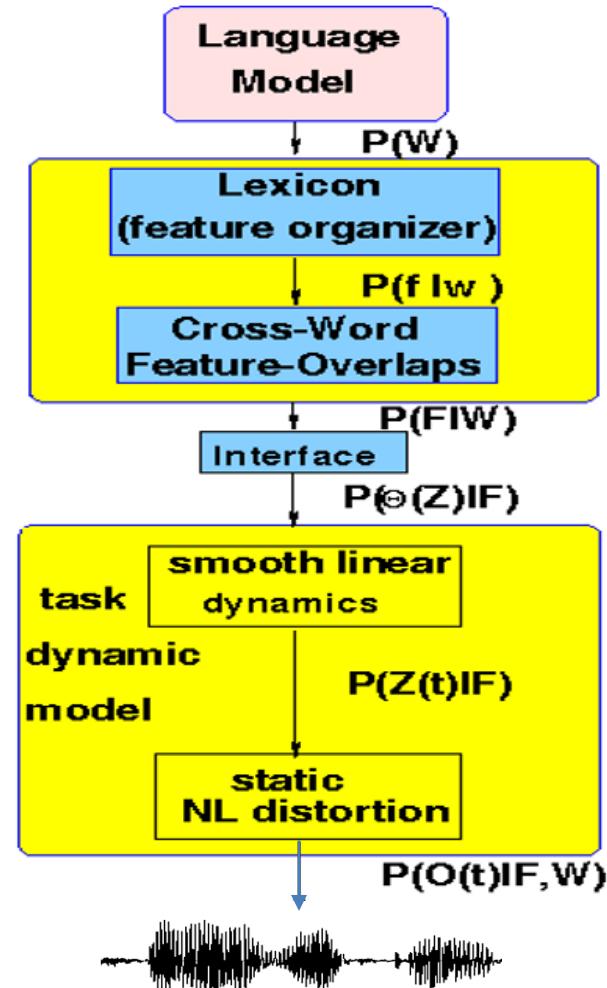
Deep Generative Model for Image-Text

(Berg-Kirkpatrick et al., 2013, 2015)



Deep Generative Model for Speech-Text

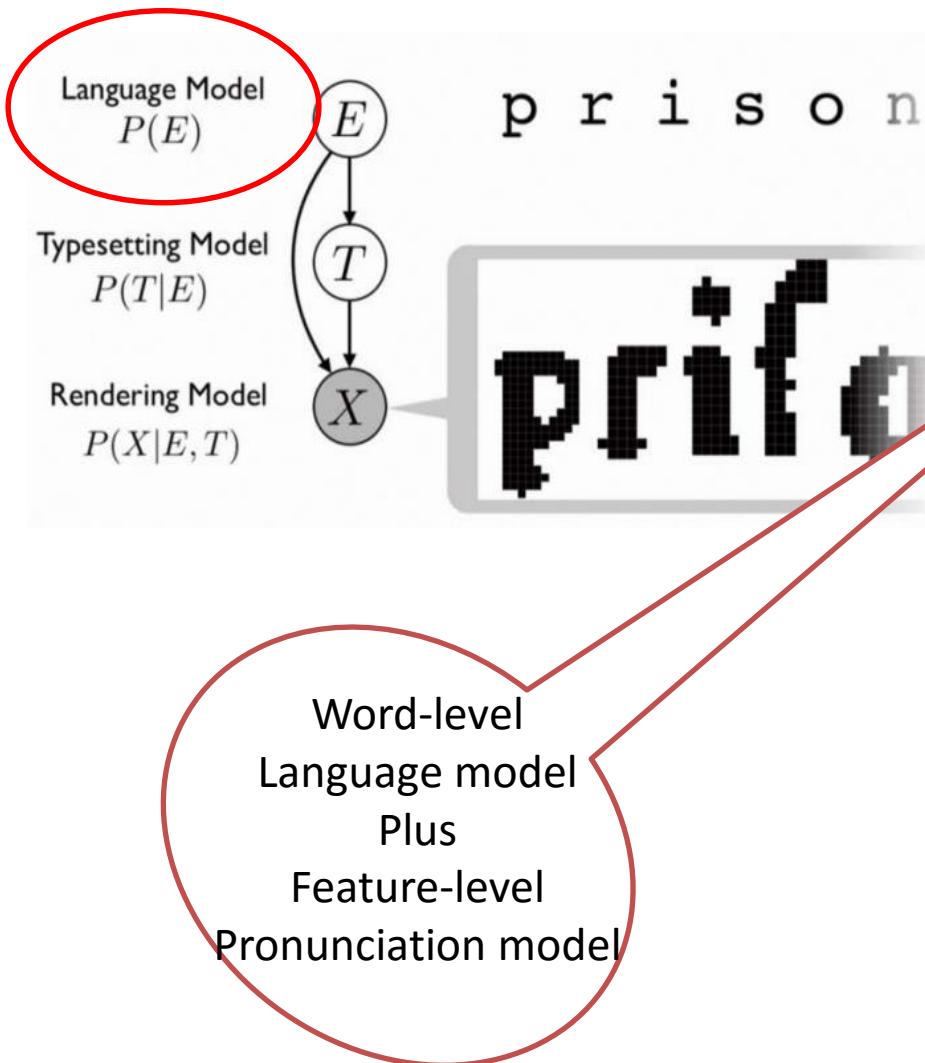
(Deng, 1998; Deng et al, 1997, 2000, 2003, 2006)



L.Deng, [A dynamic, feature-based approach to the interface between phonology & phonetics for speech modeling and recognition](#), *Speech Communication*, vol. 24, no. 4, pp. 299-323, 1998.

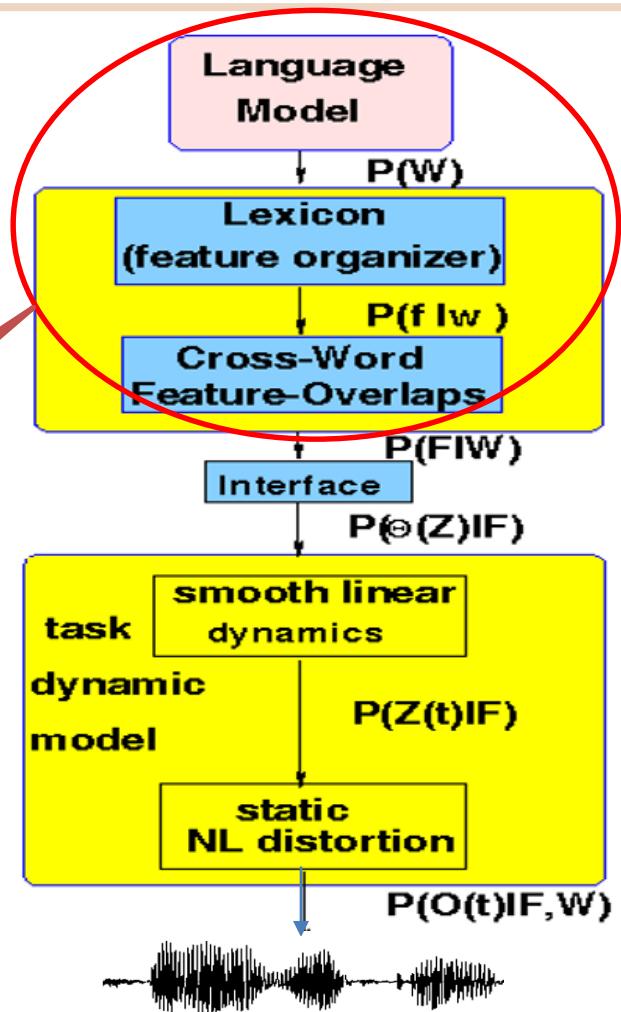
Deep Generative Model for Image-Text

(Berg-Kirkpatrick et al., 2013, 2015)



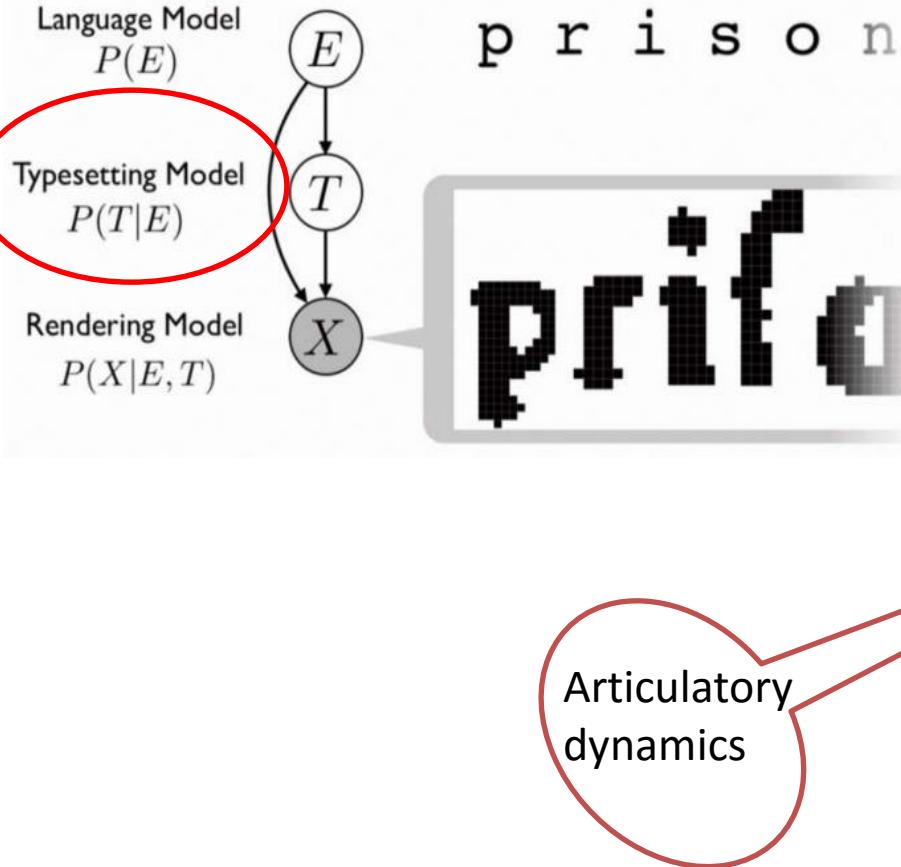
Deep Generative Model for Speech-Text

(Deng, 1998; Deng et al, 2000, 2003, 2006)



Deep Generative Model for Image-Text

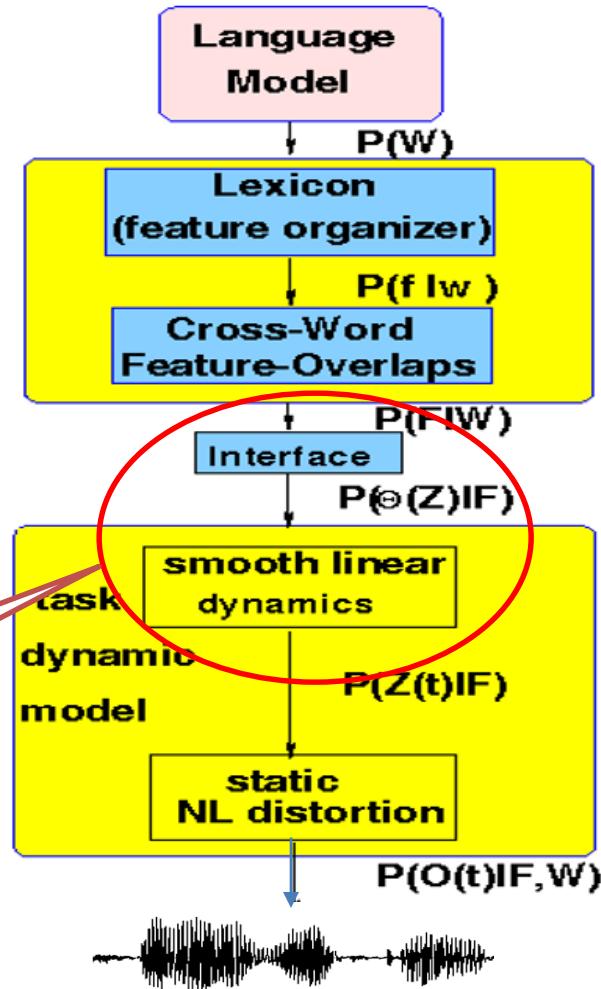
(Berg-Kirkpatrick et al., 2013, 2015)



Easy: likely no “explaining away” problem in inference and learning

Deep Generative Model for Speech-Text

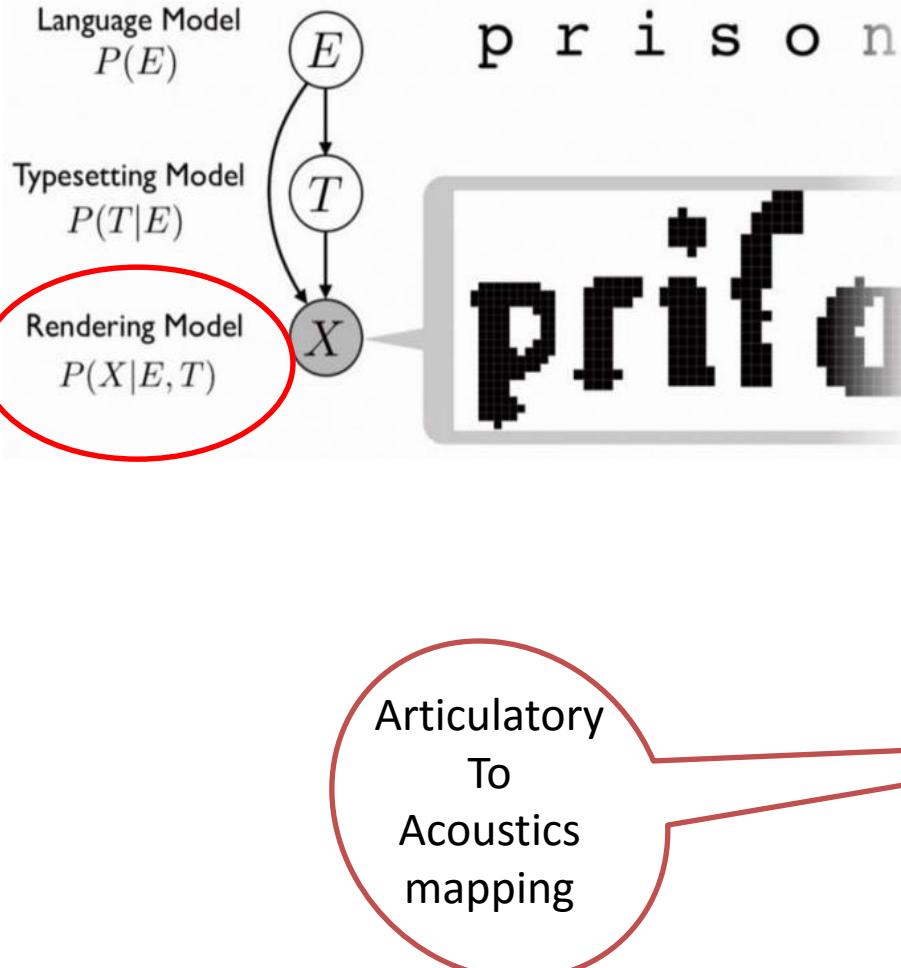
(Deng, 1998; Deng et al, 2000, 2003, 2006)



Hard: pervasive “explaining away” problem due to speech dynamics

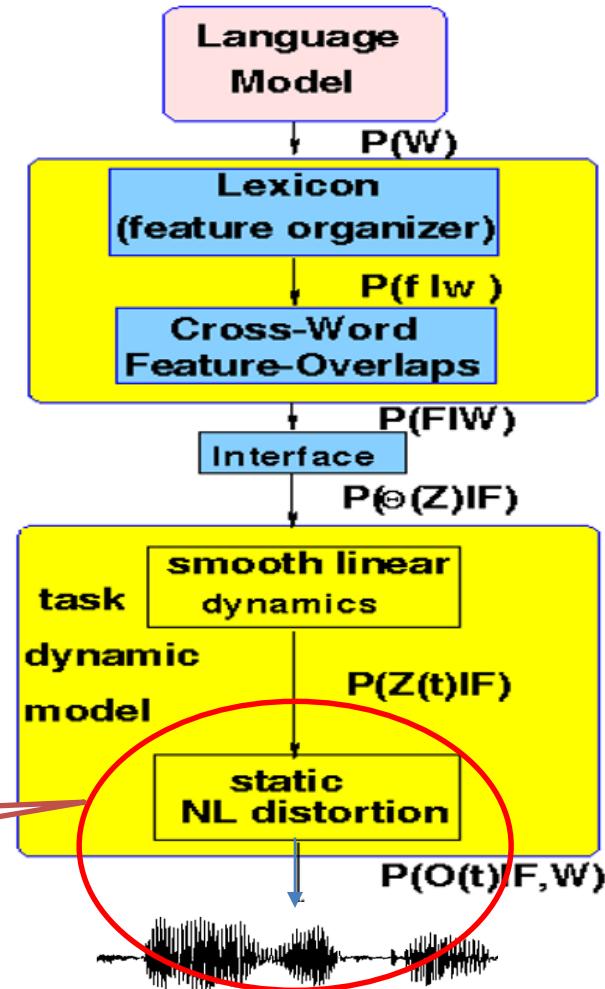
Deep Generative Model for Image-Text

(Berg-Kirkpatrick et al., 2013, 2015)

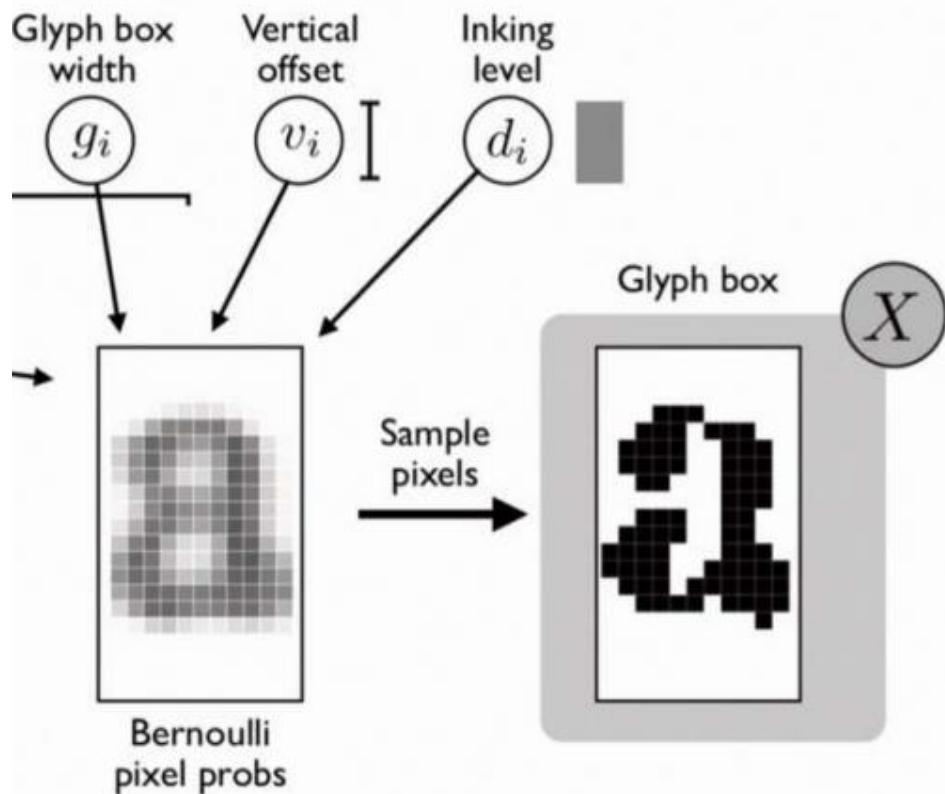


Deep Generative Model for Speech-Text

(Deng, 1998; Deng et al, 2000, 2003, 2006)



Rendering Model



Very simple, & easy to model accurately

- In contrast, articulatory-to-acoustics mapping in ASR is much more complex
- During 1997-2000, shallow NNs were used for this as “universal approximator”
- Not successful
- Now we have better DNN tool
- Even RNN/LSTM/CTC tool for dynamic modeling
- Essence: exploit the strong prior of LM: trained with billions of words/text
- No need to pair the text with acoustics; hence unsupervised learning
- Think of it as using a new objective function based on distribution matching between LM prior and the distribution of words predicted from acoustic models

Further thoughts on unsupervised ASR

- Deep generative modeling experiments not successful in 90's
 - ~~Computers were too slow~~
 - Models were too simple from text to speech waves
 - Still true → need speech scientists to work harder with technologists
 - And when generative models are not good enough, discriminative models and learning (e.g., RNN) can help a lot
 - Further, can iterate between the two, like wake-sleep (algorithm)
 - Inference/learning methods for deep generative models not mature at that time
 - Only partially true today
 - ← due to recent big advances in machine learning
 - Based on new ways of thinking about generative graphical modeling motivated by the availability of deep learning tools (e.g. DNN)
 - A brief review next

Advances in Inference Algms for Deep Generative Models

Kingma & Welling 2014, Salakhutdinov et al, 2015

ICML-2014 Talk Monday June 23, 15:20

In Track F (Deep Learning II)

**“Efficient Gradient Based Inference
through Transformations between
Bayes Nets and Neural Nets”**

Other solutions to solve the “large variance problem” in variational inference:

- Variational Bayesian Inference with Stochastic Search [D.M. Blei, M.I. Jordan and J.W. Paisley, 2012]
- Fixed-Form Variational Posterior Approximation through Stochastic Linear Regression [T. Salimans and A. Knowles, 2013].
- Black Box Variational Inference. [R. Ranganath, S. Gerrish and D.M. Blei. 2013]
- Stochastic Variational Inference [M.D. Hoffman, D. Blei, C. Wang and J. Paisley, 2013]
- Estimating or propagating gradients through stochastic neurons. [Y. Bengio, 2013].
- Neural Variational Inference and Learning in Belief Networks. [A. Mnih and K. Gregor, 2014, ICML]
- Stochastic backprop & approximation inference in deep generative models [D. Rezende, S. Mohamed, D. Wierstra, 2014]
- Semi-supervised learning with deep generative models [K. Kingma, D. Rezende, S. Mohamed, M. Welling, 2014, NIPS]
- auto-encoding variational Bayes [K. Kingma, M. Welling, 2014, ICML]
- Learning stochastic recurrent networks [Bayer and Osendorfer, 2015 ICLR]
- DRAW: A recurrent neural network for image generation. [K. Gregor, Danihelka, Rezende, Wierstra, 2015]
- Plus a number of NIPS-2015 papers, to appear.

Further thoughts on unsupervised ASR

- Deep generative modeling experiments not successful in 90's
 - ~~Computers were too slow~~
 - Models were too simple from text to speech waves
 - Still true → need speech scientists to work harder with technologists
 - **And when generative models are not good enough, discriminative models and learning (e.g., RNN) can help a lot; but to do it? A hint next**
 - Further, can iterate between the two, like wake-sleep (algorithm)
 - Inference/learning methods for deep generative models not mature at that time
 - Only partially true today
 - ← due to recent big advances in machine learning
 - Based on new ways of thinking about generative graphical modeling motivated by the availability of deep learning tools (e.g. DNN)
 - A brief review next

RNN

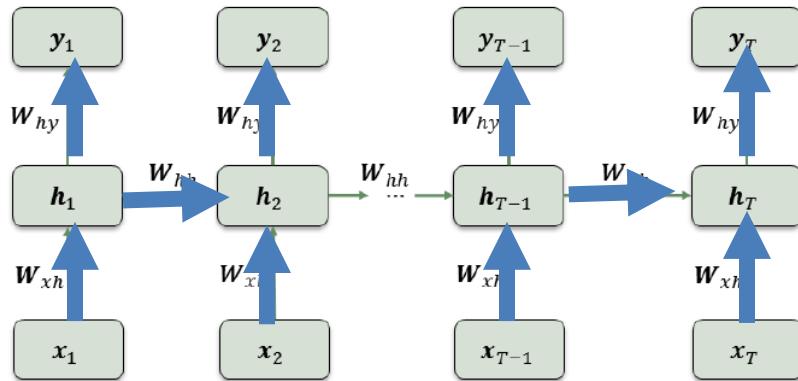
vs. Generative HDM

$$\mathbf{h}_t = f(\mathbf{h}_{t-1}; \mathbf{W}_{hh}, \mathbf{W}_{xh}, \mathbf{x}_t)$$

$$\mathbf{y}_t = g(\mathbf{h}_t; \mathbf{W}_{hy})$$

Parameterization:

- $\mathbf{W}_{hh}, \mathbf{W}_{hy}, \mathbf{W}_{xh}$: all unstructured regular matrices

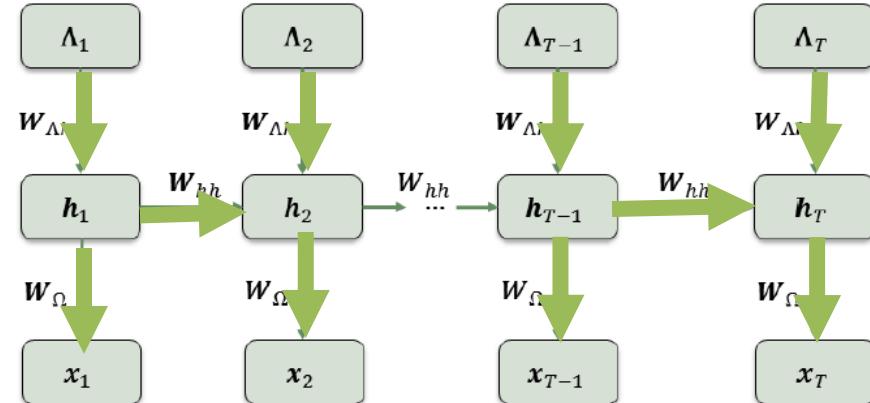


$$\mathbf{h}_t = q(\mathbf{h}_{t-1}; \mathbf{W}_{l_t}, \mathbf{t}_{l_t})$$

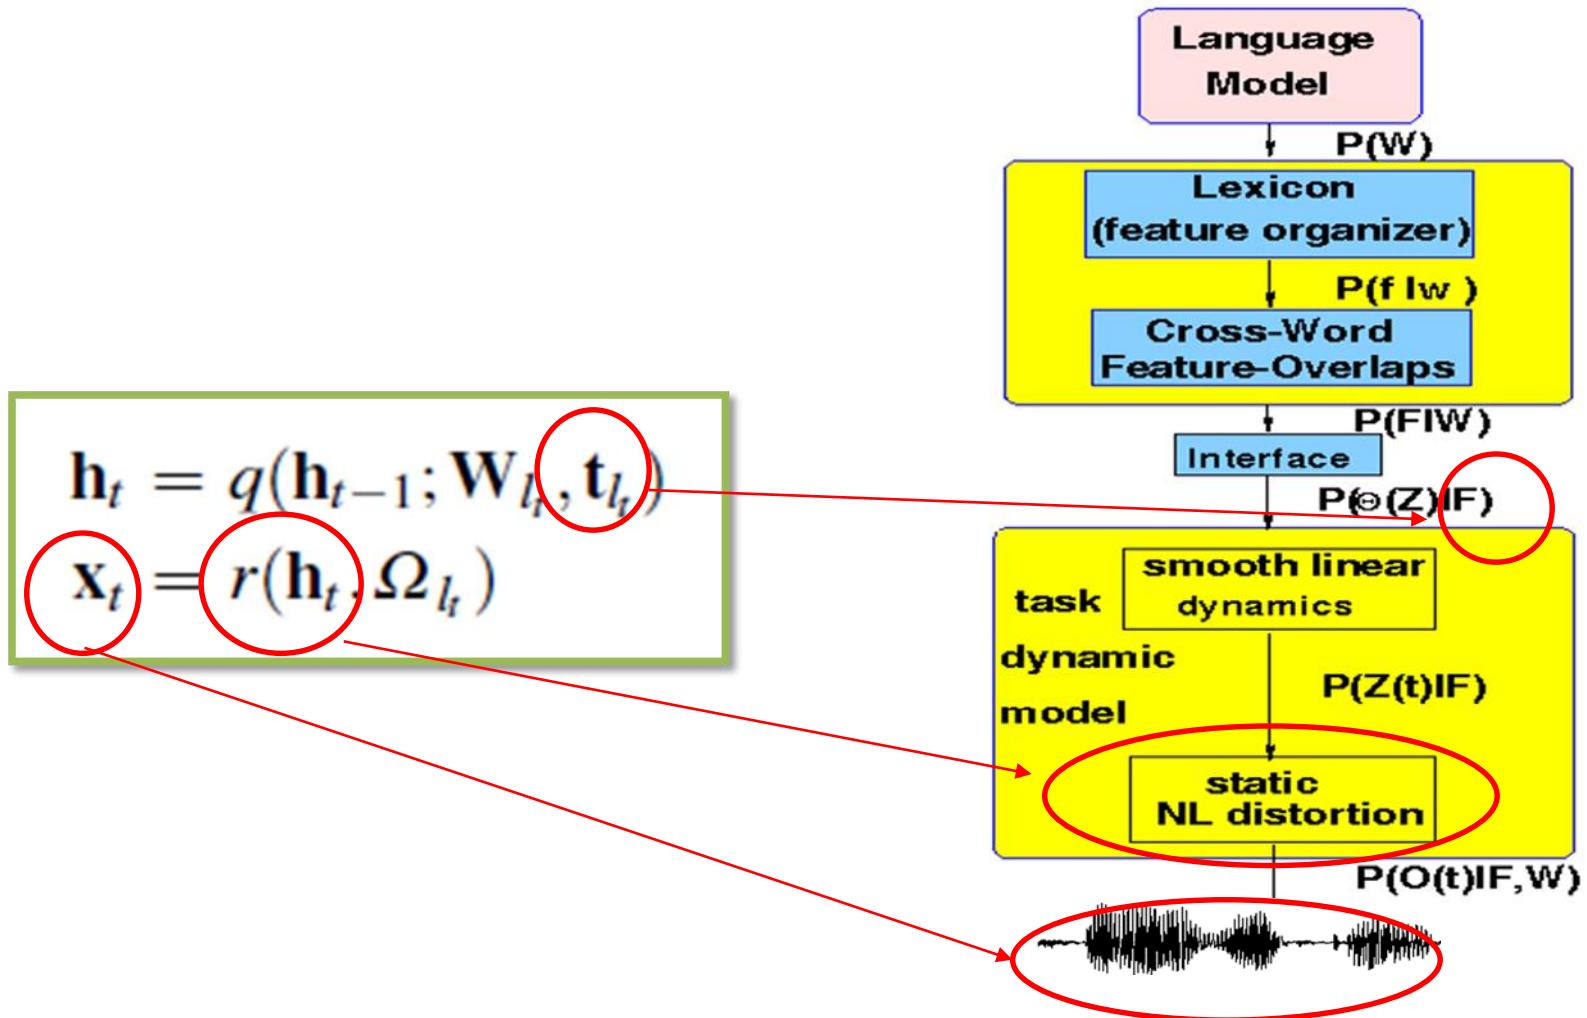
$$\mathbf{x}_t = r(\mathbf{h}_t, \Omega_{l_t})$$

Parameterization:

- $\mathbf{W}_{hh} = \mathbf{M}(\gamma_l)$; sparse system matrix
- $\mathbf{W}_{\Omega} = (\Omega_l)$; Gaussian-mix params; MLP
- $\Lambda = \mathbf{t}_l$



Generative HDM



RNN

vs. Generative HDM

$$\begin{aligned}\mathbf{h}_t &= f(\mathbf{h}_{t-1}; \mathbf{W}_{hh}, \mathbf{W}_{xh}, \mathbf{x}_t) \\ \mathbf{y}_t &= g(\mathbf{h}_t; \mathbf{W}_{hy})\end{aligned}$$

$$\begin{aligned}\mathbf{h}_t &= q(\mathbf{h}_{t-1}; \mathbf{W}_{lt}, \mathbf{t}_{lt}) \\ \mathbf{x}_t &= r(\mathbf{h}_t, \Omega_{lt})\end{aligned}$$

~DNN

~DBN

e.g. Generative pre-training

(analogous to generative DBN pretraining for DNN)

NIPS-2015 paper to appear on simpler dynamic models for a non-ASR application

Better ways of integrating deep generative/discriminative models are possible

- Hint: Example 1 where generative models are used to define the DNN architecture

end of

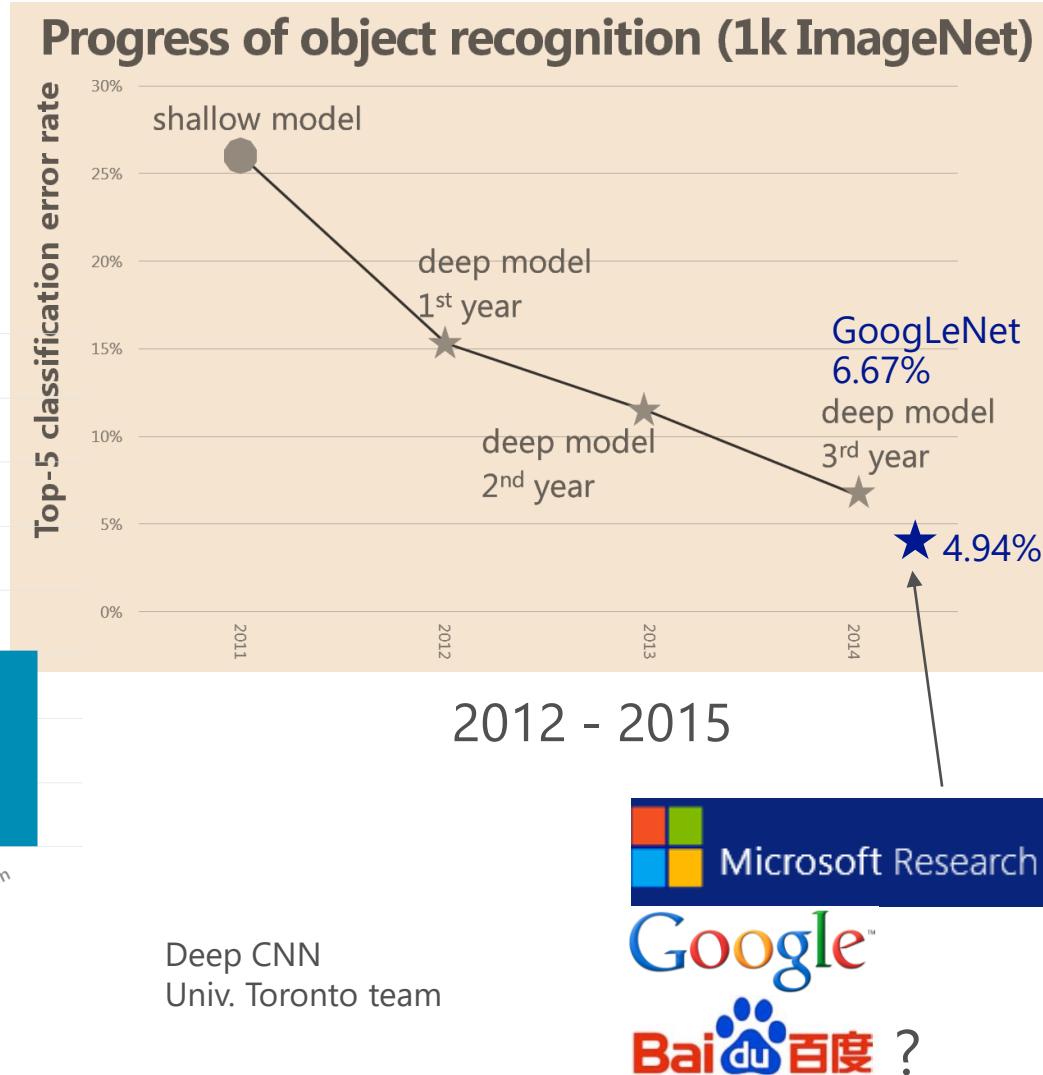
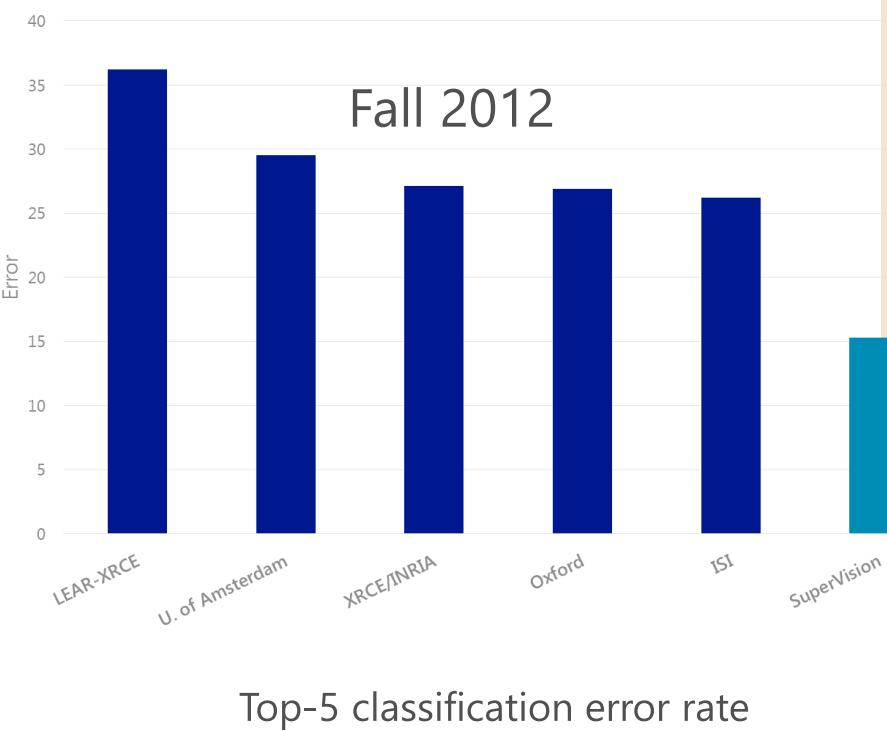
Part II: Speech

- Deep supervised learning shattered ASR via DNN/LSTM
- Deep unsupervised learning may impact more in the future
- No more low-hanging fruit

Deep Learning also Shattered the Entire Field of
Image Recognition and Computer Vision (since 2012)

ImageNet-1K Competition

Krizhevsky, Sutskever, Hinton, "ImageNet Classification with Deep Convolutional Neural Networks." *NIPS*, Dec. 2012



Part III: Language

Moving **from perception**: phonetic/word recognition, image/gesture recognition, etc
to

Cognition: **memory, attention, reasoning, Q/A, & decision making**, etc.

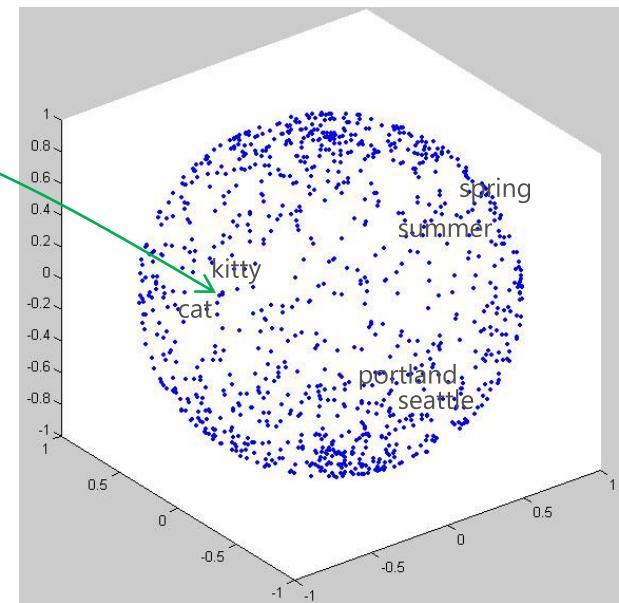
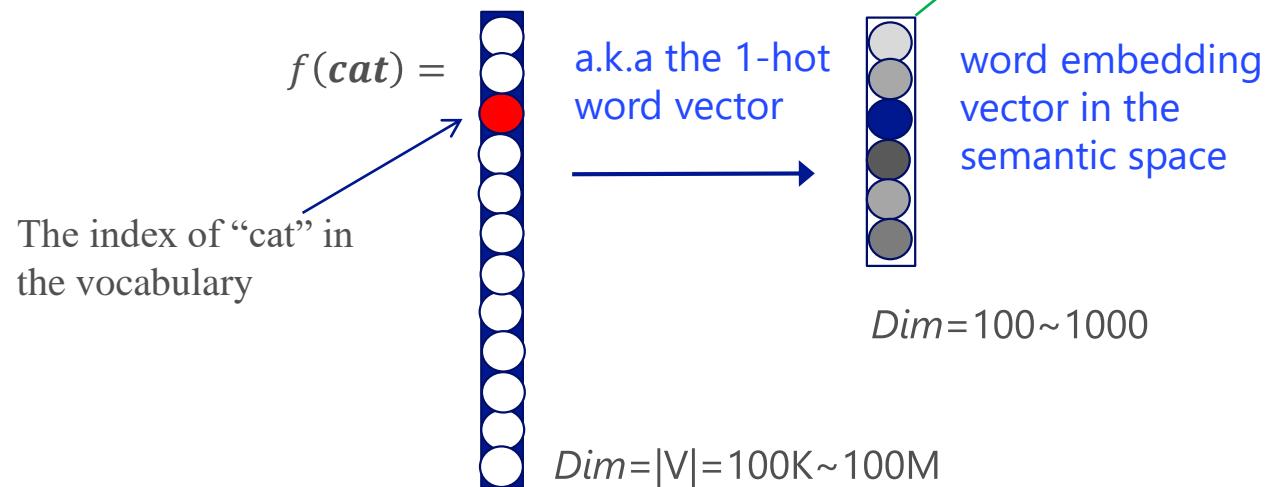
Embedding enables exploring models for these human cognitive functions

Embedding Linguistic entities in low-dimensional continuous space and Examples of NLP Applications

Semantic embedding

Project raw text into a continuous semantic space
e.g., word embedding

Captures the word meaning in a semantic space

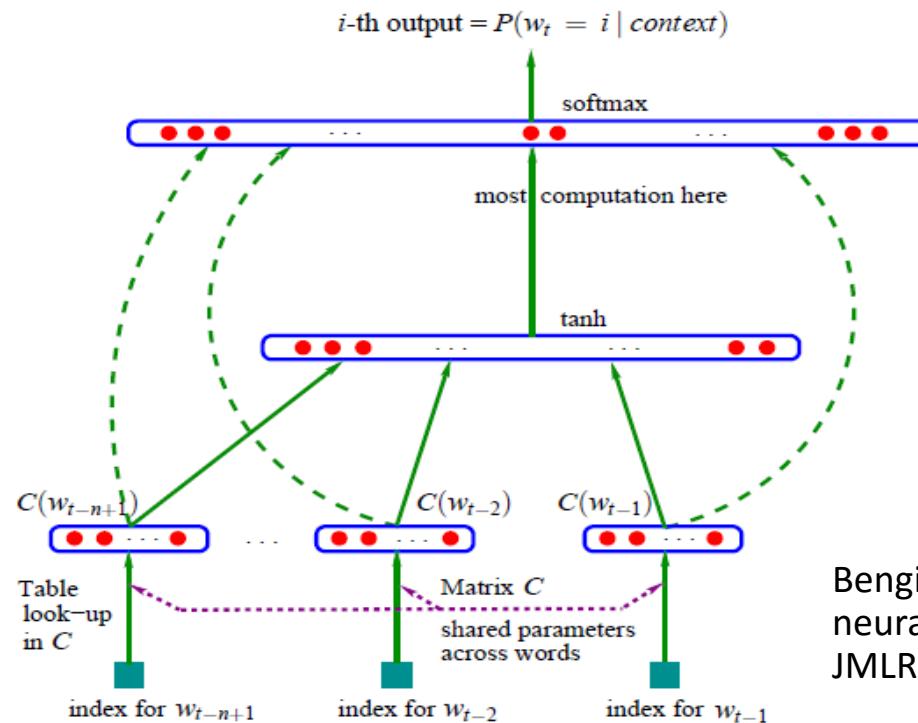


Deerwester, Dumais, Furnas, Landauer, Harshman, "Indexing by latent semantic analysis," JASIS 1990

NN Word Embedding

LM: predict the next word given the past:

e.g., $p(\text{chases}|\text{the cat}) = ?$, $p(\text{says}|\text{the cat}) = ?$



Bengio, Ducharme, Vincent, Jauvin, "A neural probabilistic language model." JMLR, 2003

SENNA word embedding

Scoring:

$$Score(w_1, w_2, w_3, w_4, w_5) = U^T \sigma(W[f_1, f_2, f_3, f_4, f_5] + b)$$

Training:

$$J = \max(0, 1 + S^- - S^+)$$

Update the model until $S^+ > 1 + S^-$

Where

$$S^+ = Score(w_1, w_2, w_3, w_4, w_5)$$

$$S^- = Score(w_1, w_2, w^-, w_4, w_5)$$

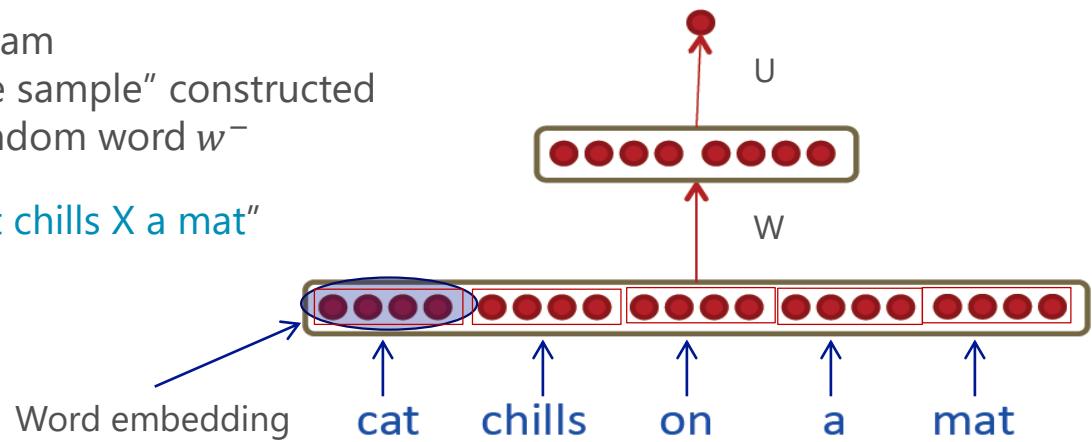
And

w_1, w_2, w_3, w_4, w_5 is a valid 5-gram

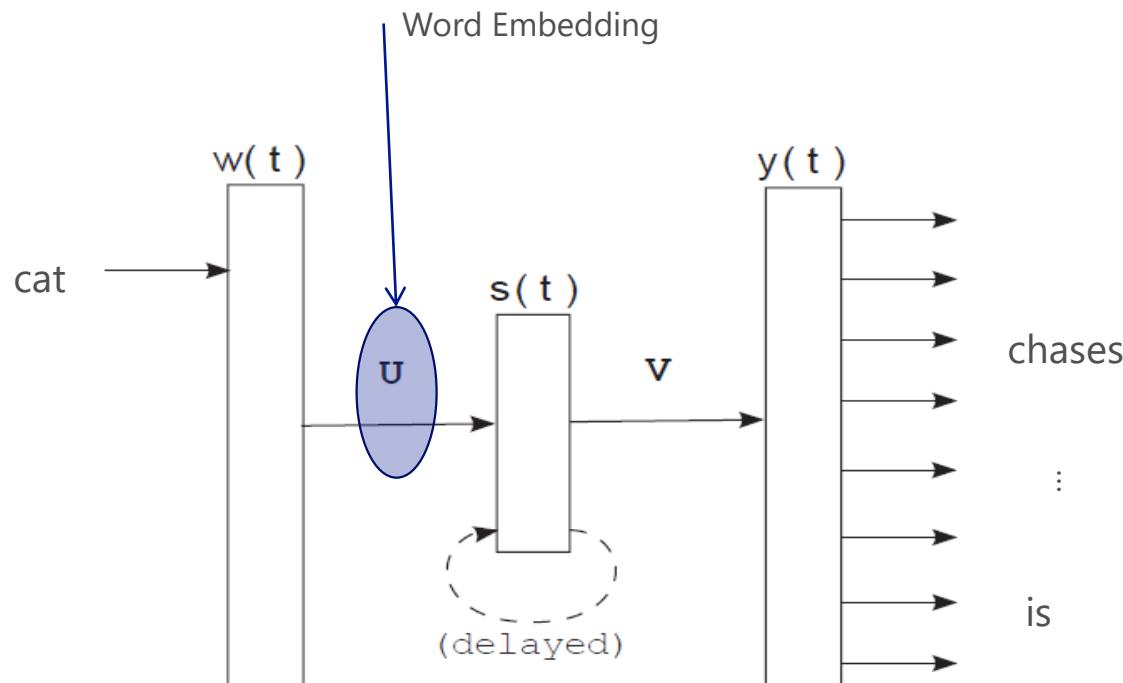
w_1, w_2, w^-, w_4, w_5 is a "negative sample" constructed by replacing the word w_3 with a random word w^-

e.g., a negative example: "cat chills X a mat"

Collobert, Weston, Bottou, Karlen,
Kavukcuoglu, Kuksa, "Natural Language
Processing (Almost) from Scratch," JMLR
2011

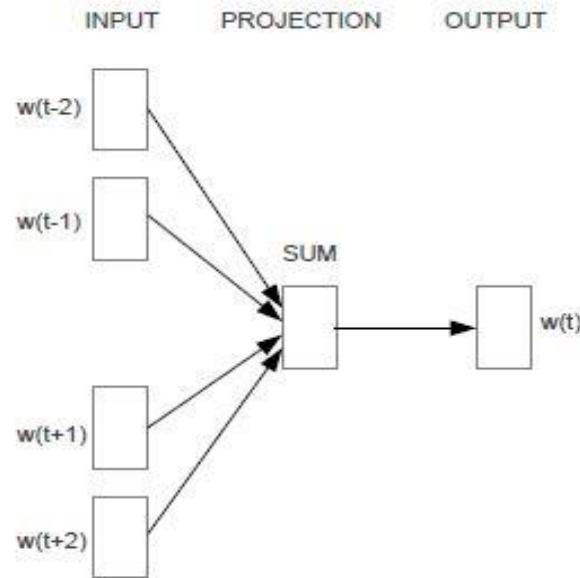


RNN-LM base word embedding



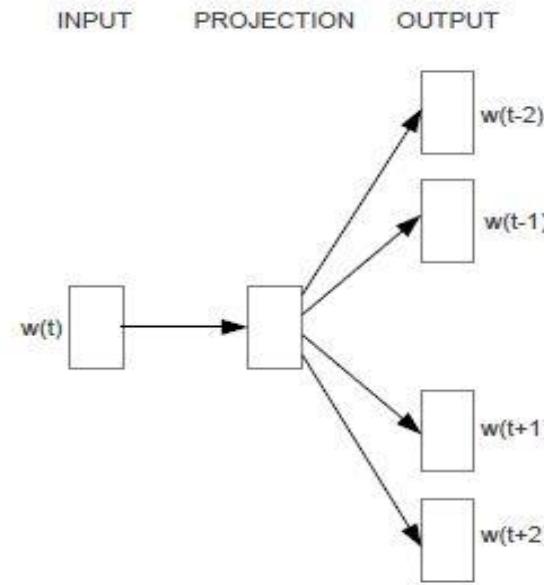
Mikolov, Yih, Zweig, "Linguistic Regularities in Continuous Space Word Representations," NAACL 2013

CBOW/Skip-gram Word Embeddings



CBOW

Continuous Bag-of-Words

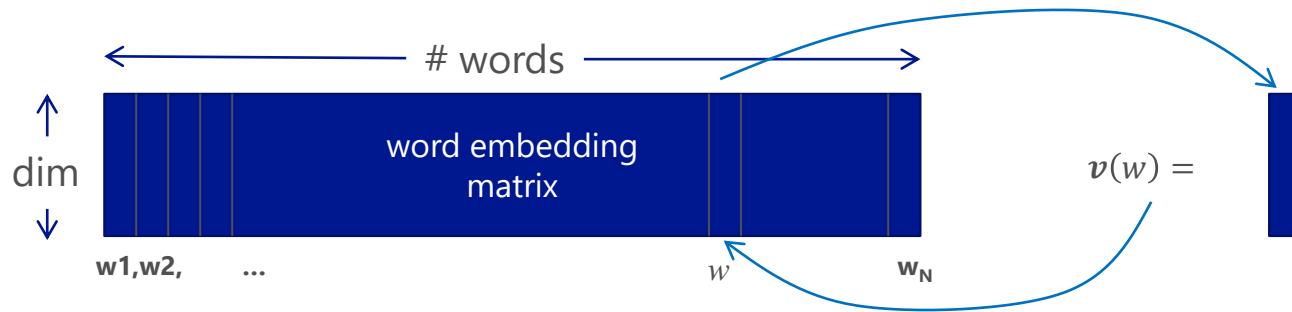


Skip-gram

The CBOW architecture (a) on the left, and the Skip-gram architecture (b) on the right.
[Mikolov et al., 2013 ICLR].

Word embedding: rethinking

- Word embedding is a neat and effective representation:

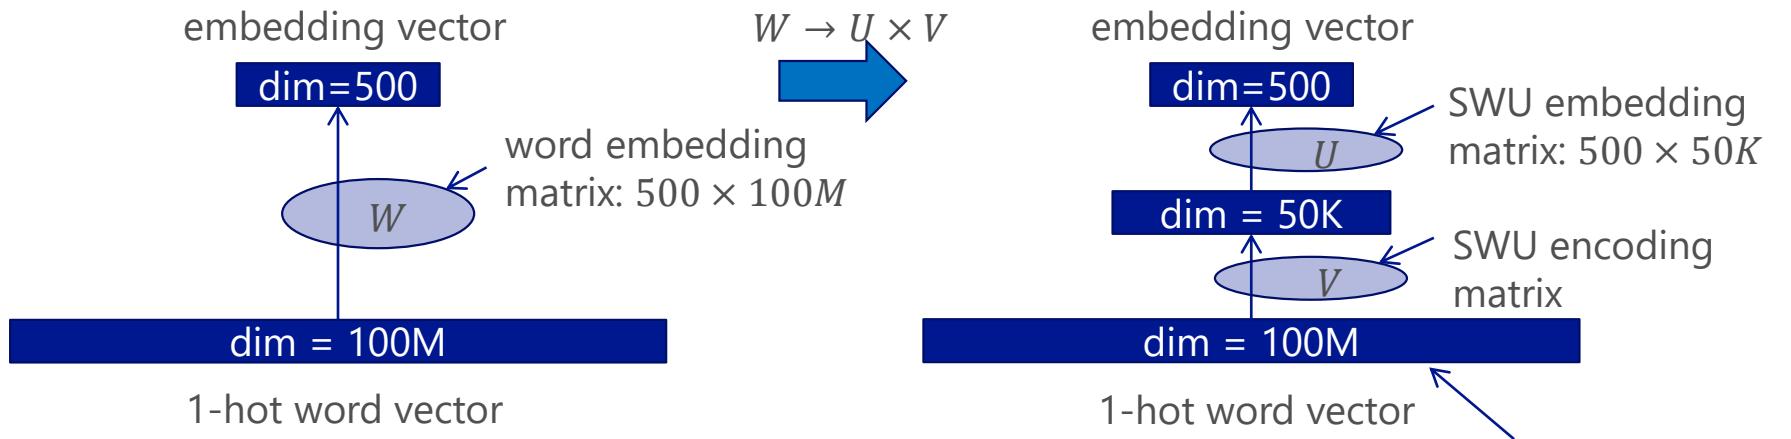


- However, for large scale NL tasks a decomposable, robust representation is preferable
 - Vocabulary of real-world big data tasks could be huge (*scalability*)
>100M unique words in a modern commercial search engine log, and keeps growing
 - New words, misspellings, and word fragments frequently occur (*generalizability*)

Build semantic embedding on top of sub-word units

Learn semantic embedding on top of sub-word units (SWU)

- Decompose *any* word into sub-word units
- *Scale* the capacity to handle almost unbounded variability (word) based on bounded variability (sub-word)



Huang, He, Gao, Deng, Acero, Heck, "Learning deep structured semantic models for web search using clickthrough data," CIKM, 2013

Sub-word unit

- Letters, context-dept letters, positioned-phones, context-dept phones, positioned-roots/morphs, context-dept morphs
- Multi-hashing approach to word input representation

Or random projection (random basis)

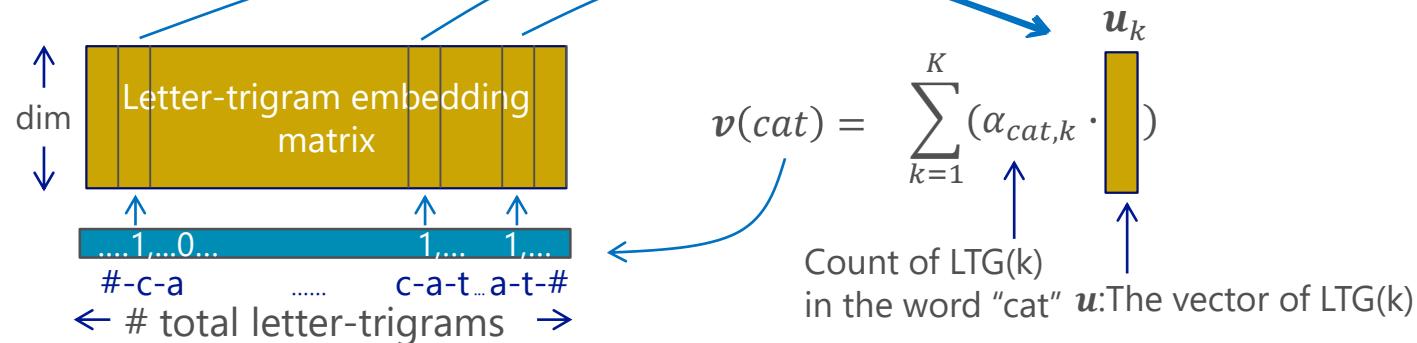
From sub-word unit embedding vectors to word vectors

SWU uses context-dependent letter, e.g., letter-trigram.

Learn one vector per letter-trigram (LTG), the encoding matrix is a fixed matrix

- Use the count of each LTG in the word for encoding

Example: cat → #cat# → #-c-a, c-a-t, a-t-#
(w/ word boundary mark #)



Two words has the same LTG:
collision rate $\approx 0.004\%$

Supervised Embeddings for Semantic Modeling with Applications

- embedding linguistic symbols by backprop
- mining distant supervision signals

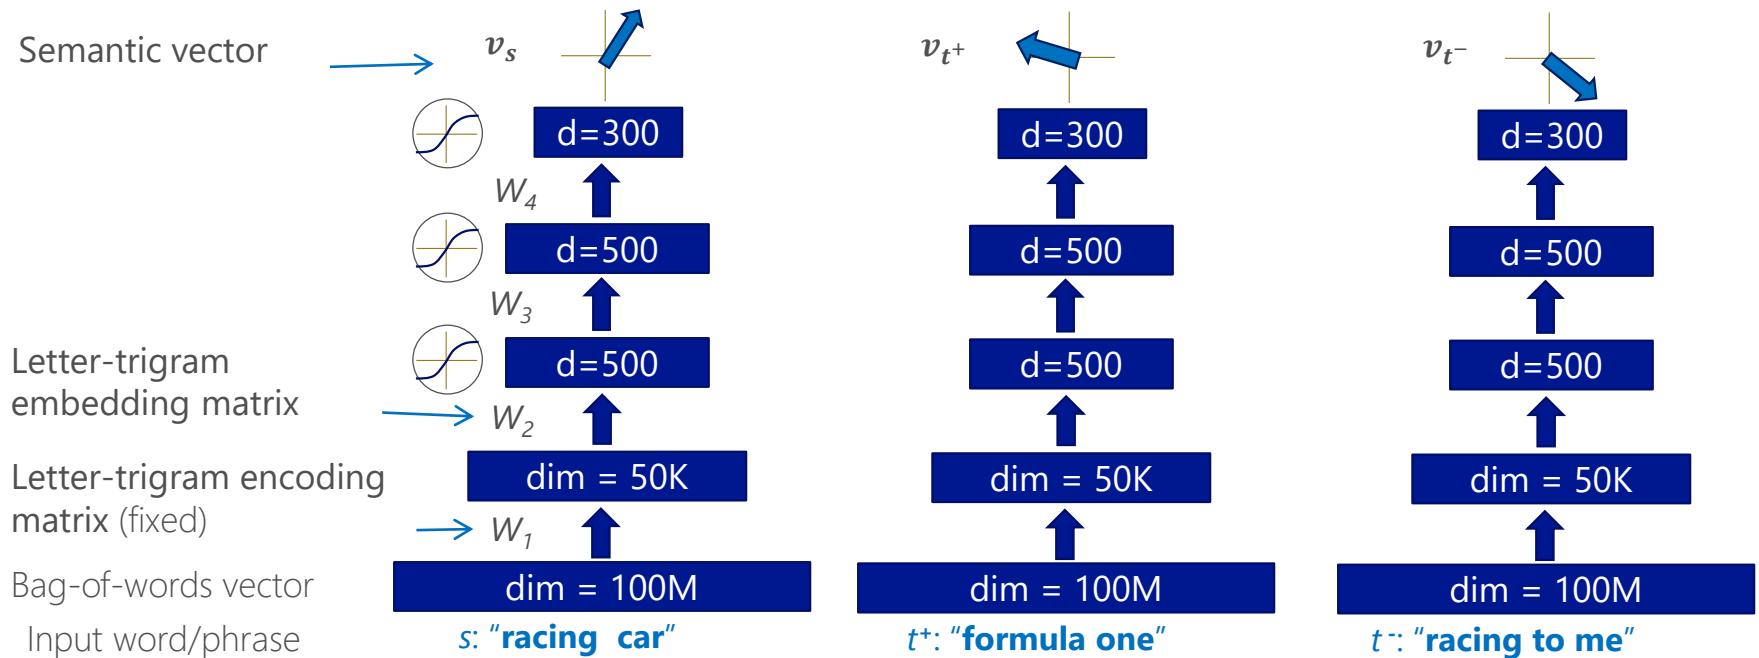
Deep Structured Semantic Model (DSSM)

- Build word/phrase, or sentence-level semantic vector representation
- Trained by a similarity-driven objective
 - projecting semantically similar phrases to vectors close to each other
 - projecting semantically different phrases to vectors far apart

DSSM for learning semantic embedding

Initialization:

Neural networks are initialized with random weights



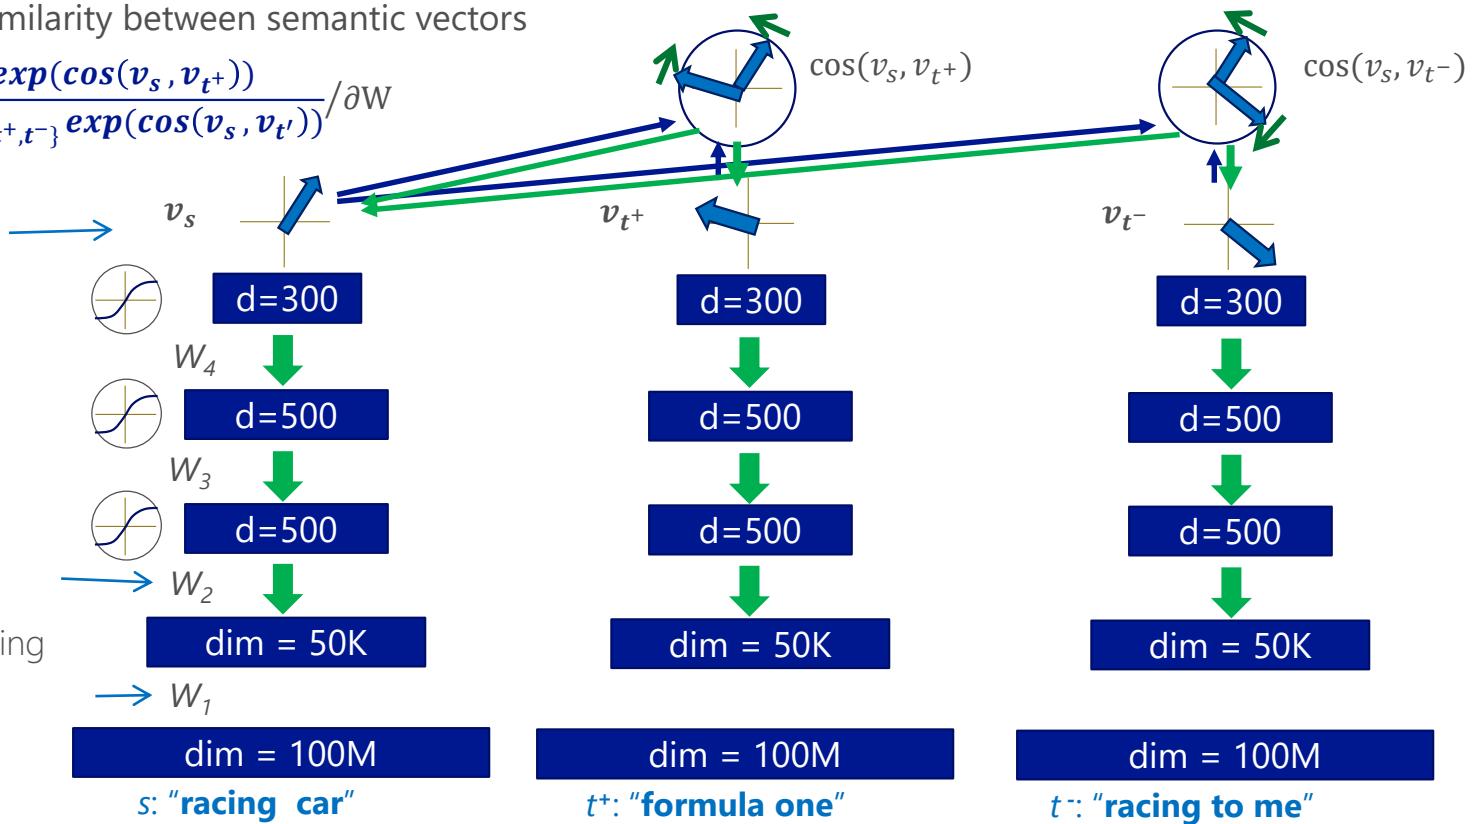
DSSM for learning semantic embedding

Training:

Compute Cosine similarity between semantic vectors

Compute gradients $\frac{\partial \exp(\cos(v_s, v_{t^+}))}{\partial W} / \sum_{t'=\{t^+, t^-\}} \exp(\cos(v_s, v_{t'}))$

Semantic vector



Letter-trigram
embedding matrix

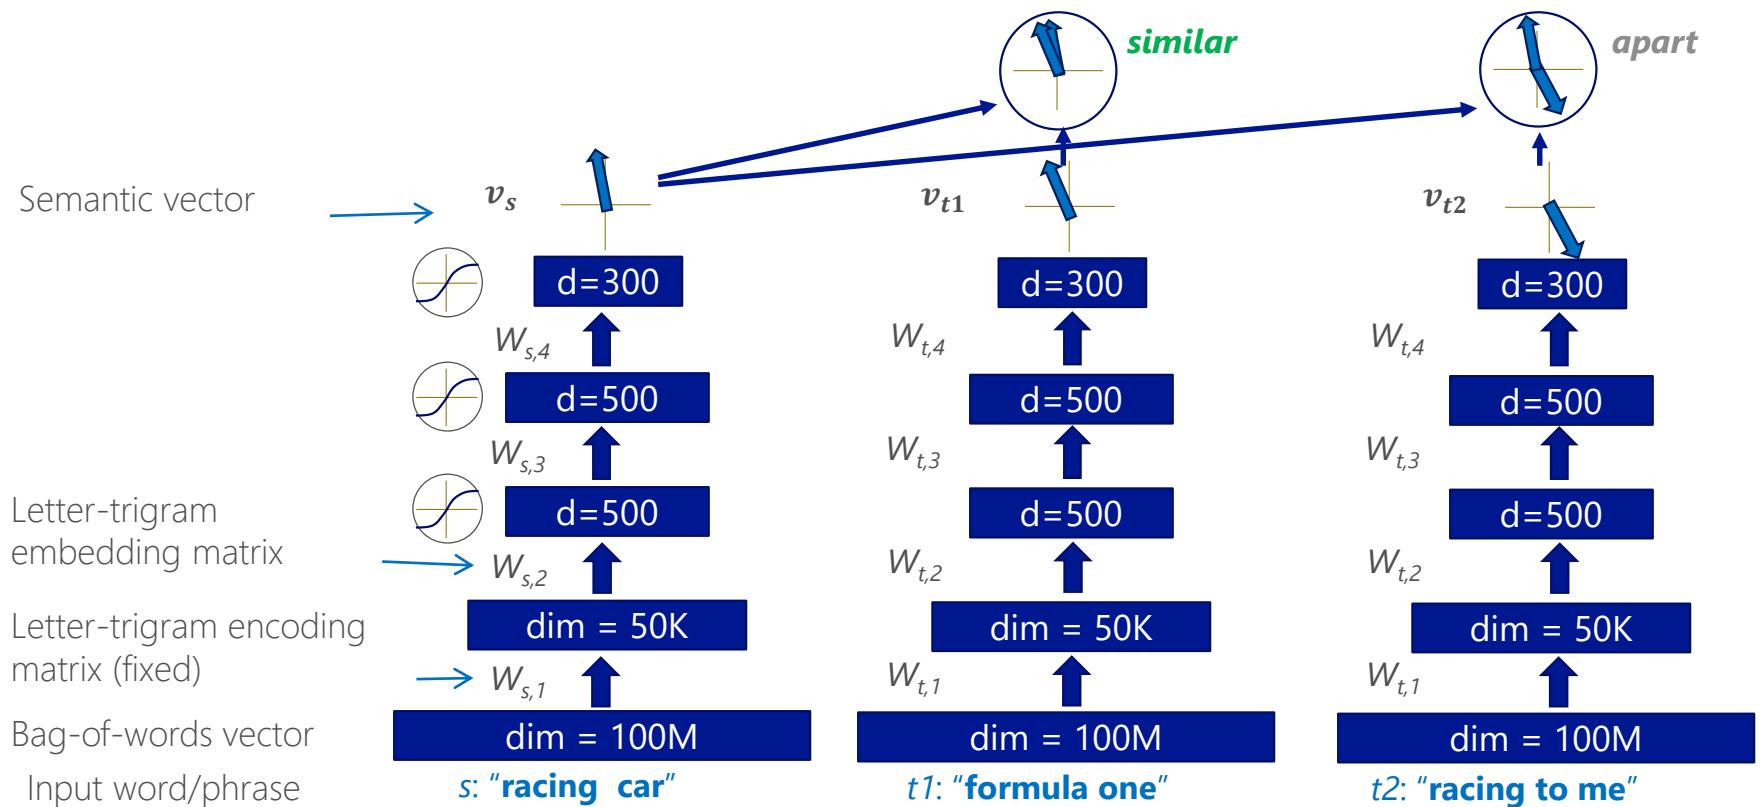
Letter-trigram encoding
matrix (fixed)

Bag-of-words vector
Input word/phrase



DSSM for learning semantic embedding

Runtime:



Training of the DSSM

Data: semantically-similar text pairs

e.g., **context <-> word** in word embedding vector learning

query <-> clicked-doc in Web Search

pattern<-> relationship in Question Answering

Objective: cosine similarity based loss

- Web search as an example: a query \mathbf{q} and a list of docs $\mathbf{D} = \{\mathbf{d}^+, \mathbf{d}_1^-, \dots \mathbf{d}_K^-\}$
 - \mathbf{d}^+ positive doc; $\mathbf{d}_1^-, \dots \mathbf{d}_K^-$ are negative docs to \mathbf{q} (e.g., sampled from not clicked docs)
- Objective: the posterior probability of clicked document given query

$$P(d^+|q) = \frac{\exp(\gamma \cos(q, d^+))}{\sum_{d \in \mathbf{D}} \exp(\gamma \cos(q, d))}$$

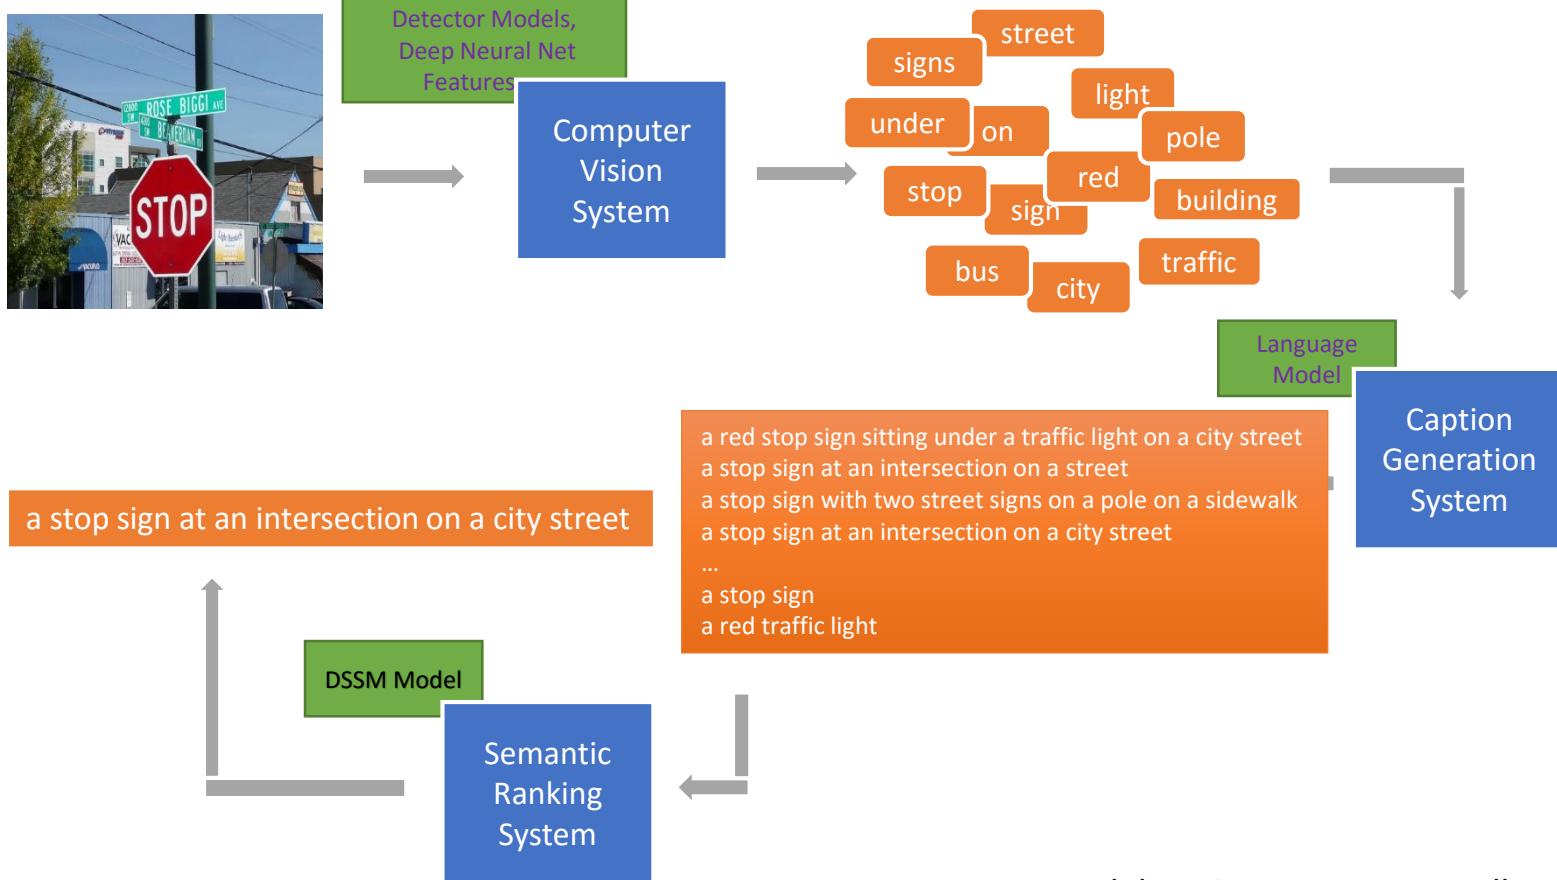
- Optimize Θ to maximize $P(d^+|q)$. SGD training on GPU (NVidia K20x)



Many applications of DSSM (many low-hanging fruits): Learning semantic similarity between X and Y

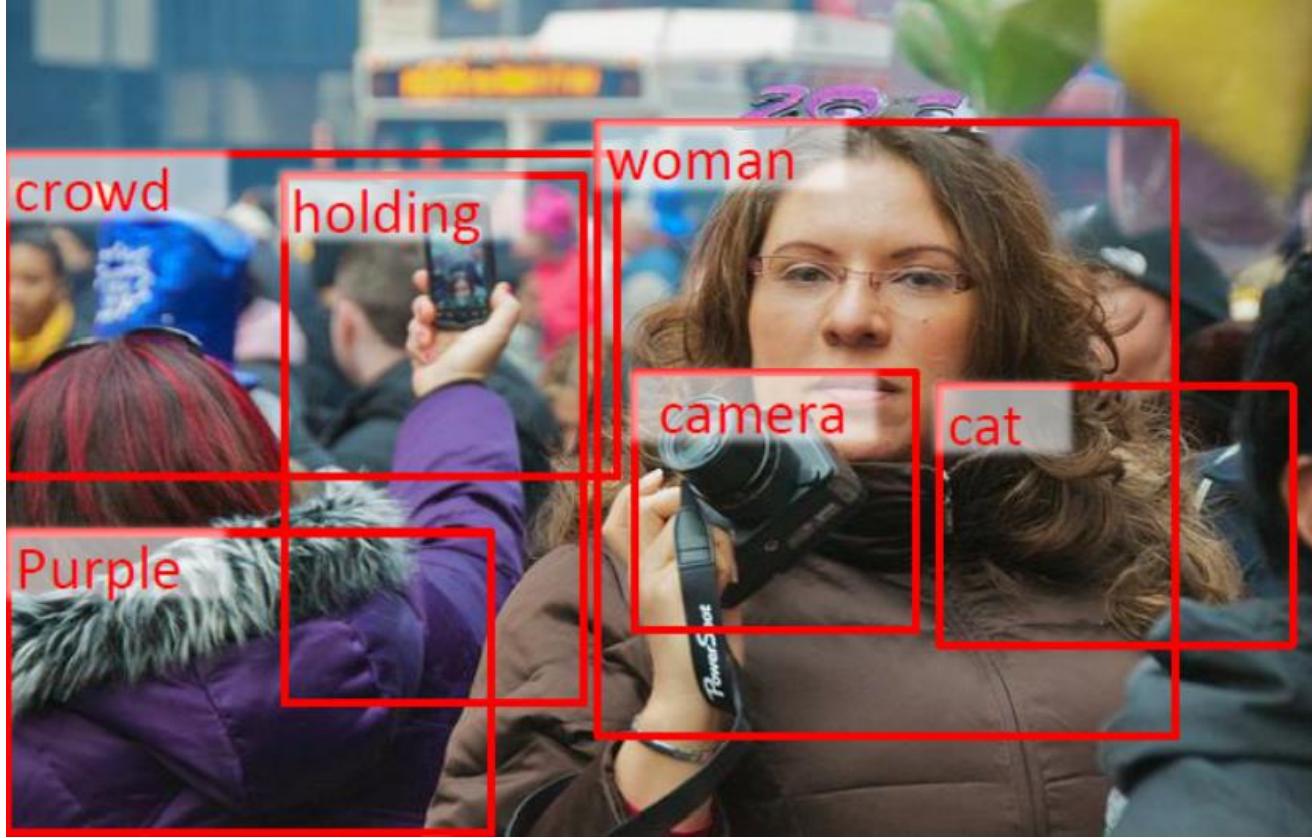
Tasks	Source X	Target Y
Word semantic embedding	<i>context</i>	<i>word</i>
Web search	<i>search query</i>	<i>web documents</i>
Query intent detection	<i>Search query</i>	<i>Use intent</i>
Question answering	<i>pattern / mention (in NL)</i>	<i>relation / entity (in KB)</i>
Machine translation	<i>sentence in language a</i>	<i>translated sentences in language b</i>
Query auto-suggestion	<i>Search query</i>	<i>Suggested query</i>
Query auto-completion	<i>Partial search query</i>	<i>Completed query</i>
Apps recommendation	<i>User profile</i>	<i>recommended Apps</i>
Distillation of survey feedbacks	<i>Feedbacks in text</i>	<i>Relevant feedbacks</i>
Automatic image captioning	<i>image</i>	<i>text caption</i>
Image retrieval	<i>text query</i>	<i>images</i>
Natural user interface	<i>command (text / speech / gesture)</i>	<i>actions</i>
Ads selection	<i>search query</i>	<i>ad keywords</i>
Ads click prediction	<i>search query</i>	<i>ad documents</i>
Email analysis: people prediction	<i>Email content</i>	<i>Recipients, senders</i>
Email search	<i>Search query</i>	<i>Email content</i>
Email decluttering	<i>Email contents</i>	<i>Email contents in similar threads</i>
Knowledge-base construction	<i>entity from source</i>	<i>entity fitting desired relationship</i>
Contextual entity search	<i>key phrase / context</i>	<i>entity / its corresponding page</i>
Automatic highlighting	<i>documents in reading</i>	<i>key phrases to be highlighted</i>
Text summarization	<i>long text</i>	<i>summarized short text</i>

Automatic image captioning (MSR system)



Fang, Gupta, Iandola, Srivastava, Deng, Dollar, Gao, He, Mitchell, Platt, Zitnick, Zweig, **“From captions to visual concepts and back,”** accepted to appear in CVPR, 2015; in arXiv 2014

Microsoft System (MSR): Use of DSSM for Global Semantic Matching



1. Word
Detection

2. Sentence
Generation

3. Sentence
Re-Ranking

woman, crowd, cat,
camera, holding,
purple

A purple camera with a woman.
A woman holding a camera in a crowd.
...
A woman holding a cat.

#1 A woman holding a
camera in a crowd.



A

a woman in a kitchen preparing food

B

woman working on counter near kitchen sink preparing a meal



Machine:

Human:

a woman in a kitchen preparing food

woman working on counter near kitchen sink preparing a meal



Machine-generated (but turker preferred)	a group of motorcycles parked next to a motorcycle	Machine-generated (but turker preferred)	a woman in a kitchen preparing food
Human-annotated (but turker not preferred)	two girls wearing short skirts and one of them sits on a motorcycle while the other stands nearby	Human-annotated (but turker not preferred)	woman working on counter near kitchen sink preparing a meal



Machine-generated (but turker preferred)	a clock tower in the middle of the street	Machine-generated (but turker preferred)	a kitchen with wooden cabinets and a sink
Human-annotated (but turker not preferred)	a statue with a clock on it near a parking lot	Human-annotated (but turker not preferred)	an ornate kitchen is designed with rustic wooden parts



Machine-generated (but turker preferred)	a man holding a tennis racquet on a tennis court
Human-annotated (but turker not preferred)	the man is on the tennis court playing a game

Potential Scenarios

take a snapshot, generate a caption,
and share the fun 😊



a dining room table and chairs
next to a window



a young boy standing on a beach



a group of people sitting at a table
with laptops



a man that is standing in front of a
store

Field test of our prototype App: on cell phone-quality photos

Competition results (CVPR-2015, June, Boston)

Won 1st Prize at MS COCO Captioning Challenge 2015!



Measure the quality of the captions by human judge (e.g., *Turing Test*).

The top teams and the state-of-the-art

	% of captions that pass the Turing Test	Official Rank
MSR	32.2%	1st(tie)
Google	31.7%	1st(tie)
MSR Captivator	30.1%	3rd(tie)
Montreal/Toronto	27.2%	3rd(tie)
Berkeley LRCN	26.8%	5th

Note: even a *Human* cannot guarantee to pass Turing Test by 100%

Other groups: Baidu/UCLA, Stanford, Tsinghua, etc.

Human	67.5%	--
-------	-------	----

Many applications of DSSM (many low-hanging fruits):

Learning semantic similarity between X and Y

Tasks	Source X	Target Y
Word semantic embedding	<i>context</i>	<i>word</i>
Web search	<i>search query</i>	<i>web documents</i>
Query intent detection	<i>Search query</i>	<i>Use intent</i>
Question answering	<i>pattern / mention (in NL)</i>	<i>relation / entity (in KB)</i>
Machine translation	<i>sentence in language a</i>	<i>translated sentences in language b</i>
Query auto-suggestion	<i>Search query</i>	<i>Suggested query</i>
Query auto-completion	<i>Partial search query</i>	<i>Completed query</i>
Apps recommendation	<i>User profile</i>	<i>recommended Apps</i>
Distillation of survey feedbacks	<i>Feedbacks in text</i>	<i>Relevant feedbacks</i>
Automatic image captioning	<i>image</i>	<i>text caption</i>
Image retrieval	<i>text query</i>	<i>images</i>
Natural user interface	<i>command (text / speech / gesture)</i>	<i>actions</i>
Ads selection	<i>search query</i>	<i>ad keywords</i>
Ads click prediction	<i>search query</i>	<i>ad documents</i>
Email analysis: people prediction	<i>Email content</i>	<i>Recipients, senders</i>
Email search	<i>Search query</i>	<i>Email content</i>
Email decluttering	<i>Email contents</i>	<i>Email contents in similar threads</i>
Knowledge-base construction	<i>entity from source</i>	<i>entity fitting desired relationship</i>
Contextual entity search	<i>key phrase / context</i>	<i>entity / its corresponding page</i>
Automatic highlighting	<i>documents in reading</i>	<i>key phrases to be highlighted</i>
Text summarization	<i>long text</i>	<i>summarized short text</i>

DECEMBER
10
2014

Bing brings the world's knowledge straight to you with Insights for Office

Today Bing and Office are introducing Insights for Office, a new, more powerful way to search for the information you need **while in Office Word Online** – available in English to all markets in the next few days. We encourage you to [try it here](#), always free.

How Bing's intelligence powers Insights for Office

Bing indexes and stores entity data from around the web representing real world people, places and things. Insights for Office utilizes Bing's ability to index the world's knowledge and our [machine learned relevance models](#) to semantically understand the most important content in a user's document and then return the most relevant results. This capability is derived largely from patterns of text analysis developed in collaboration with Microsoft Research. The results deliver the

Scenario: Contextual search in Microsoft Office/Word

The screenshot shows a Microsoft Word Online interface. In the center, the word "Lincoln" is selected in a large font. A context menu is open over this word, with "Insights" highlighted. To the right of the menu, there is an "Insights" panel. The panel title is "Abraham Lincoln". Below the title is a portrait of Abraham Lincoln and a link to "en.wikipedia.org". The main text in the document reads:

Lincoln was
the 16th president
of the United States.
He was born in

At the bottom of the screen, status bar text includes: PAGE 1 OF 1, ABOUT 14 WORDS, ENGLISH (U.S.), SAVED, and HELP IMPROVE OFFICE.

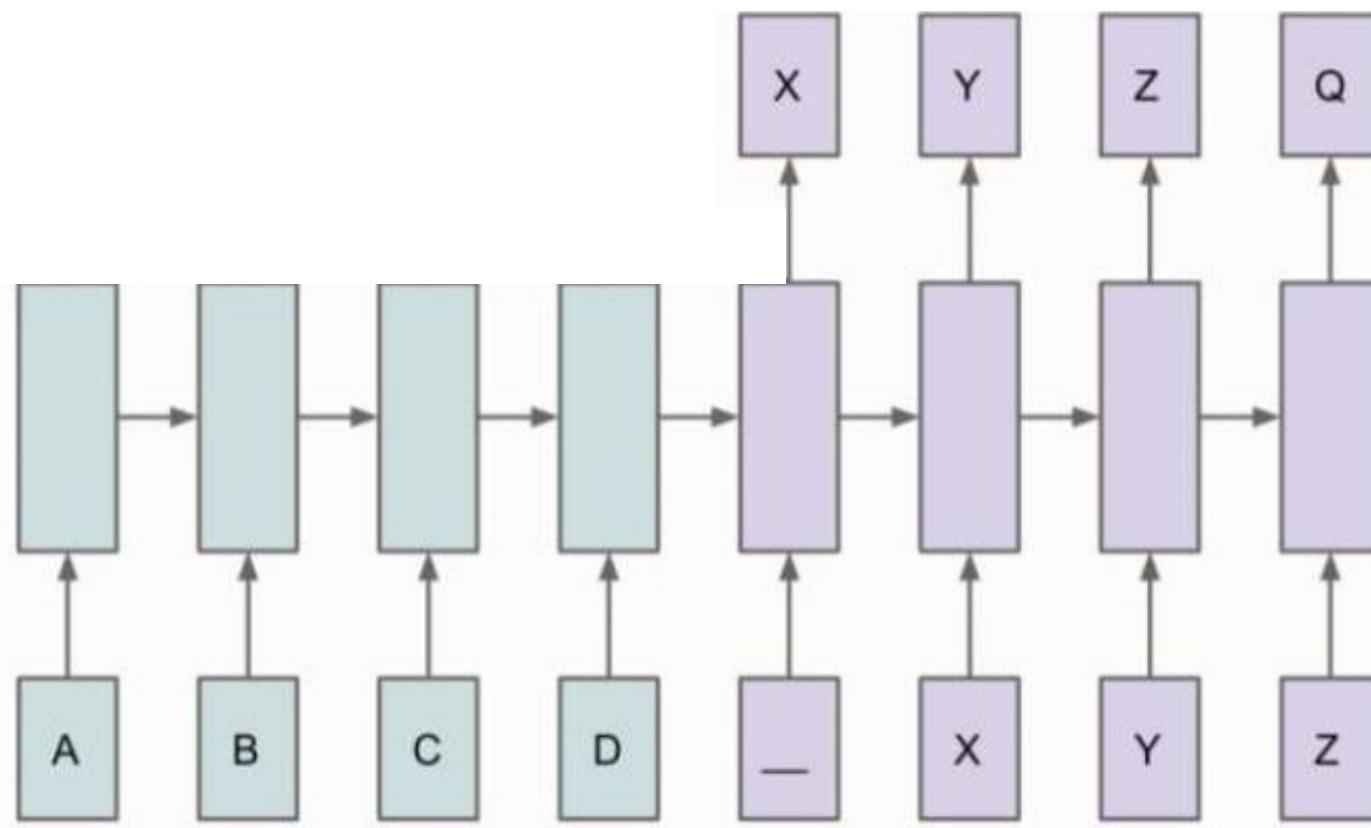
When “Lincoln” is selected, pages of a car company, movie, or the town in Nebraska will not appear¹⁸³

Towards Modeling Cognitive Functions: Memory and Attention

- seq-to-seq learning via LSTM with attention mechanism
- memory nets and neural Turing machines
- dynamic memory nets
- from seq2seq to seq2struct and to struct2struct

Seq-2-Seq Learning for Machine Translation

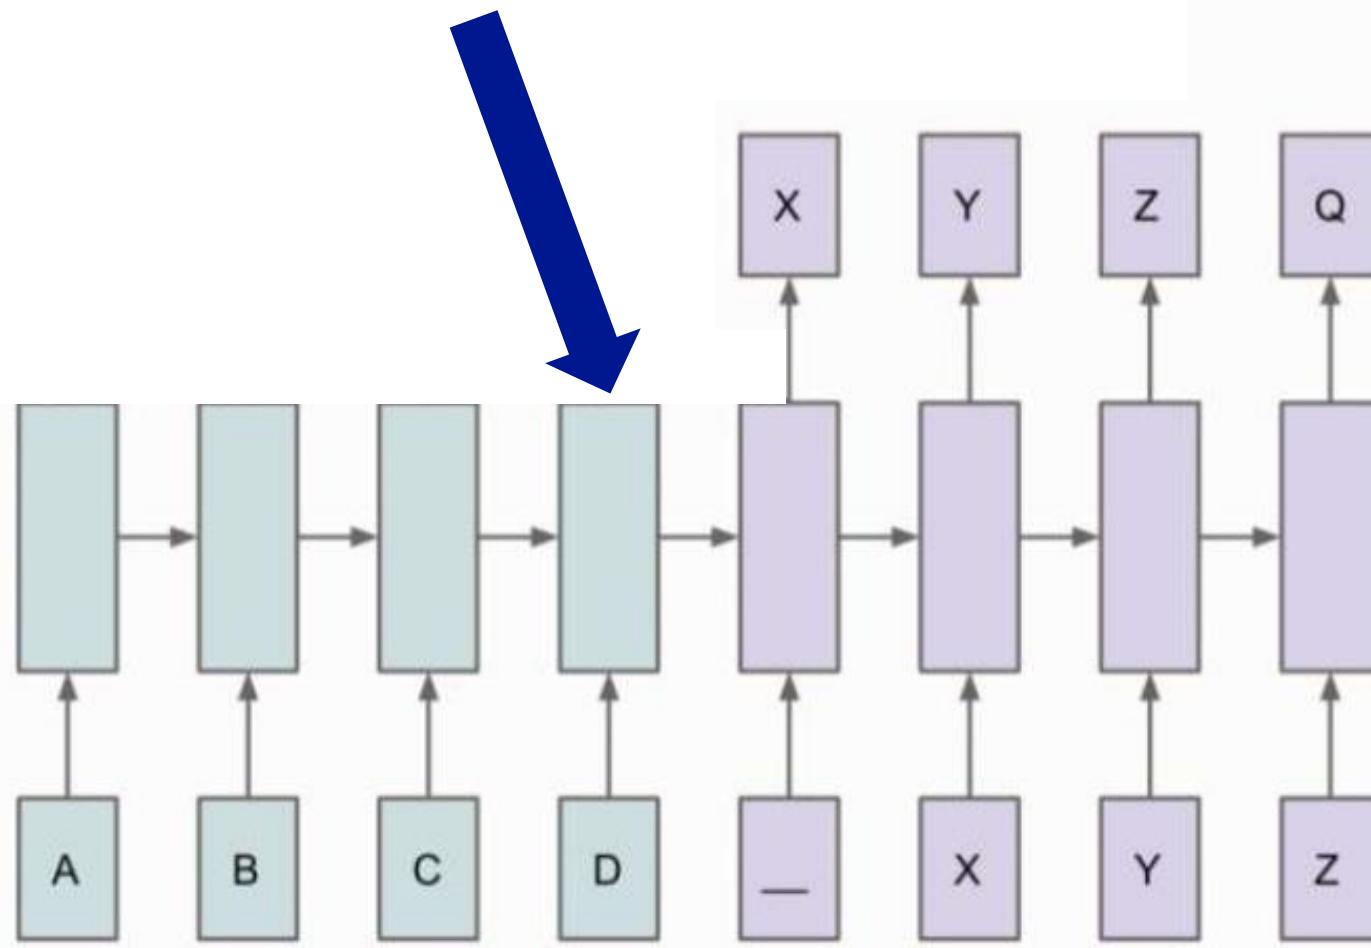
[Sutskever, Vinyals, Le, NIPS, 2014]



Seq-2-Seq Learning for Machine Translation

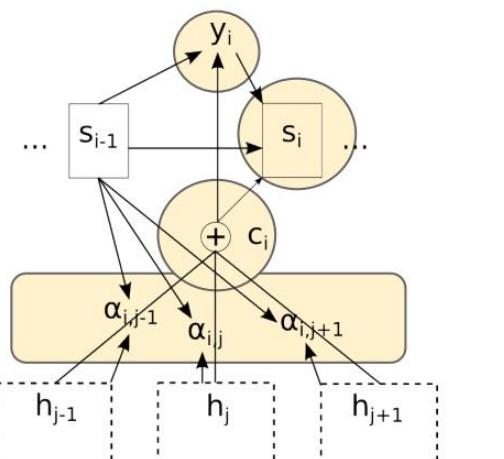
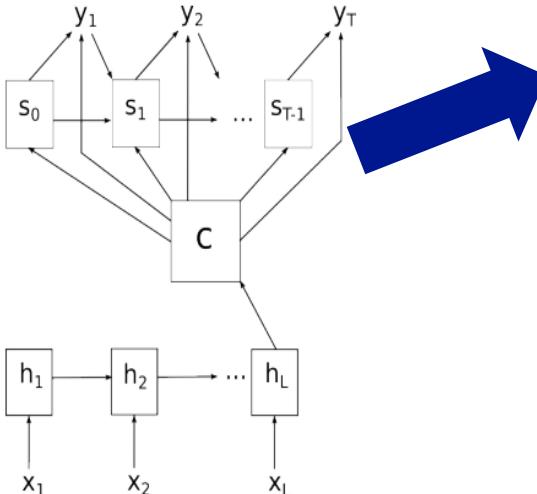
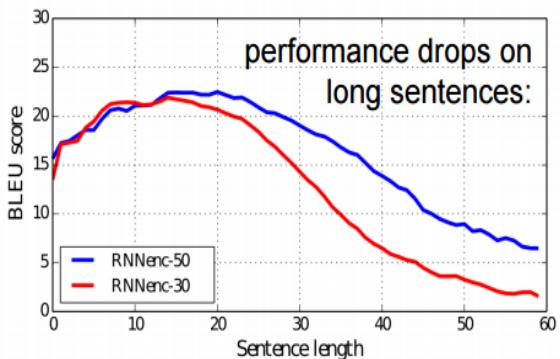
[Sutskever, Vinyals, Le, NIPS, 2014]

“thought vector” for src language



RNN Encoder-Decoder: Issues

- has to remember the whole sentence
- fixed size representation can be the bottleneck
- humans do it differently



"attention" mechanism
used for "softly" selecting
relevant input portion from
memory

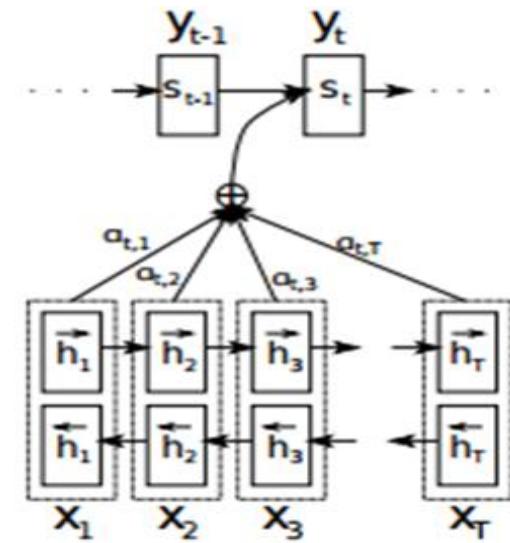


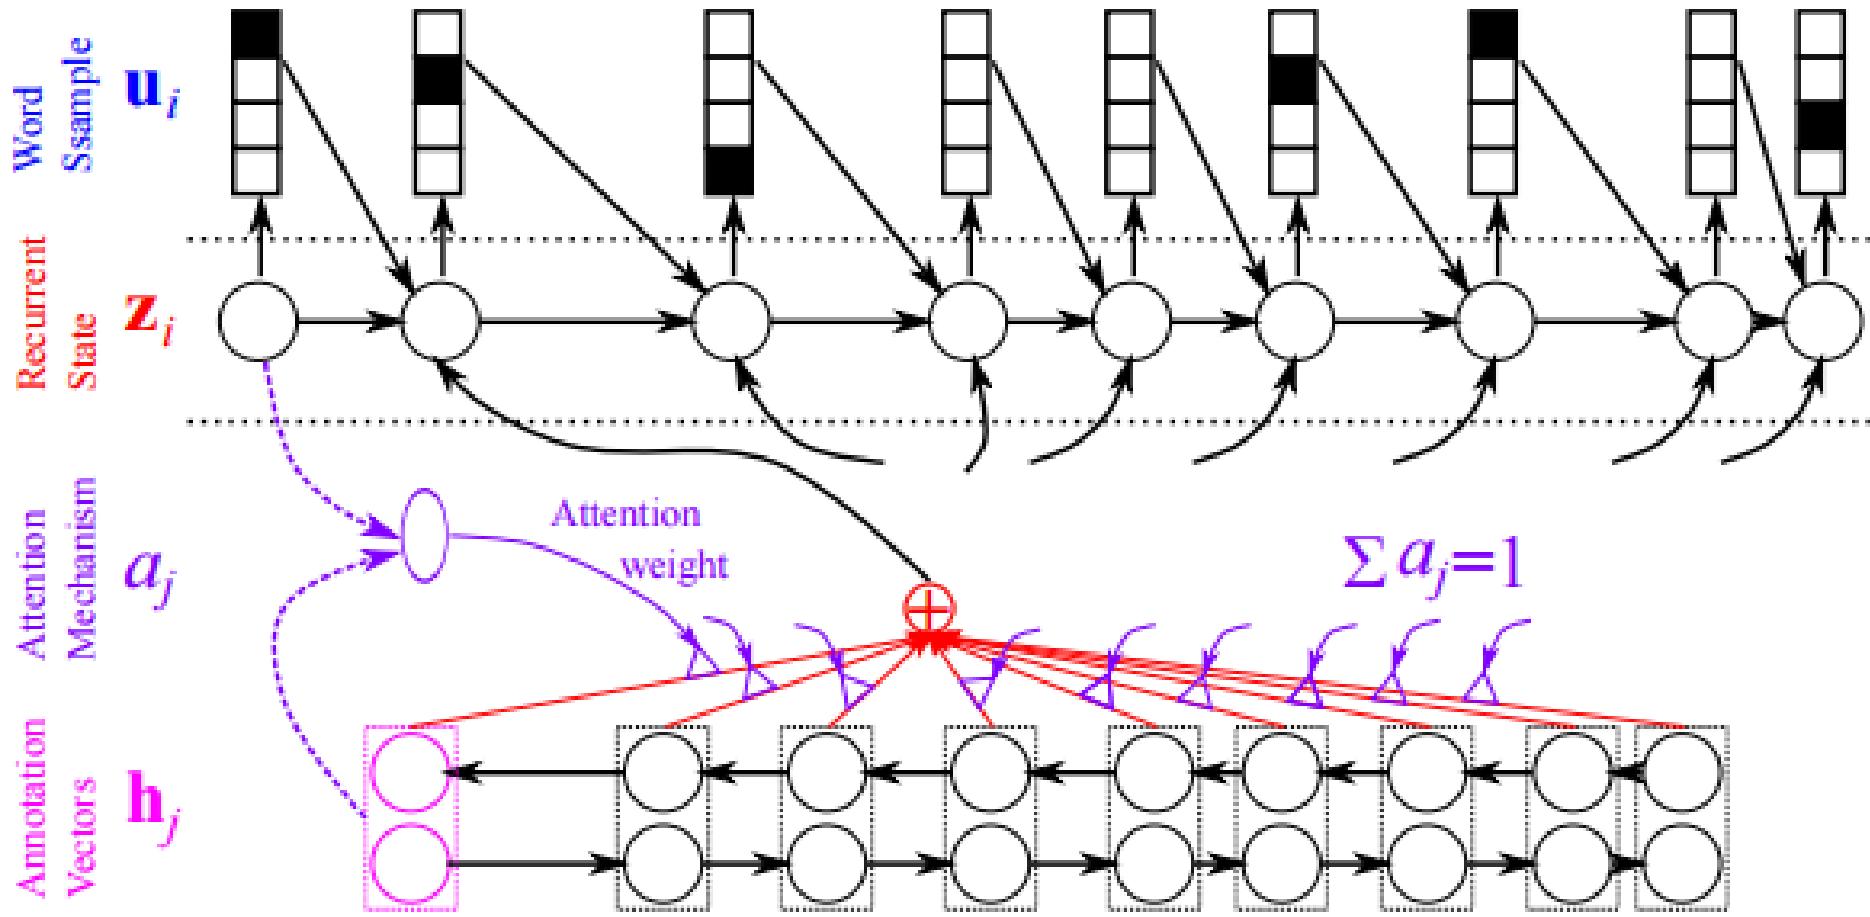
Figure 1: The graphical illustration of the proposed model trying to generate the t -th target word y_t given a source sentence (x_1, x_2, \dots, x_T) .



The full picture where “attention” is situated

(slide from Y. Bengio)

$$f = (\text{La, croissance, économique, s'est, ralenti, ces, dernières, années, .})$$

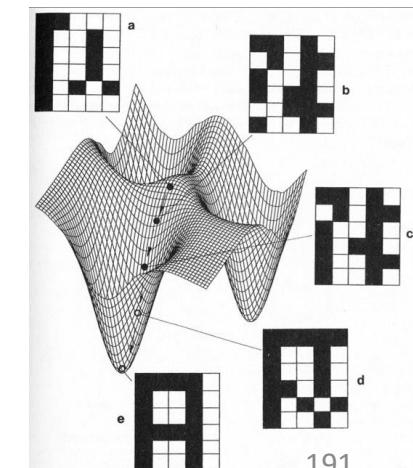
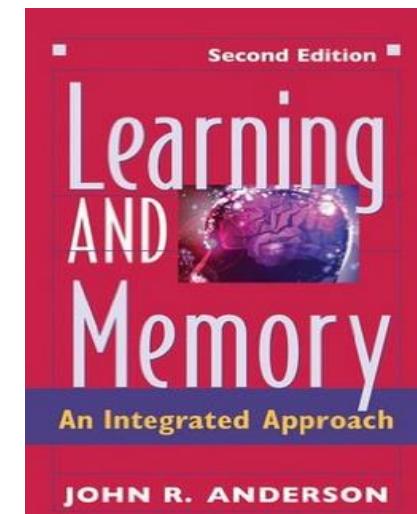
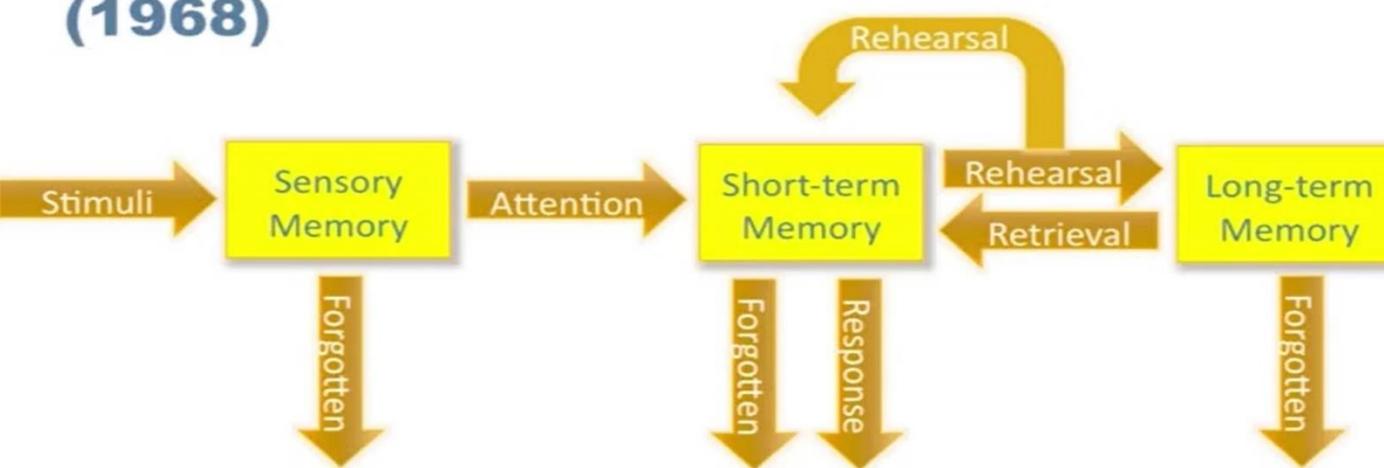


$$e = (\text{Economic, growth, has, slowed, down, in, recent, years, .})$$

Popular theories of human memory/attention

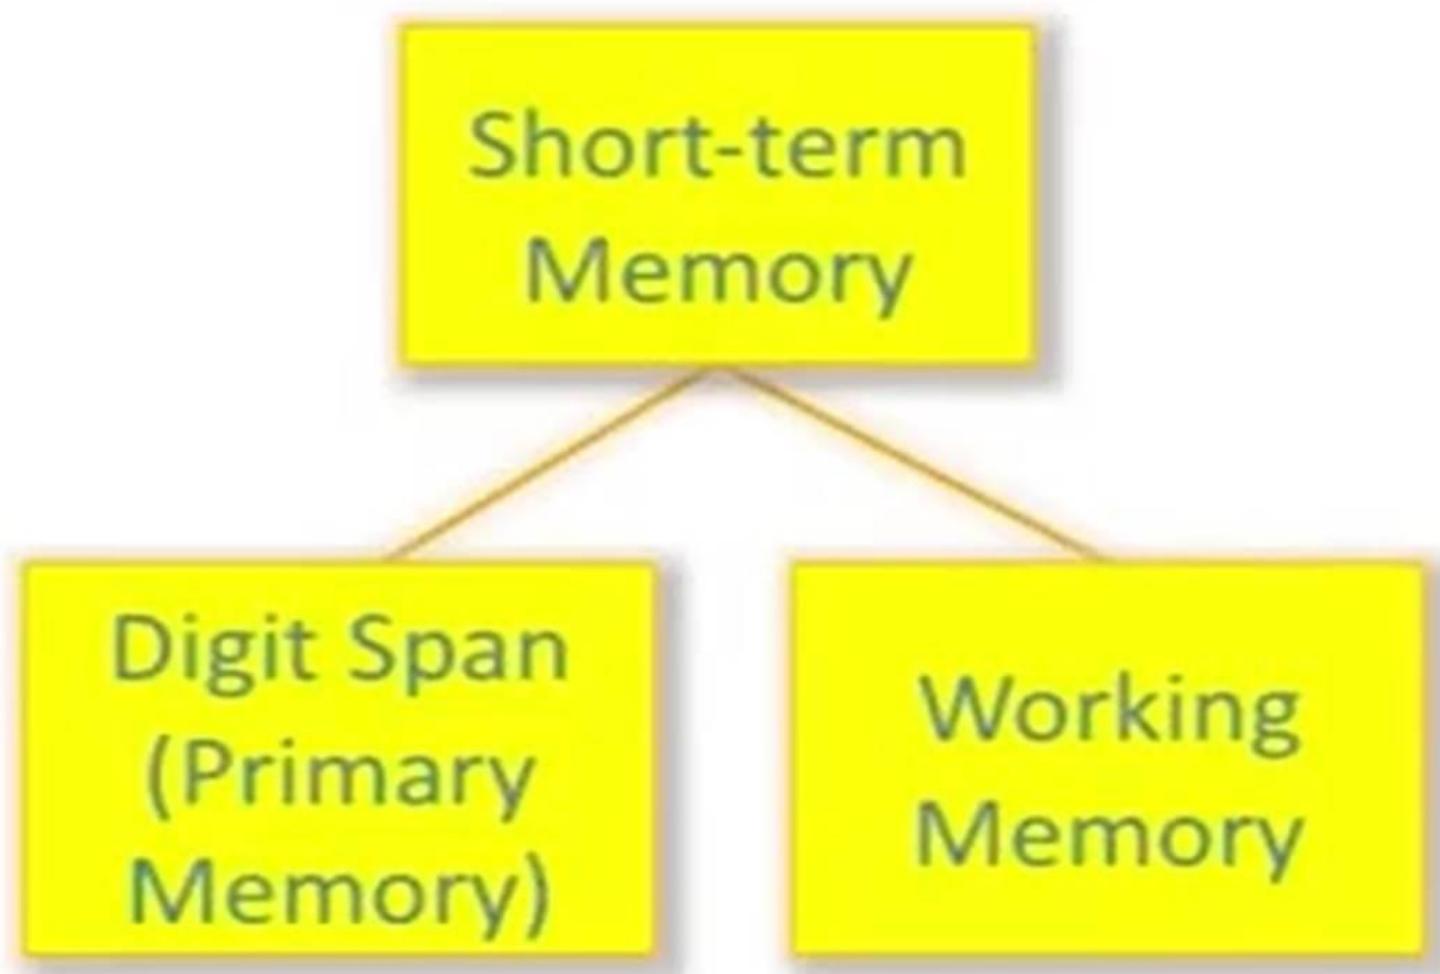
The attention and memory models discussed so far are far from human memory/attention Mechanisms (https://en.wikipedia.org/wiki/Atkinson%20Shiffrin_memory_model):

Atkinson and Shiffrin (1968)



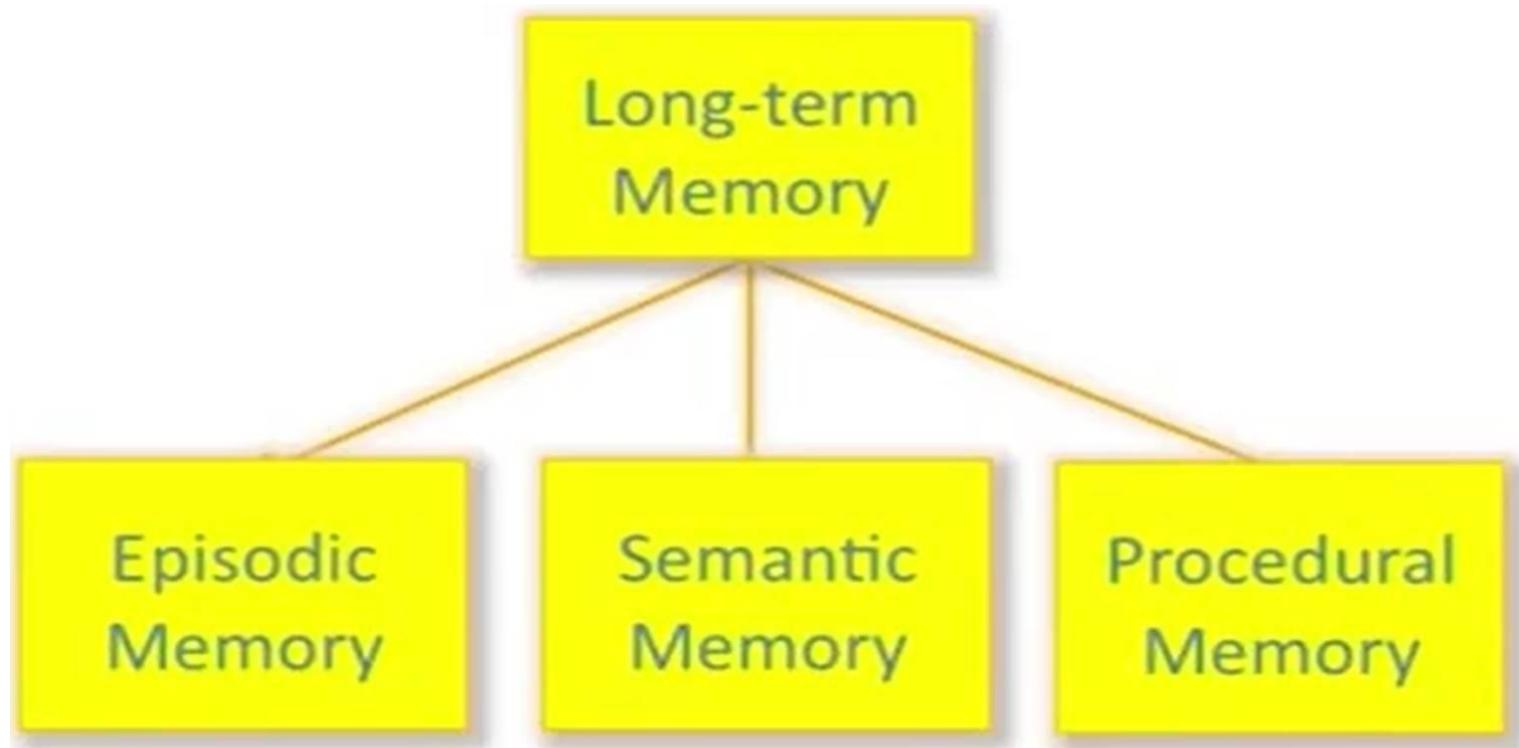
Hopfield nets store (associative) memories as attractors of the dynamic network

LSTM mainly models short-term memory



LSTM does not model long-term memory well

- LSTM makes short-term memory lasting via a simple “unit-loop” mechanism, very different from long-term memory in human cognition
- Review of a very recent modeling study on episodic and semantic memories, extending the basic LSTM formulation





TYPES OF LONG-TERM MEMORY

Episodic Memory	"When was the last time you rode a bicycle?"	Internal Diary	Retrieval by conscious contextual cues
Semantic Memory	"What is a bicycle?"	Internal Encyclopedia	Retrieval by conscious conceptual cues
Procedural Memory	Bike-riding Skill	Internal Computer Program	Unconscious automatic play-back

Slides from: Coursera, 2014

Ask Me Anything: Dynamic Memory Networks for Natural Language Processing

Ankit Kumar Ozan Irsoy Jonathan Su James Bradbury Robert English
Brian Pierce Peter Ondruska Mohit Iyyer Ishaan Gulrajani Richard Socher
firstname@metamind.io
MetaMind
Palo Alto, CA

Abstract

Most tasks in natural language processing can be cast into question answering (QA) problems over language input. We introduce the dynamic memory network (DMN), a unified neural network framework which processes input sequences and questions, forms semantic and episodic memories, and generates relevant answers. Questions trigger an iterative attention process which allows the model to condition its attention on the result of previous iterations. These results are then reasoned over in a hierarchical recurrent sequence model to generate answers.

Going beyond L-STM --- towards more realist long-term memory (episodic & semantic)

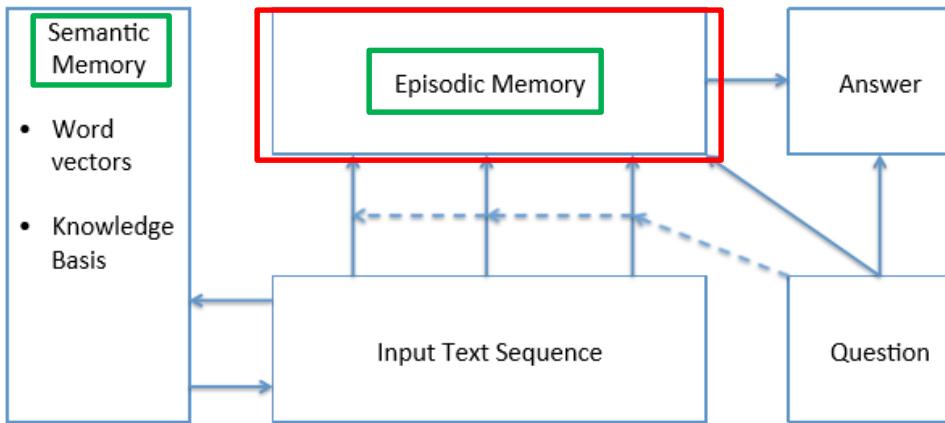


Figure 2: Overview of DMN modules. Communication between them is indicated by arrows and uses only vector representations. Questions trigger gates which allow vectors for certain input words or sentences to be given to the episodic memory module. The final state of the episodic memory is the input to the answer module.

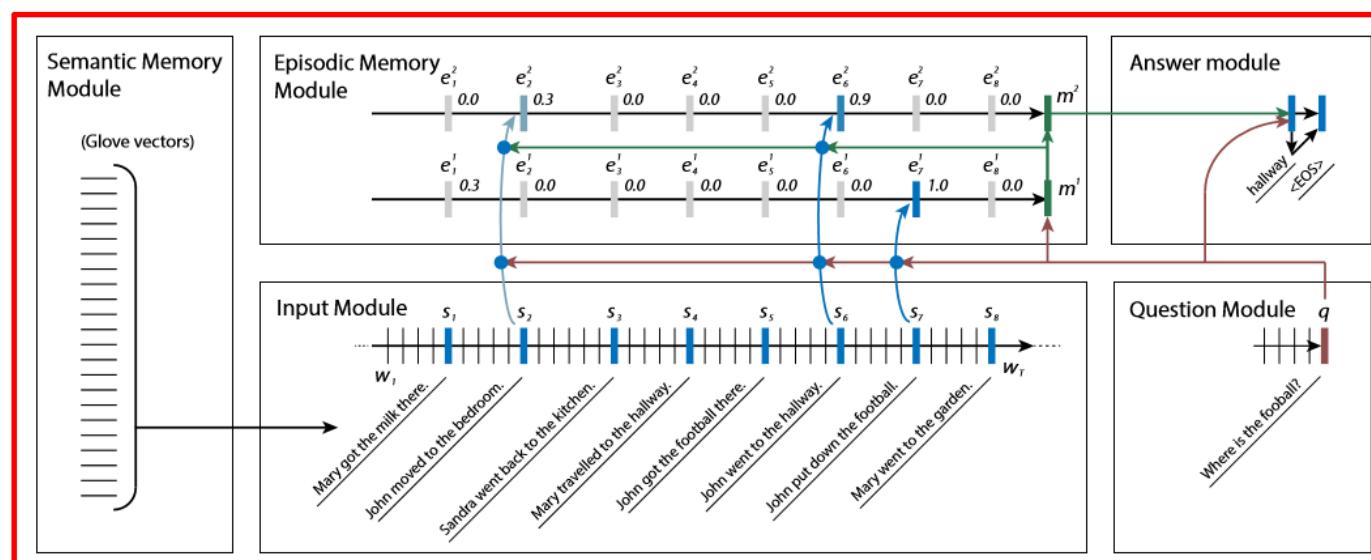


Figure 3: Real example of an input sentence sequence and the attention gates that are triggered by a specific question. Gate values g_t^i are shown above the corresponding vectors. The gates change with each search over inputs. We do not draw connections for gates that are close to zero. See Section 4.1 for details on the dataset that this example comes from.

Towards Modeling Cognitive Functions: Memory and Attention

- seq-to-seq learning via LSTM with attention mechanism
- memory nets and neural Turing machines
- dynamic memory nets
- from seq2seq to seq2struct & to struct2struct**

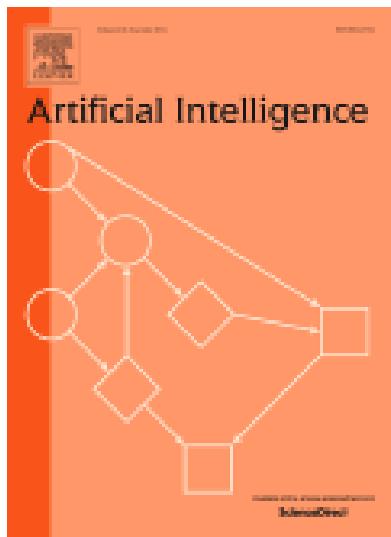
Review: embedding in the form of “flat” vectors

- A linguistic or physical entity or a simple “relation”

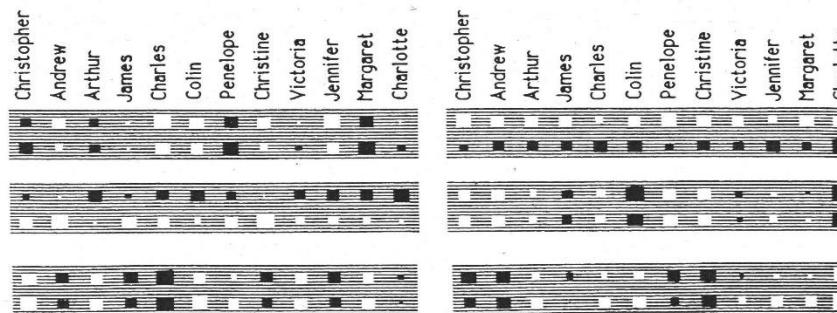


mapping via distributed representations by NN

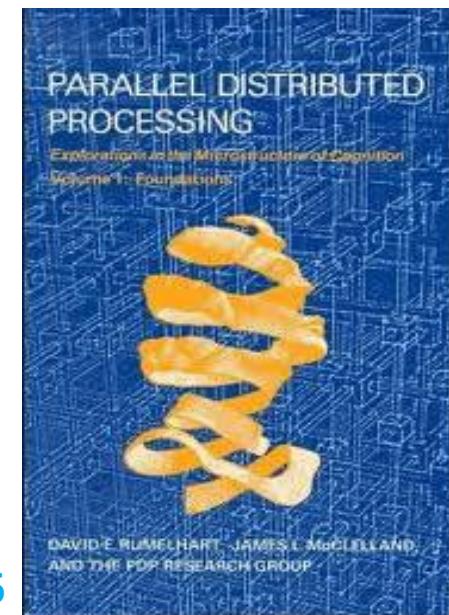
A low-dim continuous-space vector or **embedding**



Special Issue, vol. 46 (1990)
Connectionist Symbol Processing
(4 articles)



PDP book, 1986



Extension: “flat” vectors → structures (tree/graph)

- Structured embedding vectors via tensor-product rep.



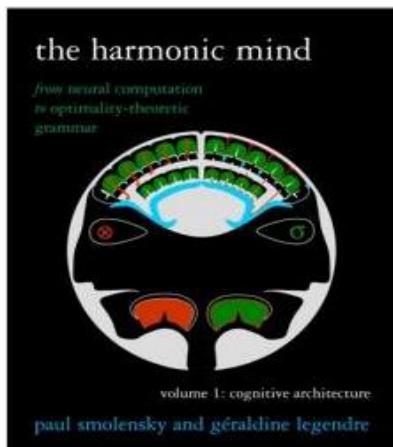
symbolic semantic parse tree (complex relation)

Then, reasoning in symbolic-space (traditional AI) can be beautifully carried out in the continuous-space in human cognitive and neural-net (i.e., connectionist) terms

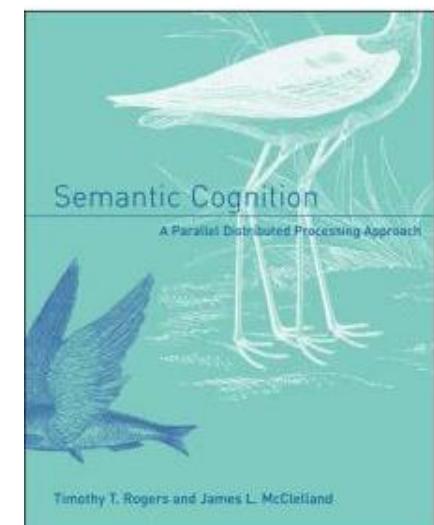
Smolensky & Legendre: The Harmonic Mind, MIT Press, 2006

From Neural Computation to Optimality-Theoretic Grammar

Volume I: Cognitive Architecture; Volume 2: Linguistic Implications



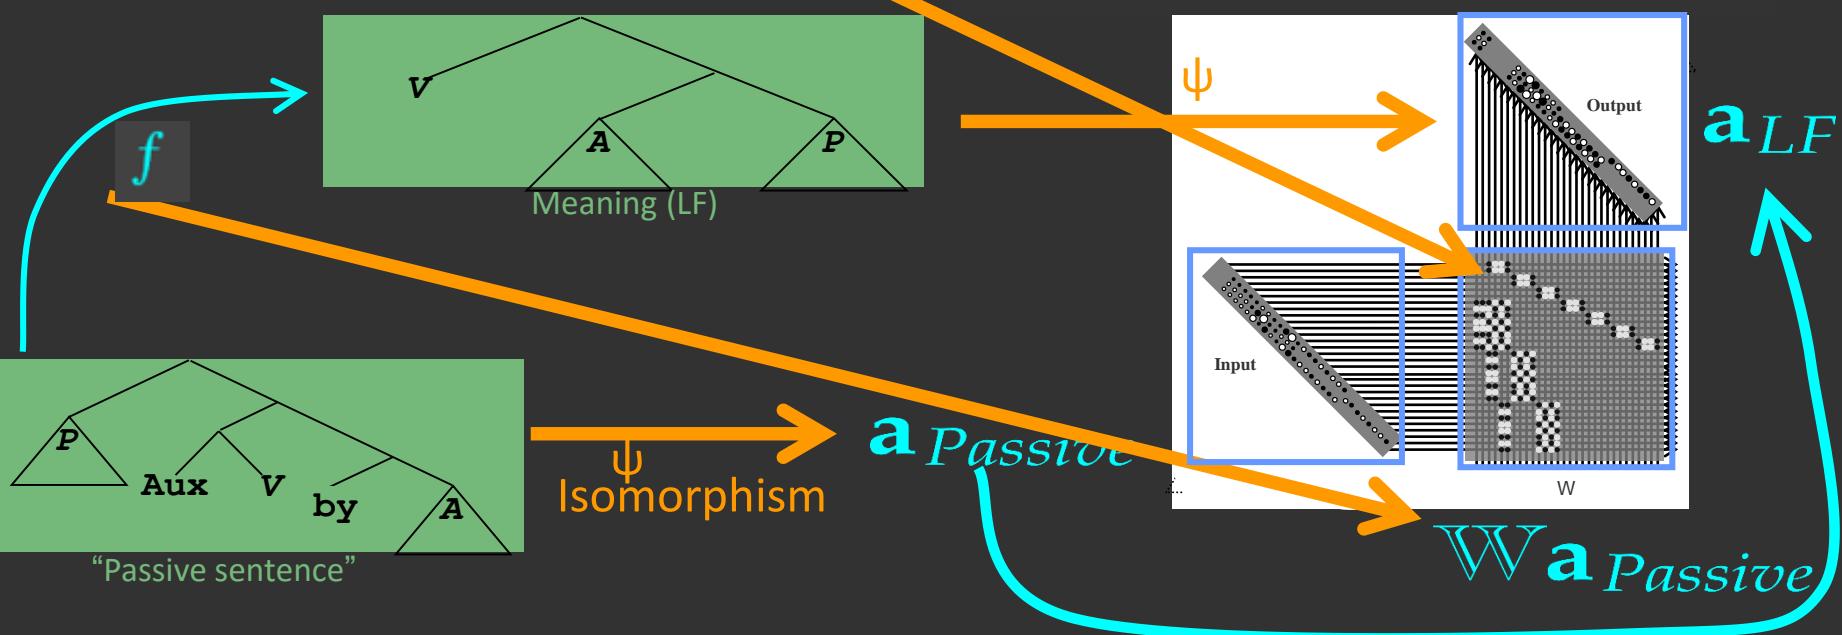
Rogers & McClelland
Semantic Cognition
MIT Press, 2006



Few leaders are admired by George Bush \xrightarrow{f} admire(George Bush, few leaders)

$$f(s) = \text{cons}(\text{ex}_1(\text{ex}_0(\text{ex}_1(s))), \text{cons}(\text{ex}_1(\text{ex}_1(\text{ex}_1(s))), \text{ex}_0(s)))$$

$$W = W_{\text{cons}_0}[W_{\text{ex}_1}W_{\text{ex}_0}W_{\text{ex}_1}] + \\ W_{\text{cons}_1}[W_{\text{cons}_0}(W_{\text{ex}_1}W_{\text{ex}_1}W_{\text{ex}_1}) + W_{\text{cons}_1}(W_{\text{ex}_0})]$$



Summary & Perspective

- Speech recognition is the first success example of deep learning at industry scale
- Deep learning is very effective in speech recognition, speech translation ([Skype Translator](#)), image recognition ([Onedrive Image tagging](#)), image captioning, language understanding ([Cortana](#)), semantic intelligence, multimodal and multitask learning, web search, advertising, entity search ([Insights for MS Office](#)), user and business activity prediction, etc.
- Enabling factors:
 - Big datasets for training deep models
 - Powerful GPGPU computing
 - Innovations in deep learning architectures and algorithms
 - How to discover distant supervision signals free from human labeling
 - How to build deep learning systems grounded on exploiting such “smart” signals (example: DSSM)
 - ...

Summary & Perspective

- Speech recognition: **all** low-hanging fruits are taken
 - i.e. **more** innovation and hard work needed than before
- Image recognition: **most** low-hanging fruits are taken
- Natural Language: does not seem there is much low-hanging fruit there
 - i.e. **even more** innovation and hard work needed than before
- Big data analytics (e.g. user behavior, business activities, etc):
 - A new frontier
- Small data: deep learning may still win (e.g. 2012 Kaggle's drug discovery)
- Perceptual data: deep learning methods always win, and win big
- Be careful: data with adversarial nature; data with odd variability

Issues for “Near” Future of Deep Learning

- For perceptual tasks (e.g. speech, image/video, gesture, etc.)
 - With supervised data: what will be the limit for growing accuracy wrt increasing amounts of labeled data?
 - Beyond this limit or when labeled data are exhausted or non-economical to collect, will novel and effective unsupervised deep learning emerge and what will they be (e.g. deep generative models)?
 - Many new innovations are to come, likely in the area of unsupervised learning
- For cognitive tasks (e.g. natural language, reasoning, knowledge, decision making, etc.)
 - Will supervised deep learning (e.g. MT) beat the non-deep-learning state of the art like speech/image recognition?
 - How to distill/exploit “distant” supervision signals for supervised deep learning?
 - Will dense vector embedding be sufficient for language? Do we really need to directly encode and recover syntactic/semantic structure of language?
 - Even more new innovations are to come, likely in the area of new architectures and learning methods pertaining to distant supervised learning

The Future of Deep Learning

- Continued rapid progress in language processing methods and applications by both industry and academia
 - From image to video processing/understanding
 - From supervised learning (huge success already) to **unsupervised** learning (not much success yet but ideas abound)
 - From perception to cognition
 - More exploration of **attention** modeling
 - Combine representation learning with complex **knowledge** extraction & **reasoning**
 - Modeling human **memory** functions more faithfully
 - Learning to **act and control** (deep reinforcement learning)
 - Successes in business applications will propel more rapid advances in deep learning (positive feedbacks)

Additional References

- Auli, M., Galley, M., Quirk, C. and Zweig, G., 2013. Joint language and translation modeling with recurrent neural networks. In EMNLP.
- Auli, M., and Gao, J., 2014. Decoder integration and expected bleu training for recurrent neural network language models. In ACL.
- Baker, J., Li Deng, Jim Glass, S. Khudanpur, C.-H. Lee, N. Morgan, and D. O'Shughnessy, Research Developments and Directions in Speech Recognition and Understanding, Part 1, in *IEEE Signal Processing Magazine*, vol. 26, no. 3, pp. 75-80, 2009.
- Baker, J., Li Deng, S. Khudanpur, C.-H. Lee, J. Glass, and N. Morgan, Updated MINDS Report on Speech Recognition and Understanding, in *IEEE Signal Processing Magazine*, vol. 26, no. 4, July 2009.
- Bengio, Y., 2009. Learning deep architectures for AI. *Foundamental Trends Machine Learning*, vol. 2.
- Bengio, Y., Courville, A., and Vincent, P. 2013. Representation learning: A review and new perspectives. *IEEE Trans. PAMI*, vol. 38, pp. 1798-1828.
- Bengio, Y., Ducharme, R., and Vincent, P., 2000. A Neural Probabilistic Language Model, in NIPS.
- Bridle, J., L. Deng, J. Picone, H. Richards, J. Ma, T. Kamm, M. Schuster, S. Pike, R. Reagan, An investigation of segmental hidden dynamic models of speech coarticulation for automatic speech recognition. *Final Report for 1998 Workshop on Langauge Engineering, CLSP* (Johns Hopkins, 1998).
- Collobert, R., Weston, J., Bottou, L., Karlen, M., Kavukcuoglu, K., and Kuksa, P., 2011. Natural language processing (almost) from scratch. in *JMLR*, vol. 12.
- Dahl, G., Yu, D., Deng, L., and Acero, 2012. A. Context-dependent, pre-trained deep neural networks for large vocabulary speech recognition, *IEEE Trans. Audio, Speech, & Language Proc.*, Vol. 20 (1), pp. 30-42.
- Deerwester, S., Dumais, S. T., Furnas, G. W., Landauer, T., and Harshman, R. 1990. Indexing by latent semantic analysis. *J. American Society for Information Science*, 41(6): 391-407
- Deng, L. A dynamic, feature-based approach to the interface between phonology & phonetics for speech modeling and recognition, *Speech Communication*, vol. 24, no. 4, pp. 299-323, 1998.
- Deng, L. Computational Models for Speech Production, in *Computational Models of Speech Pattern Processing*, pp. 199-213, Springer Verlag, 1999.
- Deng, L. Switching Dynamic System Models for Speech Articulation and Acoustics, in *Mathematical Foundations of Speech and Language Processing*, vol. 138, pp. 115-134, Springer Verlag, 2003.
- Deng, L., K. Hassanein, M. Elmasry, Analysis of the correlation structure for a neural predictive model with application to speech recognition. *Neural Networks*. 7(2), 331-339, 1994.
- Deng L., G. Ramsay, and D. Sun, Production models as a structural basis for automatic speech recognition," Speech Communication (special issue on speech production modeling), in *Speech Communication*, vol. 22, no. 2, pp. 93-112, August 1997.
- Deng L. and J. Ma, Spontaneous Speech Recognition Using a Statistical Coarticulatory Model for the Vocal Tract Resonance Dynamics, *Journal of the Acoustical Society of America*, 2000.
- Deng, L. and O'Shaughnessy, O. *Speech Processing—A Dynamic and Optimization-Oriented Approach*, Marcel Dekker, New York, 2003
- Deng L. and Yu, L. DEEP LEARNING: Methods and Applications, NOW Publishings, 2014.
- Deng L. and Xiao Li, Machine Learning Paradigms for Speech Recognition: An Overview, in *IEEE Transactions on Audio, Speech, and Language Processing*, vol. 21, no. 5, pp. 1060-1089, May 2013.
- Deng L. and Xuedong Huang, Challenges in Adopting Speech Recognition, in *Communications of the ACM*, vol. 47, no. 1, pp. 11-13, January 2004.
- Deng, L., Seltzer, M., Yu, D., Acero, A., Mohamed, A., and Hinton, G., Binary Coding of Speech Spectrograms Using a Deep Auto-encoder, *Interspeech*, 2010.
- Deng, L., Tur, G., He, X., and Hakkani-Tur, D. 2012. Use of kernel deep convex networks and end-to-end learning for spoken language understanding, *Proc. IEEE Workshop on Spoken Language Technologies*.
- Deng, L., Yu, D. and Acero, A. 2006. Structured speech modeling, *IEEE Trans. on Audio, Speech and Language Processing*, vol. 14, no. 5, pp. 1492-1504.
- Deng, L., Geoffrey Hinton, and Brian Kingsbury, New types of deep neural network learning for speech recognition and related applications: An overview, *ICASSP*, 2013.
- Deng L. and Dong Yu, Use of Differential Cepstra as Acoustic Features in Hidden Trajectory Modeling for Phonetic Recognition, *ICASSP*, 2007.
- Deng, L., Xiaodong He, and Jianfeng Gao, Deep Stacking Networks for Information Retrieval, *ICASSP*, 2013.
- Deng L. and Dong Yu, Deep Convex Network: A Scalable Architecture for Speech Pattern Classification, in *Interspeech*, 2011.
- Deng, L., Yu, D., and Platt, J. 2012. Scalable stacking and learning for building deep architectures, *Proc. ICASSP*.
- Deoras, A., and Sarikaya, R., 2013. Deep belief network based semantic taggers for spoken language understanding, in *INTERSPEECH*.
- Devlin, J., Zbib, R., Huang, Z., Lamar, T., Schwartz, R., and Makhoul, J., 2014. Fast and Robust Neural Network Joint Models for Statistical Machine Translation, *ACL*.

Additional References

- Frome, A., Corrado, G., Shlens, J., Bengio, S., Dean, J., Ranzato, M., and Mikolov, T., 2013. DeViSE: A Deep Visual-Semantic Embedding Model, Proc. NIPS.
- Gao, J., He, X., Yih, W-t., and Deng, L. 2014a. Learning continuous phrase representations for translation modeling. In ACL.
- Gao, J., He, X., and Nie, J-Y. 2010. Clickthrough-based translation models for web search: from word models to phrase models. In CIKM.
- Gao, J., Pantel, P., Gamon, M., He, X., and Deng, L. 2014. Modeling interestingness with deep neural networks. In EMNLP
- Gao, J., Toutanova, K., Yih., W-T. 2011. Clickthrough-based latent semantic models for web search. In SIGIR.
- Gao, J., Yuan, W., Li, X., Deng, K., and Nie, J-Y. 2009. Smoothing clickthrough data for web search ranking. In SIGIR.
- Gao, J., and He, X. 2013. Training MRF-based translation models using gradient ascent. In NAACL-HLT.
- Graves, A., Jaitly, N., and Mohamed, A., 2013a. Hybrid speech recognition with deep bidirectional LSTM, Proc. ASRU.
- Graves, A., Mohamed, A., and Hinton, G., 2013. Speech recognition with deep recurrent neural networks, Proc. ICASSP.
- He, X. and Deng, L., 2013. Speech-Centric Information Processing: An Optimization-Oriented Approach, in Proceedings of the IEEE.
- He, X. and Deng, L., 2012. Maximum Expected BLEU Training of Phrase and Lexicon Translation Models , ACL.
- He, X., Deng, L., and Chou, W., 2008. Discriminative learning in sequential pattern recognition, Sept. IEEE Sig. Proc. Mag.
- Hinton, G., Deng, L., Yu, D., Dahl, G., Mohamed, A., Jaitly, N., Senior, A., Vanhoucke, V., Nguyen, P., Sainath, T., and Kingsbury, B., 2012. Deep Neural Networks for Acoustic Modeling in Speech Recognition, IEEE Signal Processing Magazine, vol. 29, no. 6, pp. 82-97.
- Hinton, G., and Salakhutdinov, R., 2010. Discovering binary codes for documents by learning deep generative models. Topics in Cognitive Science.
- Hu, Y., Auli, M., Gao, Q., and Gao, J. 2014. Minimum translation modeling with recurrent neural networks. In EACL.
- Huang, E., Socher, R., Manning, C, and Ng, A. 2012. Improving word representations via global context and multiple word prototypes, Proc. ACL.
- Huang, P., He, X., Gao, J., Deng, L., Acero, A., and Heck, L. 2013. Learning deep structured semantic models for web search using clickthrough data. In CIKM.
- Hutchinson, B., Deng, L., and Yu, D., 2012. A deep architecture with bilinear modeling of hidden representations: Applications to phonetic recognition, Proc. ICASSP.
- Hutchinson, B., Deng, L., and Yu, D., 2013. Tensor deep stacking networks, IEEE Trans. Pattern Analysis and Machine Intelligence, vol. 35, pp. 1944 - 1957.
- Kiros, R., Zemel, R., and Salakhutdinov, R. 2013. Multimodal Neural Language Models, Proc. NIPS Deep Learning Workshop.
- Koehn, P. 2009. Statistical Machine Translation. Cambridge University Press.
- Krizhevsky, A., Sutskever, I, and Hinton, G., 2012. ImageNet Classification with Deep Convolutional Neural Networks, NIPS.
- Le, H-S., Oparin, I., Allauzen, A., Gauvain, J-L., Yvon, F., 2013. Structured output layer neural network language models for speech recognition, IEEE Transactions on Audio, Speech and Language Processing.
- LeCun, Y., Bottou, L., Bengio, Y., and Haffner, P. 1998. Gradient-based learning applied to document recognition, Proceedings of the IEEE, Vol. 86, pp. 2278-2324.
- Lee, L., H. Attias, Li Deng, and P. Fieguth, A Multimodal Variational Approach to Learning and Inference in Switching State Space Models, ICASSP, 2004.
- Ma, J. and L. Deng, Target-directed mixture dynamic models for spontaneous speech recognition. IEEE Trans. Audio Speech Process. 12(1), 47–58, 2004
- Mesnil, G., He, X., Deng, L., and Bengio, Y., 2013. Investigation of Recurrent-Neural-Network Architectures and Learning Methods for Spoken Language Understanding, in Interspeech.
- Mikolov, T. 2012. Statistical Language Models based on Neural Networks, Ph.D. thesis, Brno University of Technology.
- Mikolov, T., Chen, K., Corrado, G., and Dean, J. 2013. Efficient estimation of word representations in vector space, Proc. ICLR.
- Mikolov, T., Kombrink, S., Burget, L., Cernocky, J., Khudanpur, S., 2011. Extensions of Recurrent Neural Network LM. ICASSP.
- Mikolov, T., Yih, W., Zweig, G., 2013. Linguistic Regularities in Continuous Space Word Representations. In NAACL-HLT.
- Mohamed, A., G. Dahl, G. Hinton, Acoustic modeling using deep belief networks. IEEE Trans. Audio Speech Lang. Process. 20(1), 14–22, 2012 & 2009
- Mohamed, A., Yu, D., and Deng, L. 2010. Investigation of full-sequence training of deep belief networks for speech recognition, Proc. Interspeech.
- Ngiam, J., Khosla, A., Kim, M., Nam, J., Lee, H., and Ng, A. 2011. Multimodal deep learning, Proc. ICML.

Additional References

- Picone, P., S. Pike, R. Regan, T. Kamm, J. Bridle, L. Deng, Z. Ma, H. Richards, M. Schuster, 1999. Initial evaluation of hidden dynamic models on conversational speech, ICASSP.
- Robinson, T. An application of recurrent nets to phone probability estimation. IEEE Trans. Neural Networks. **5**(2), 298–305, 1994.
- Sainath, T., Mohamed, A., Kingsbury, B., and Ramabhadran, B. 2013. Convolutional neural networks for LVCSR, Proc. ICASSP.
- Salakhutdinov R., and Hinton, G., 2007 Semantic hashing. in Proc. SIGIR Workshop Information Retrieval and Applications of Graphical Models
- Schwenk, H., Dchelotte, D., Gauvain, J-L., 2006. Continuous space language models for statistical machine translation, in COLING-ACL
- Seide, F., Li, G., and Yu, D. 2011. Conversational speech transcription using context-dependent deep neural networks, Proc. Interspeech
- Shen, Y., He, X., Gao, J., Deng, L., and Mesnil, G. 2014. Learning Semantic Representations Using Convolutional Neural Networks for Web Search, in Proceedings of WWW.
- Shen, Y., He, X., Gao, J., Deng, L., and Mesnil, G. 2014. A convolutional latent semantic model for web search. CIKM
- Shen, X., L. Deng, Maximum likelihood in statistical estimation of dynamical systems: decomposition algorithm and simulation results. Signal Process. **57**, 65–79, 1997.
- Socher, R., Huval, B., Manning, C., Ng, A., 2012. Semantic compositionality through recursive matrix-vector spaces. In EMNLP.
- Socher, R., Lin, C., Ng, A., and Manning, C. 2011. Learning continuous phrase representations and syntactic parsing with recursive neural networks, Proc. ICML.
- Socher, R., Perelygin, A., Wu, J., Chuang, J., Manning, C., Ng A., and Potts. C. 2013. Recursive Deep Models for Semantic Compositionality Over a Sentiment Treebank, Proc. EMNLP
- Song, X. He, X., Gao. J., and Deng, L. 2014. Learning Word Embedding Using the DSSM. MSR Tech Report.
- Song, Y., Wang, H., and He, X., 2014. Adapting Deep RankNet for Personalized Search. Proc. WSDM.
- Tur, G., Deng, L., Hakkani-Tur, D., and He, X., 2012. Towards Deeper Understanding Deep Convex Networks for Semantic Utterance Classification, in ICASSP.
- Togneri, R., L. Deng, Joint state and parameter estimation for a target-directed nonlinear dynamic system model. IEEE Trans. Signal Process. **51**(12), 3061–3070, 2003.
- Vinyals, O., Y. Jia, Li Deng, and Trevor Darrell, Learning with Recursive Perceptual Representations, NIPS, 2012.
- Wright, S., Kanevsky, D., Deng, L., He, X., Heigold, G., and Li, H. Optimization Algorithms and Applications for Speech and Language Processing, in IEEE Transactions on Audio, Speech, and Language Processing, vol. 21, no. 11. 2013.
- Yao, K., Zweig, G., Hwang, M-Y. , Shi, Y., Yu, D., 2013. Recurrent neural networks for language understanding, submitted to Interspeech.
- Yann, D., Tur, G., Hakkani-Tur, D., Heck, L., 2014. Zero-Shot Learning and Clustering for Semantic Utterance Classification Using Deep Learning, in ICLR.
- Yu, D. and Deng, L. Discriminative pretraining of deep neural networks, US Patent, 2013.
- Yu, D., Deng, L., G. Dahl, Roles of pre-training and fine-tuning in context-dependent DBN-HMMs for real-world speech recognition, *NIPS Workshop* 2010.

Thank You